

ELLA: Exploration through Learned Language Abstraction

Suvir Mirchandani Siddharth Karamcheti Dorsa Sadigh

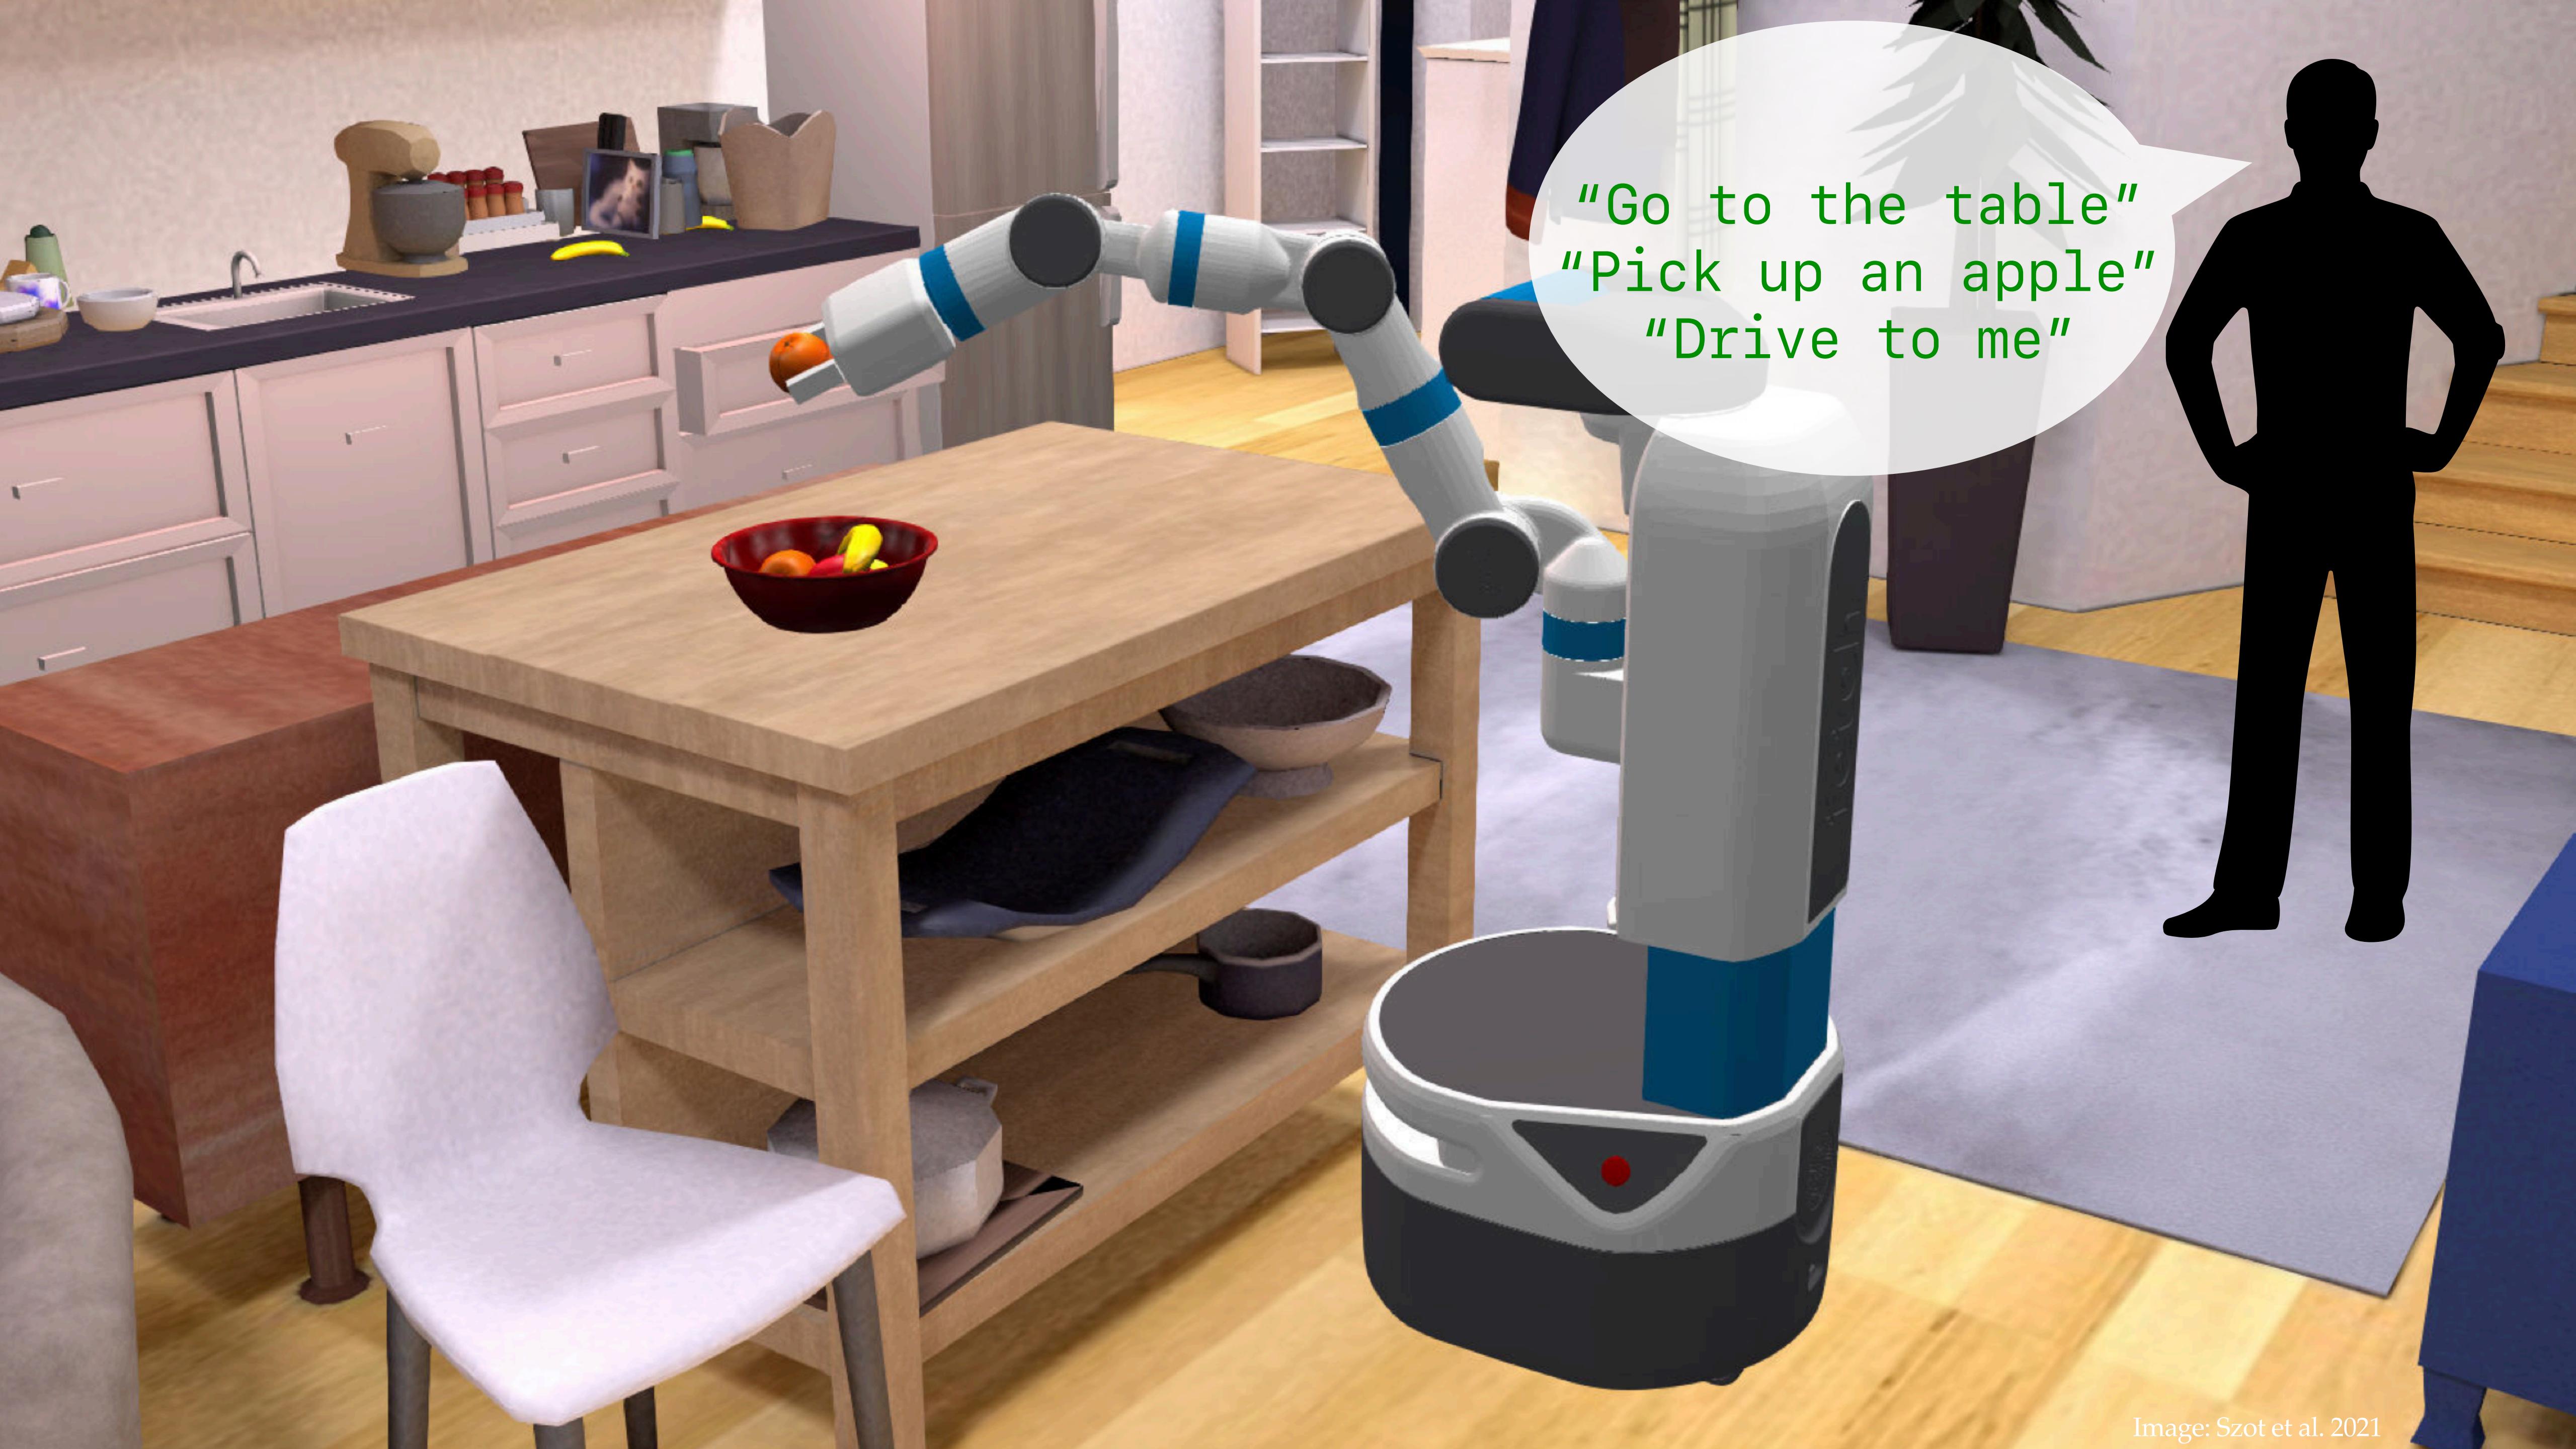




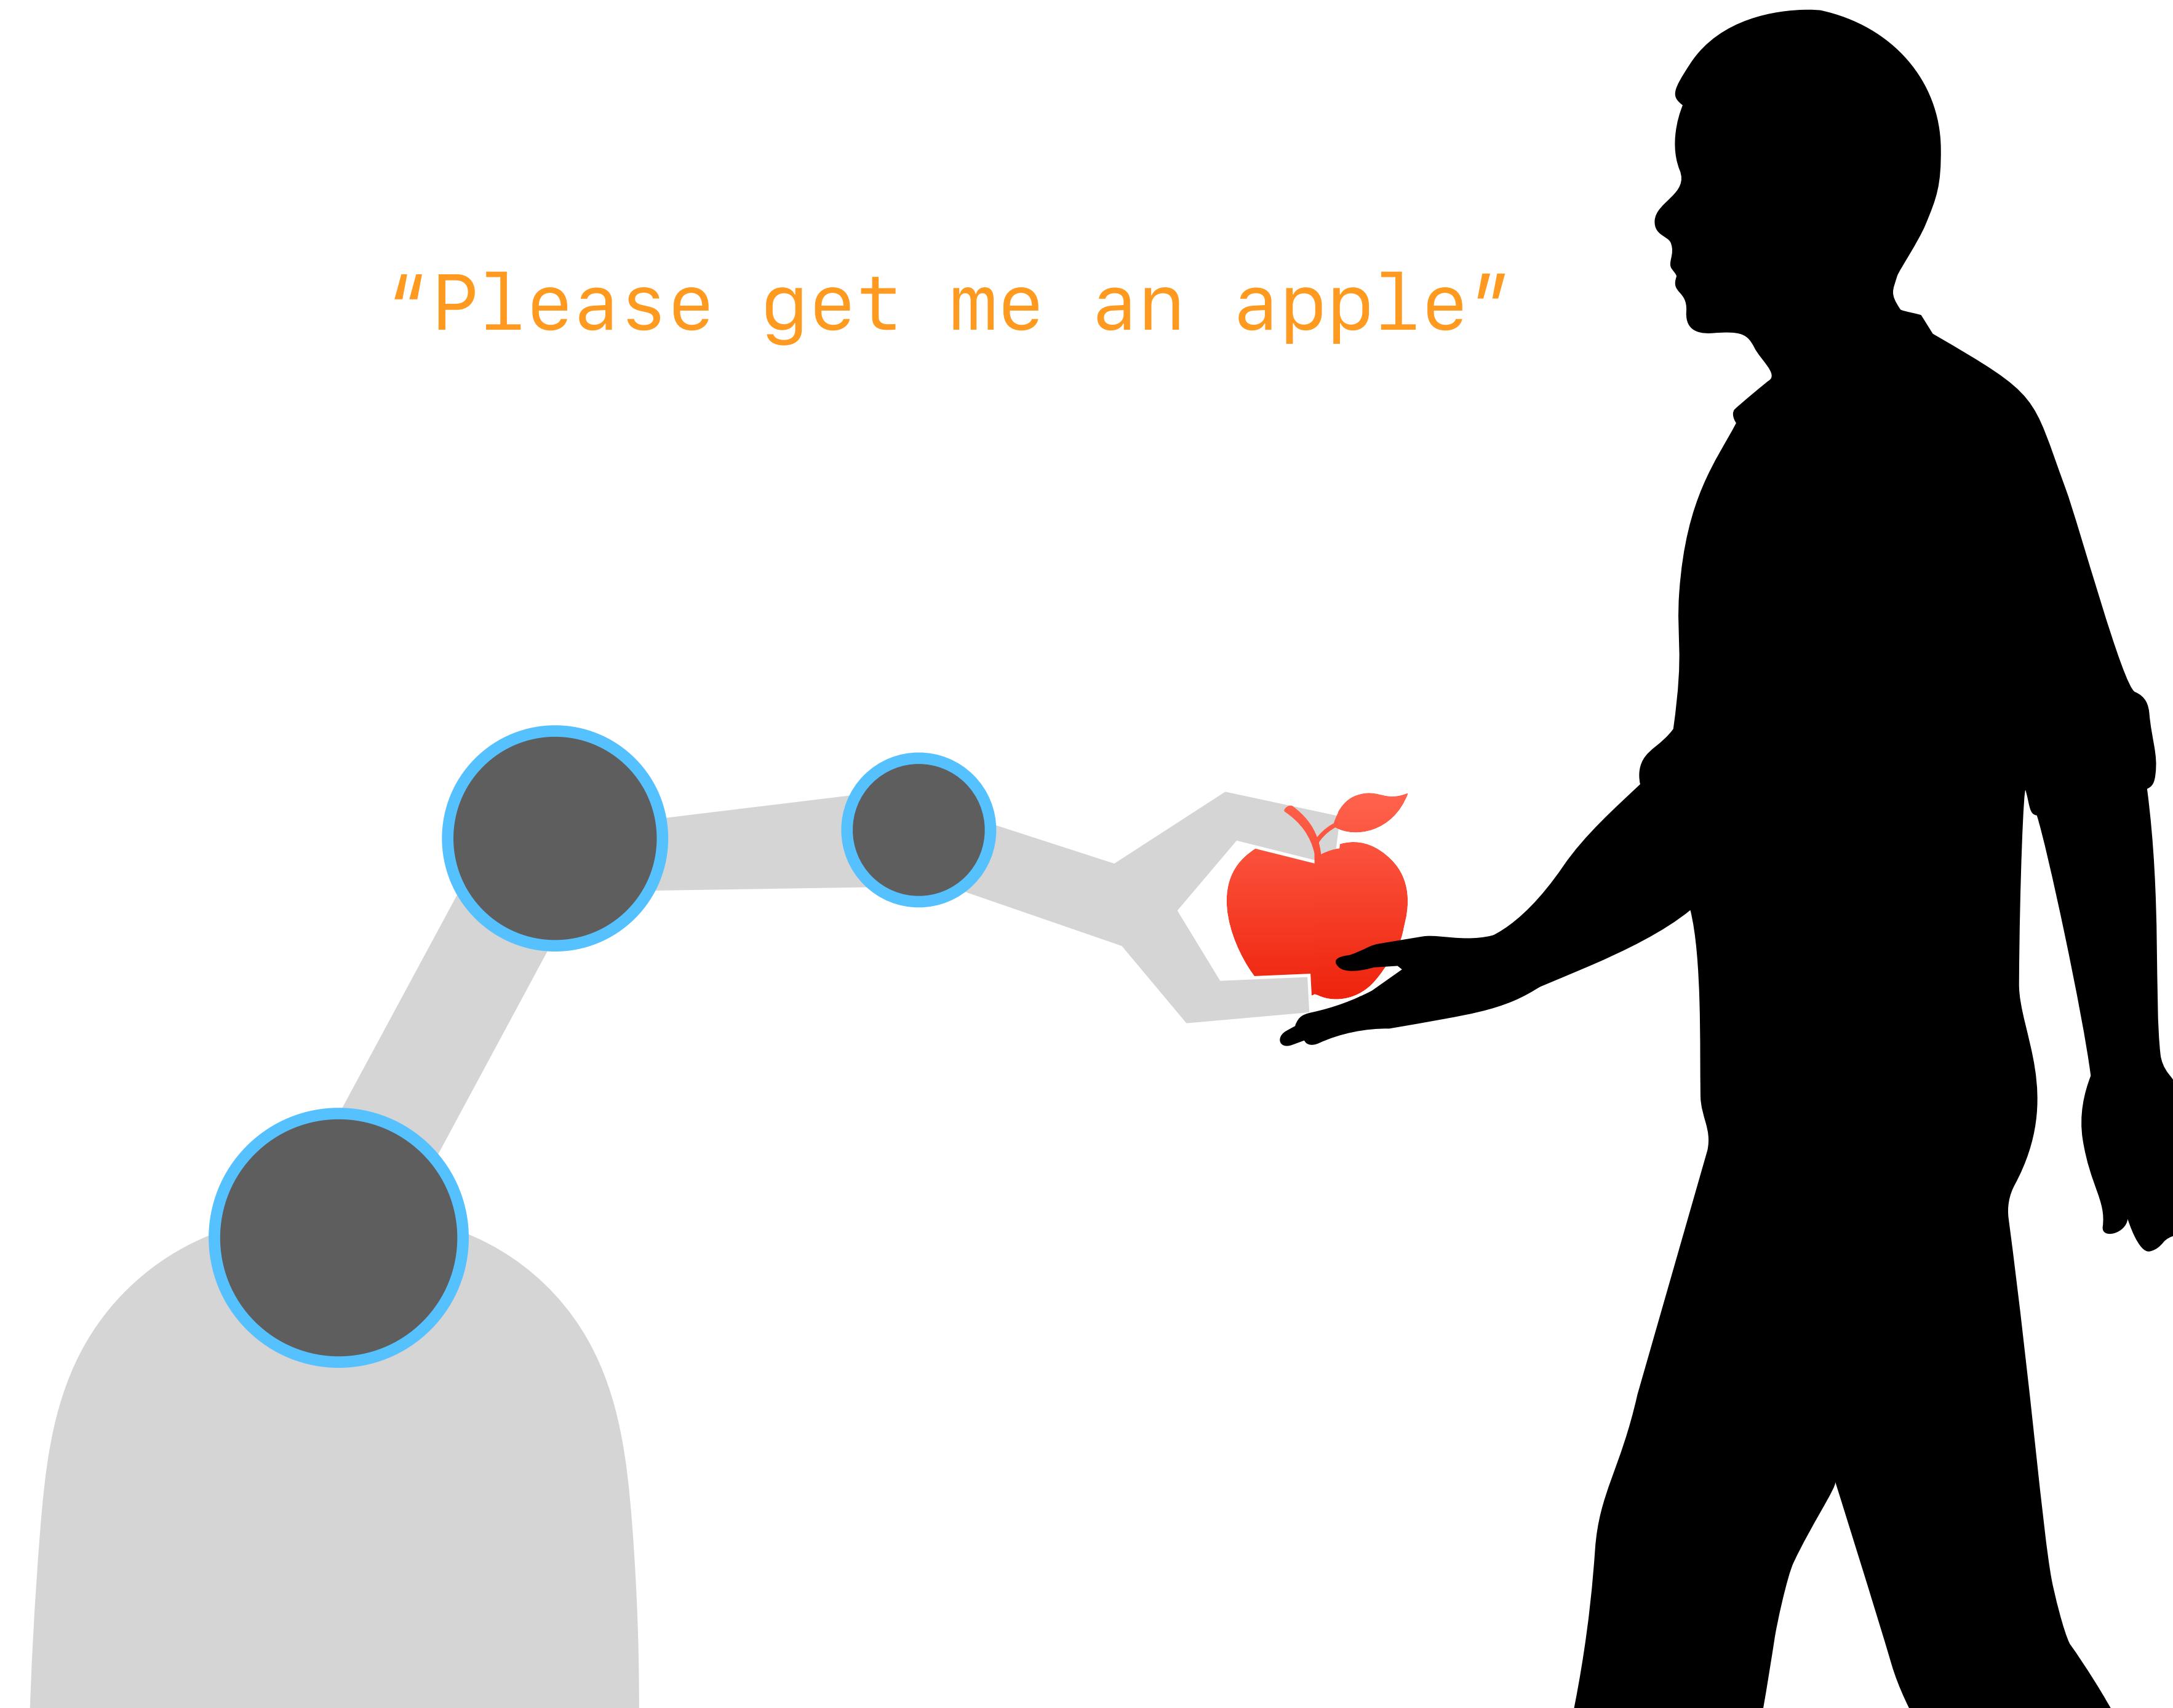
Image: Szot et al. 2021



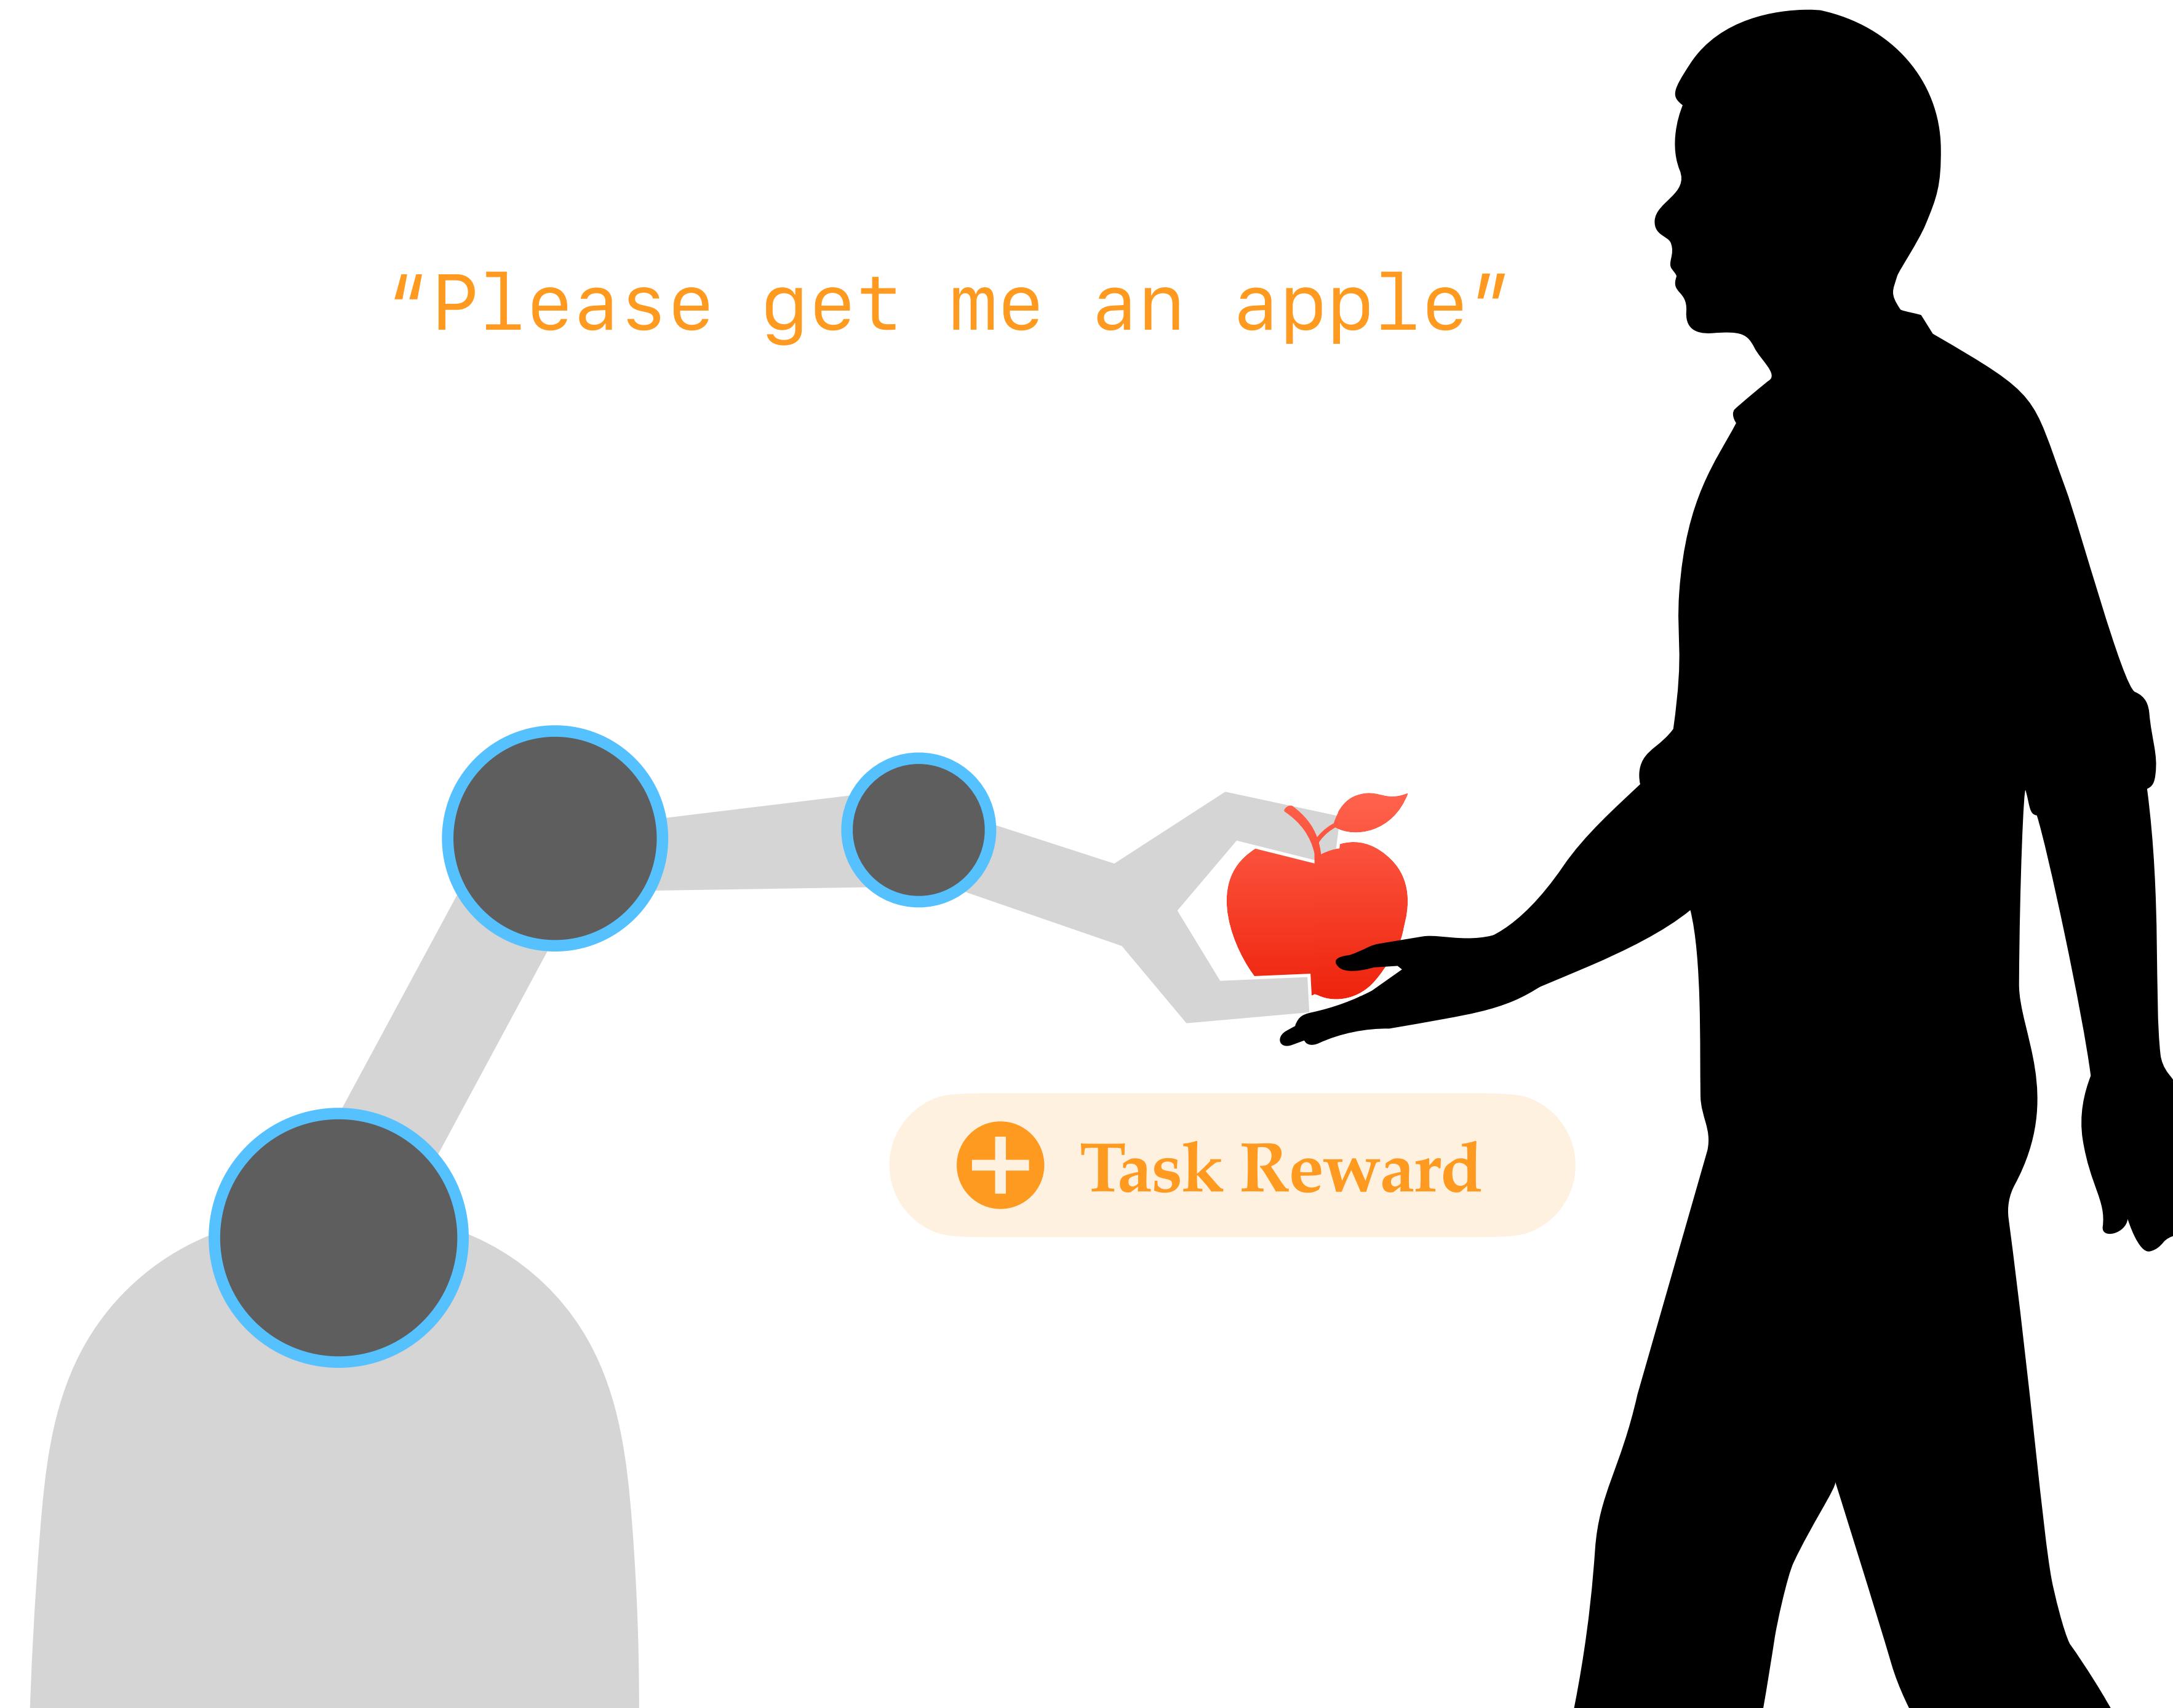
"Please get me
an apple"



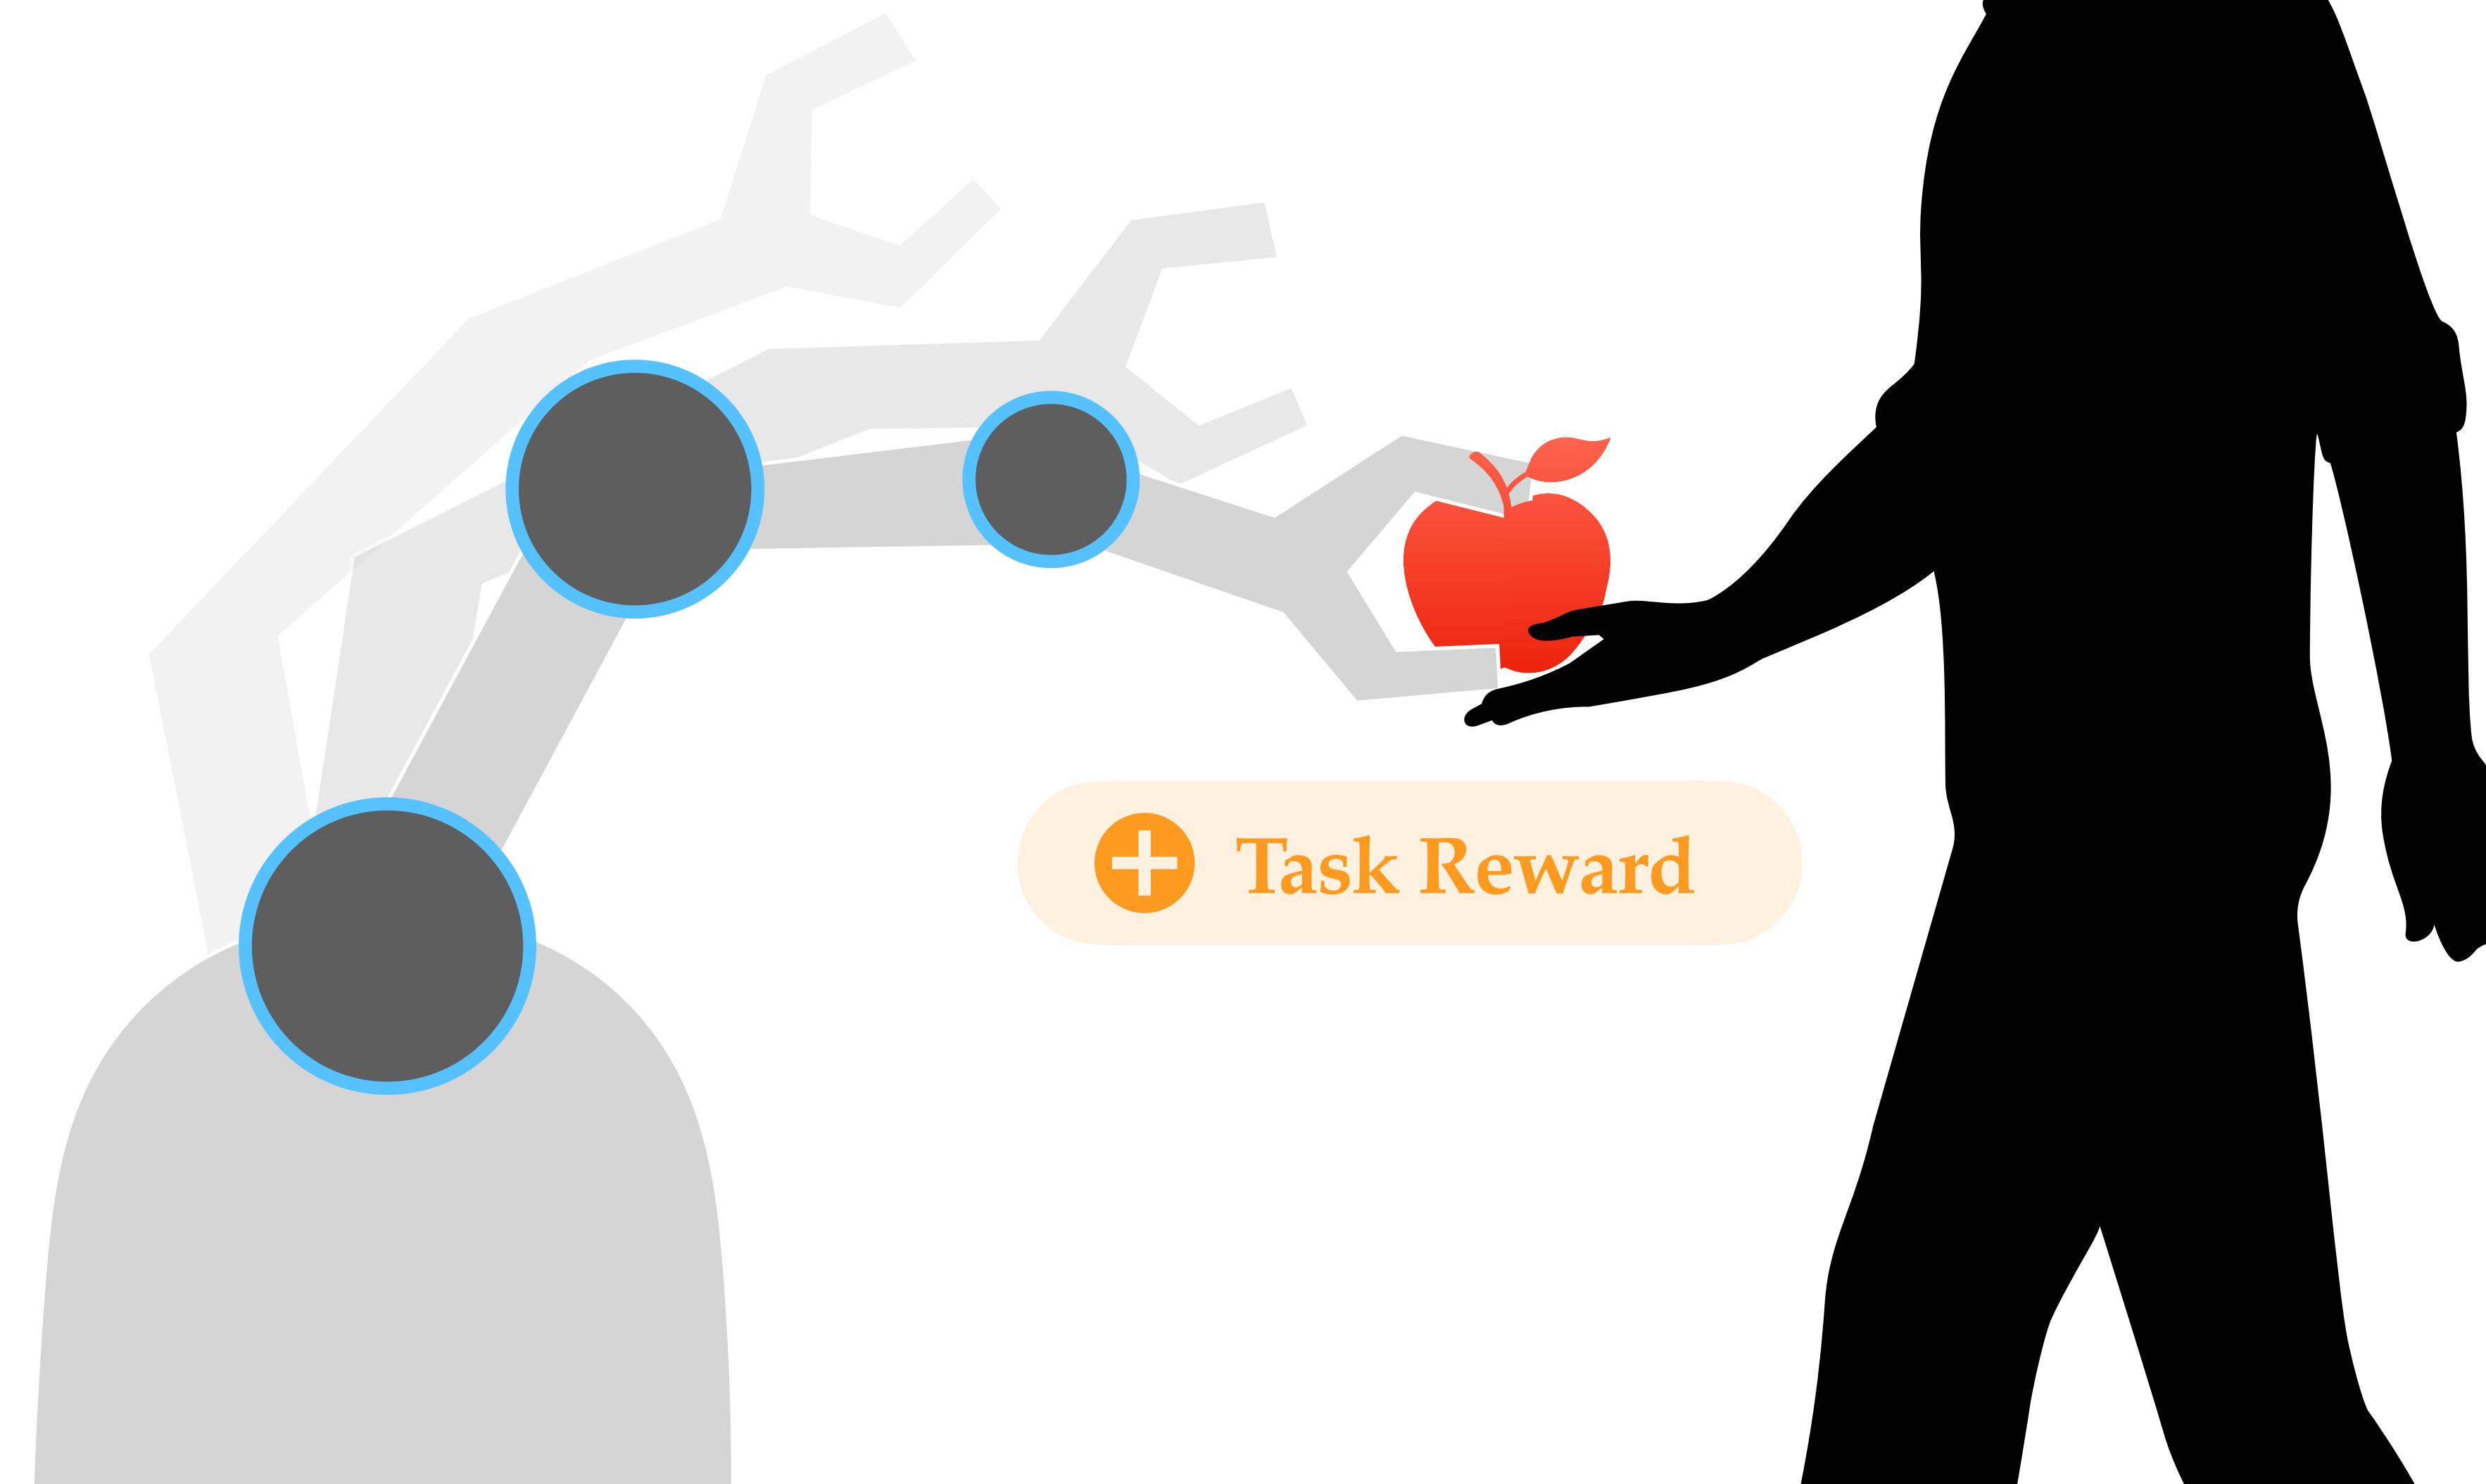
"Please get me an apple"



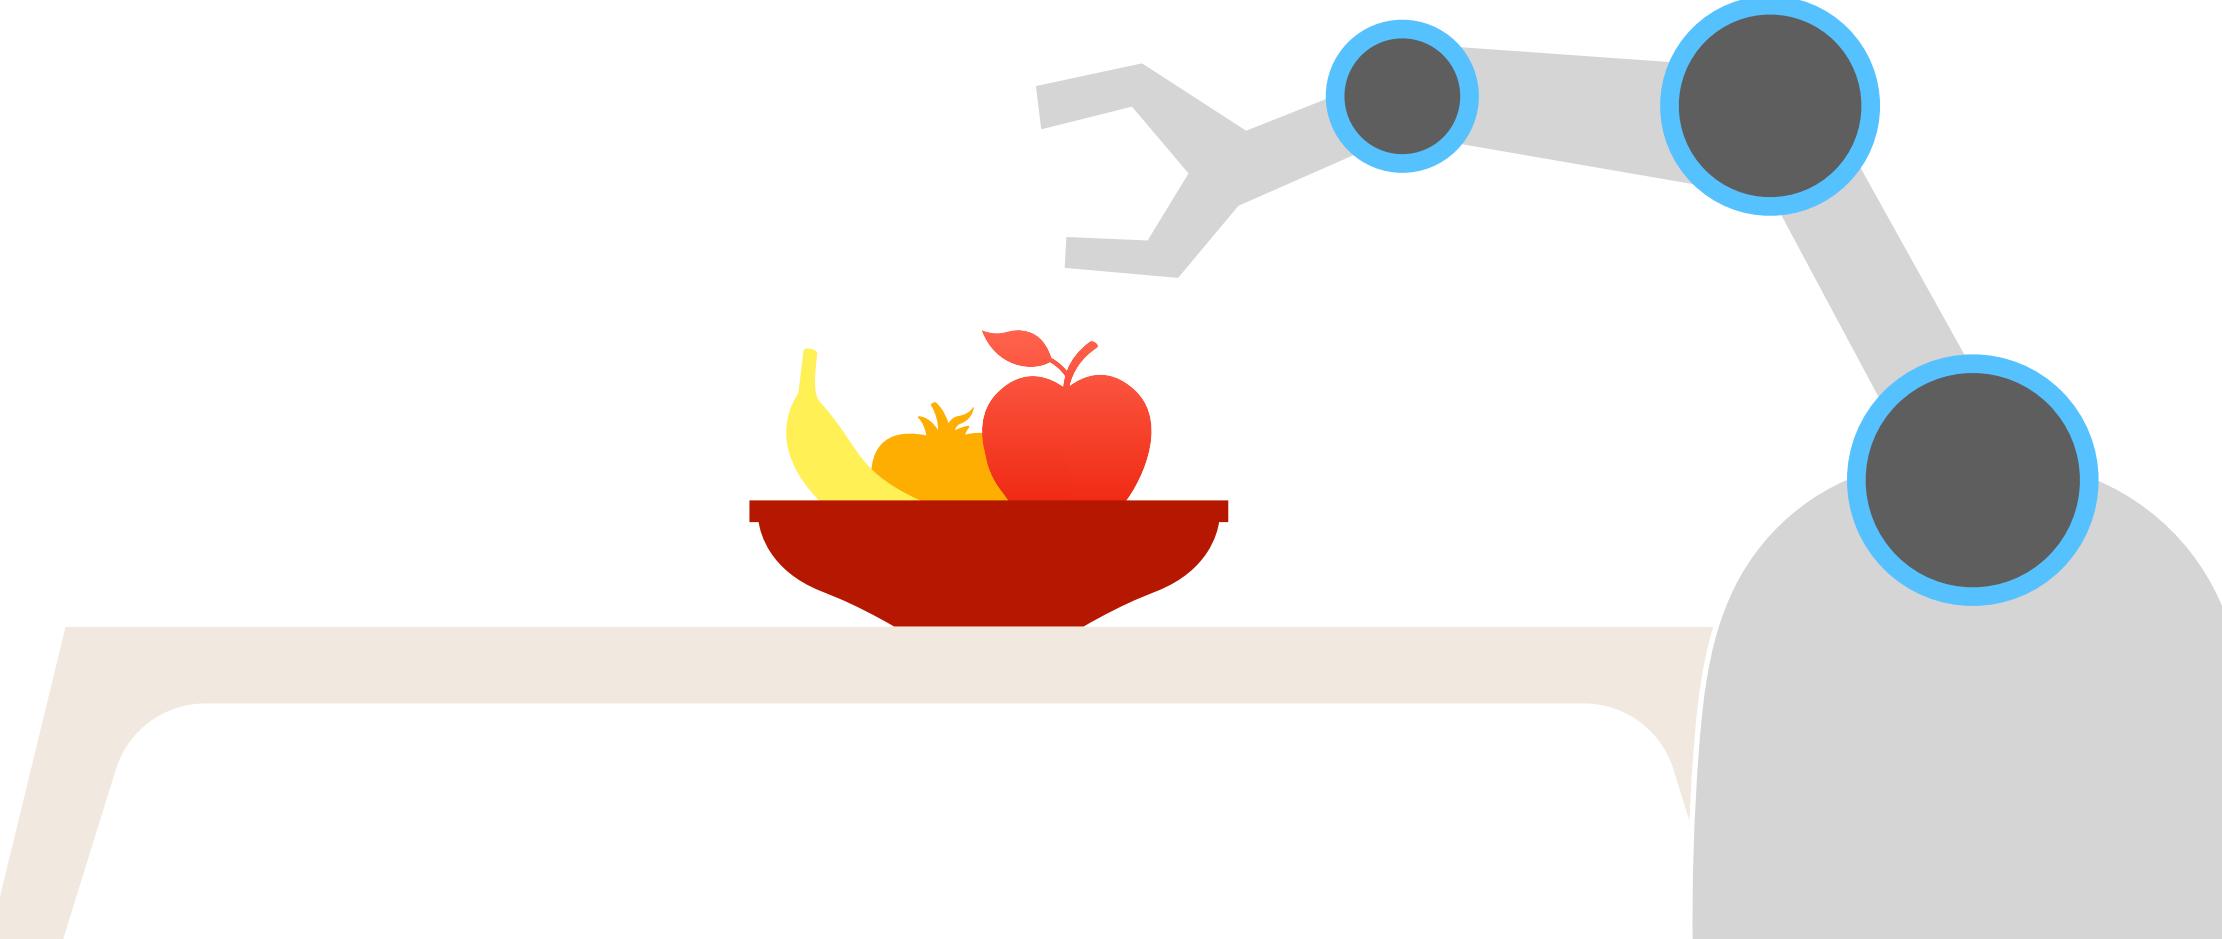
"Please get me an apple"



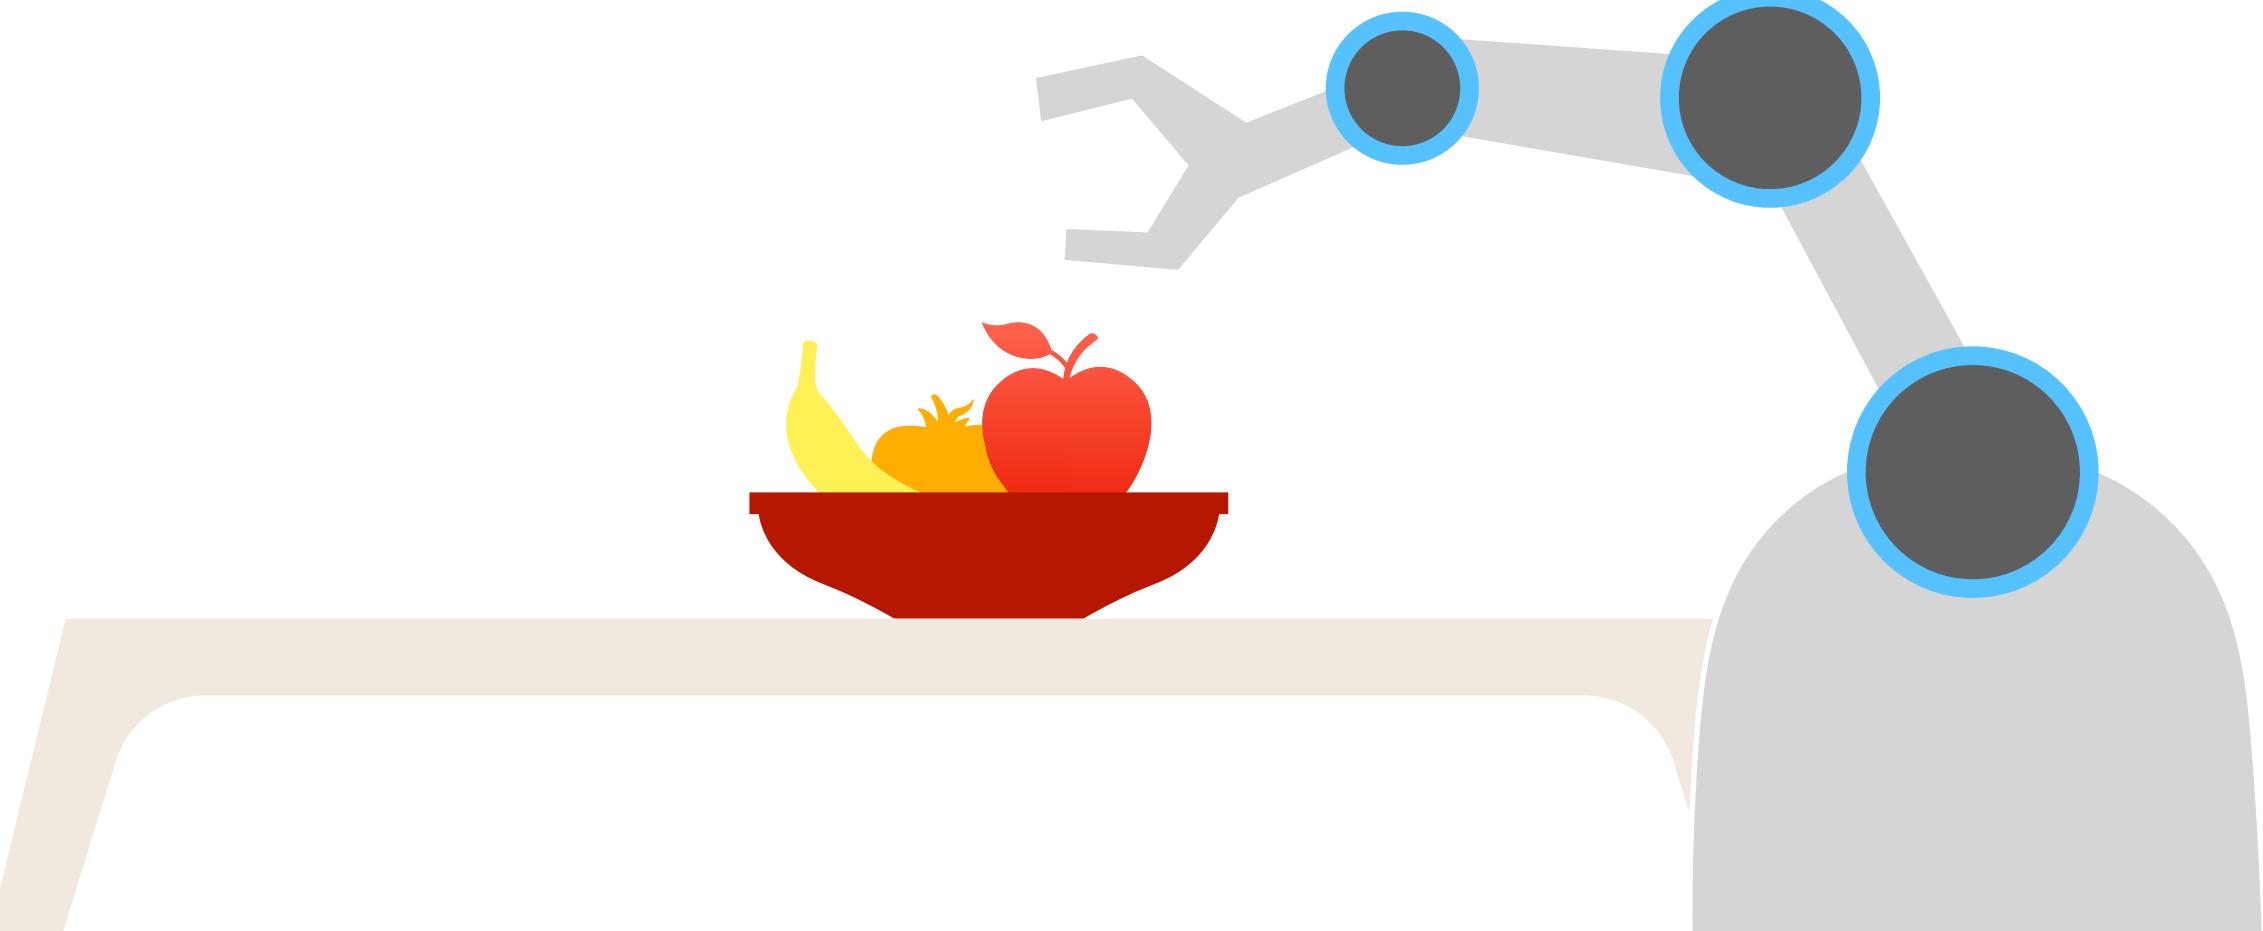
"Please get me an apple"



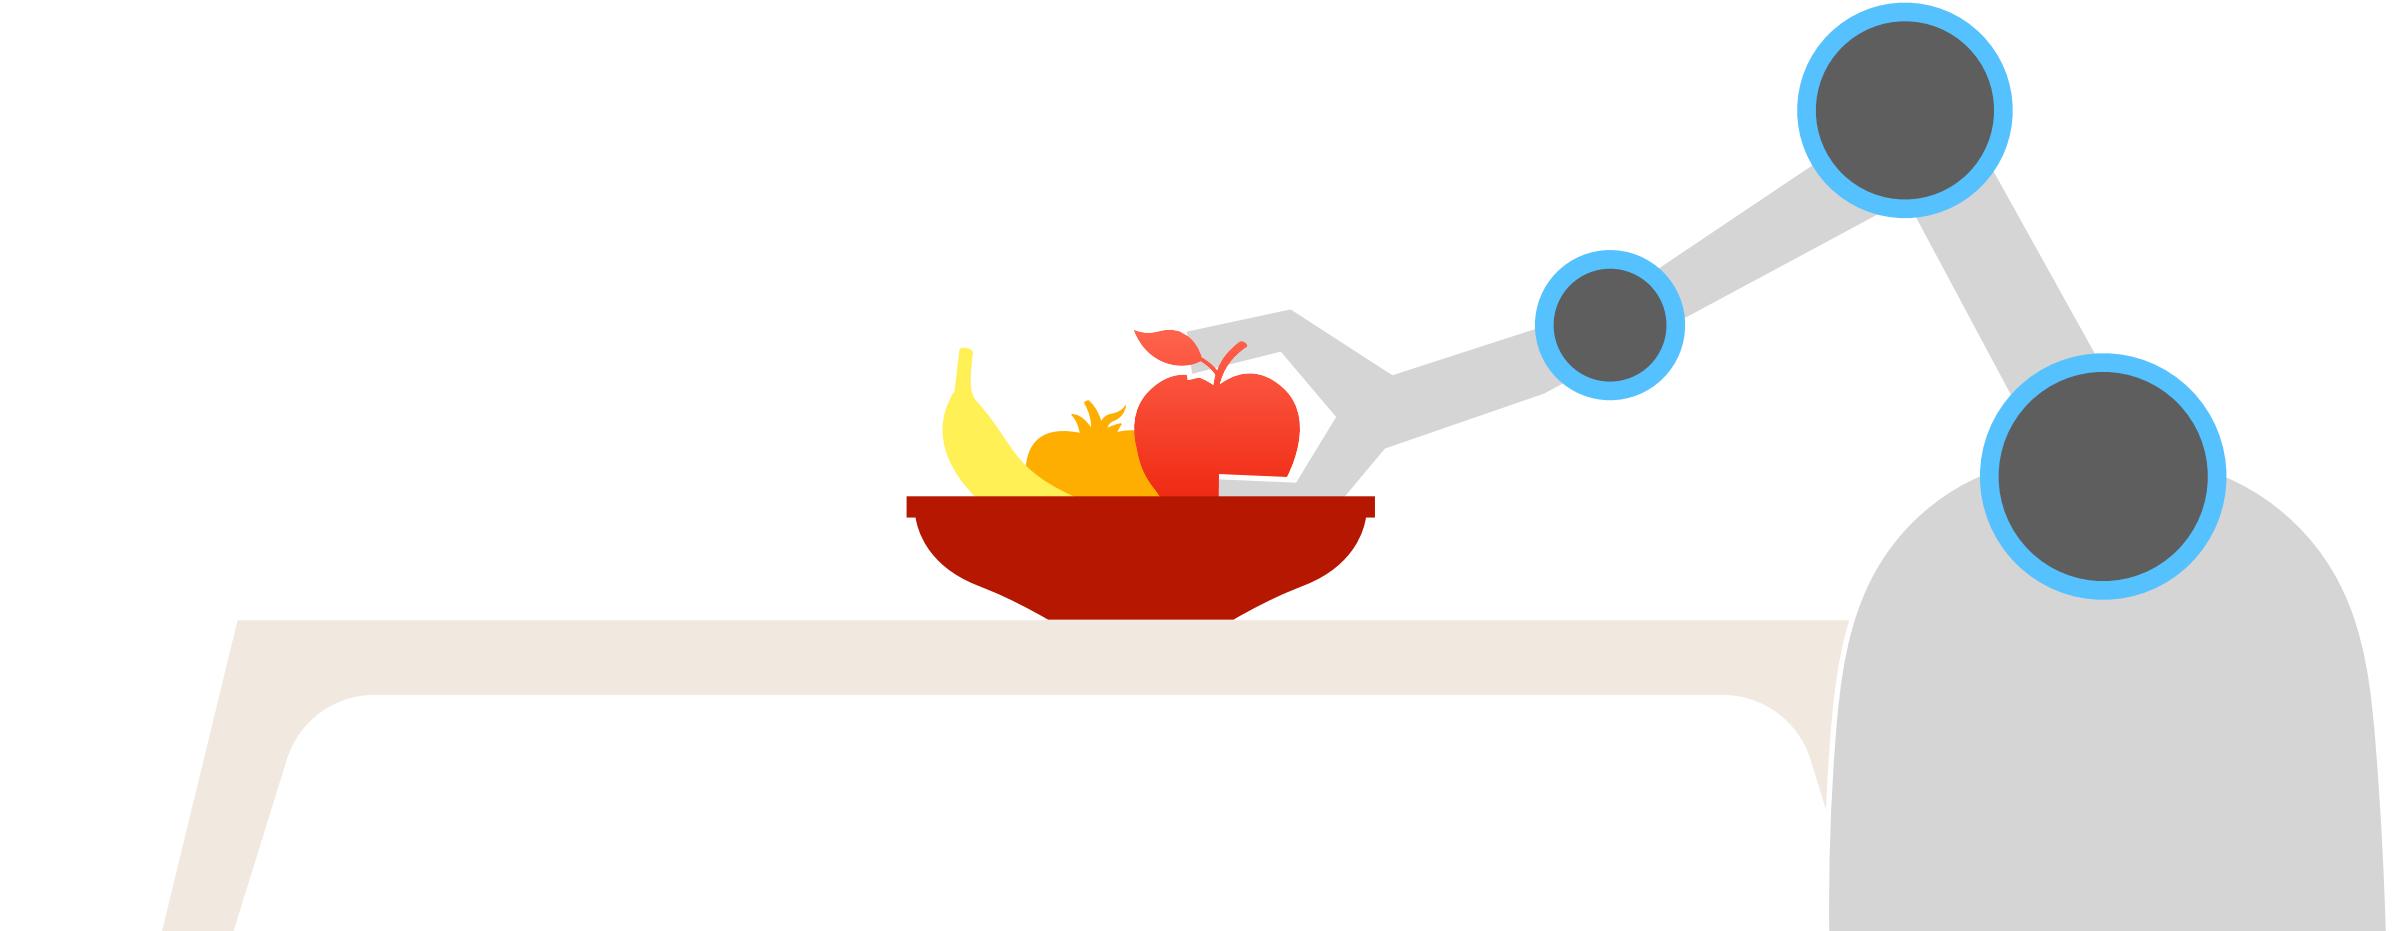
"Go to the table"



"Go to the table"

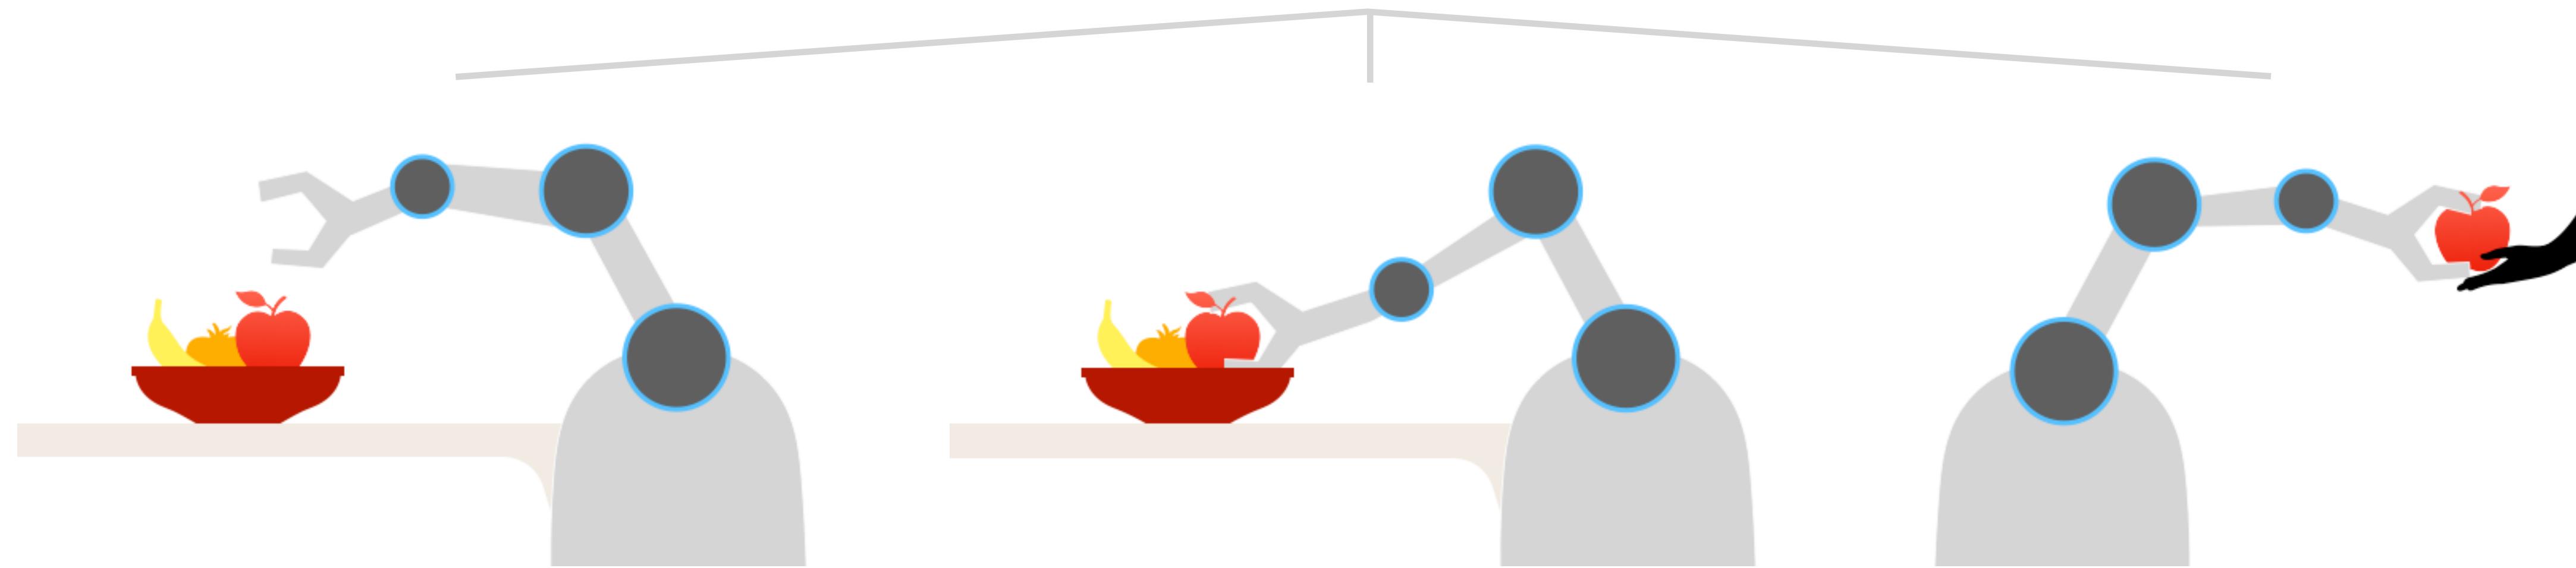


"Pick up an apple"



Learning from sparse rewards for
complex tasks is **not sample-efficient**.

"Please get me an apple"

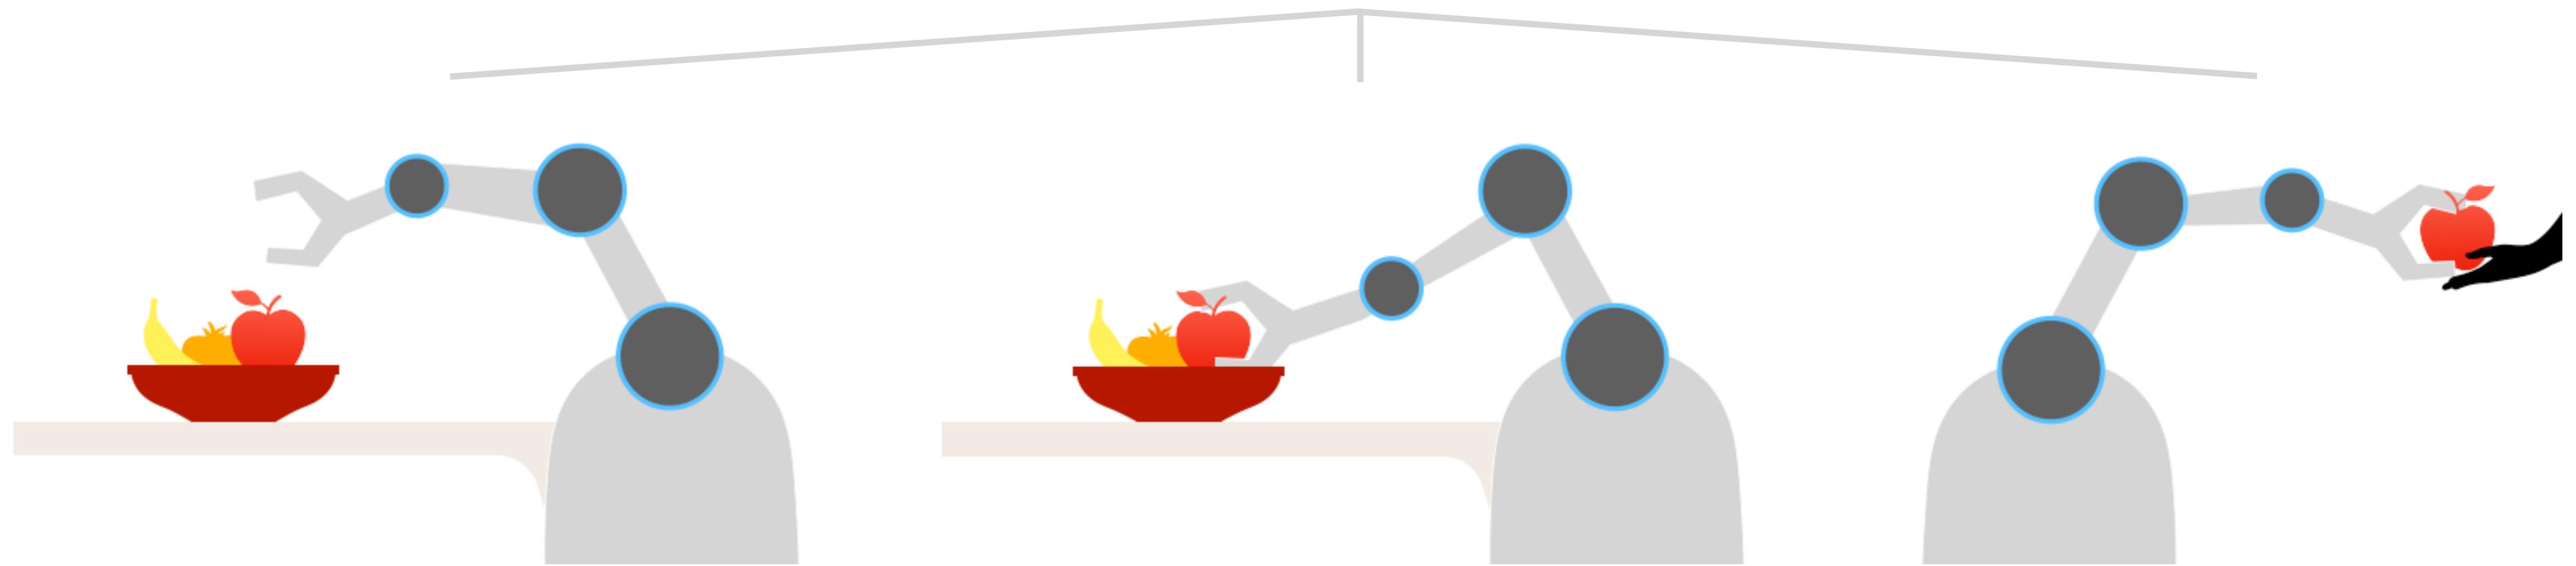


"Go to the table"

"Pick up an apple"

"Bring it to me"

"Please get me an apple"



"Go to the table"

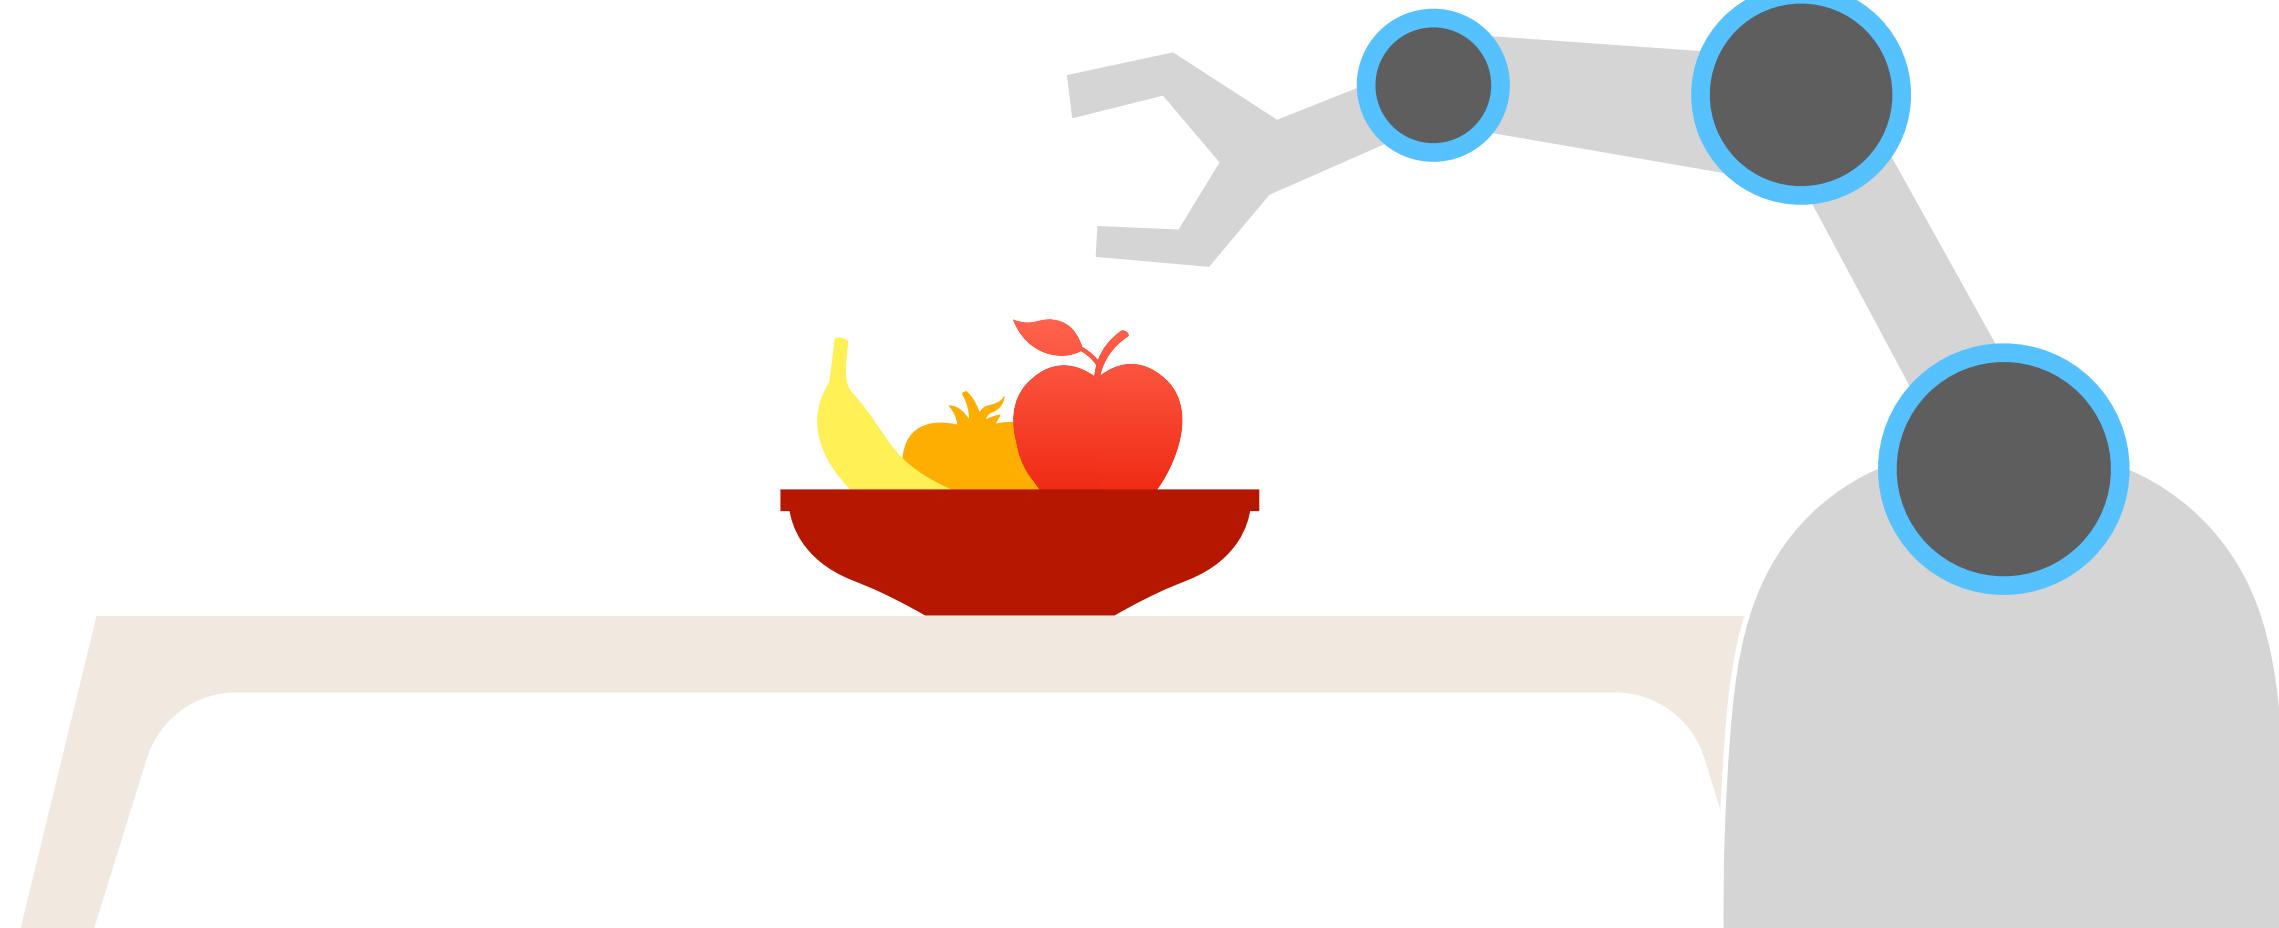
"Pick up an apple"

"Bring it to me"

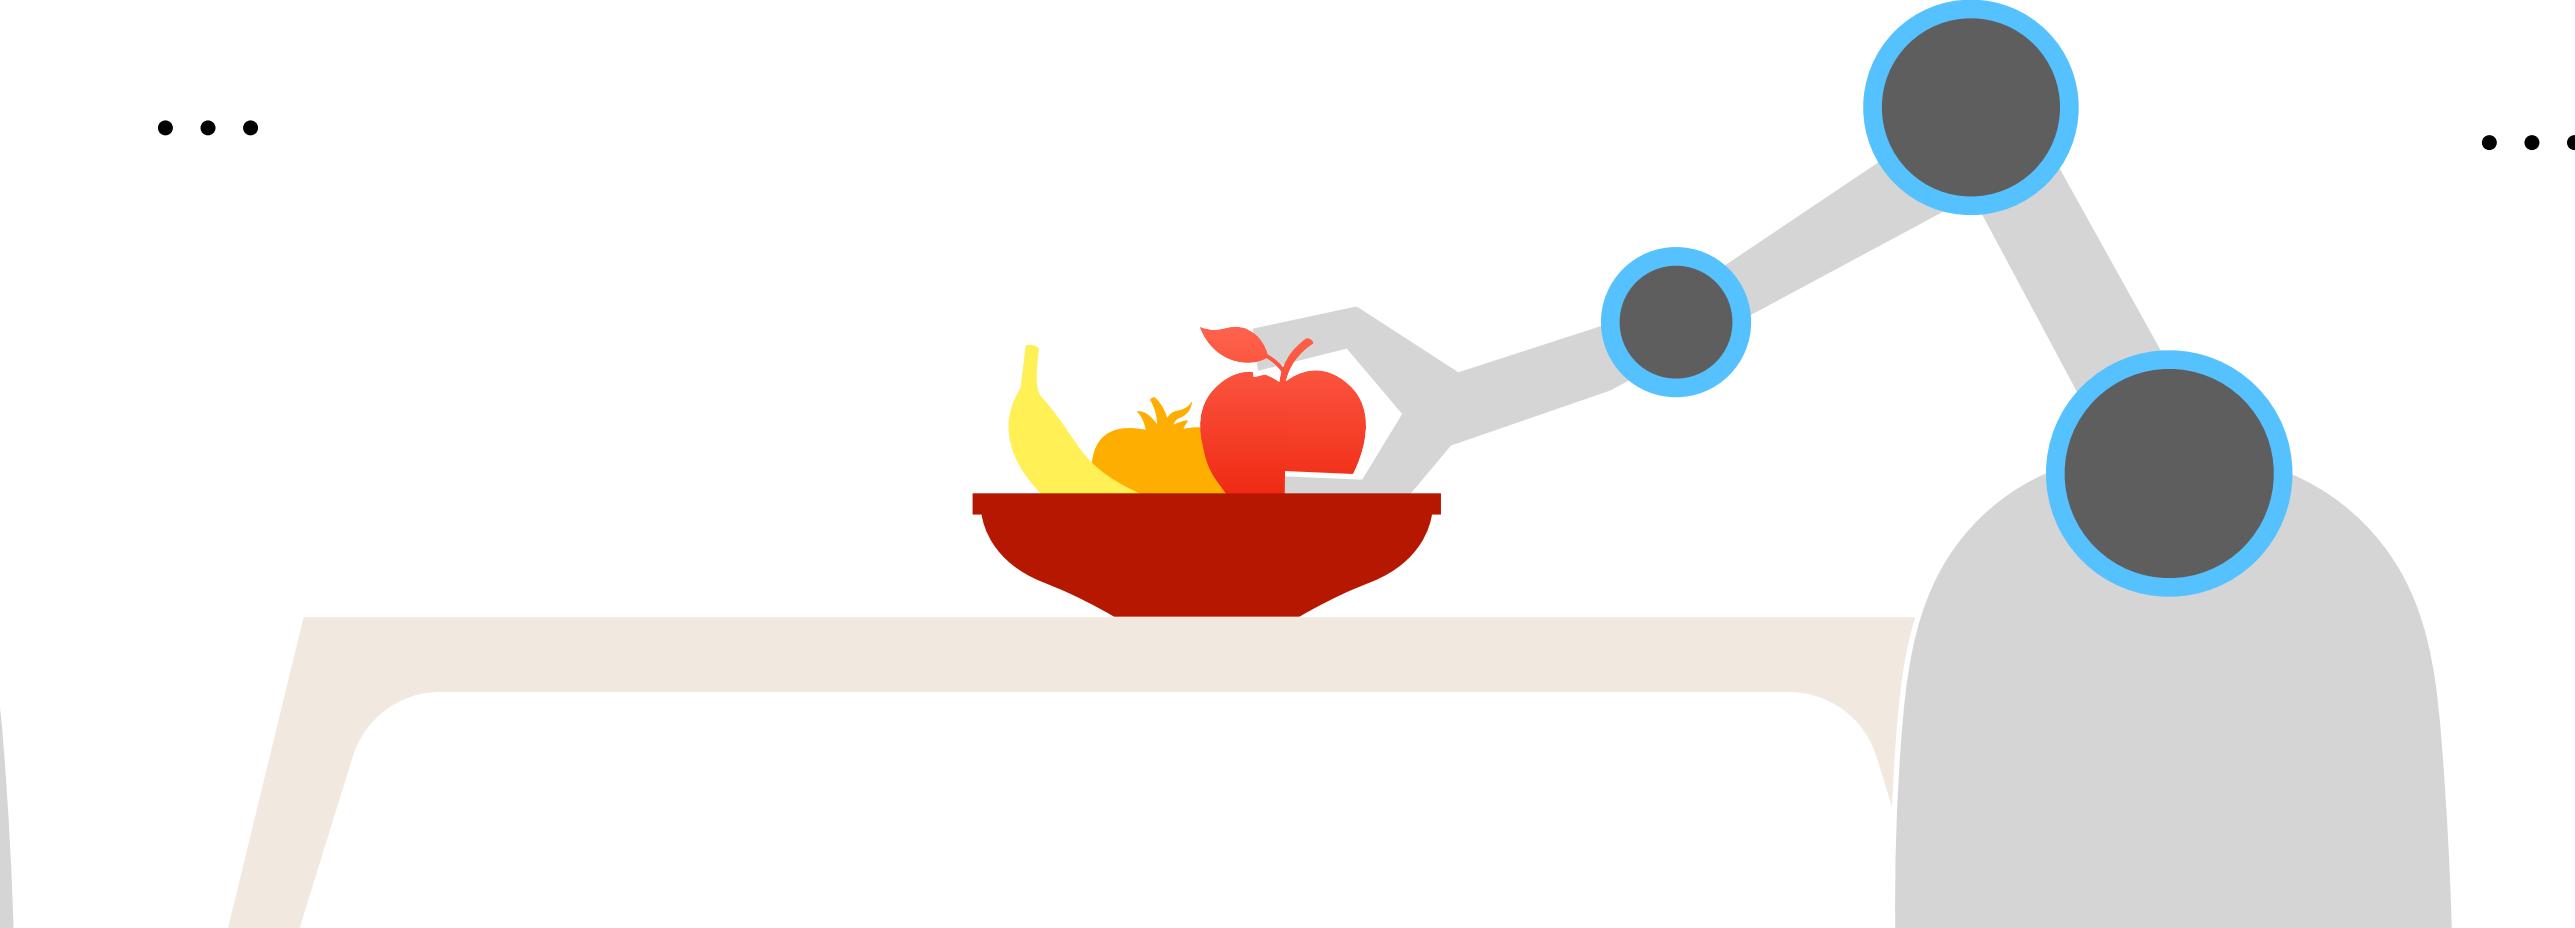


Can we use the principle of abstraction in
language to guide exploration?

"Go to the table"

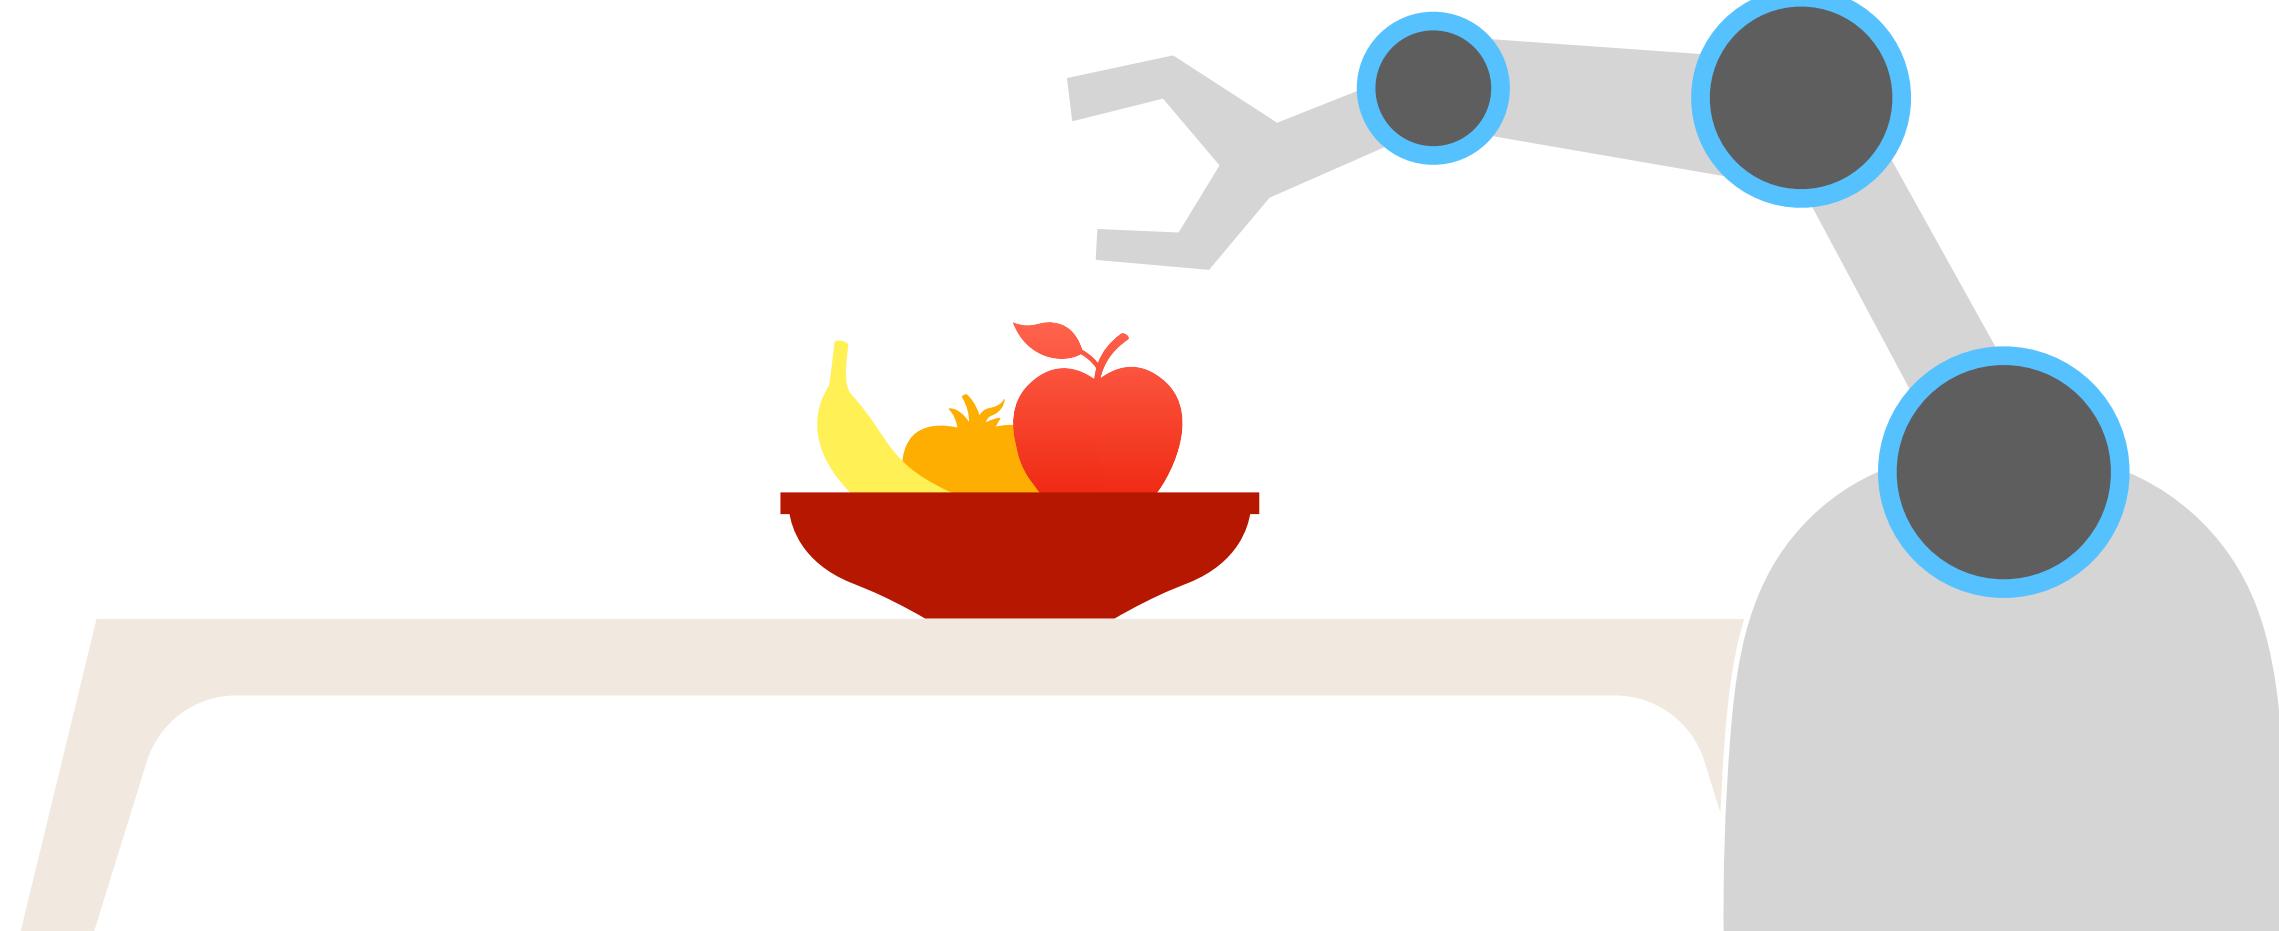


"Pick up an apple"

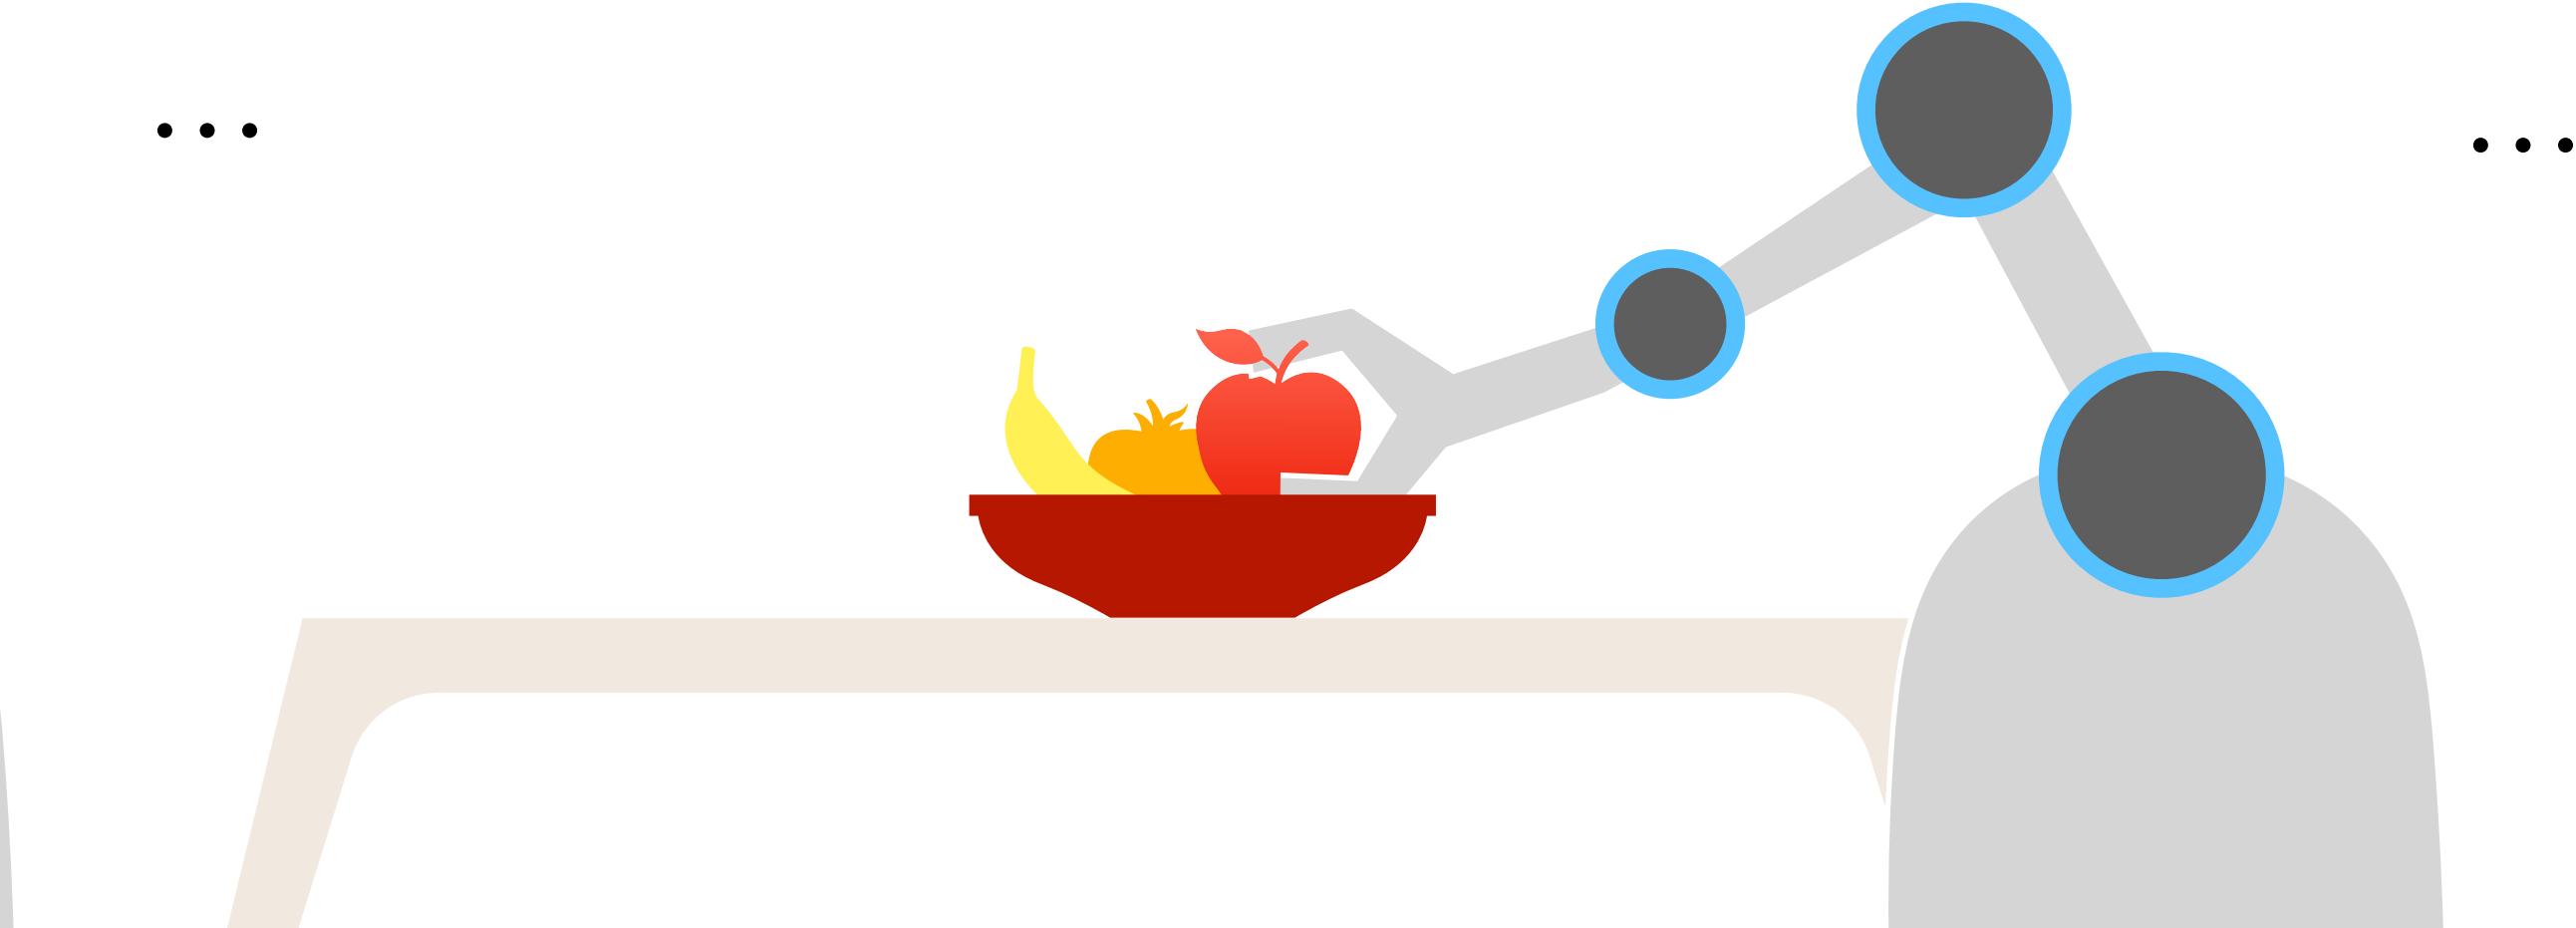


"Go to the table"

 Subtask Reward

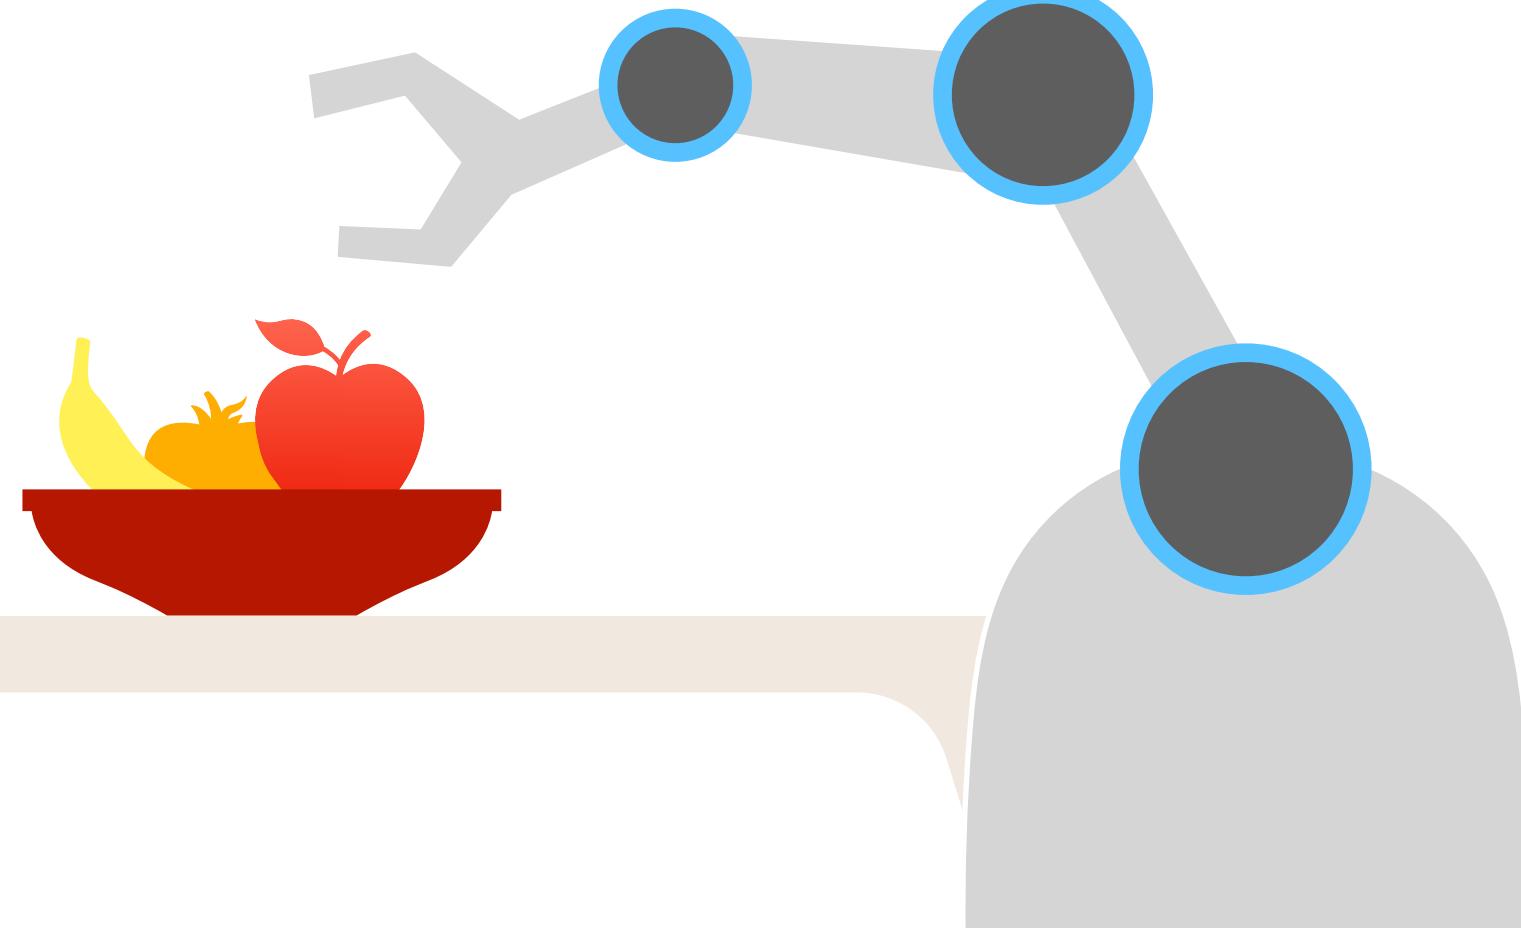


"Pick up an apple"



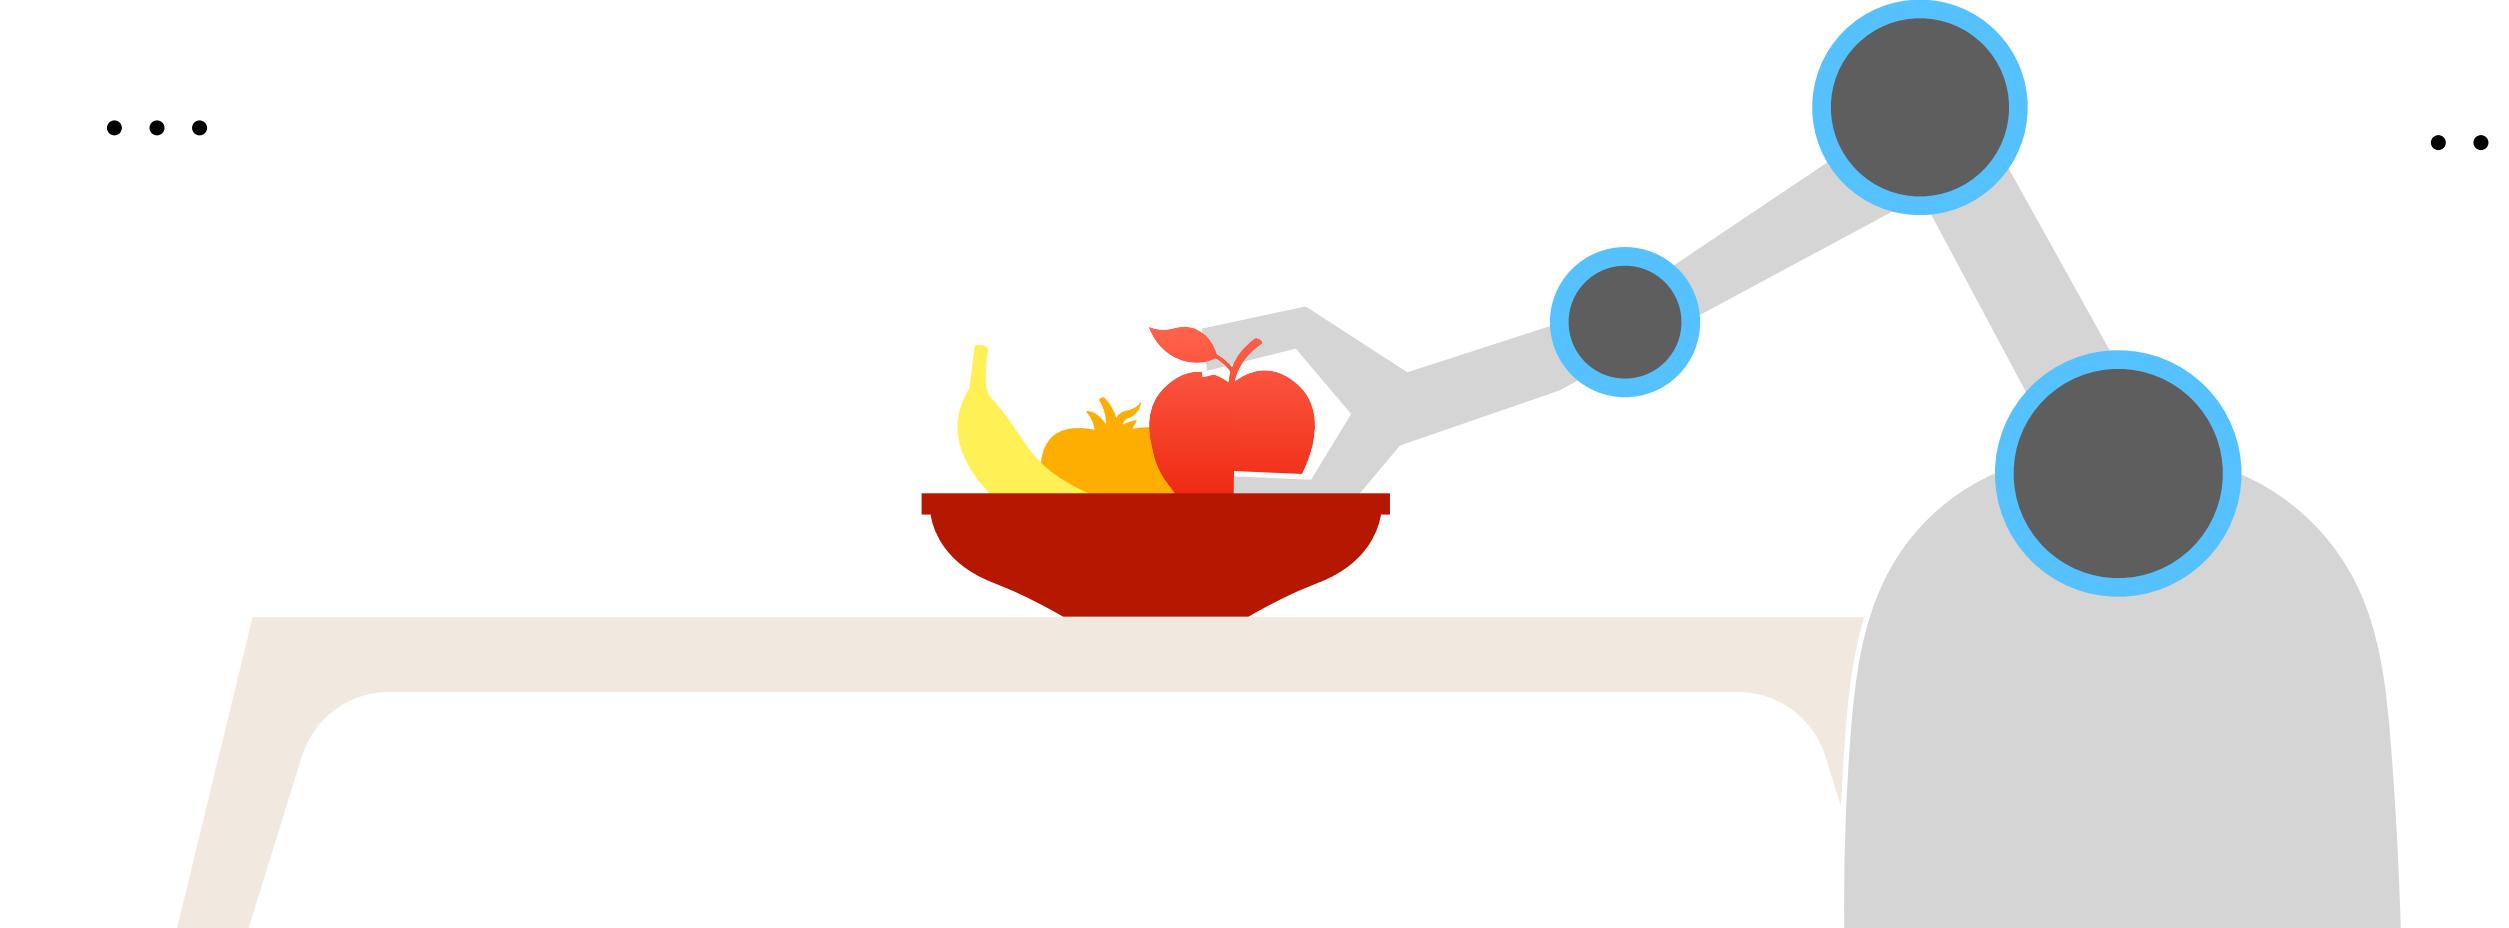
"Go to the table"

+ Subtask Reward

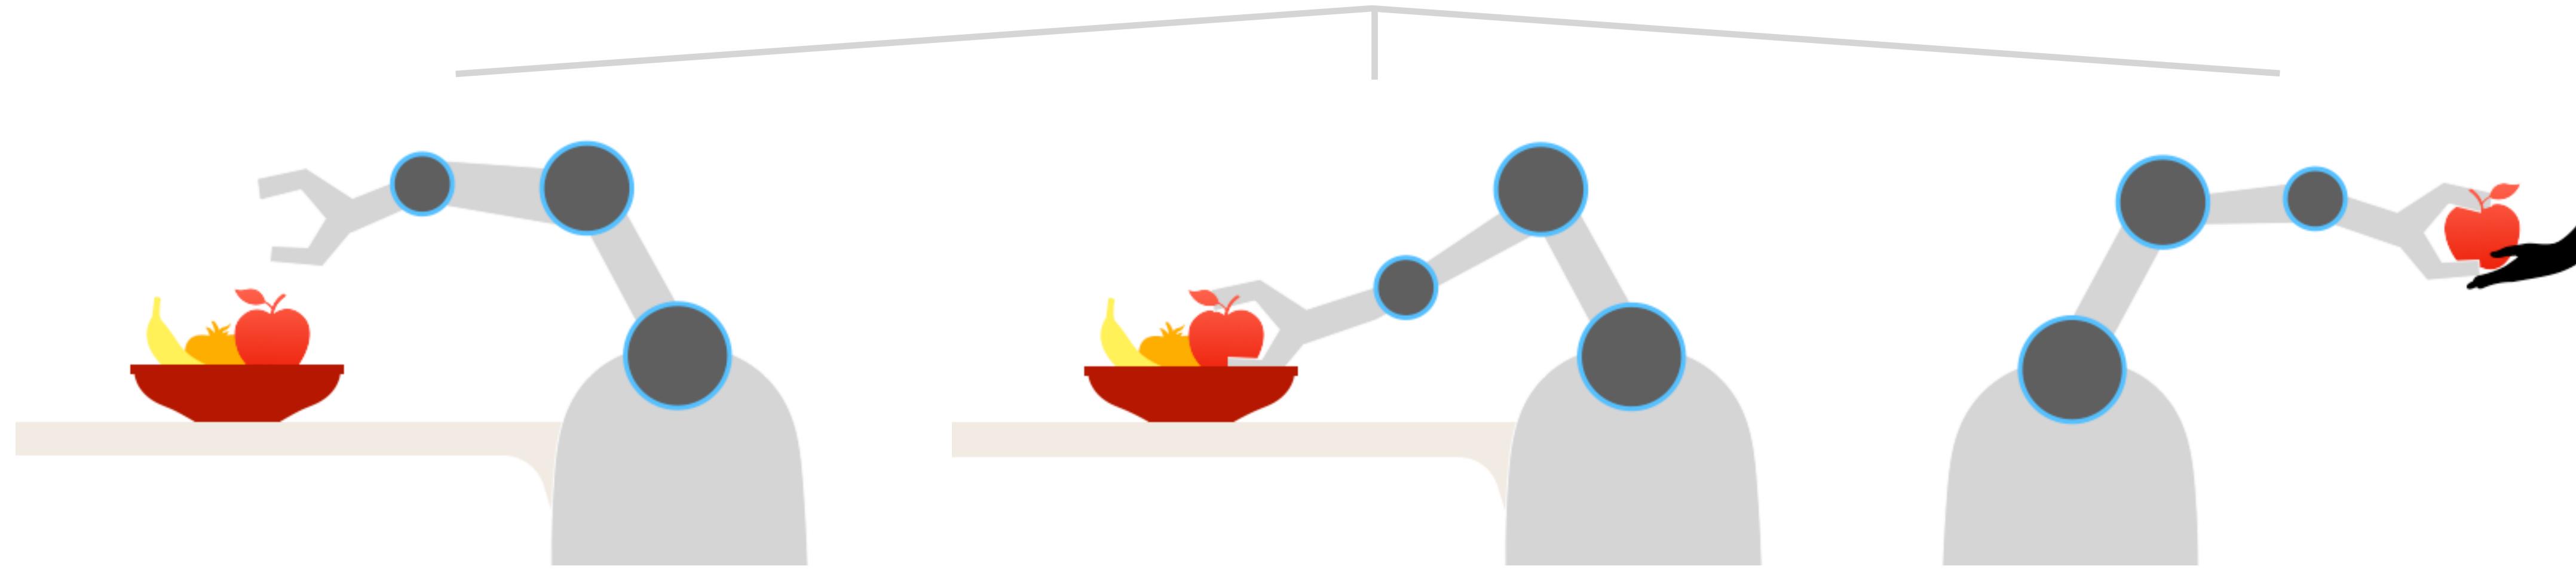


"Pick up an apple"

+ Subtask Reward



"Bring me an apple"



"Go to the table"

"Pick up an apple"

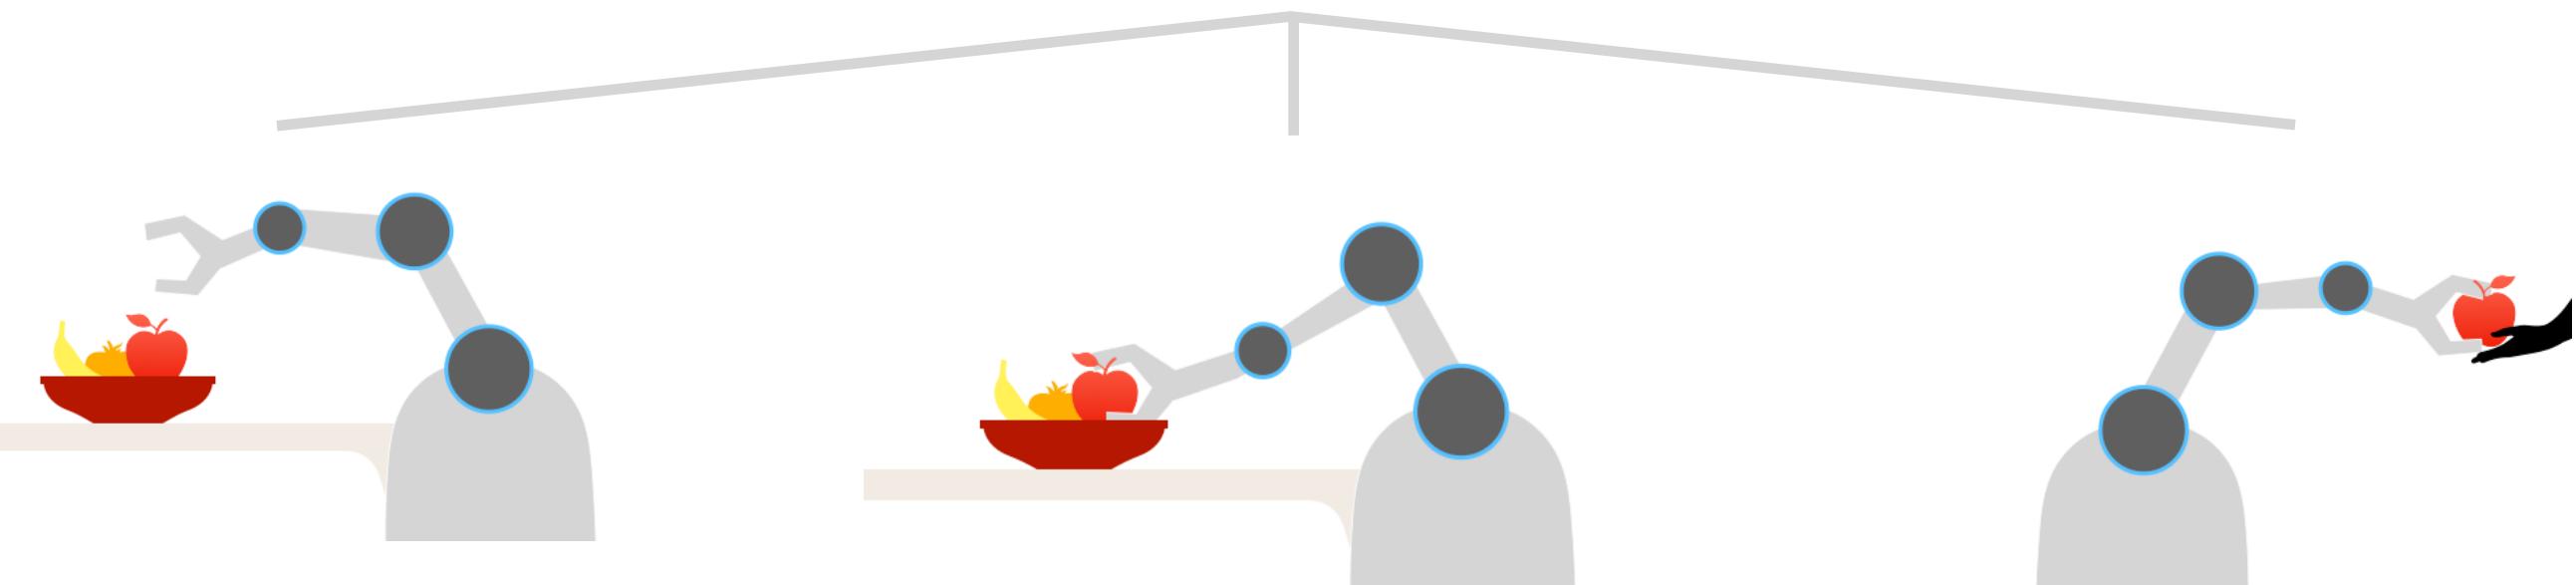
"Bring it to me"

<turn left> . . .

<lift up> . . .

<turn right> . . .

"Bring me an apple"



"Go to the table"



<turn left> . . .

"Pick up an apple"



<lift up> . . .

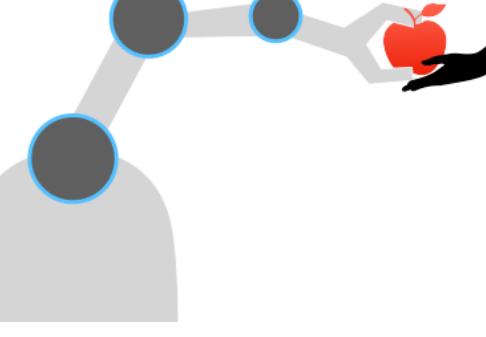
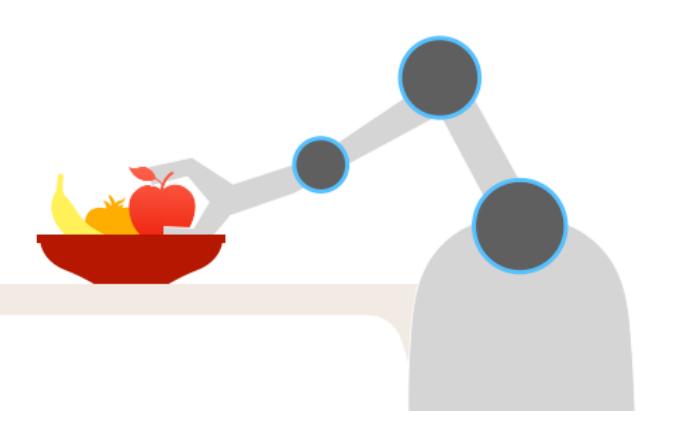
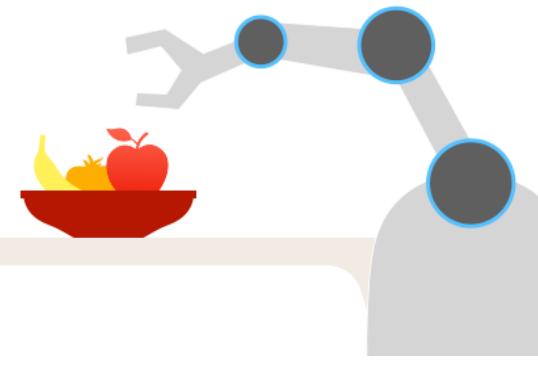
"Bring it to me"



<turn right> . . .

"Bring me an apple"

g_h



"Go to the table"

g_l

<turn left> . . .

"Pick up an apple"

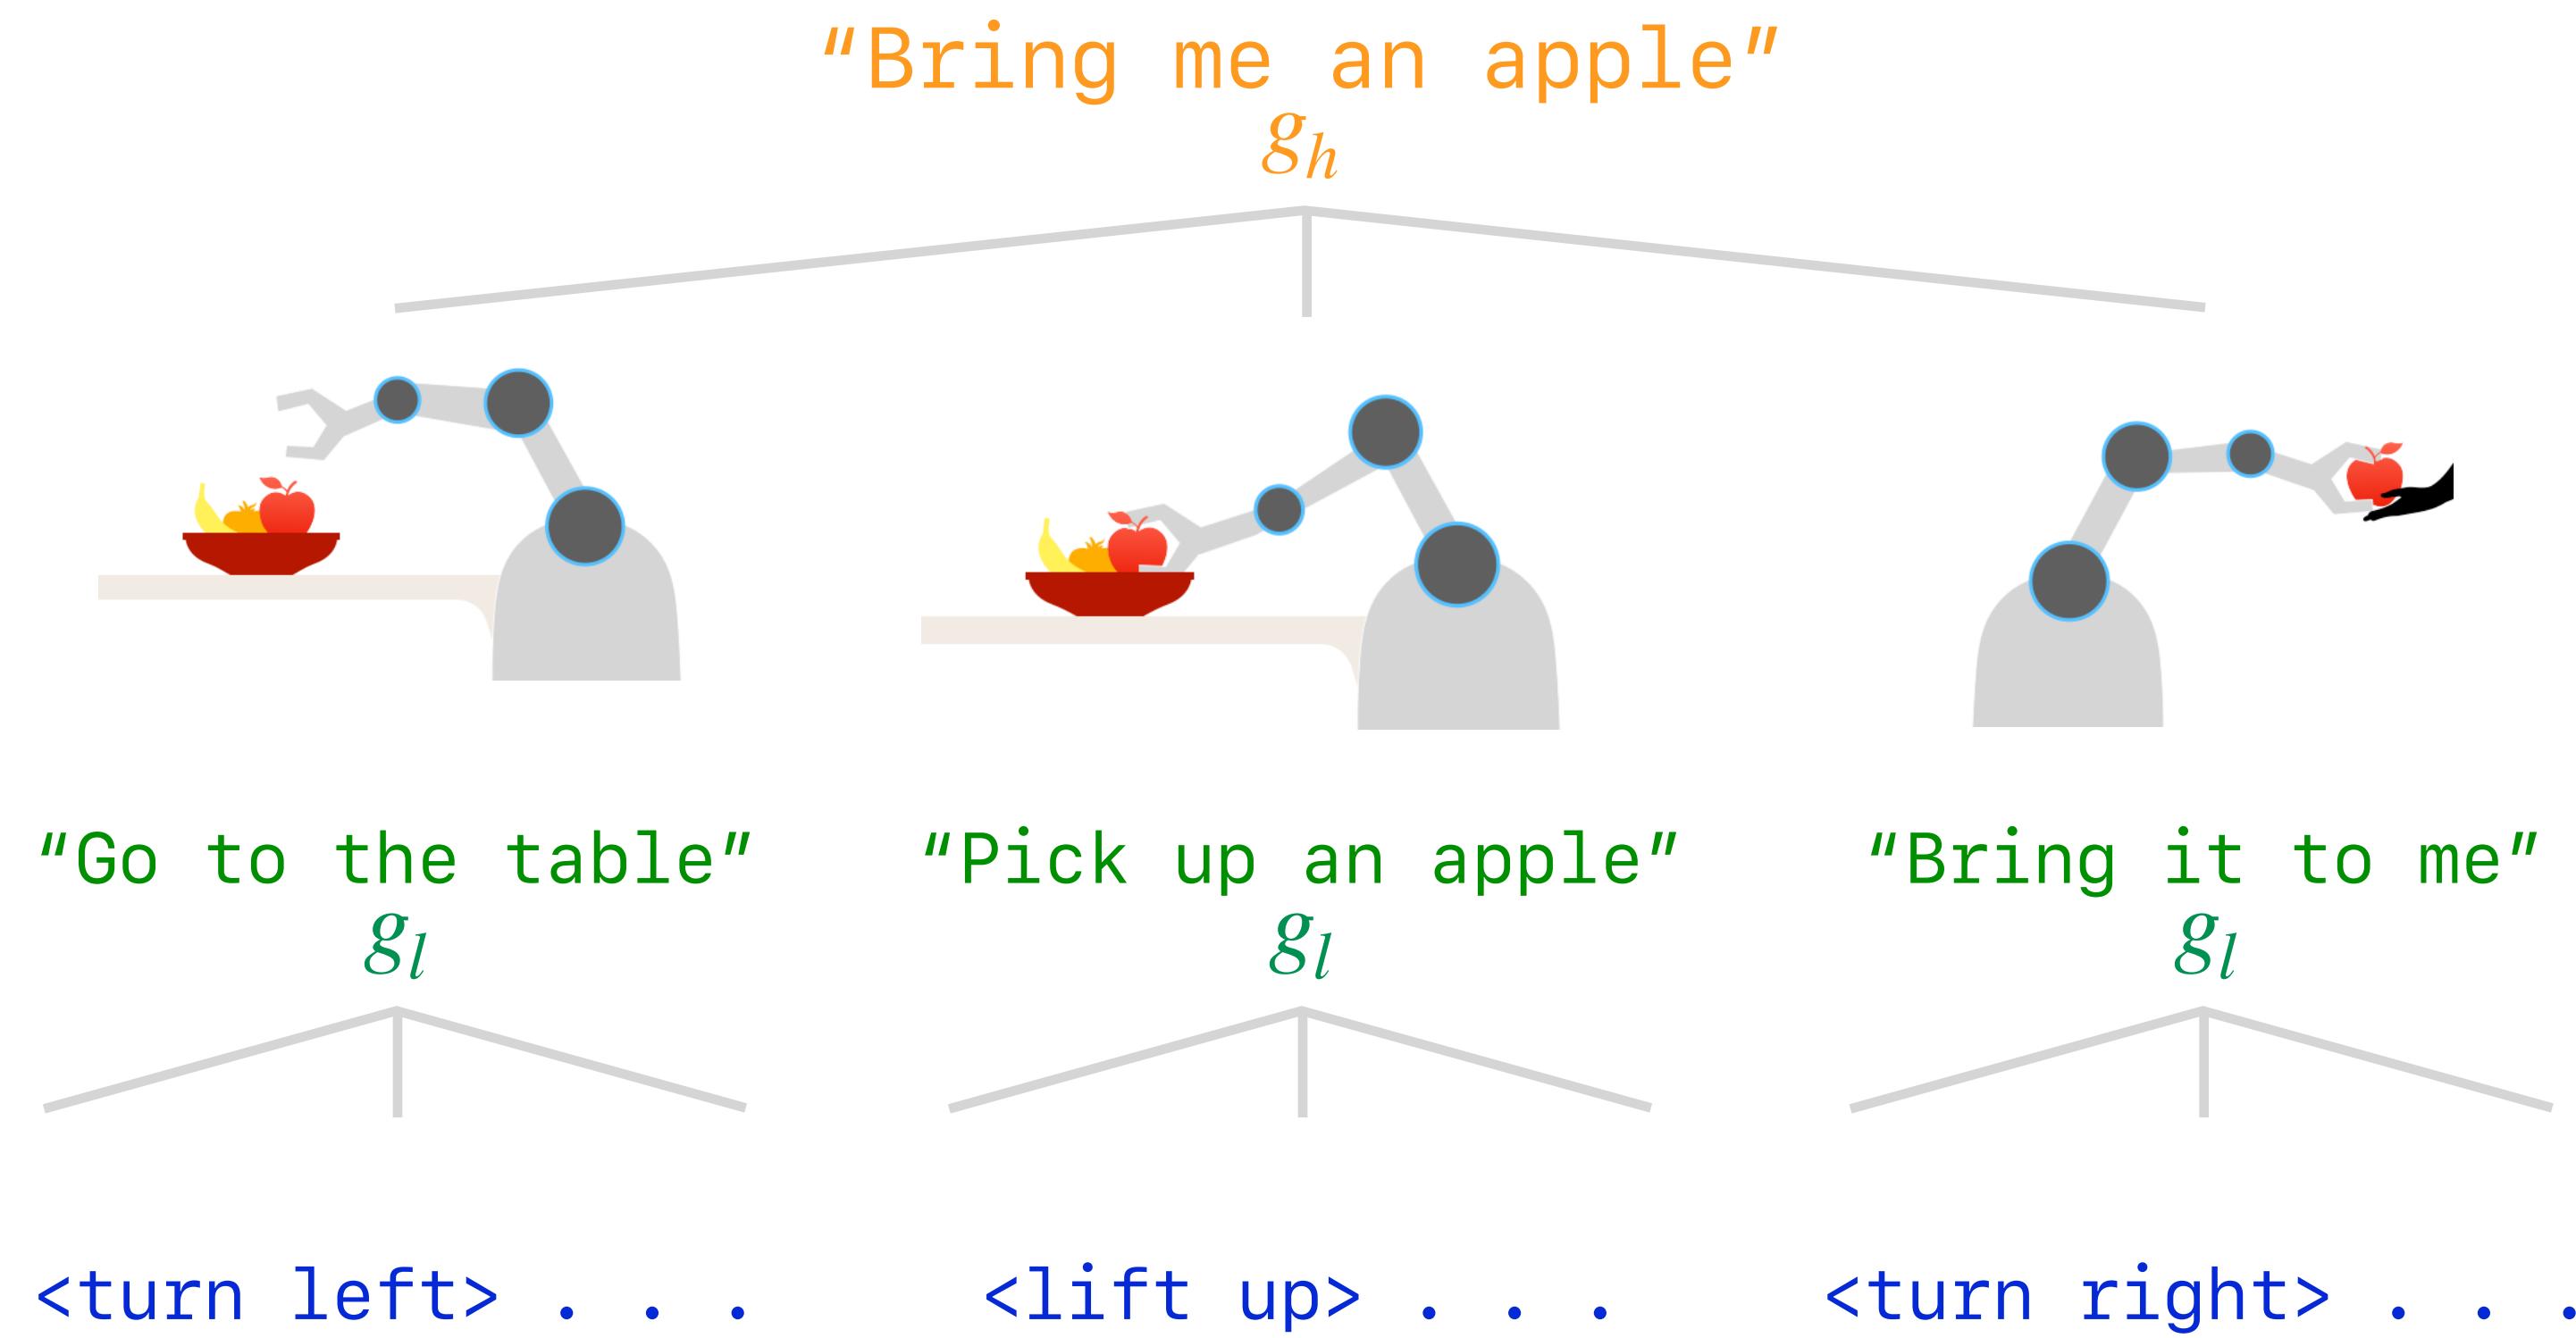
g_l

<lift up> . . .

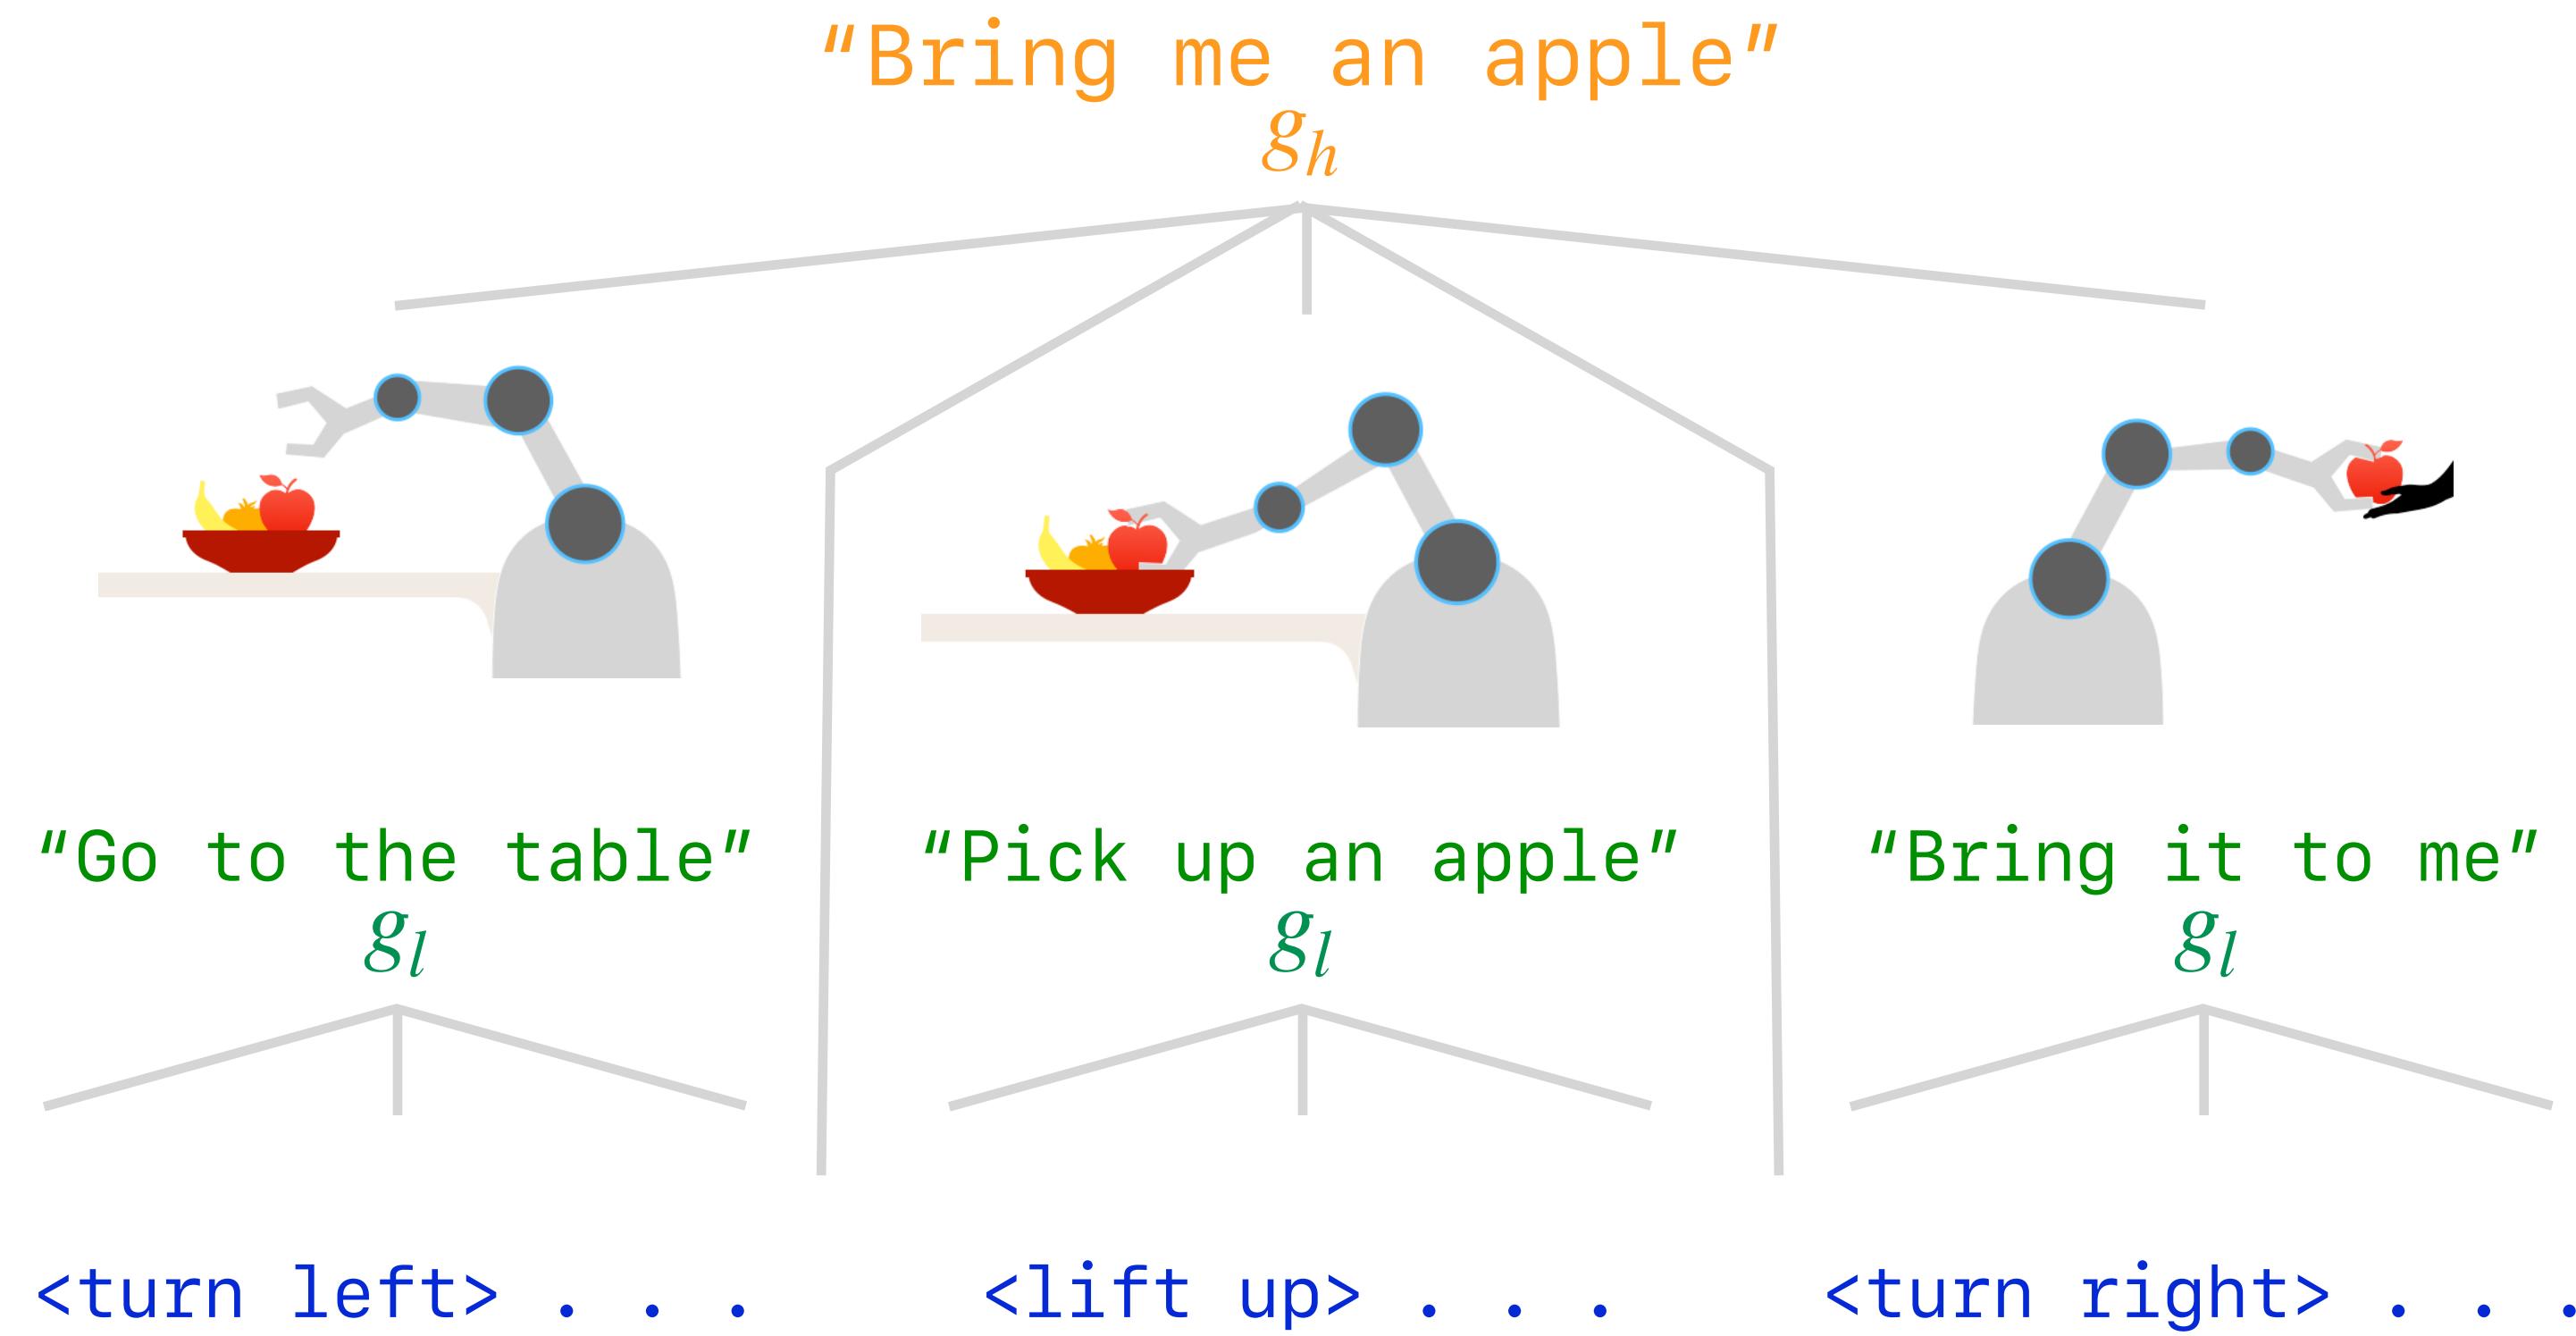
"Bring it to me"

g_l

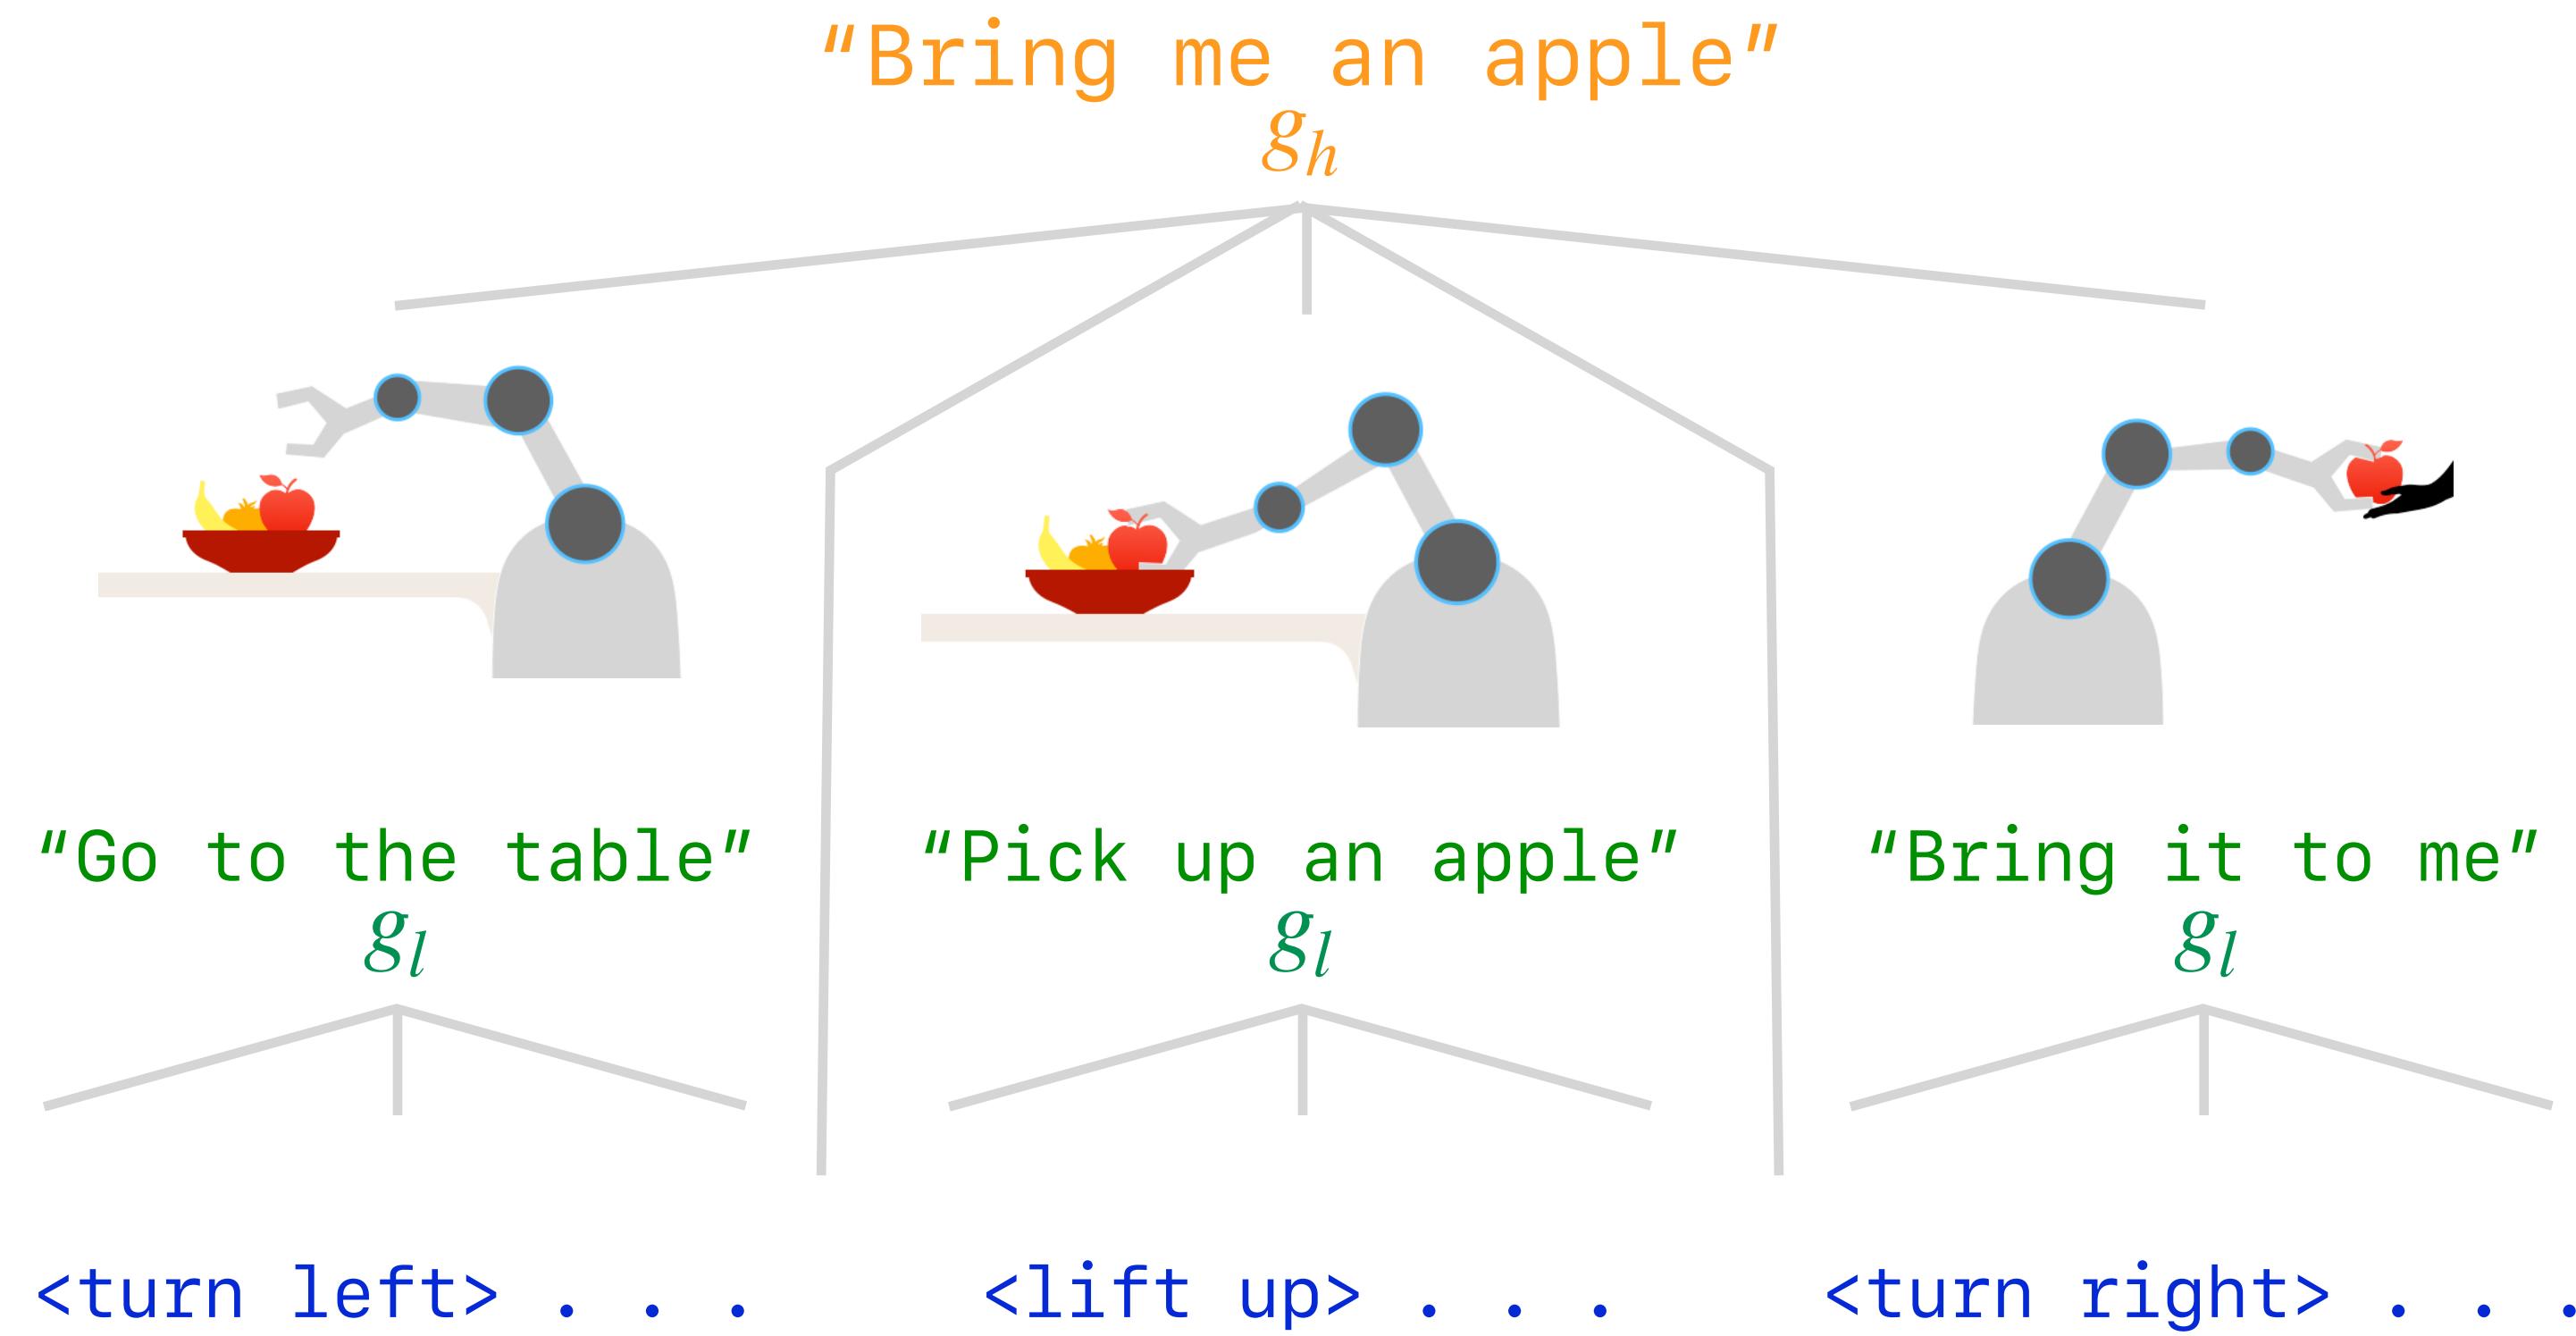
<turn right> . . .



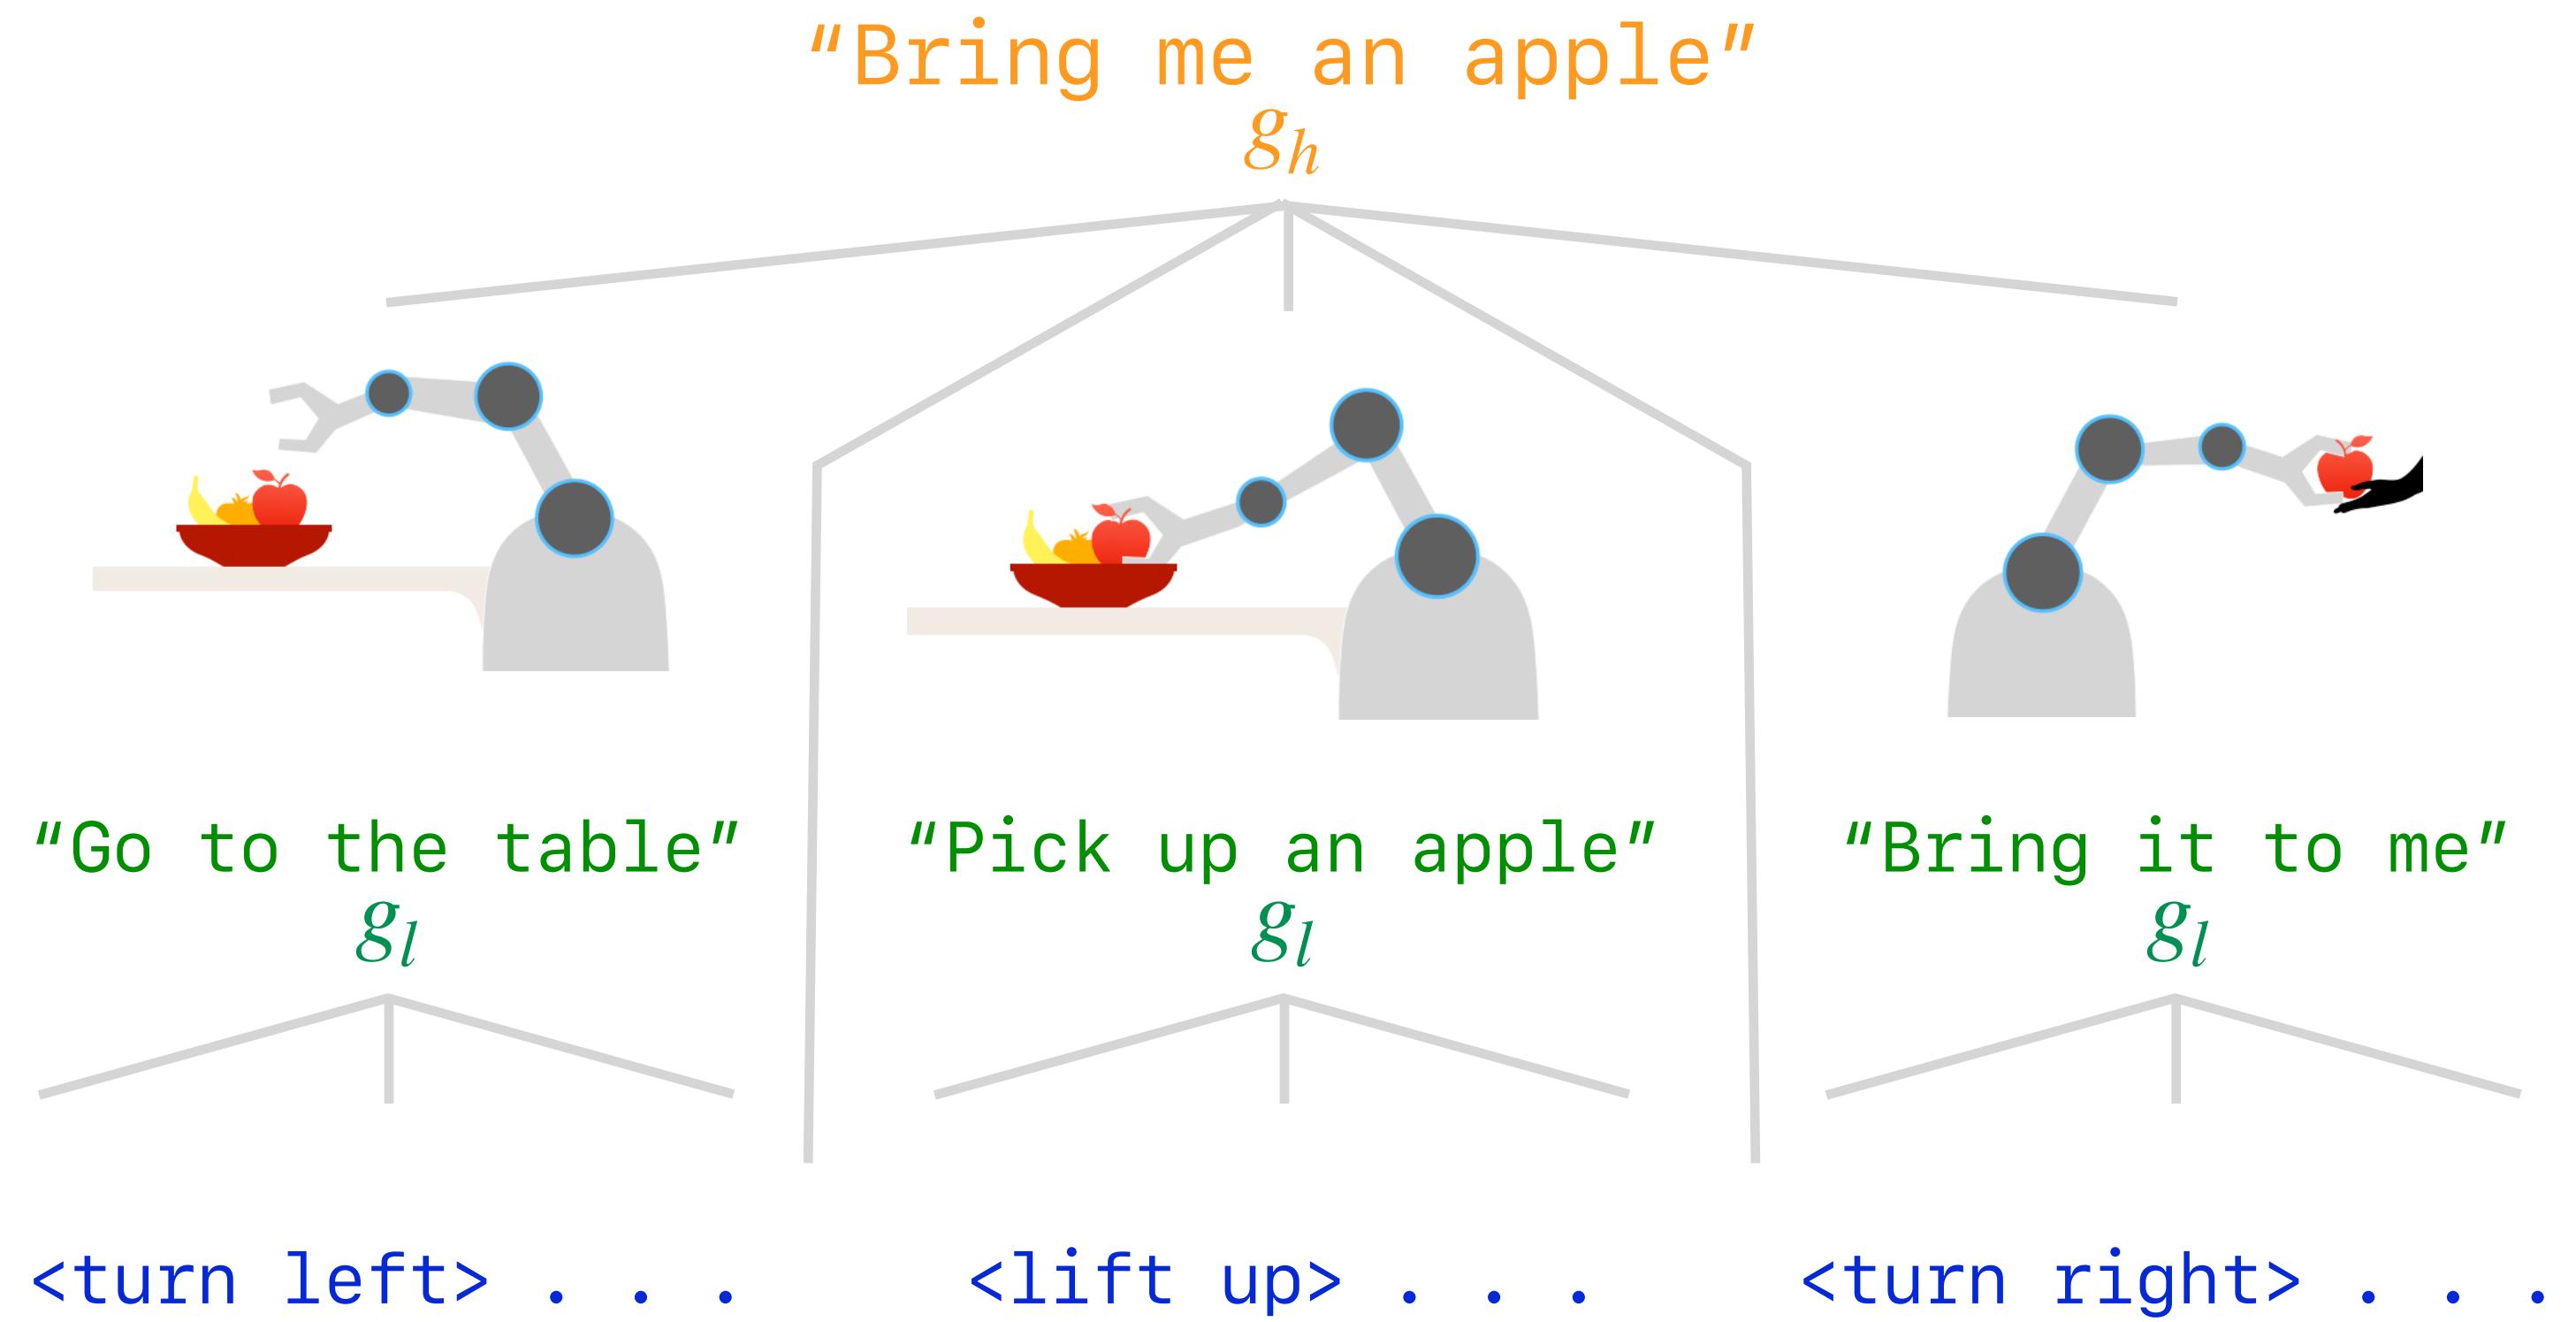
- **Strict factorization**
 - Jiang et al. 2019
 - Das et al. 2018
 - + **Explicit decomposition**
 - Andreas et al. 2017
 - Waytowich et al. 2019



- **Strict factorization**
 - Jiang et al. 2019
 - Das et al. 2018
- + **Explicit decomposition**
 - Andreas et al. 2017
 - Waytowich et al. 2019

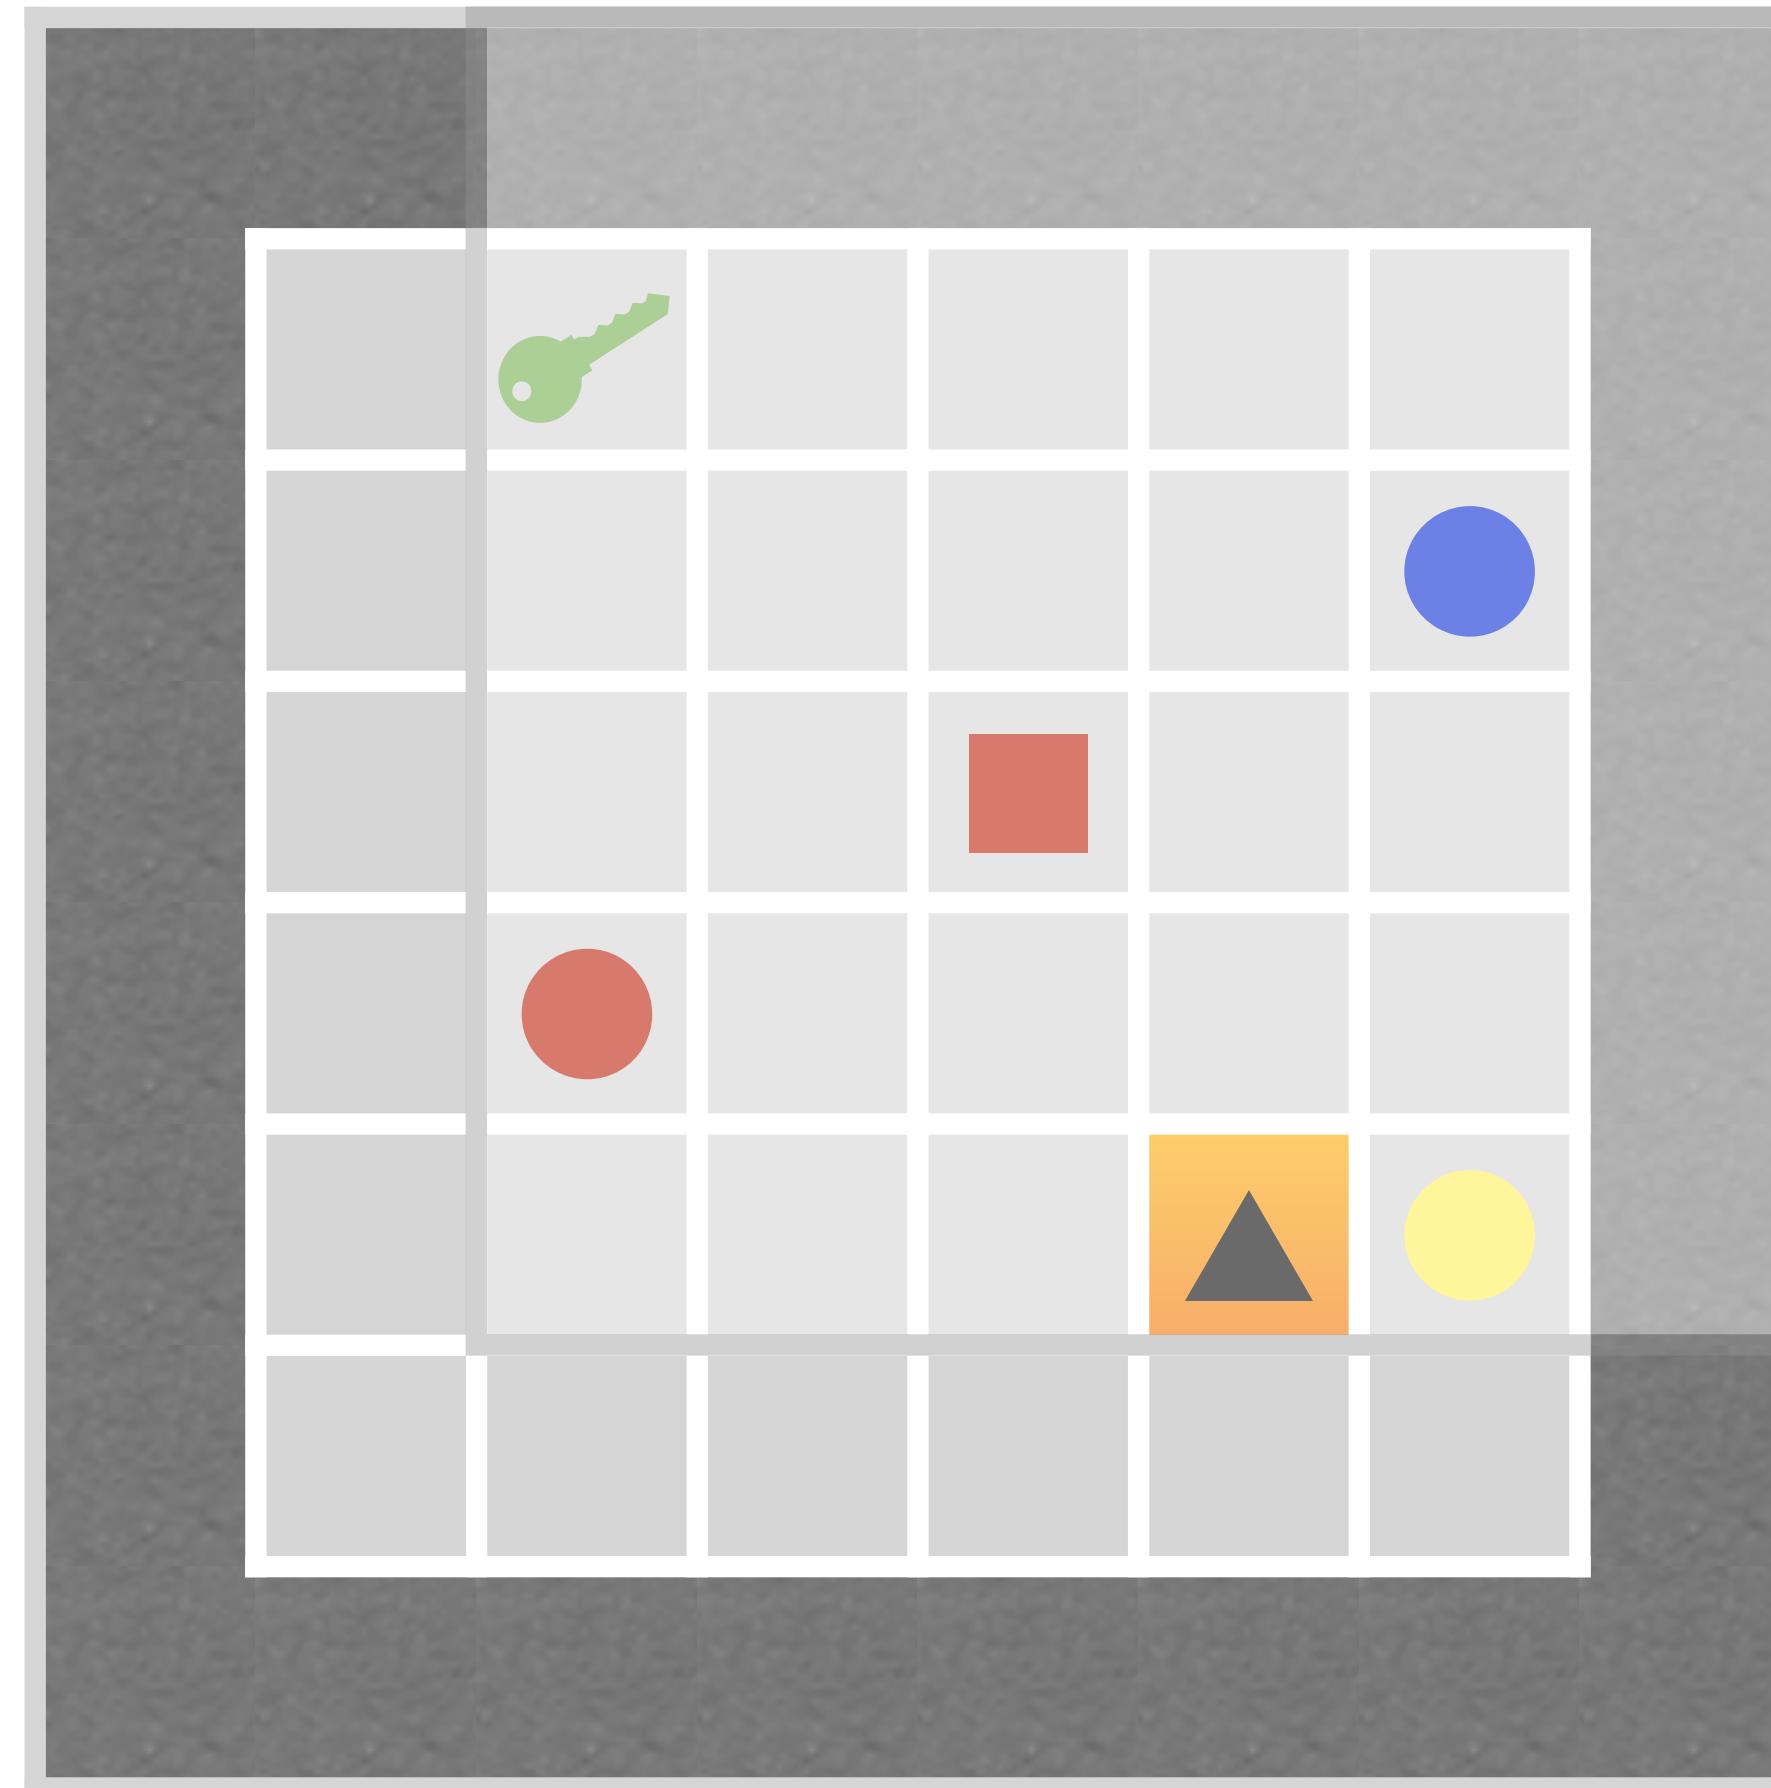


- **Strict factorization**
 - Jiang et al. 2019
 - Das et al. 2018
 - + **Explicit decomposition**
 - Andreas et al. 2017
 - Waytowich et al. 2019
- **Non-abstractive**
 - Goyal et al. 2019



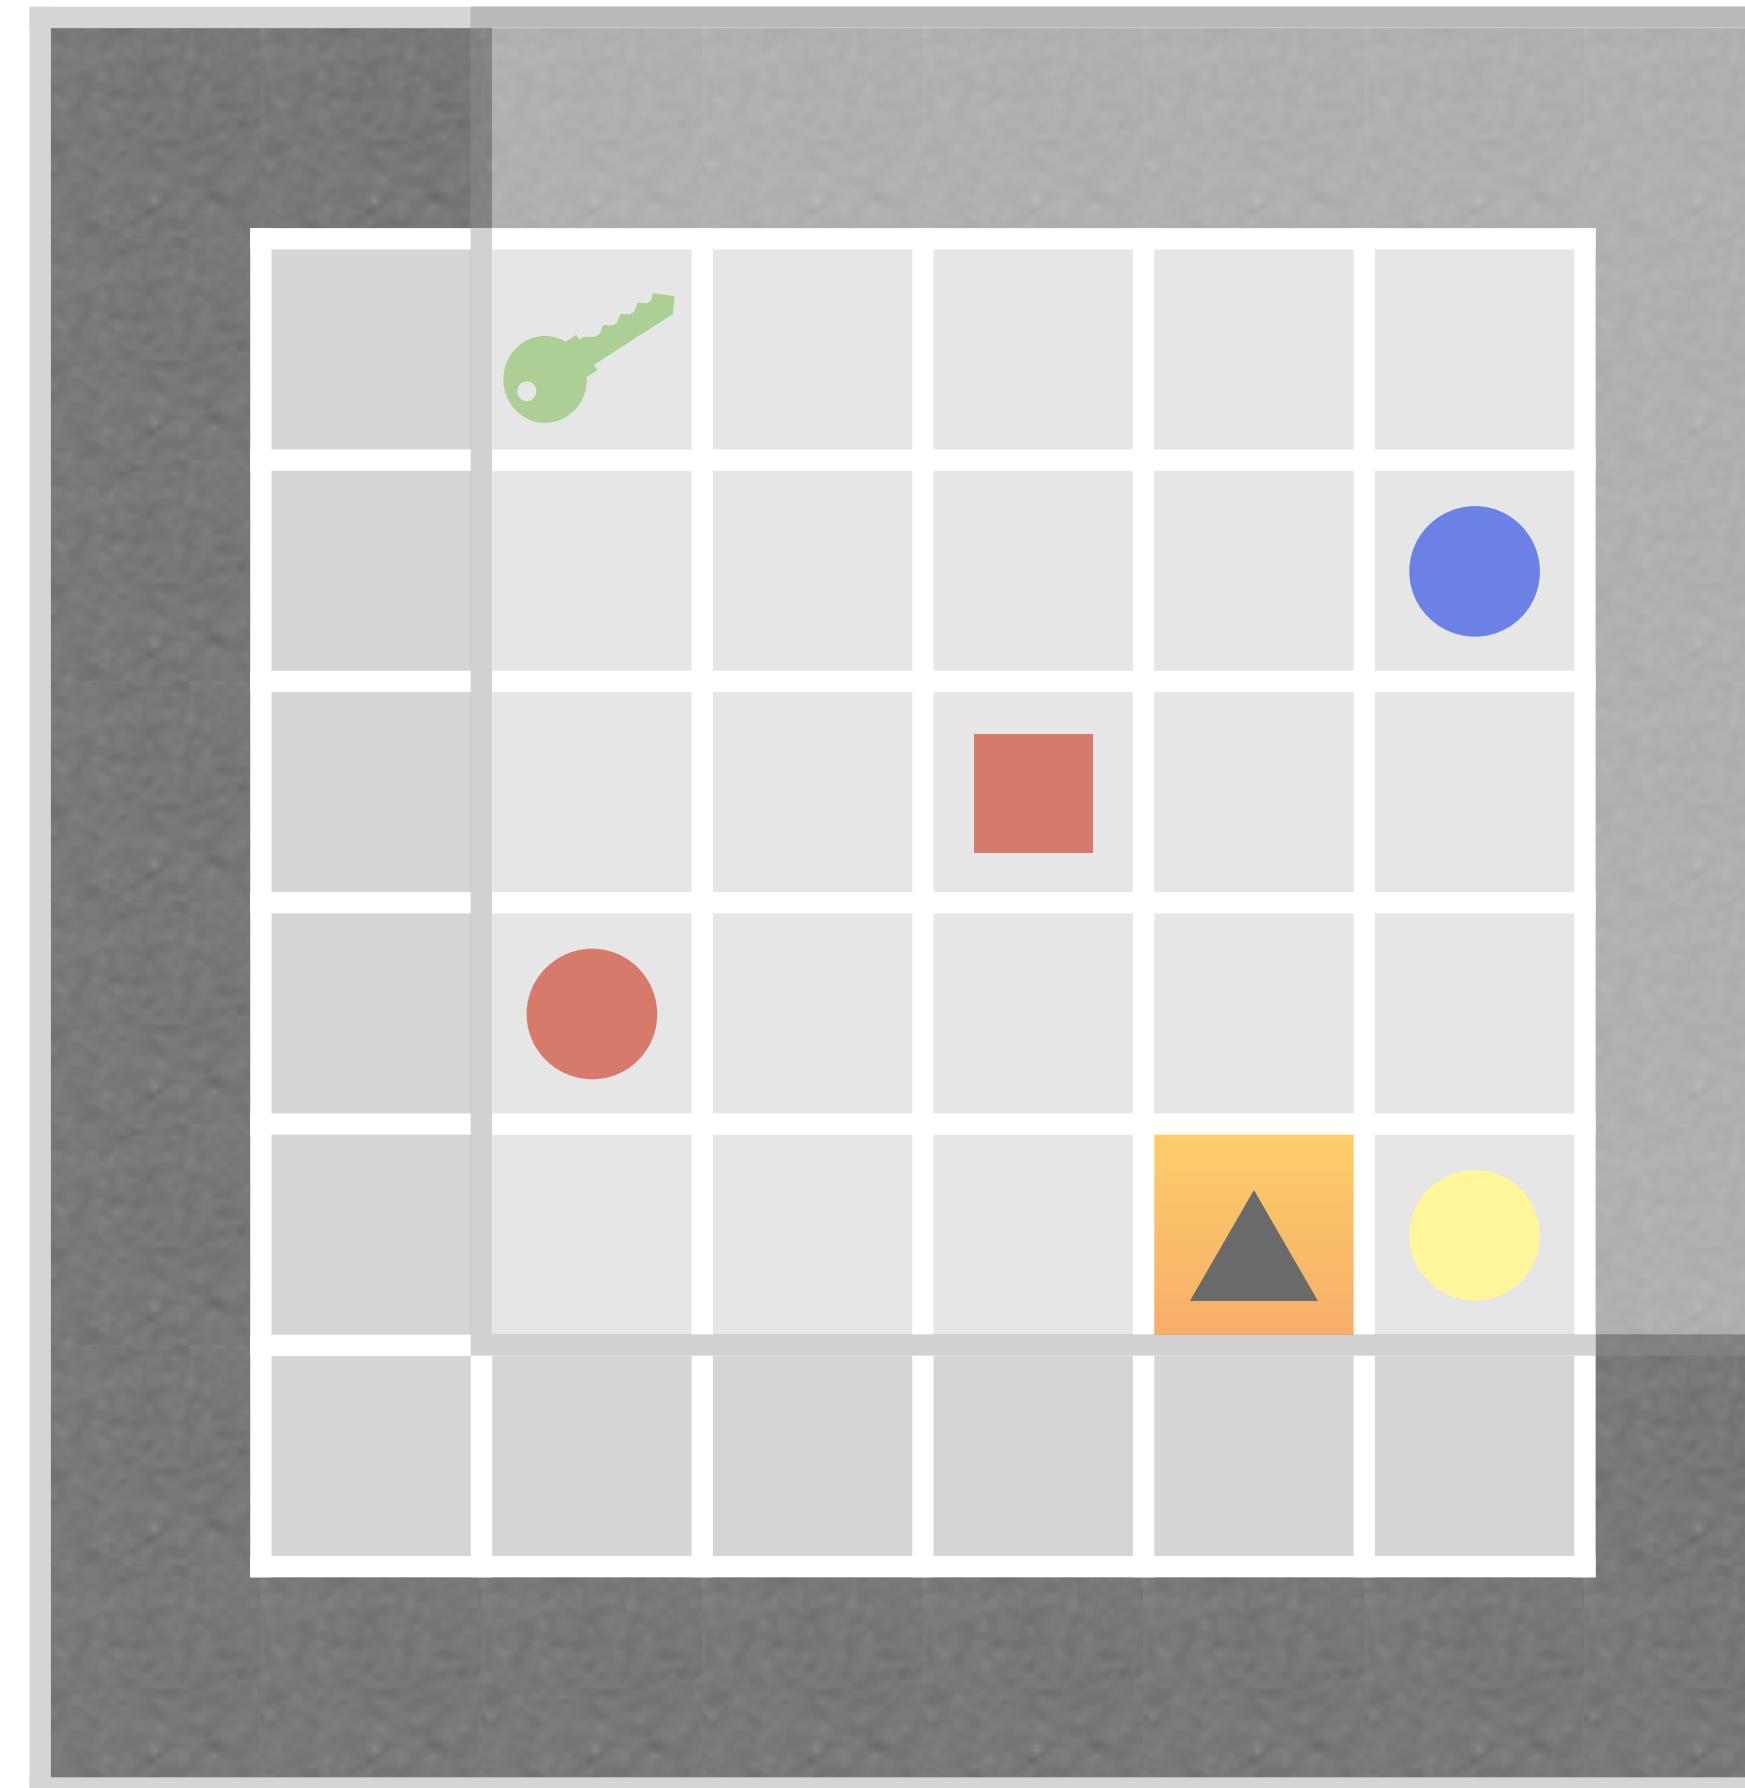
- **Strict factorization**
 - Jiang et al. 2019
 - Das et al. 2018
 - + **Explicit decomposition**
 - Andreas et al. 2017
 - Waytowich et al. 2019
- **Non-abstractive**
 - Goyal et al. 2019
- **Intrinsic motivation & curiosity**
 - Burda et al. 2019

Formulation



$$(S, A, T, R, G, G_\ell, \gamma)$$

Formulation



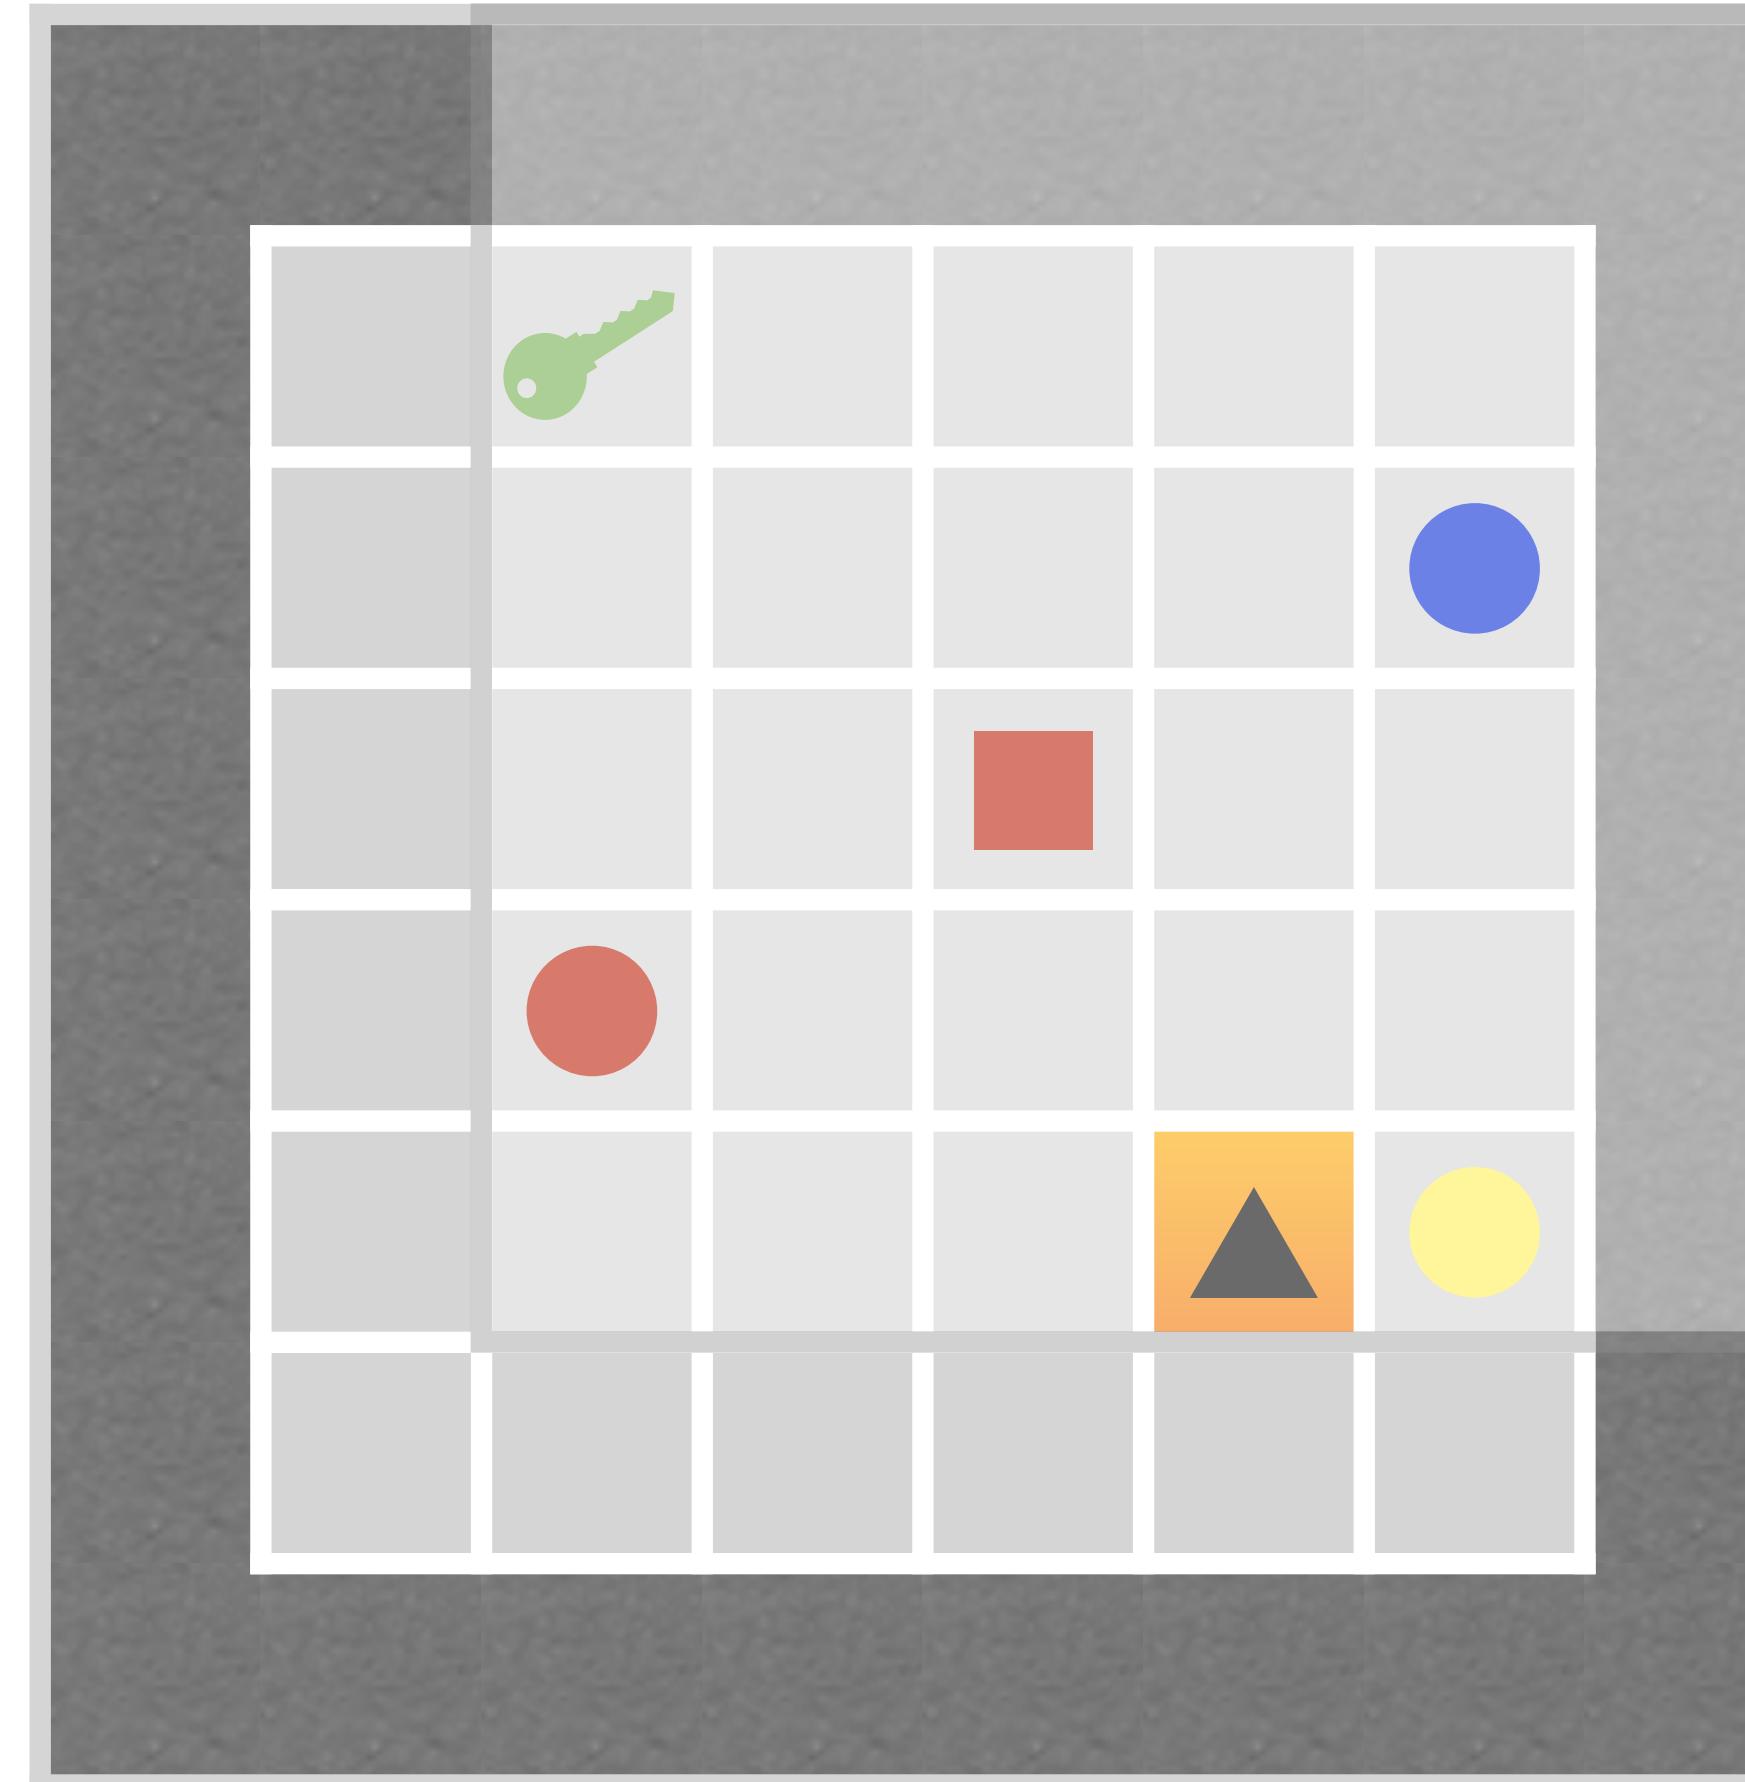
$$(S, A, T, R, G, G_\ell, \gamma)$$

A: Primitive Actions

<turn left>
<turn right>
<move forward>

<pick up>
<put down>
<toggle>
<done>

Formulation



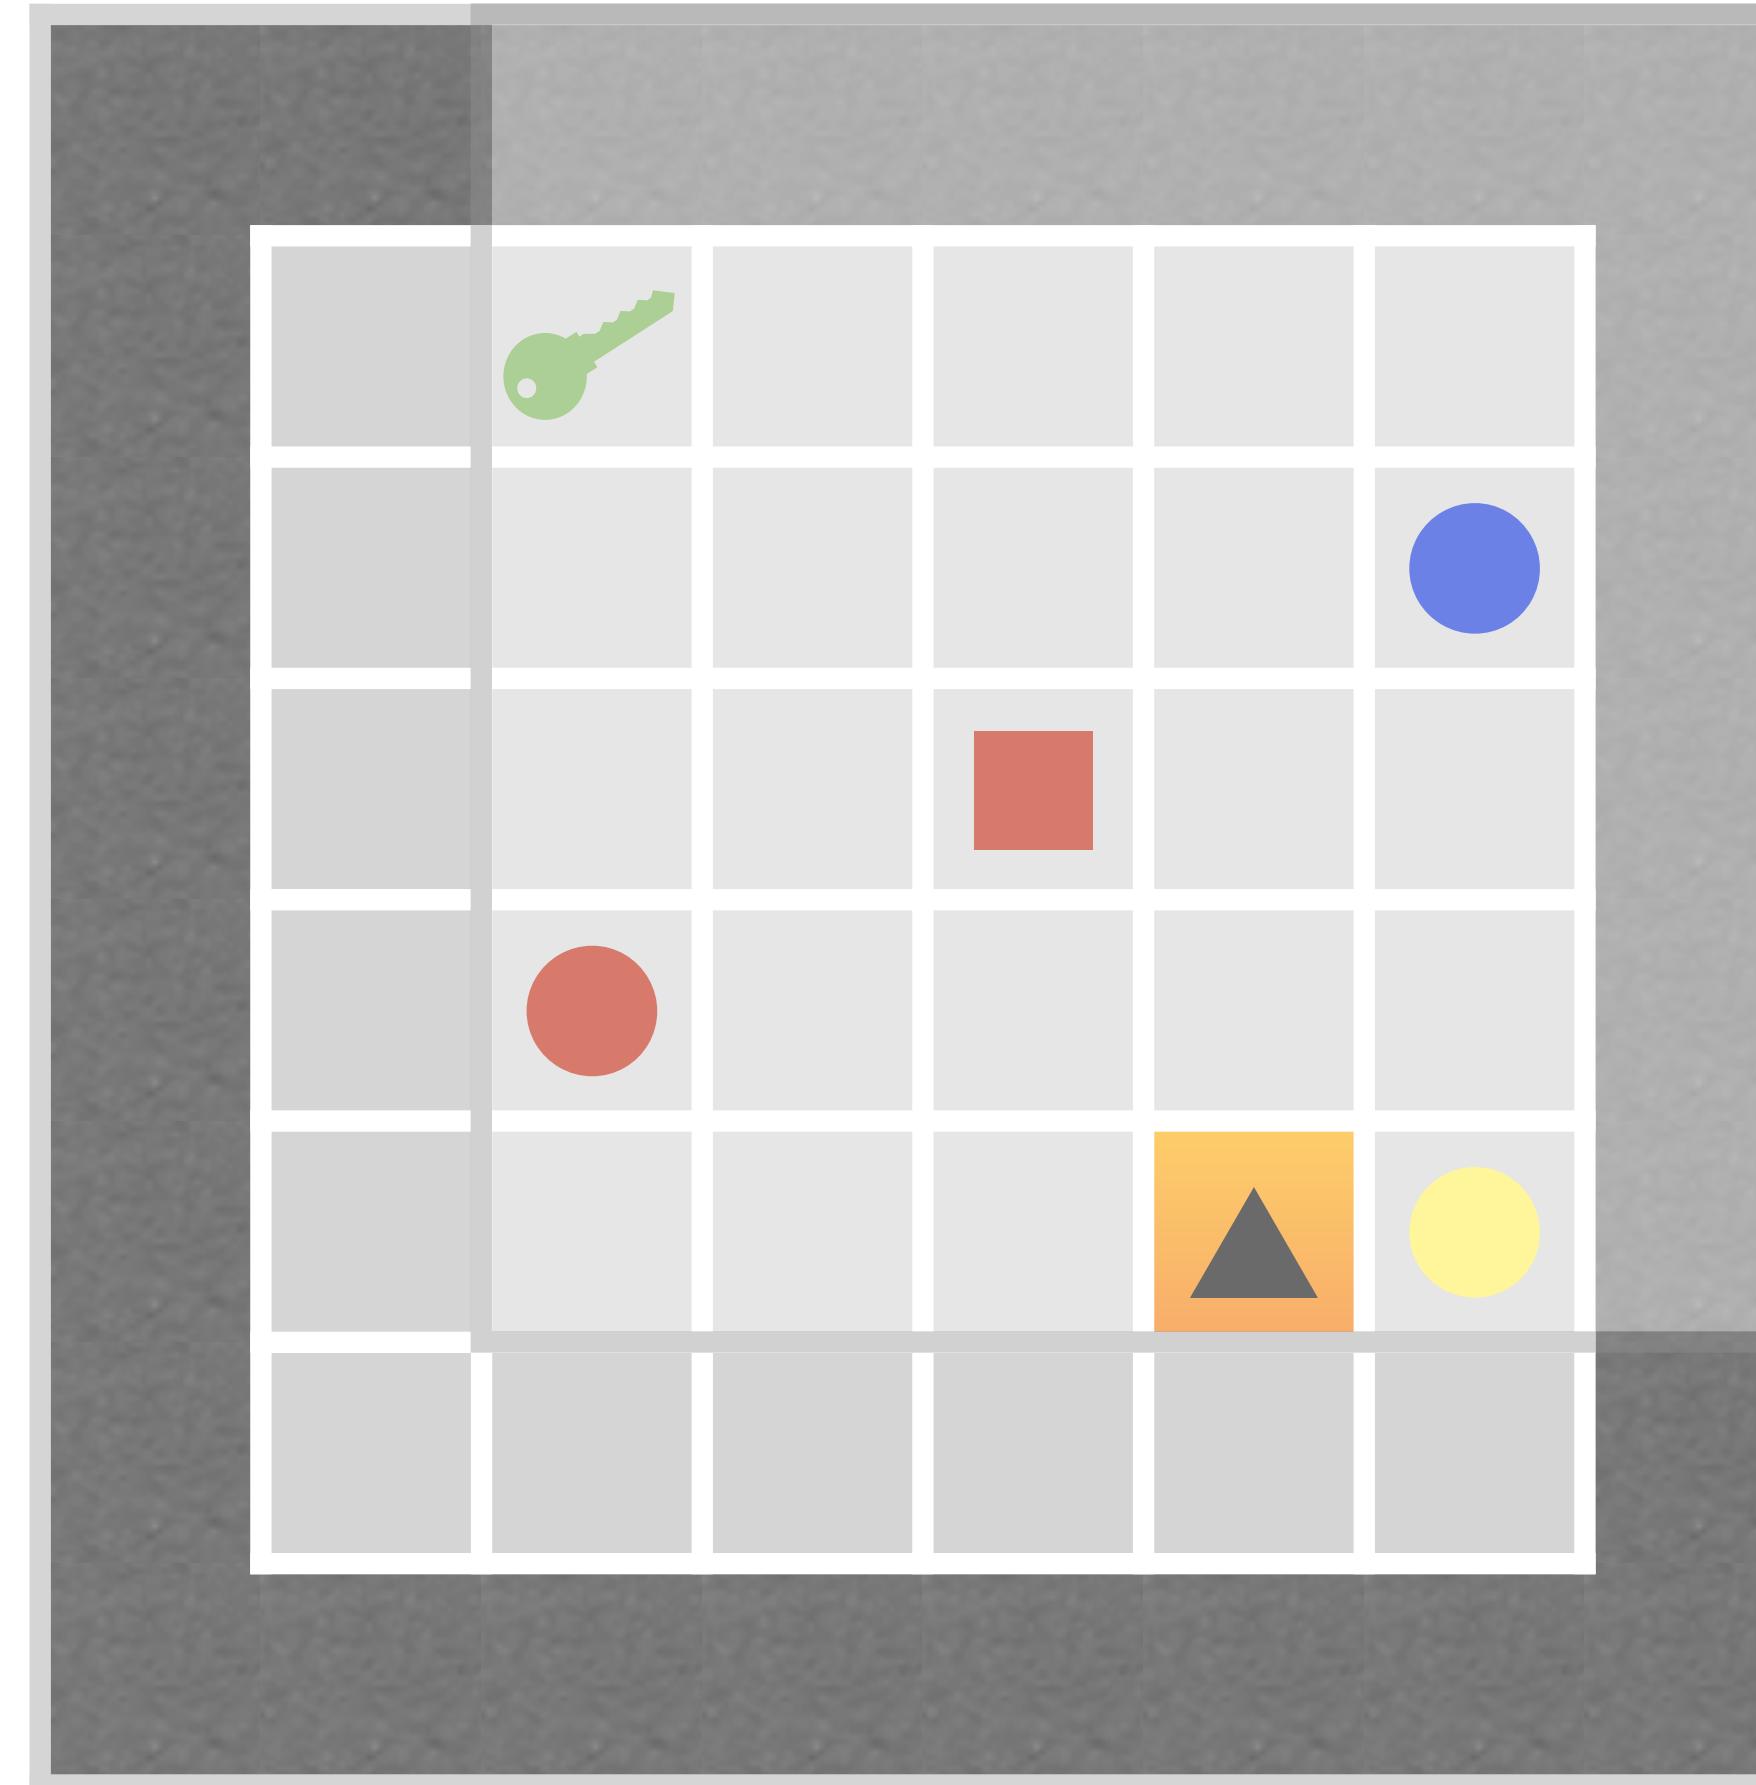
$$(S, A, T, R, G, G_\ell, \gamma)$$

A: Primitive Actions

<turn left>
<turn right>
<move forward>

<pick up>
<put down>
<toggle>
<done>

Formulation



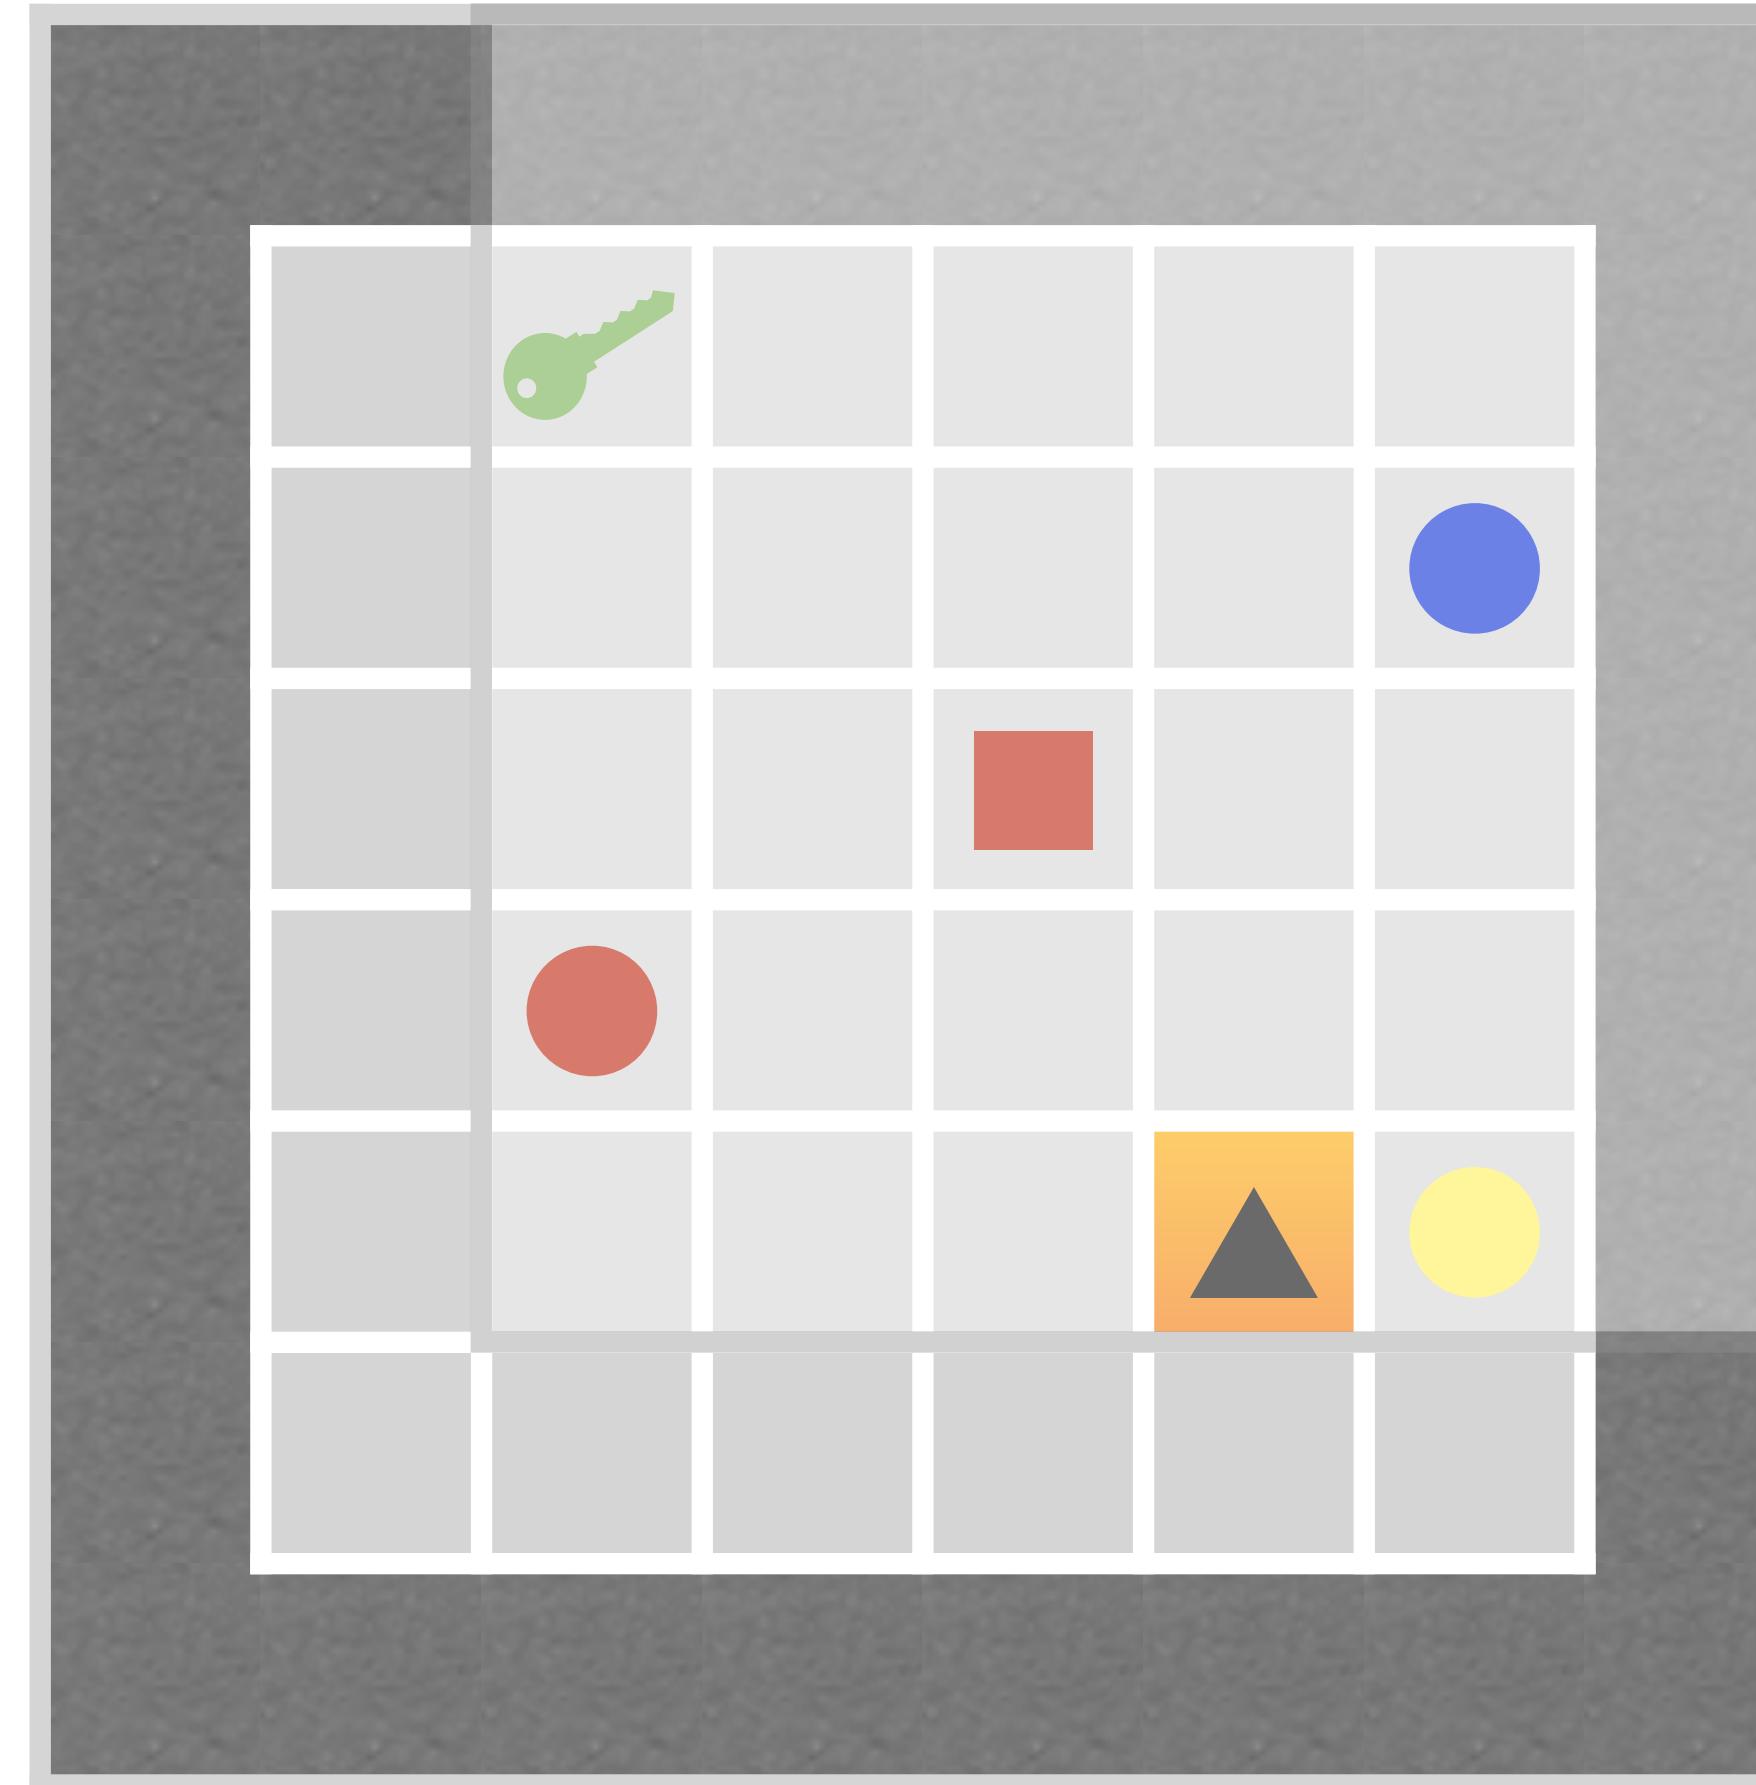
$$(S, A, T, R, G, G_\ell, \gamma)$$

A: Primitive Actions

<turn left>
<turn right>
<move forward>

<pick up>
<put down>
<toggle>
<done>

Formulation



$$(S, A, T, R, G, G_\ell, \gamma)$$

G: High-Level Instructions

"go to the red ball and
then to the blue ball"

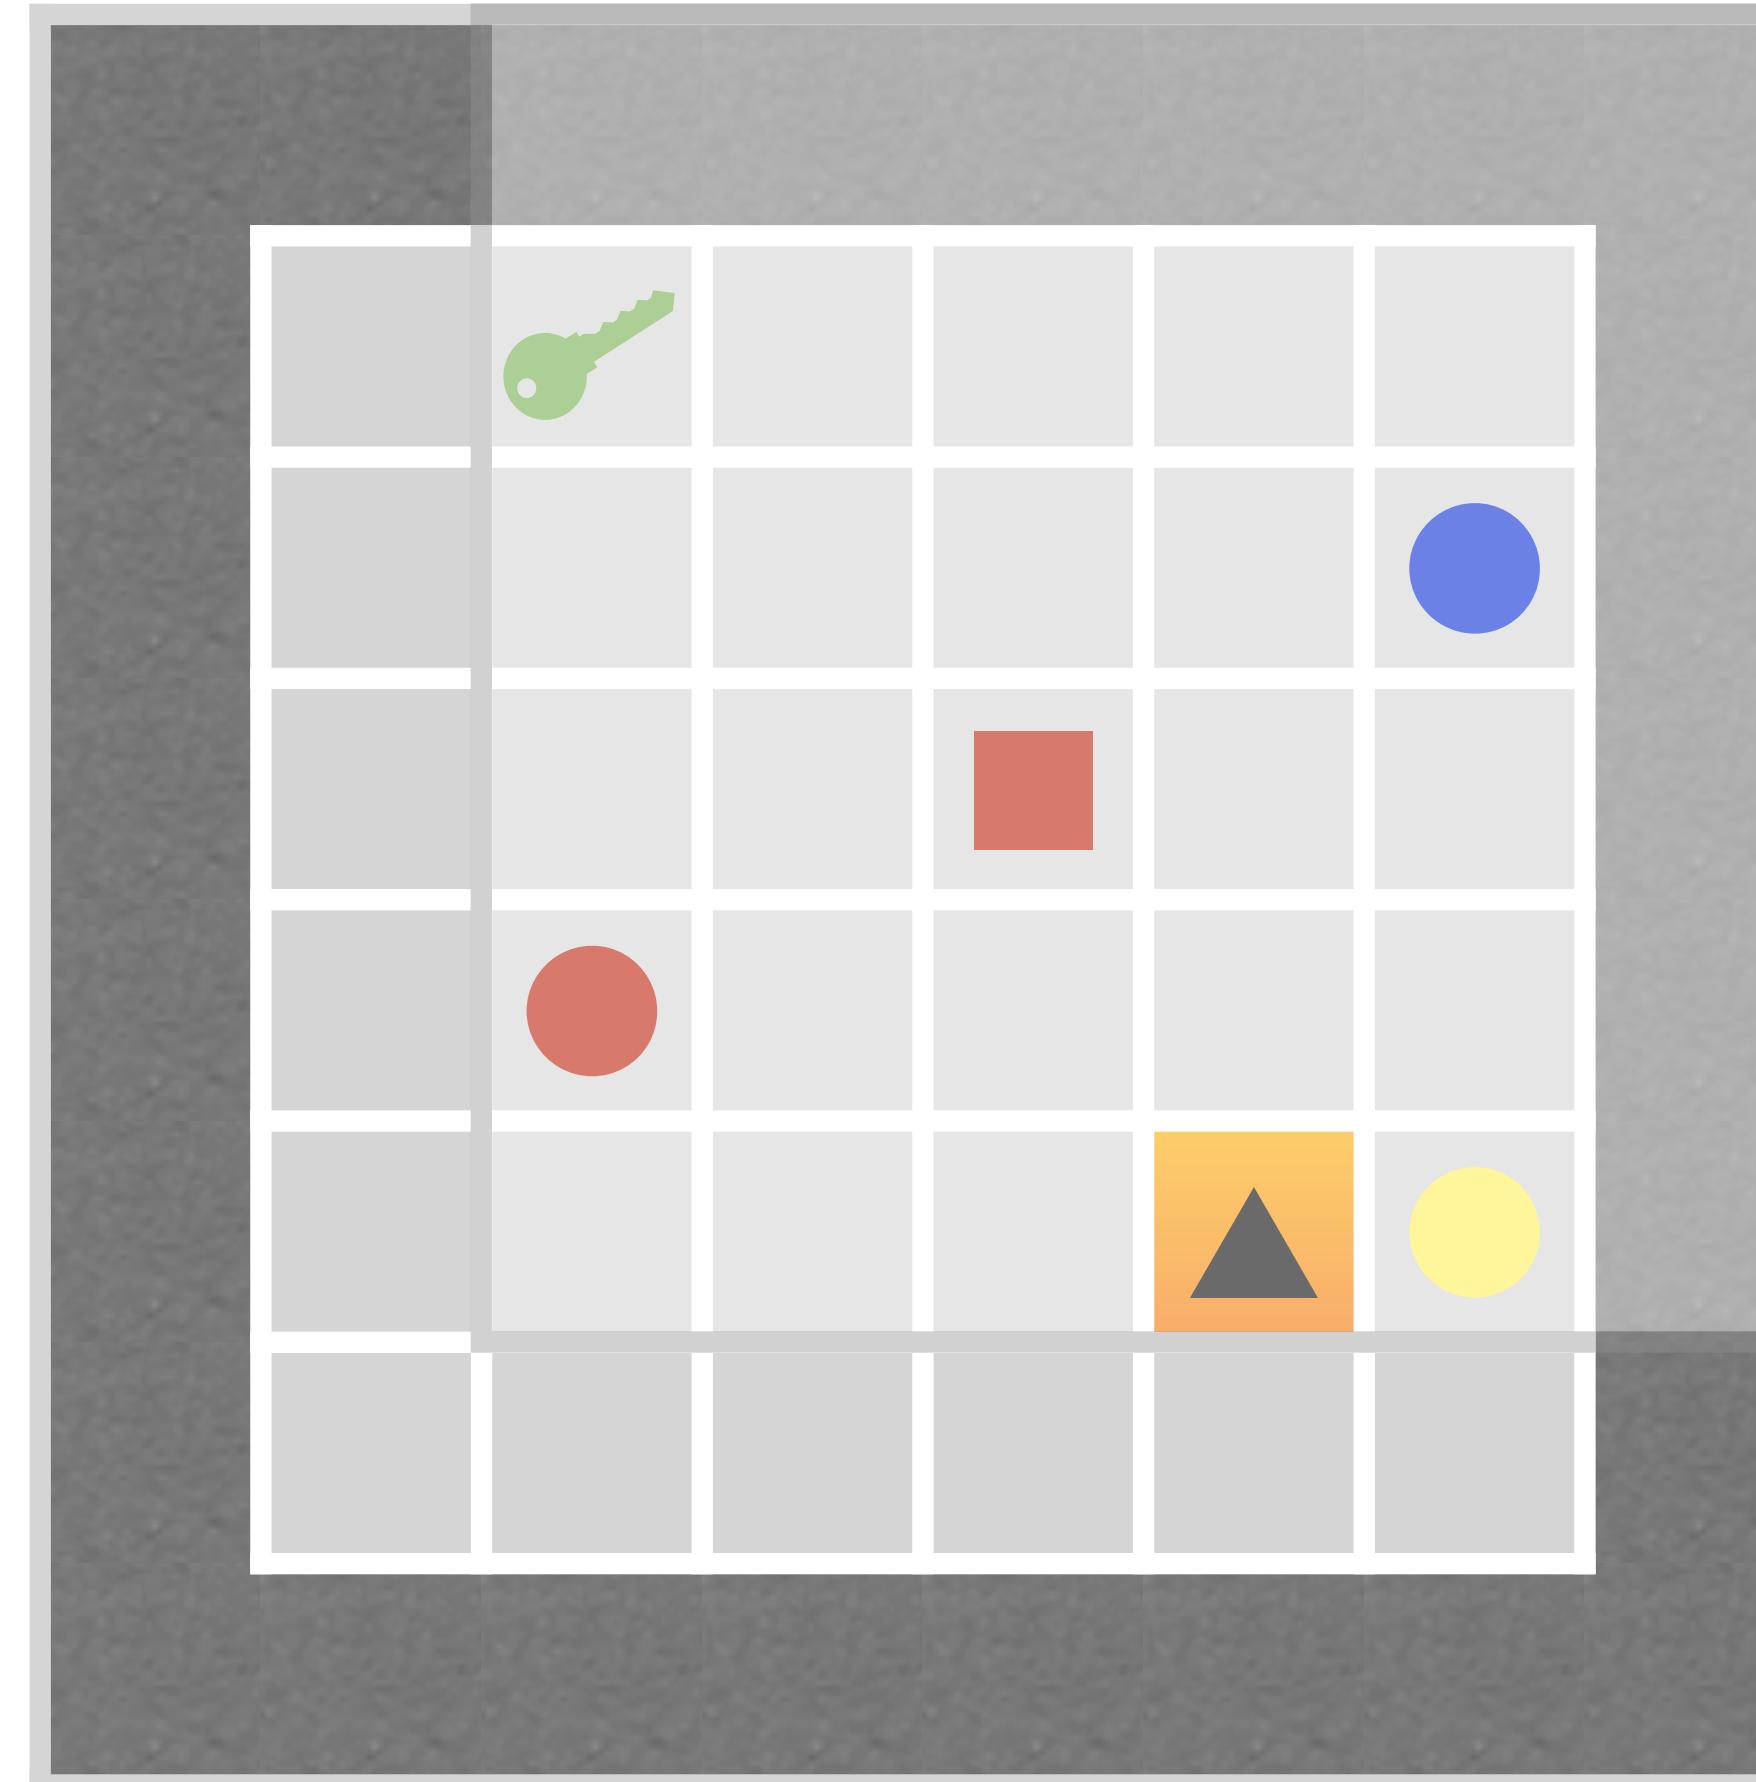
"put the red ball next to
the blue ball"

A: Primitive Actions

<turn left>
<turn right>
<move forward>

<pick up>
<put down>
<toggle>
<done>

Formulation



$$(S, A, T, R, G, G_\ell, \gamma)$$

G : High-Level Instructions

"go to the red ball and
then to the blue ball"

"put the red ball next to
the blue ball"

G_ℓ : Low-Level Instructions

"go to the red ball"

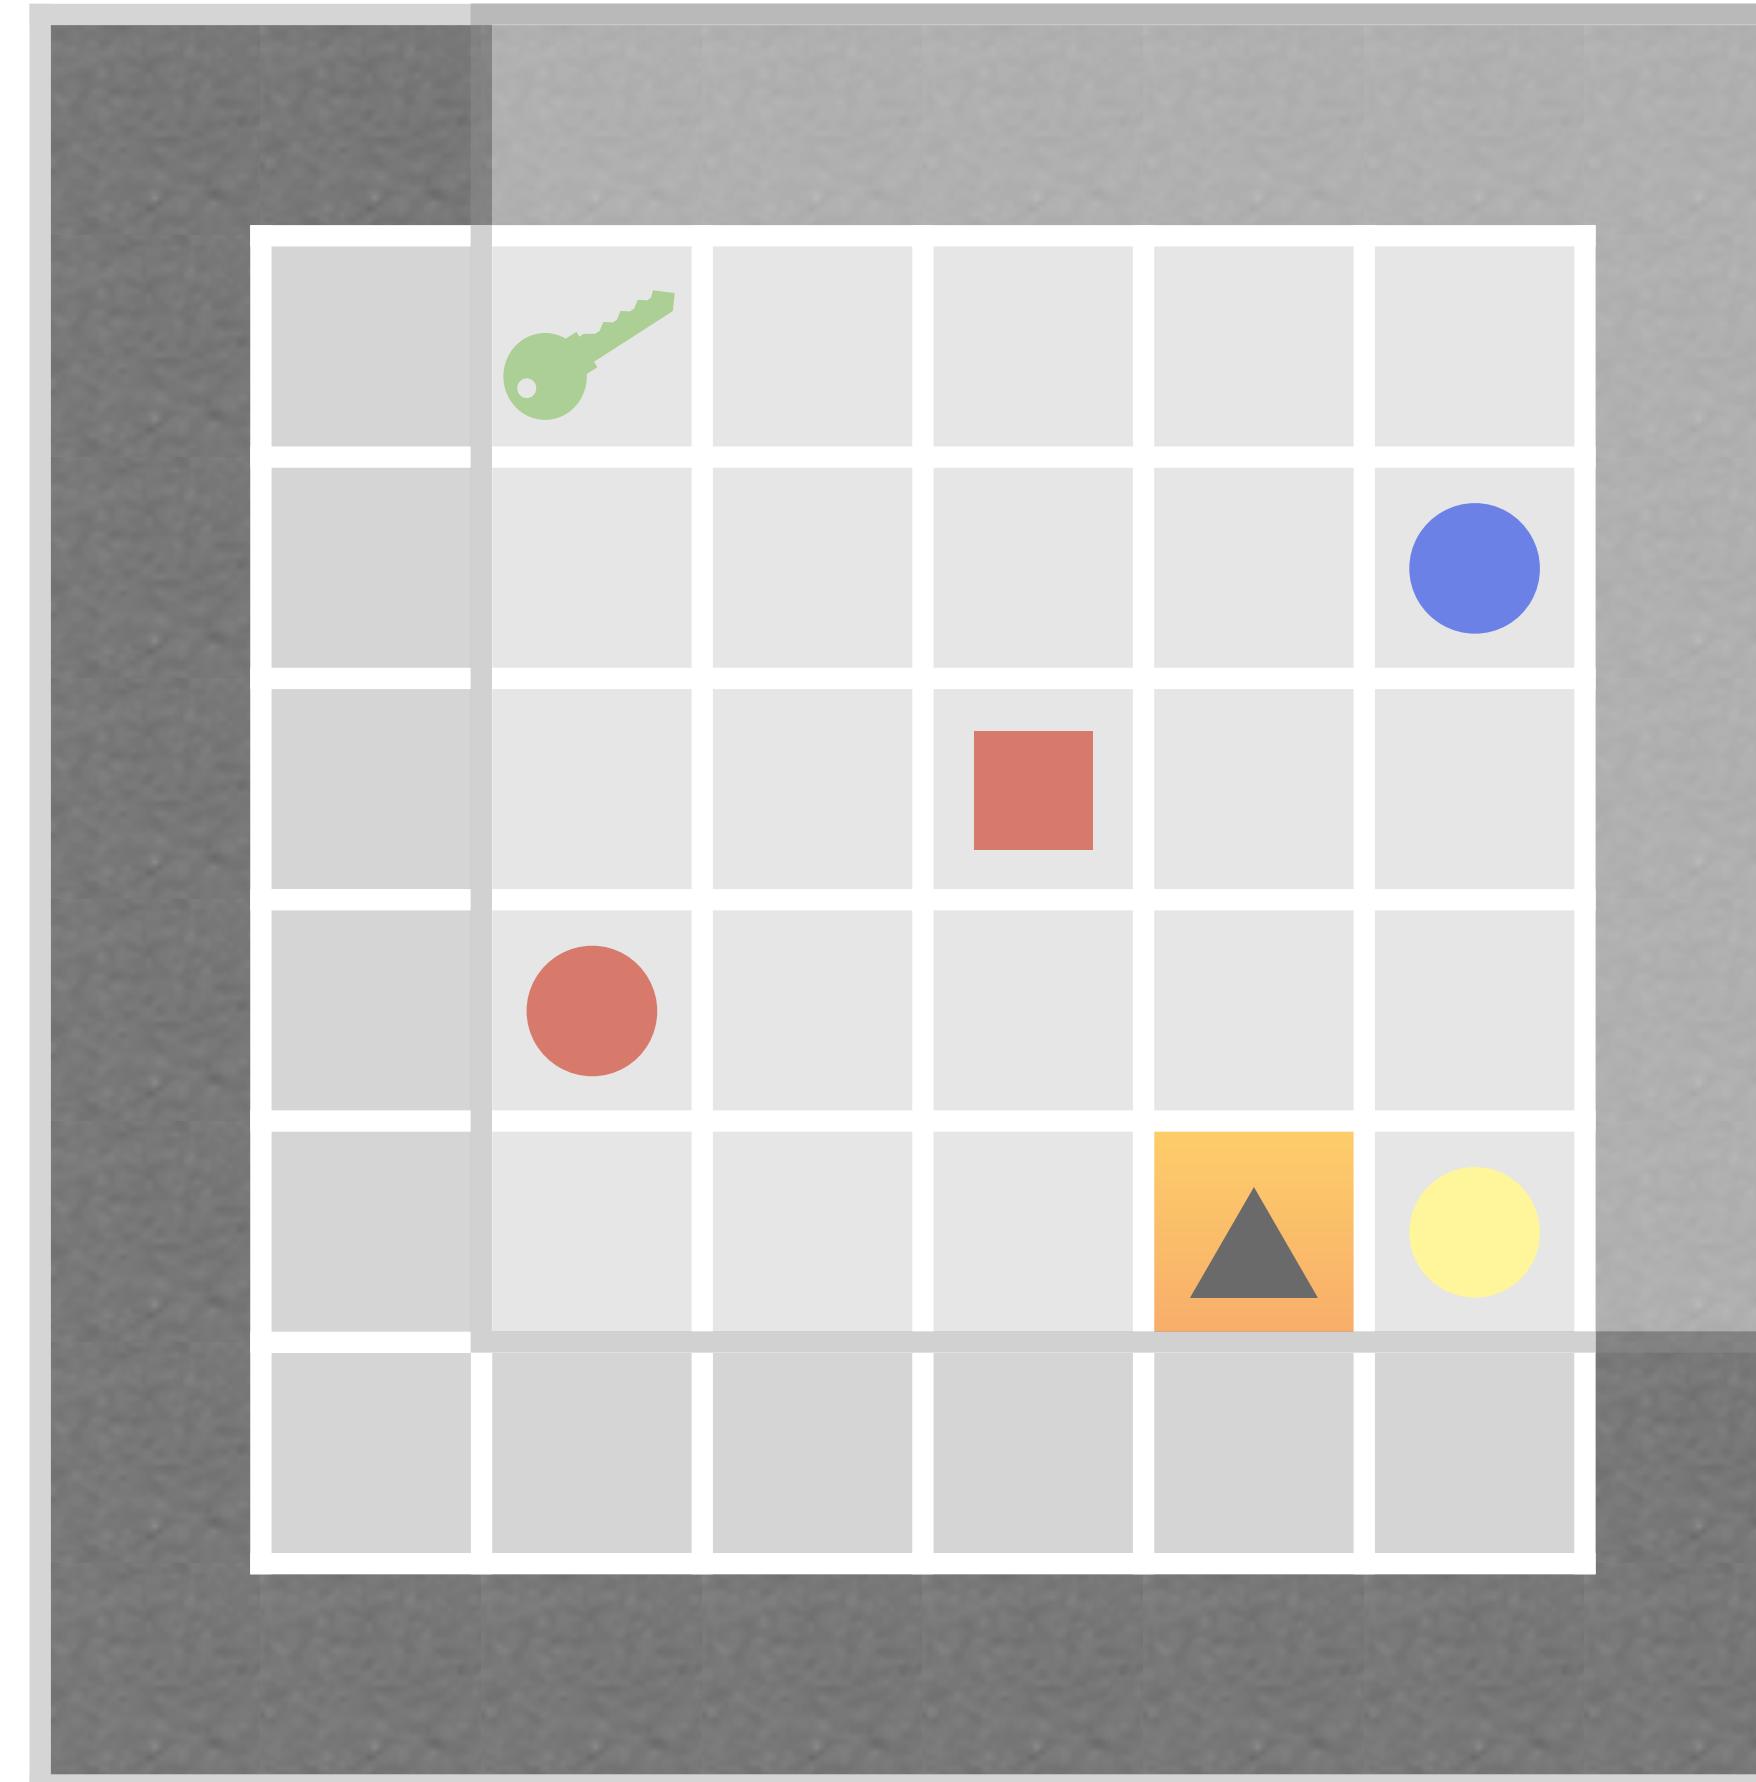
"go to the blue ball"

A : Primitive Actions

<turn left>
<turn right>
<move forward>

<pick up>
<put down>
<toggle>
<done>

Formulation



$$(S, A, T, R, G, G_\ell, \gamma)$$

G : High-Level Instructions

"go to the red ball and
then to the blue ball"

"put the red ball next to
the blue ball"

G_ℓ : Low-Level Instructions

"go to the red ball"

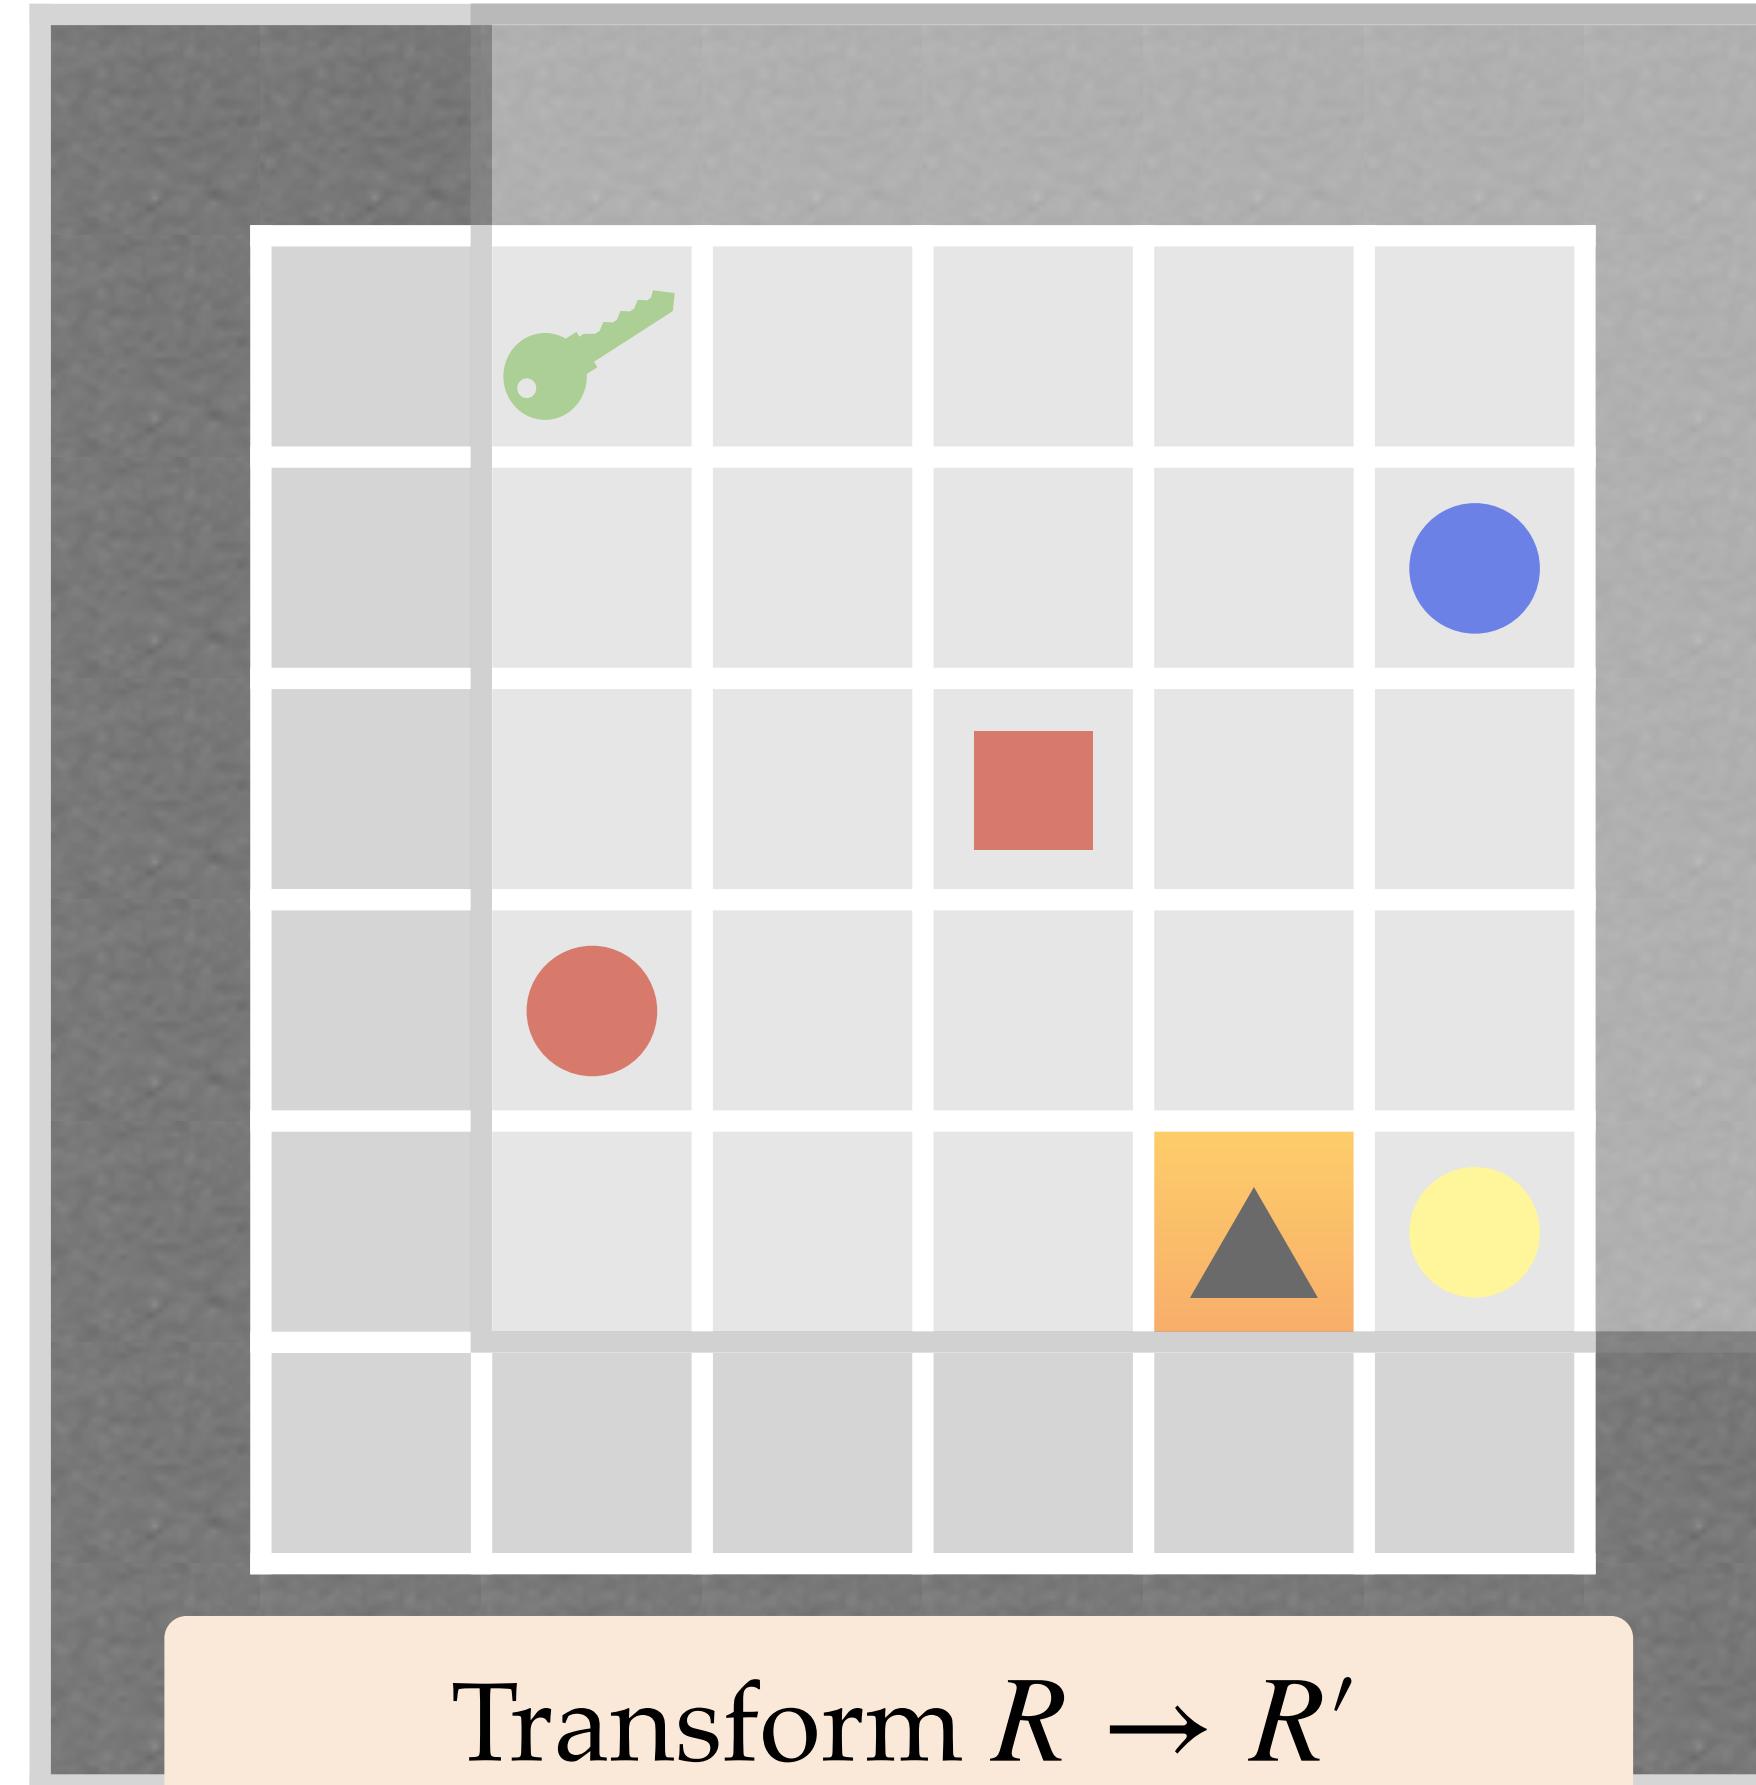
"go to the blue ball"

A : Primitive Actions

<turn left>
<turn right>
<move forward>

<pick up>
<put down>
<toggle>
<done>

Formulation



$$(S, A, T, R, G, G_\ell, \gamma)$$

G : High-Level Instructions

"go to the red ball and
then to the blue ball"

"put the red ball next to
the blue ball"

G_ℓ : Low-Level Instructions

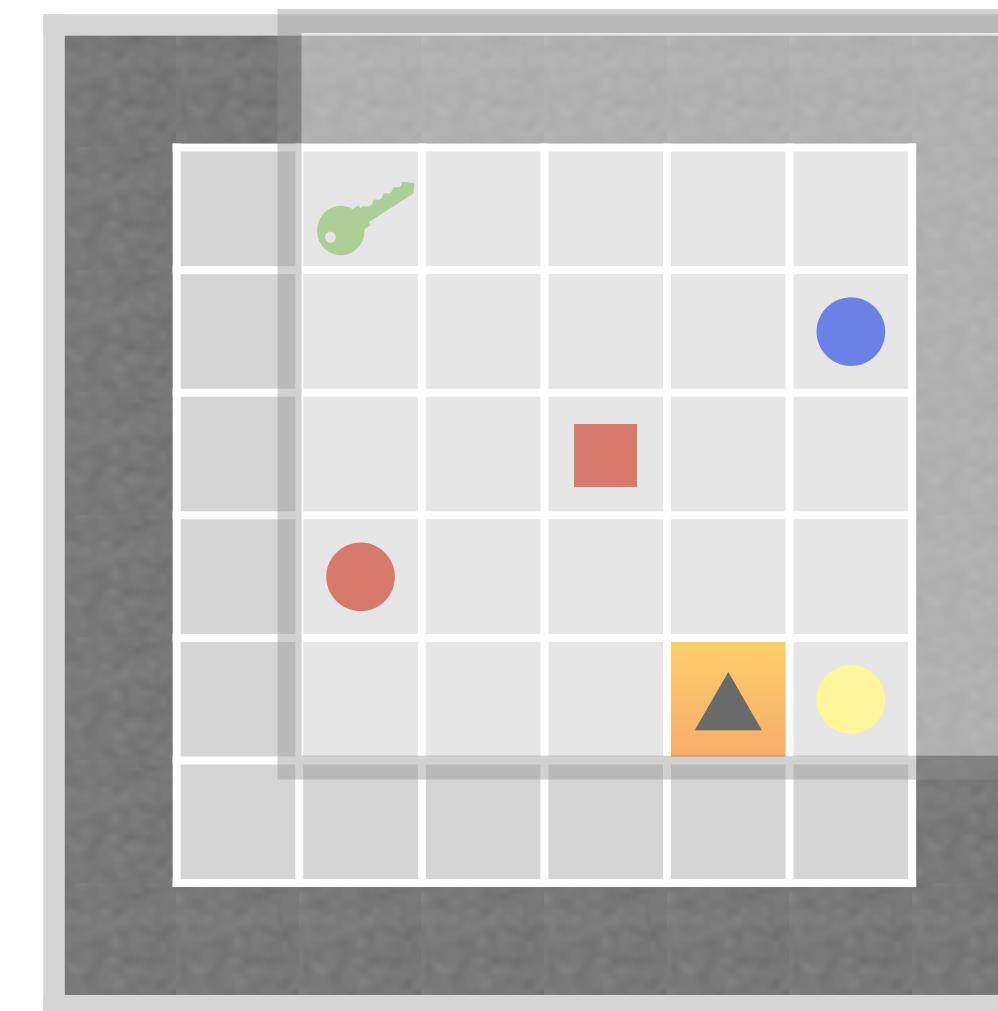
"go to the red ball"

"go to the blue ball"

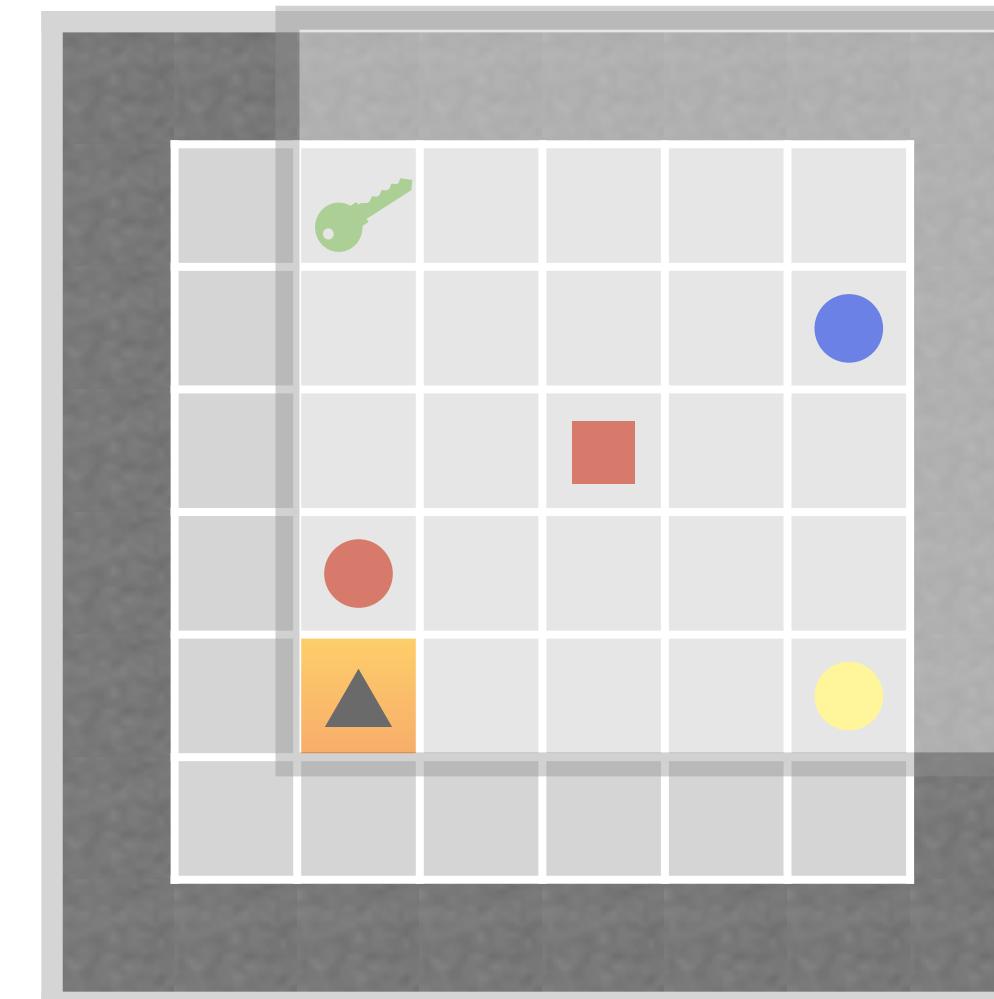
A : Primitive Actions

<turn left>
<turn right>
<move forward>

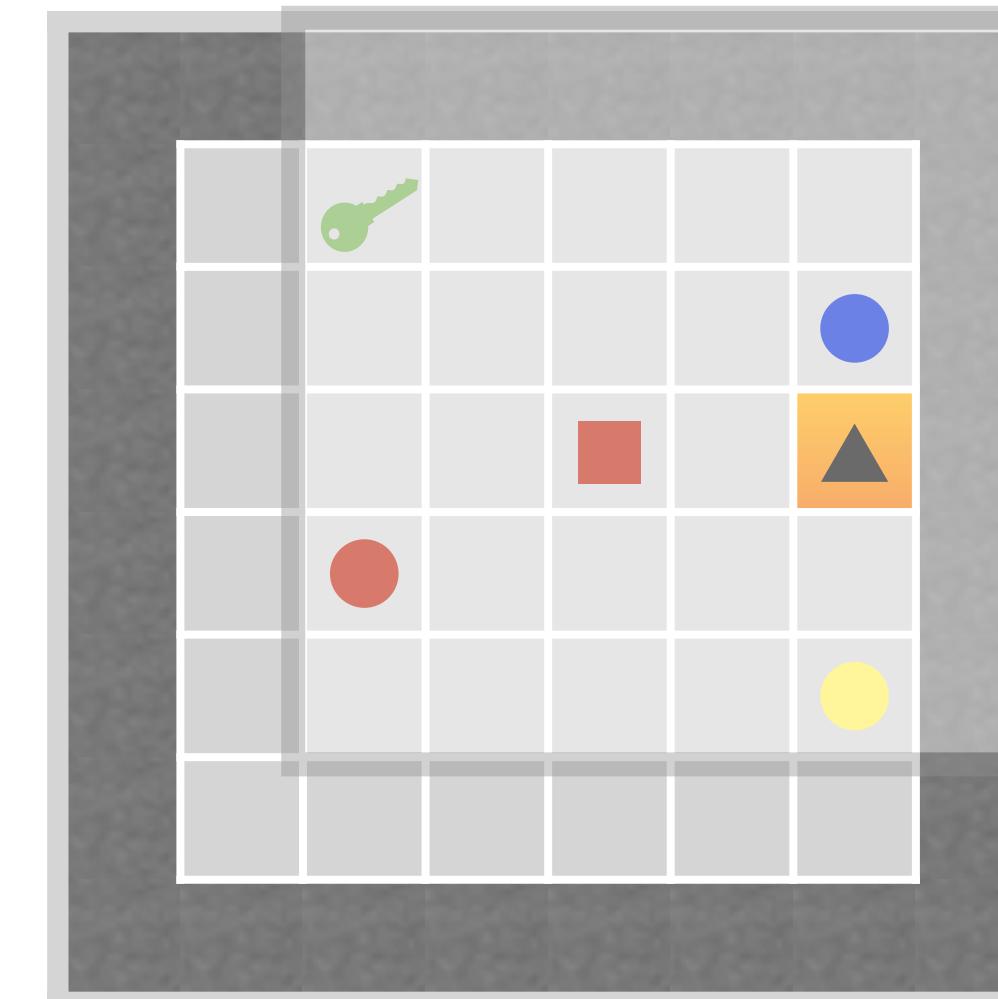
<pick up>
<put down>
<toggle>
<done>



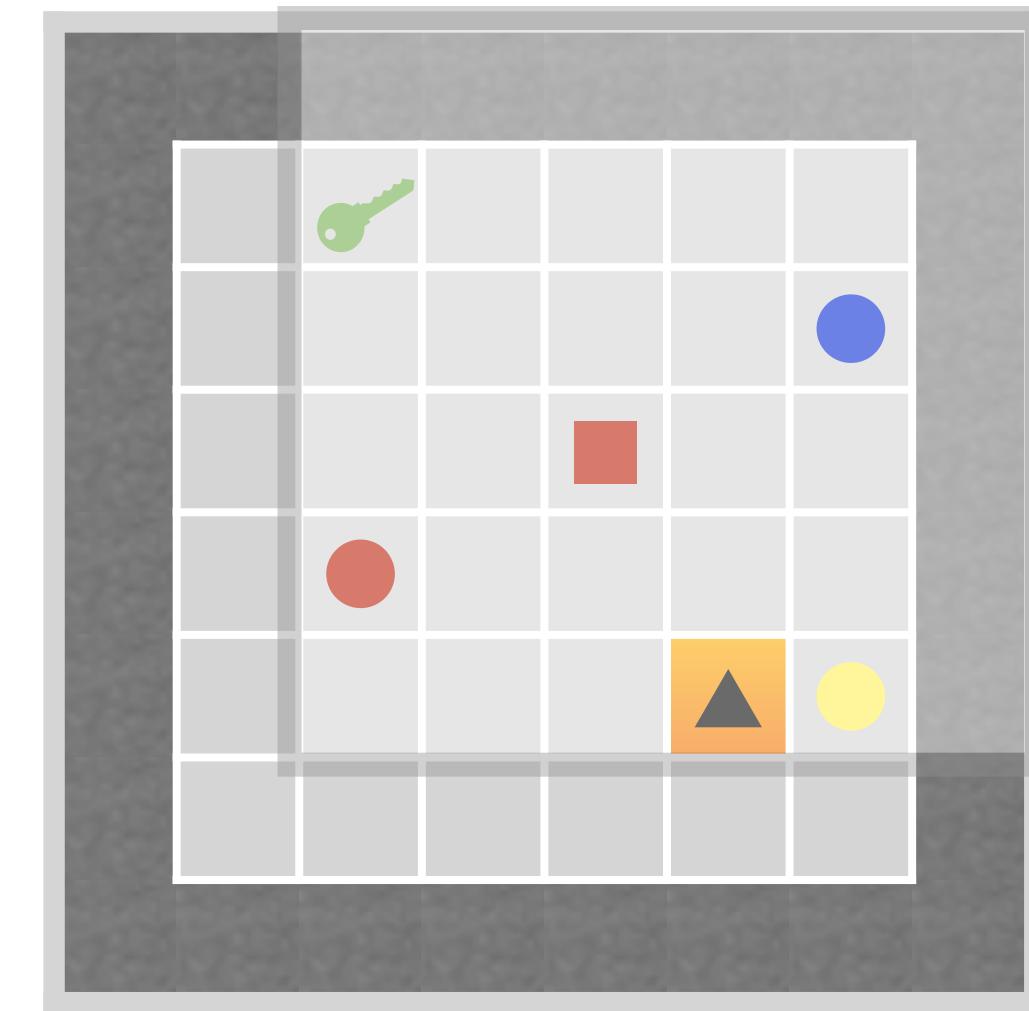
"go to the red ball and then to the blue ball"



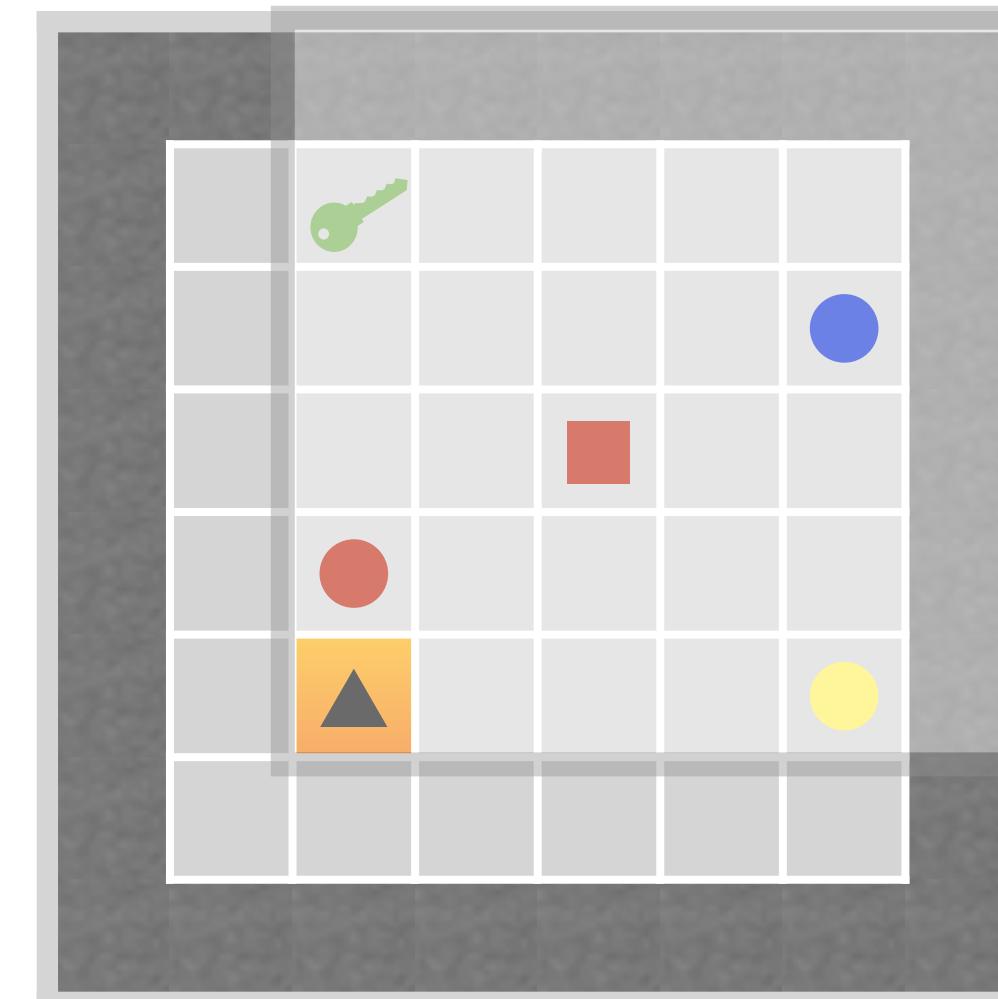
"go to the red ball"



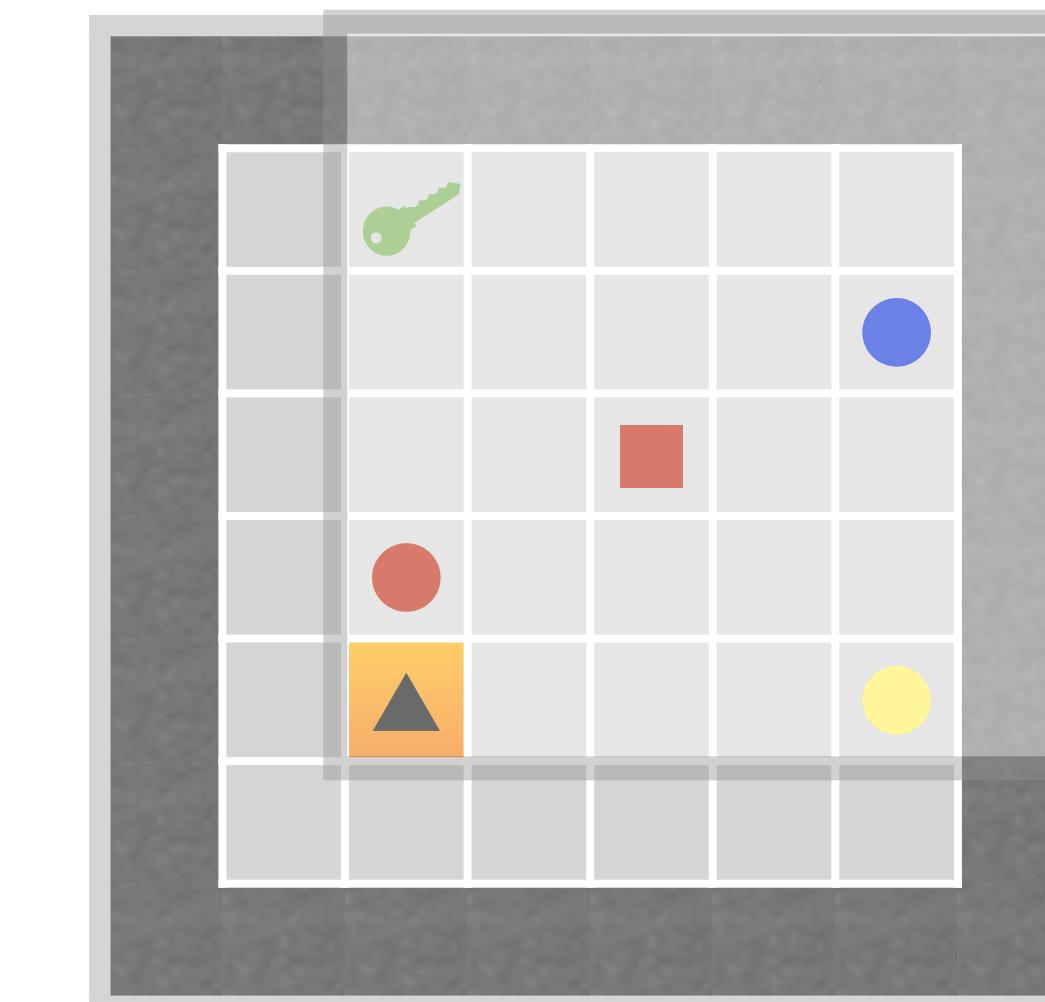
"go to the blue ball"



put the red ball next to the blue ball



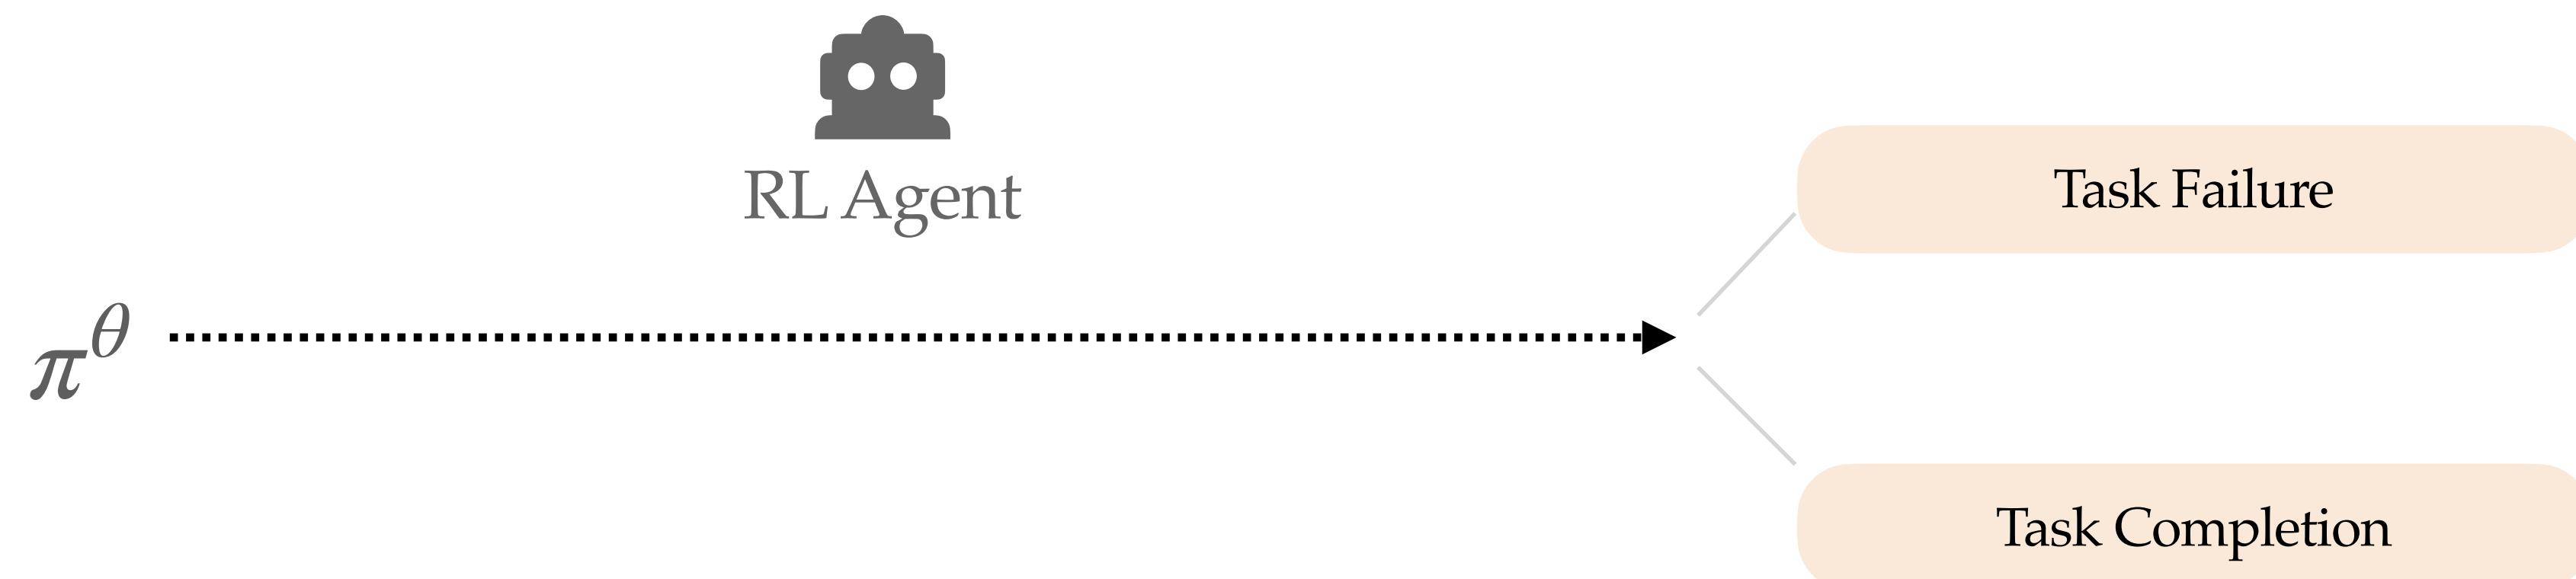
"go to the red ball"



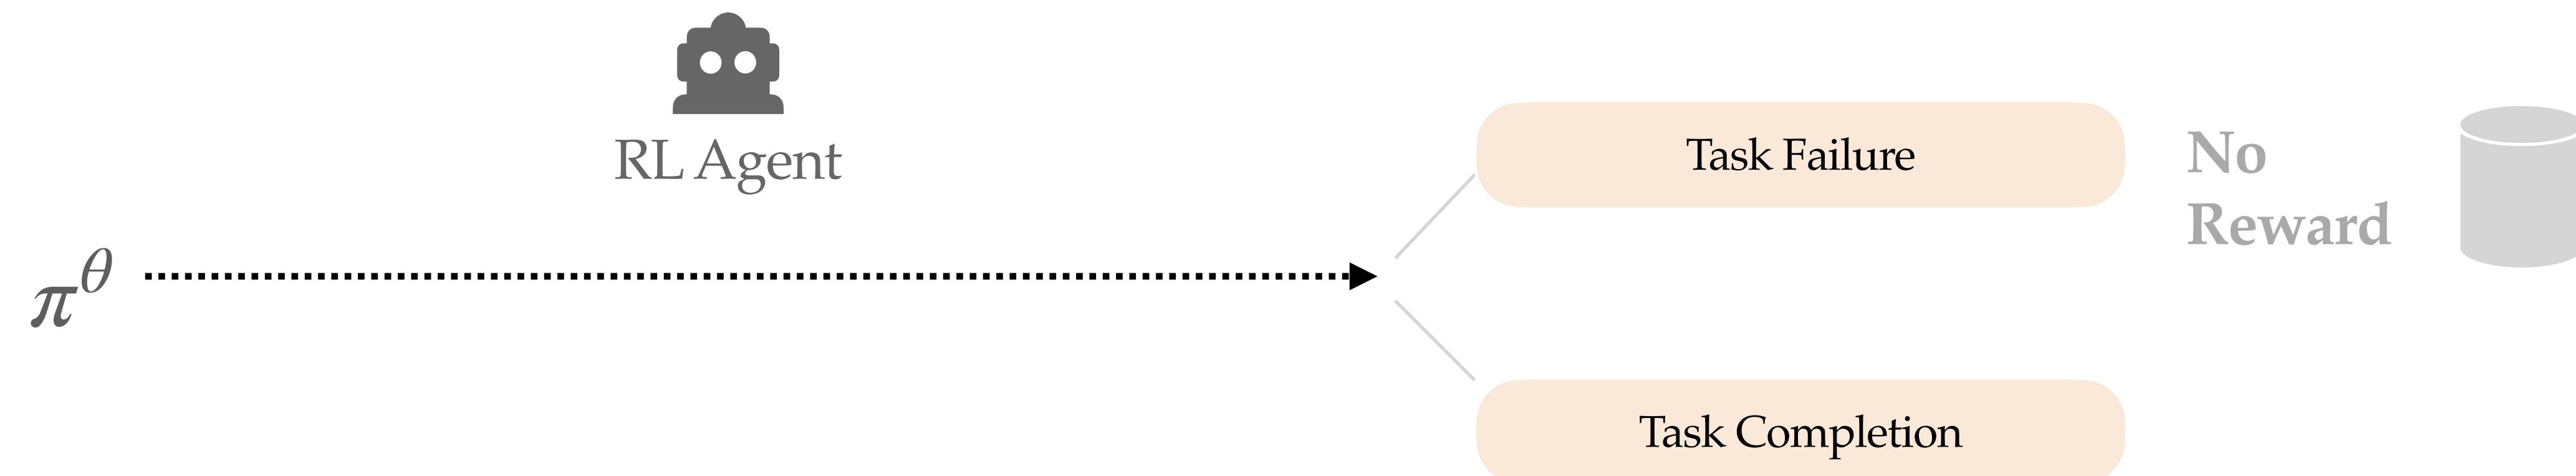
<pick up>

...

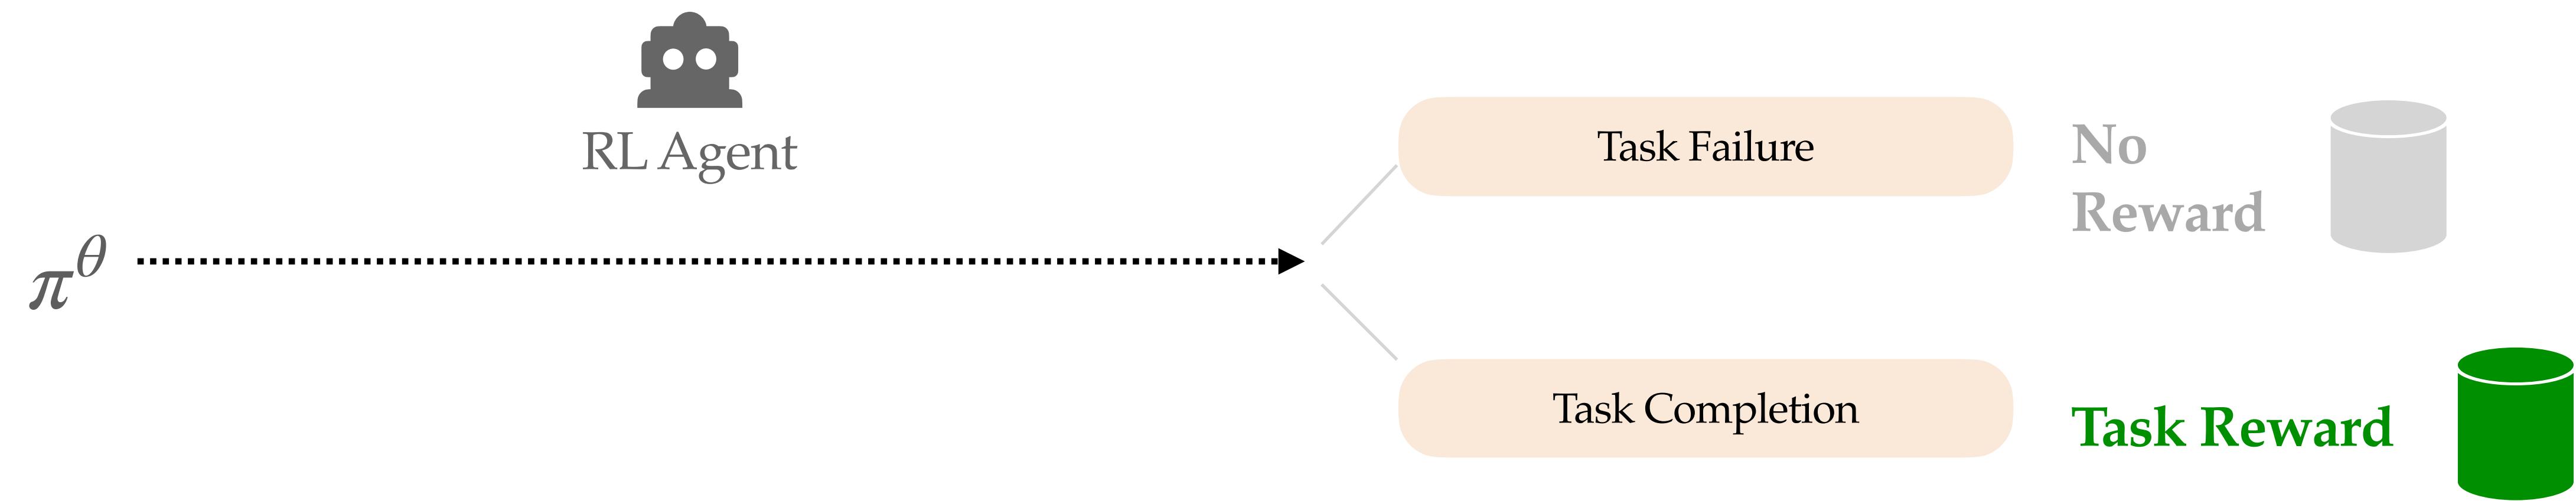
ELLA



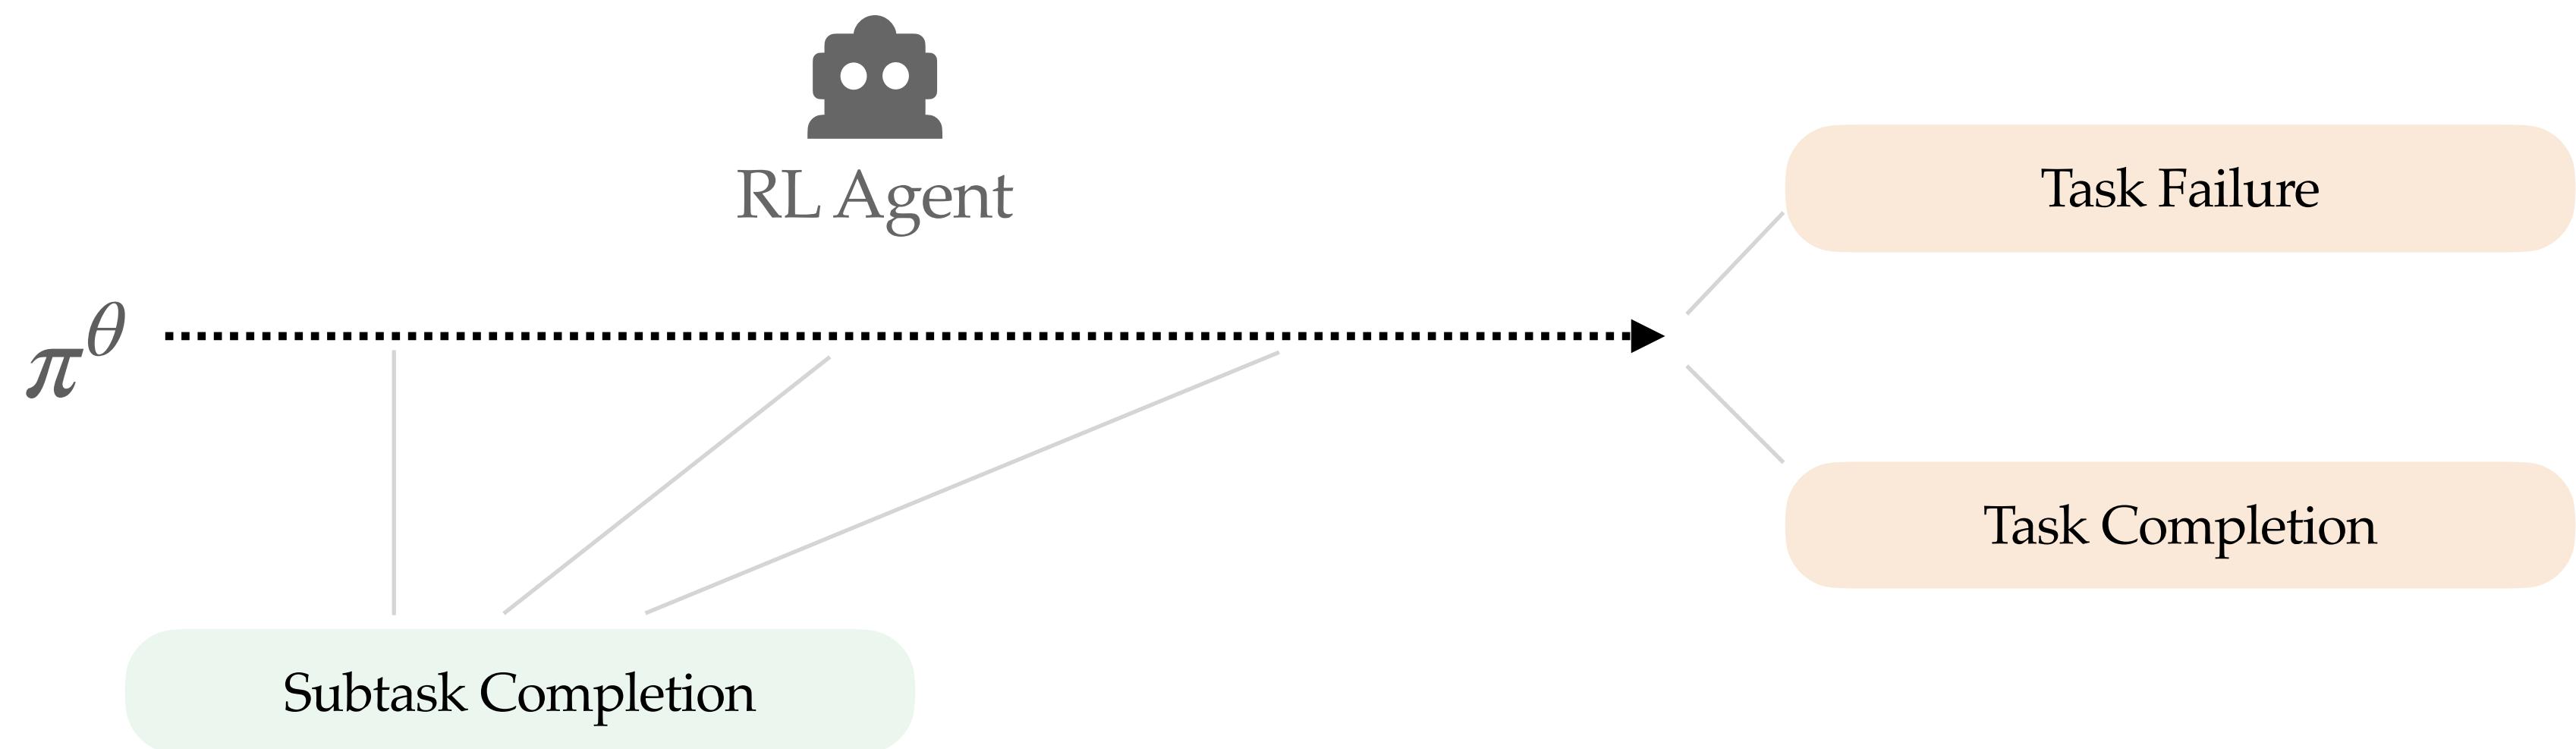
ELLA



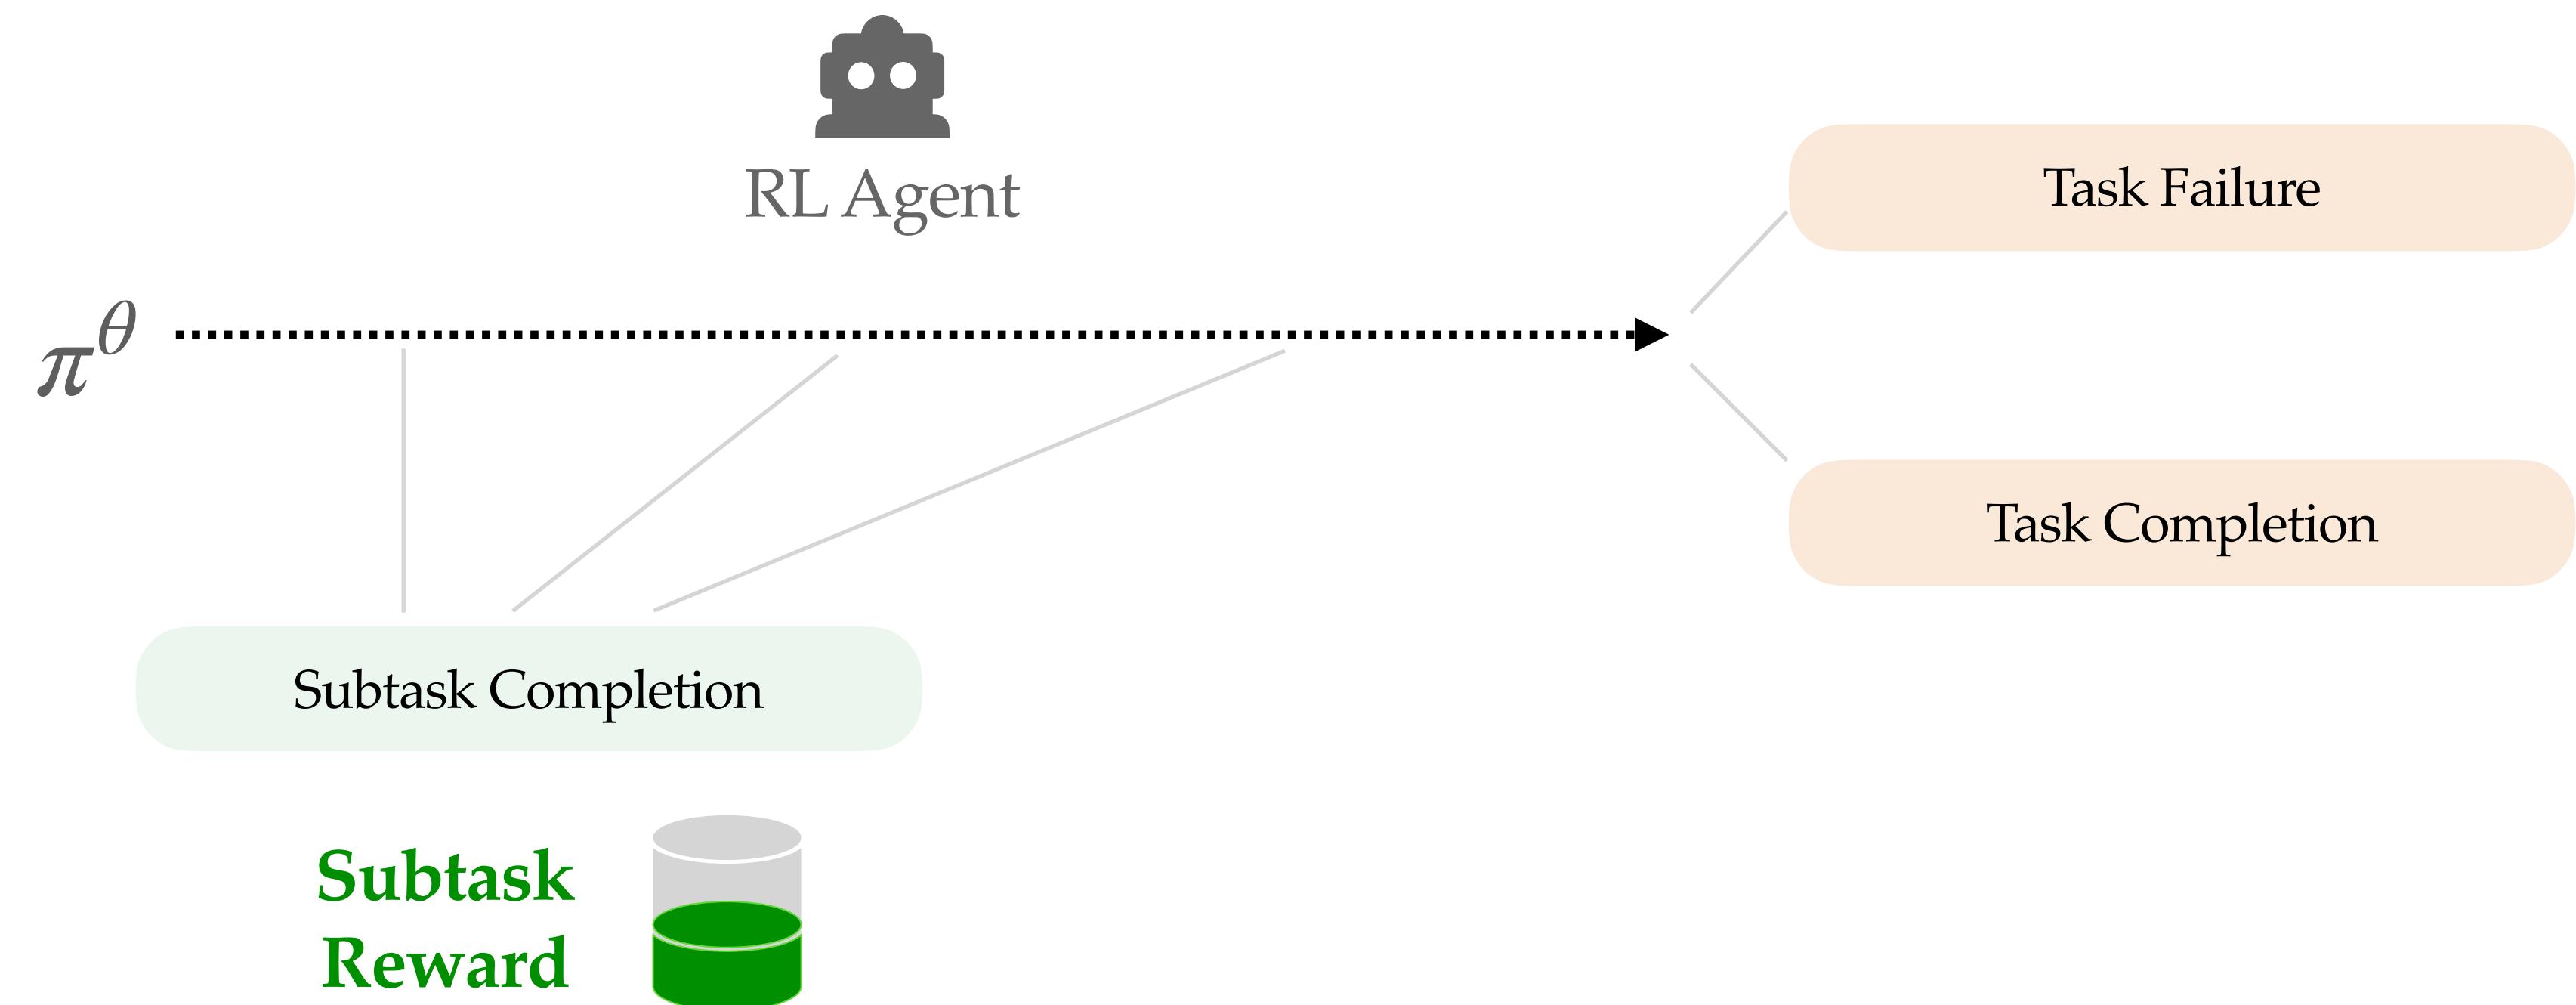
ELLA



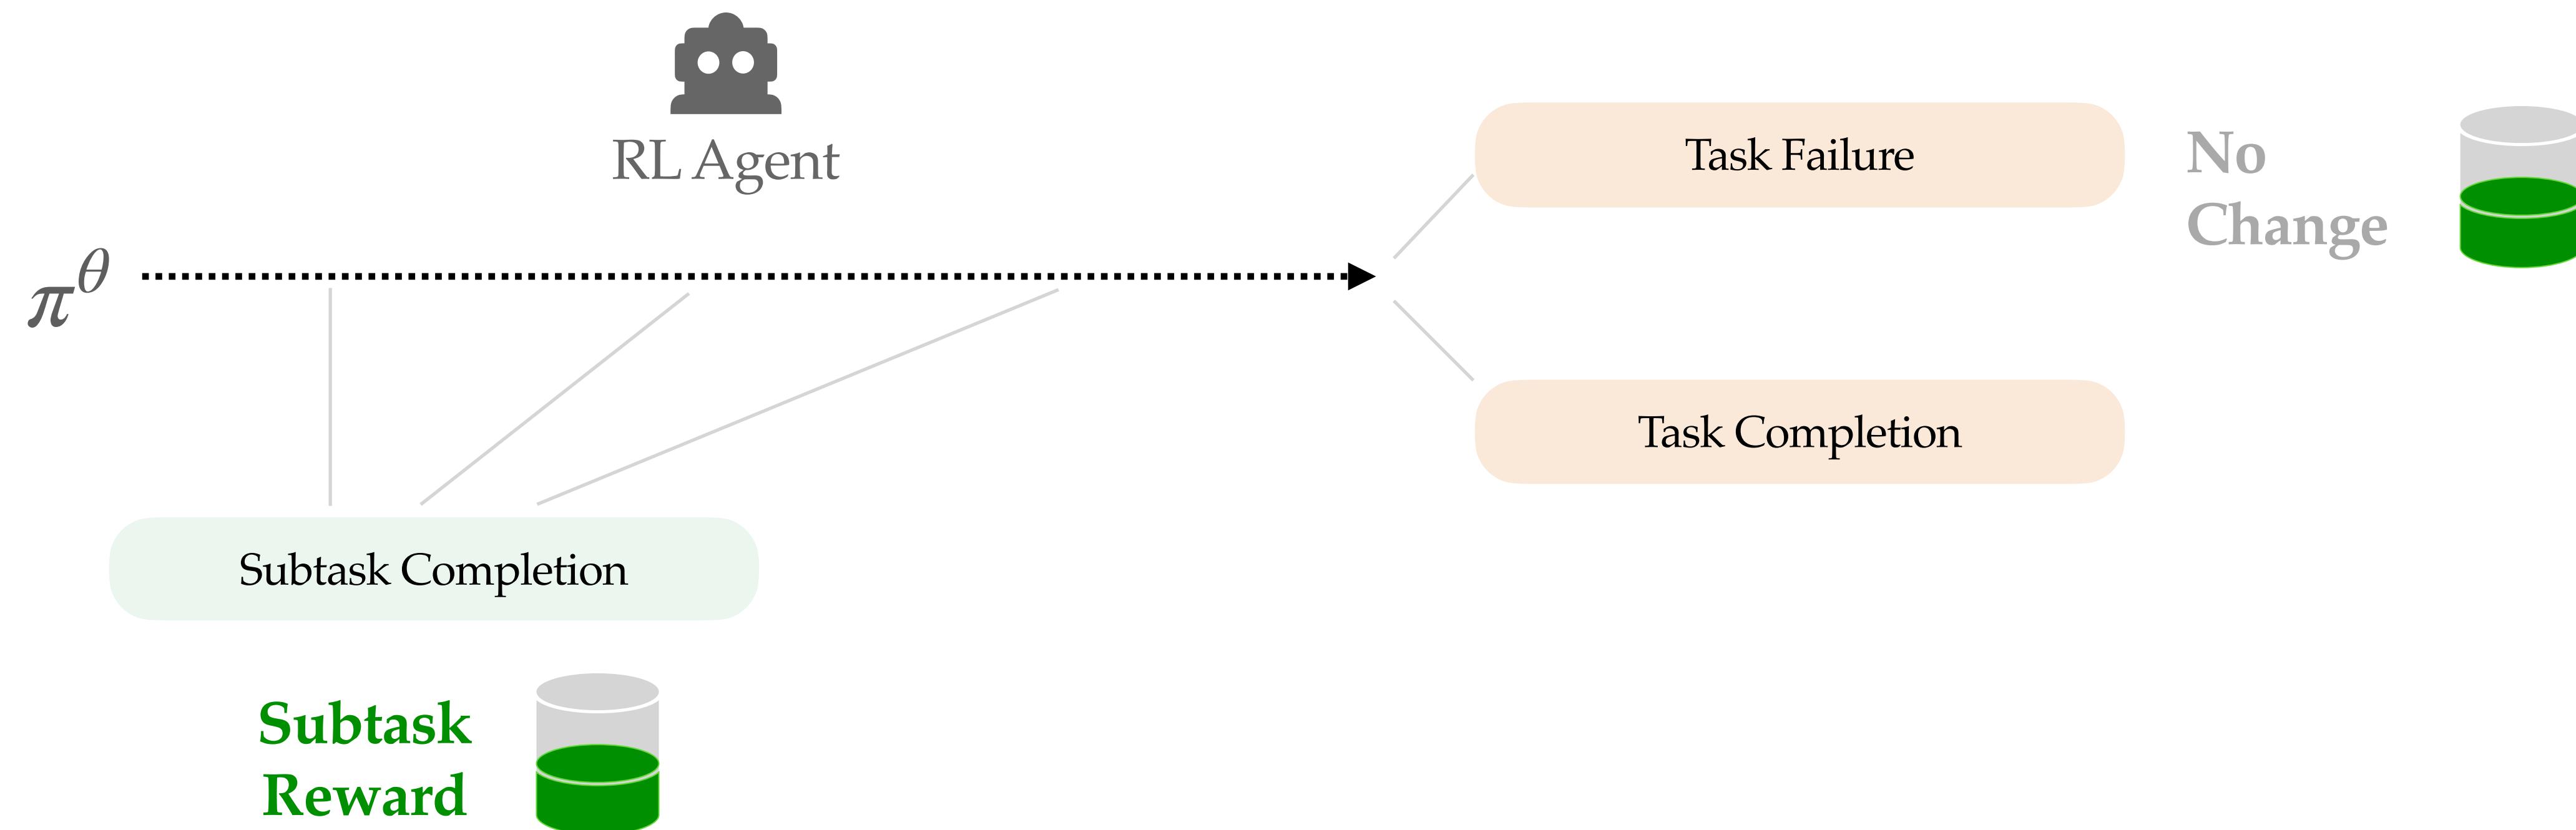
ELLA



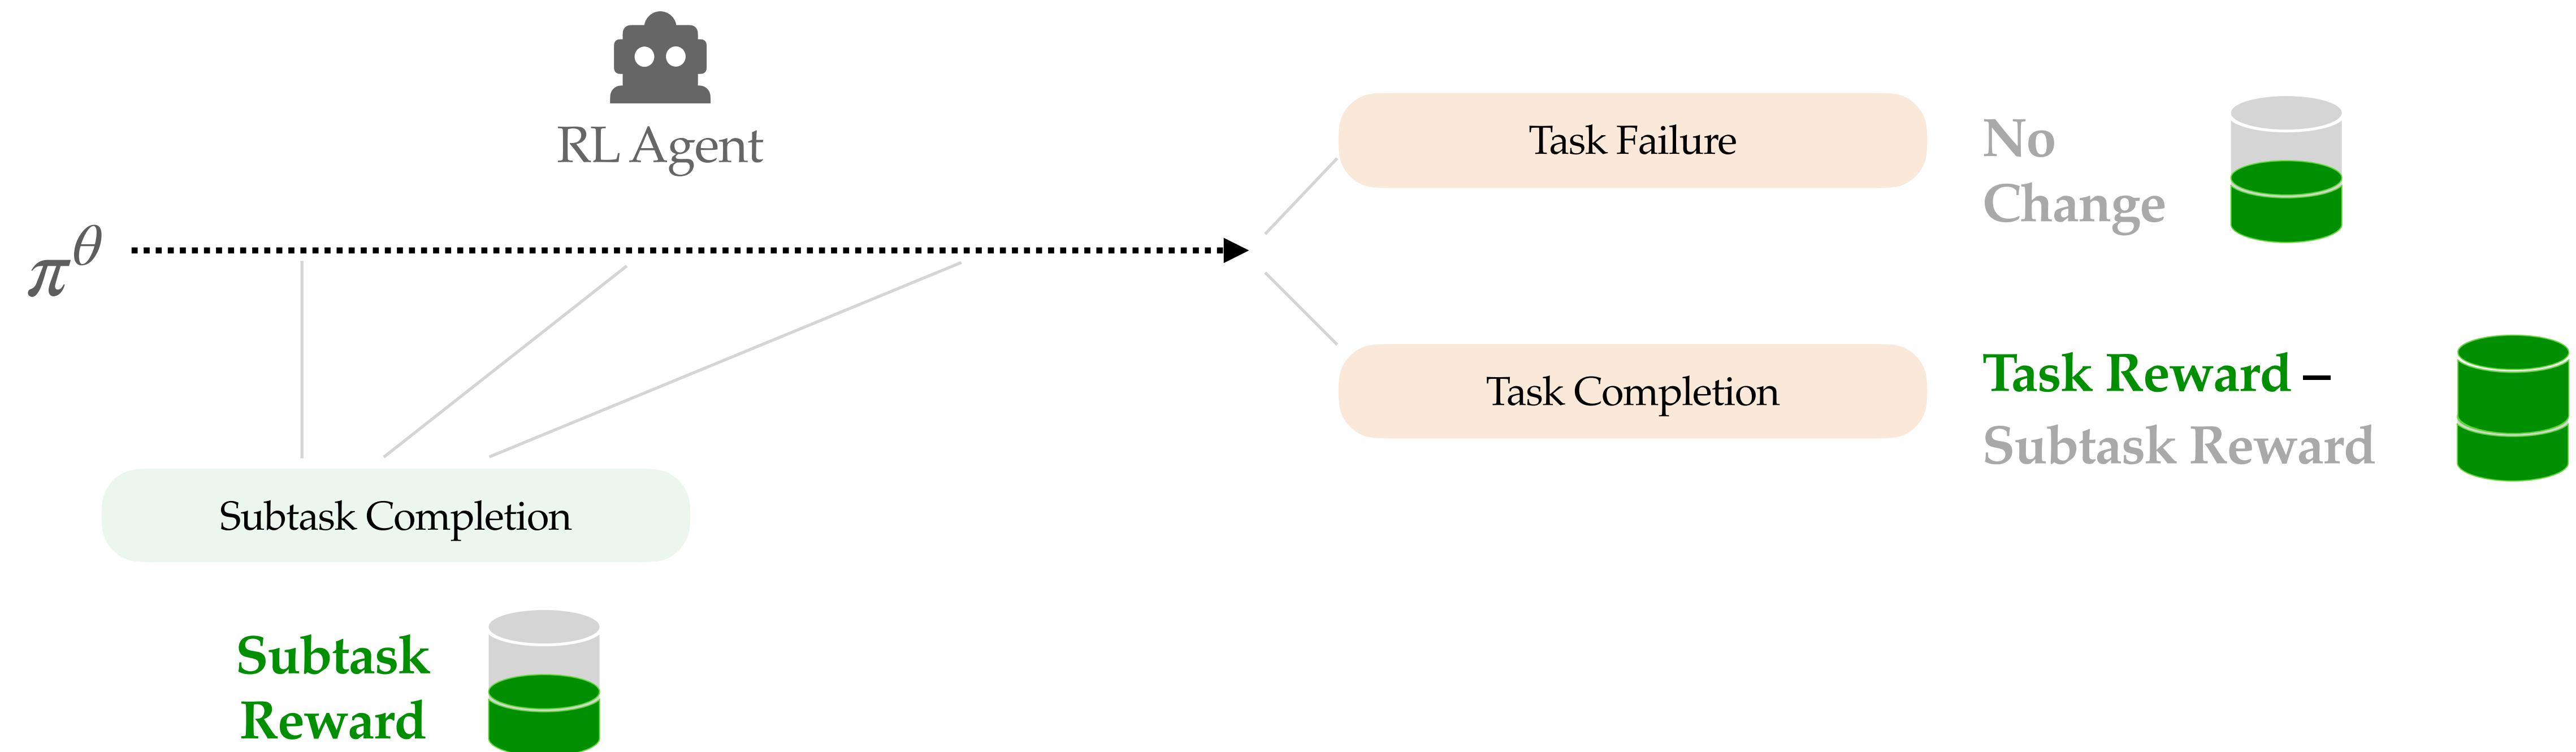
ELLA



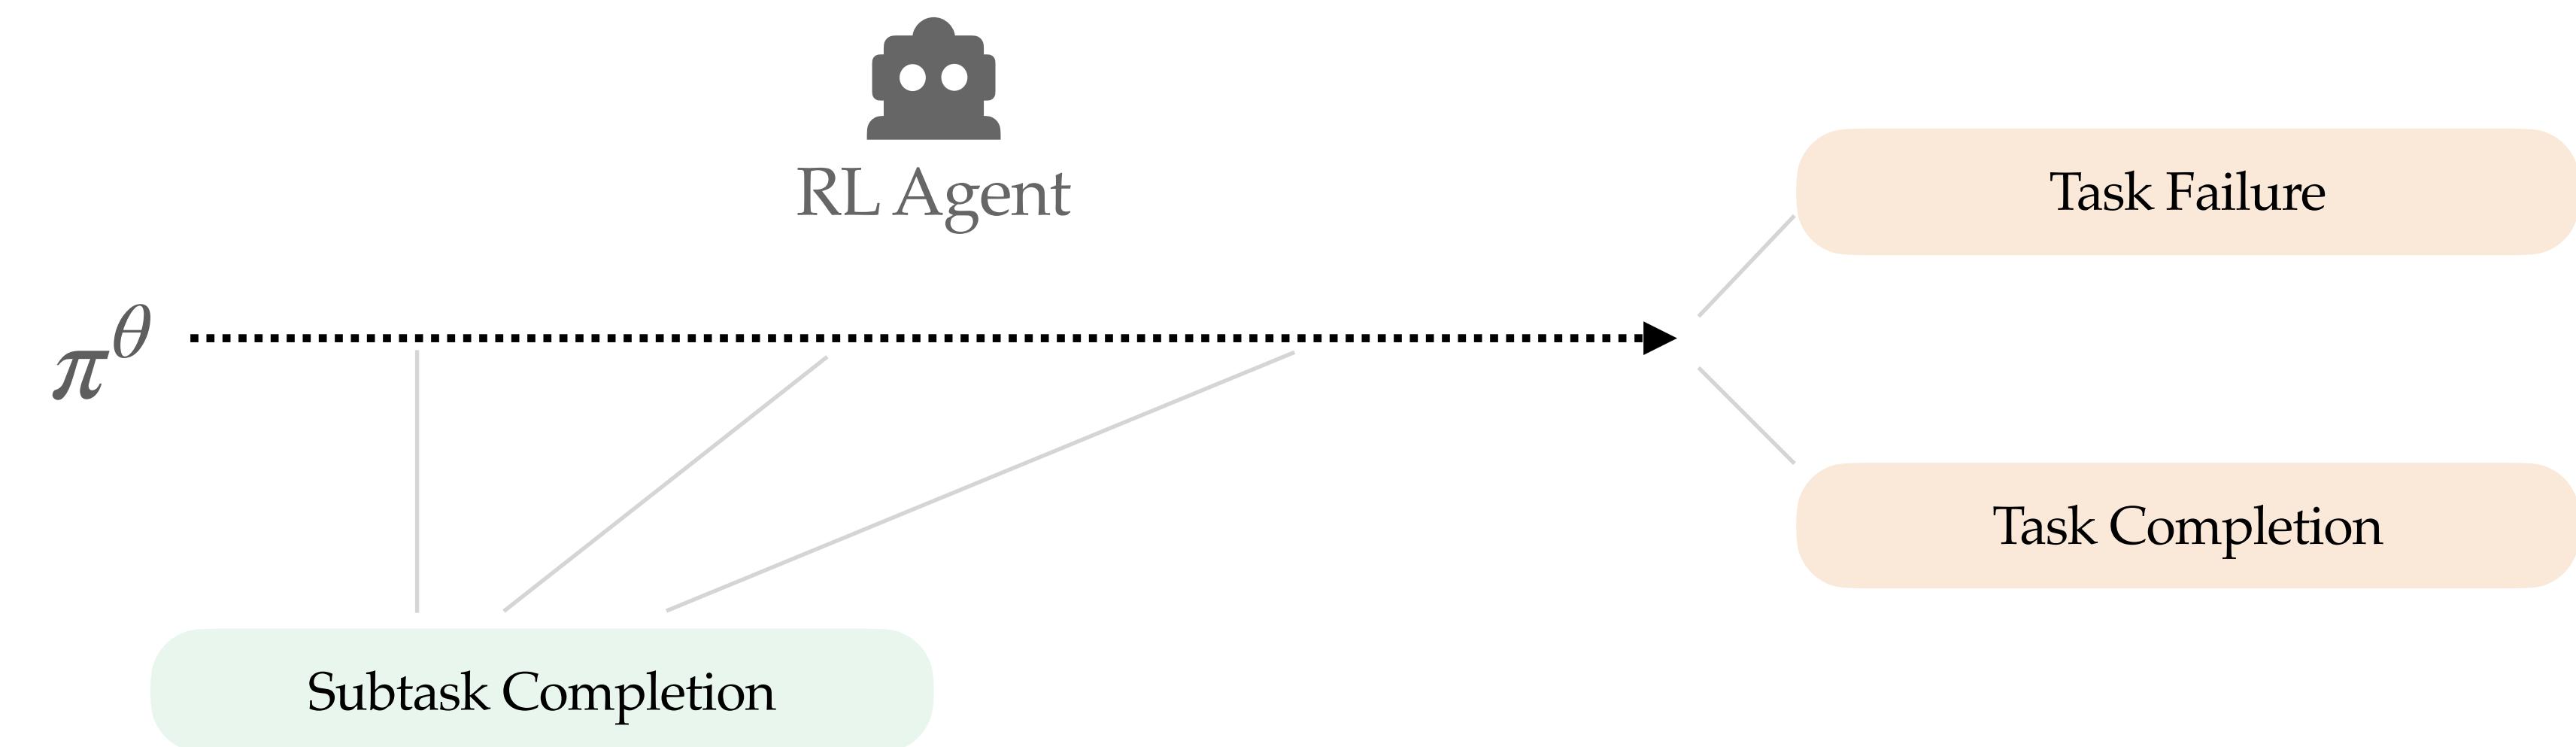
ELLA



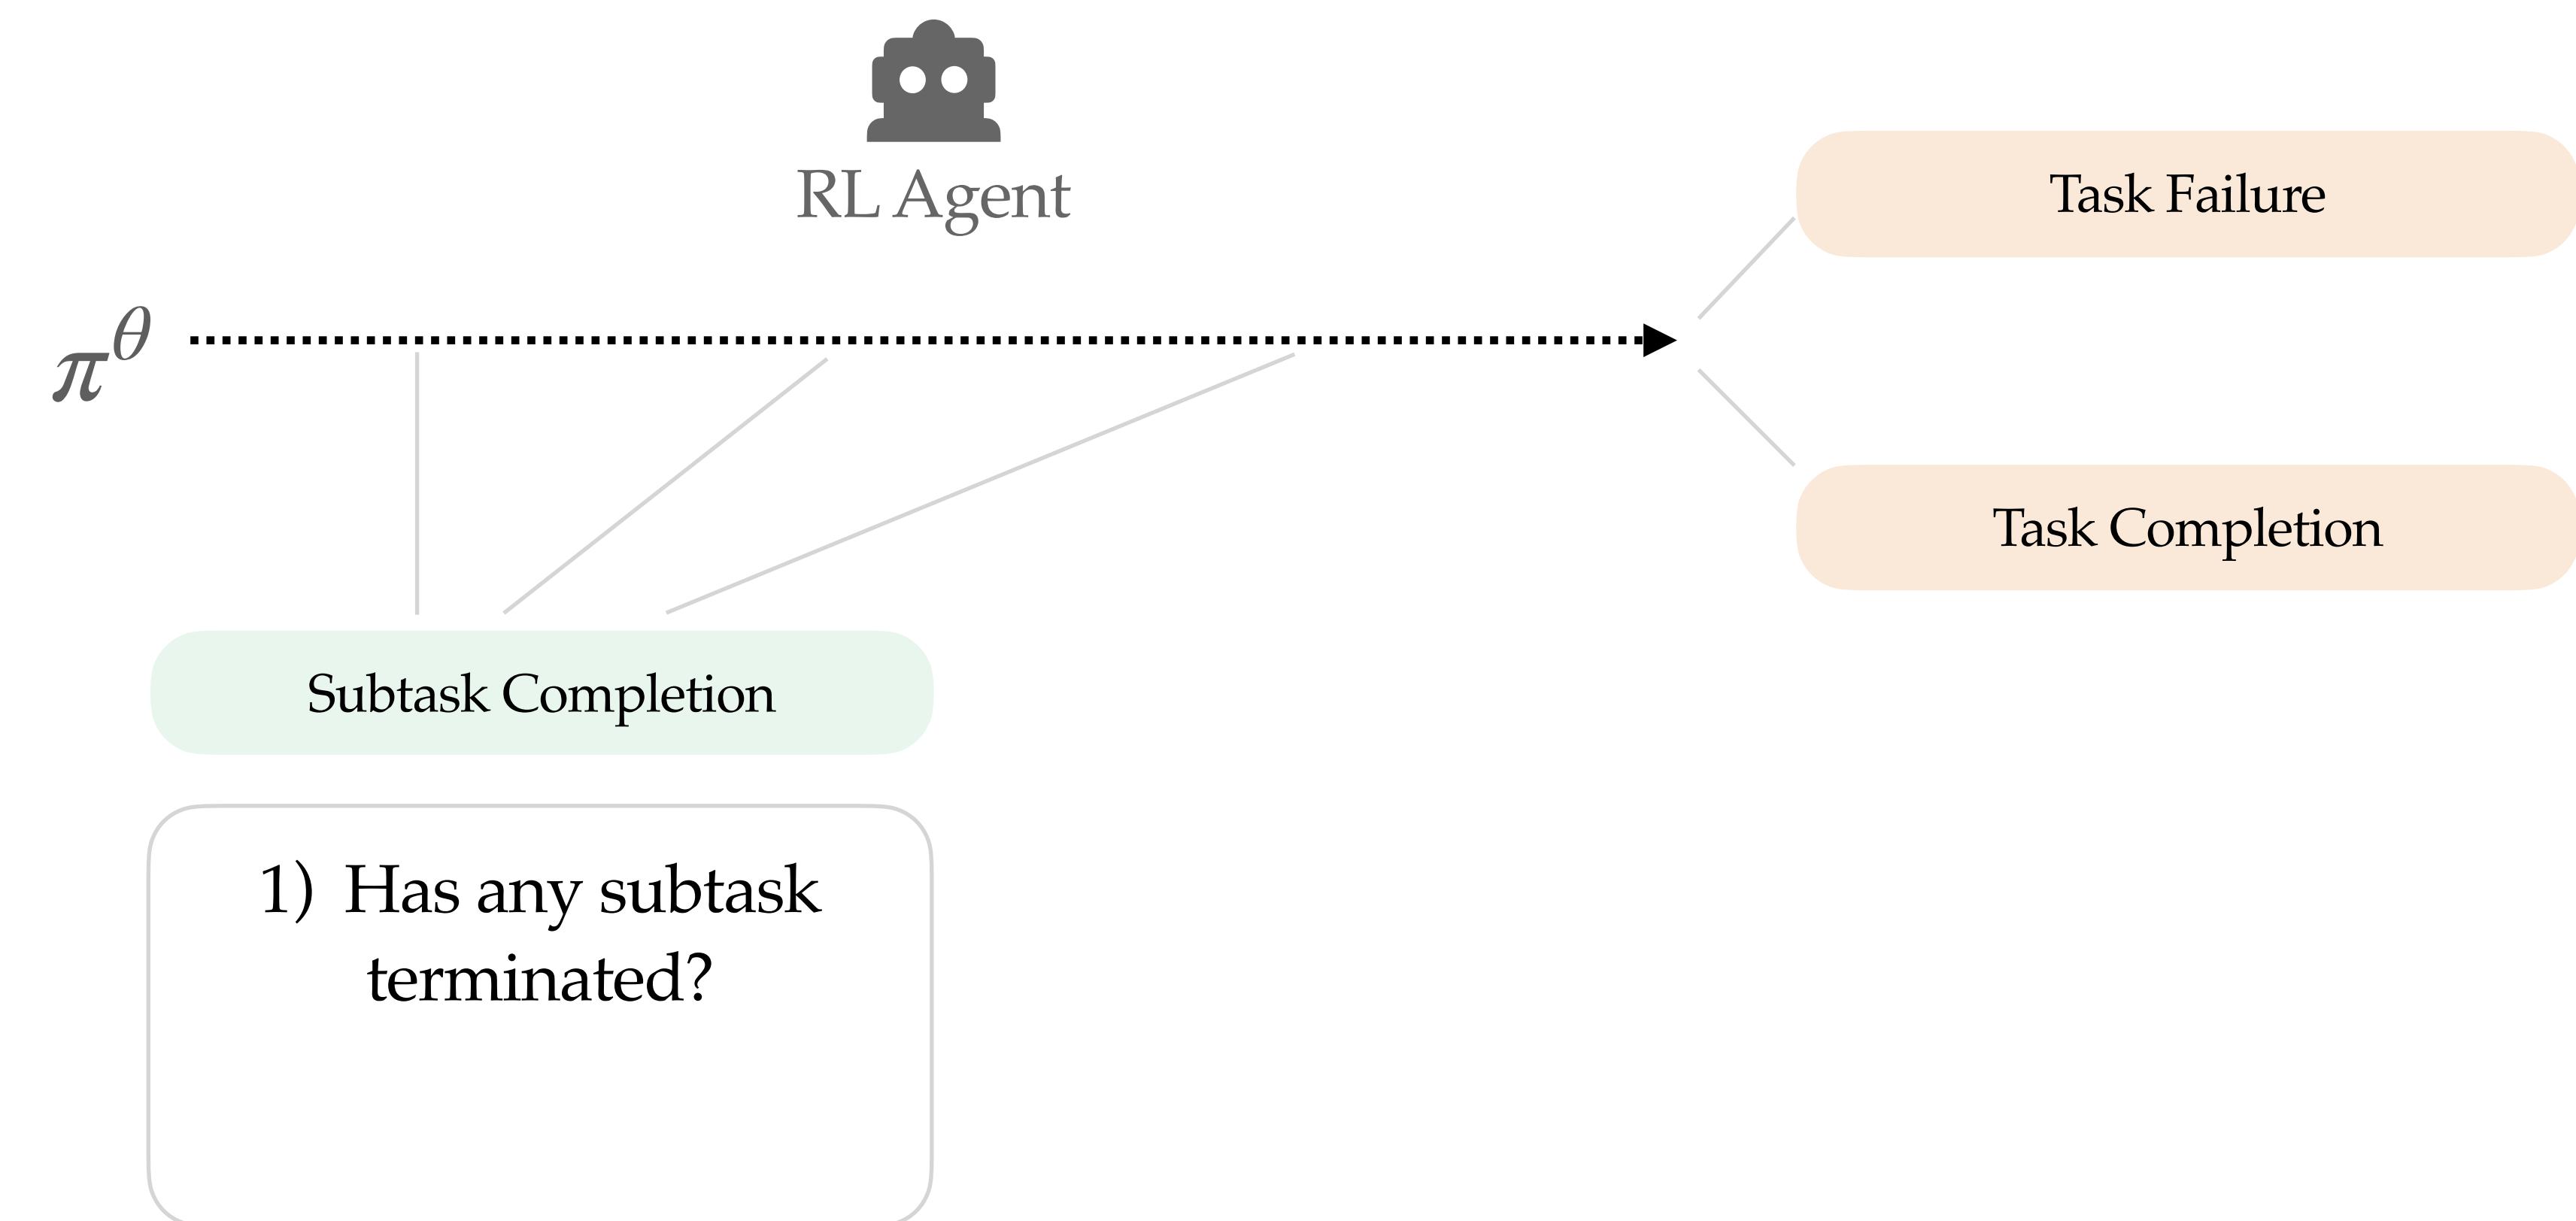
ELLA



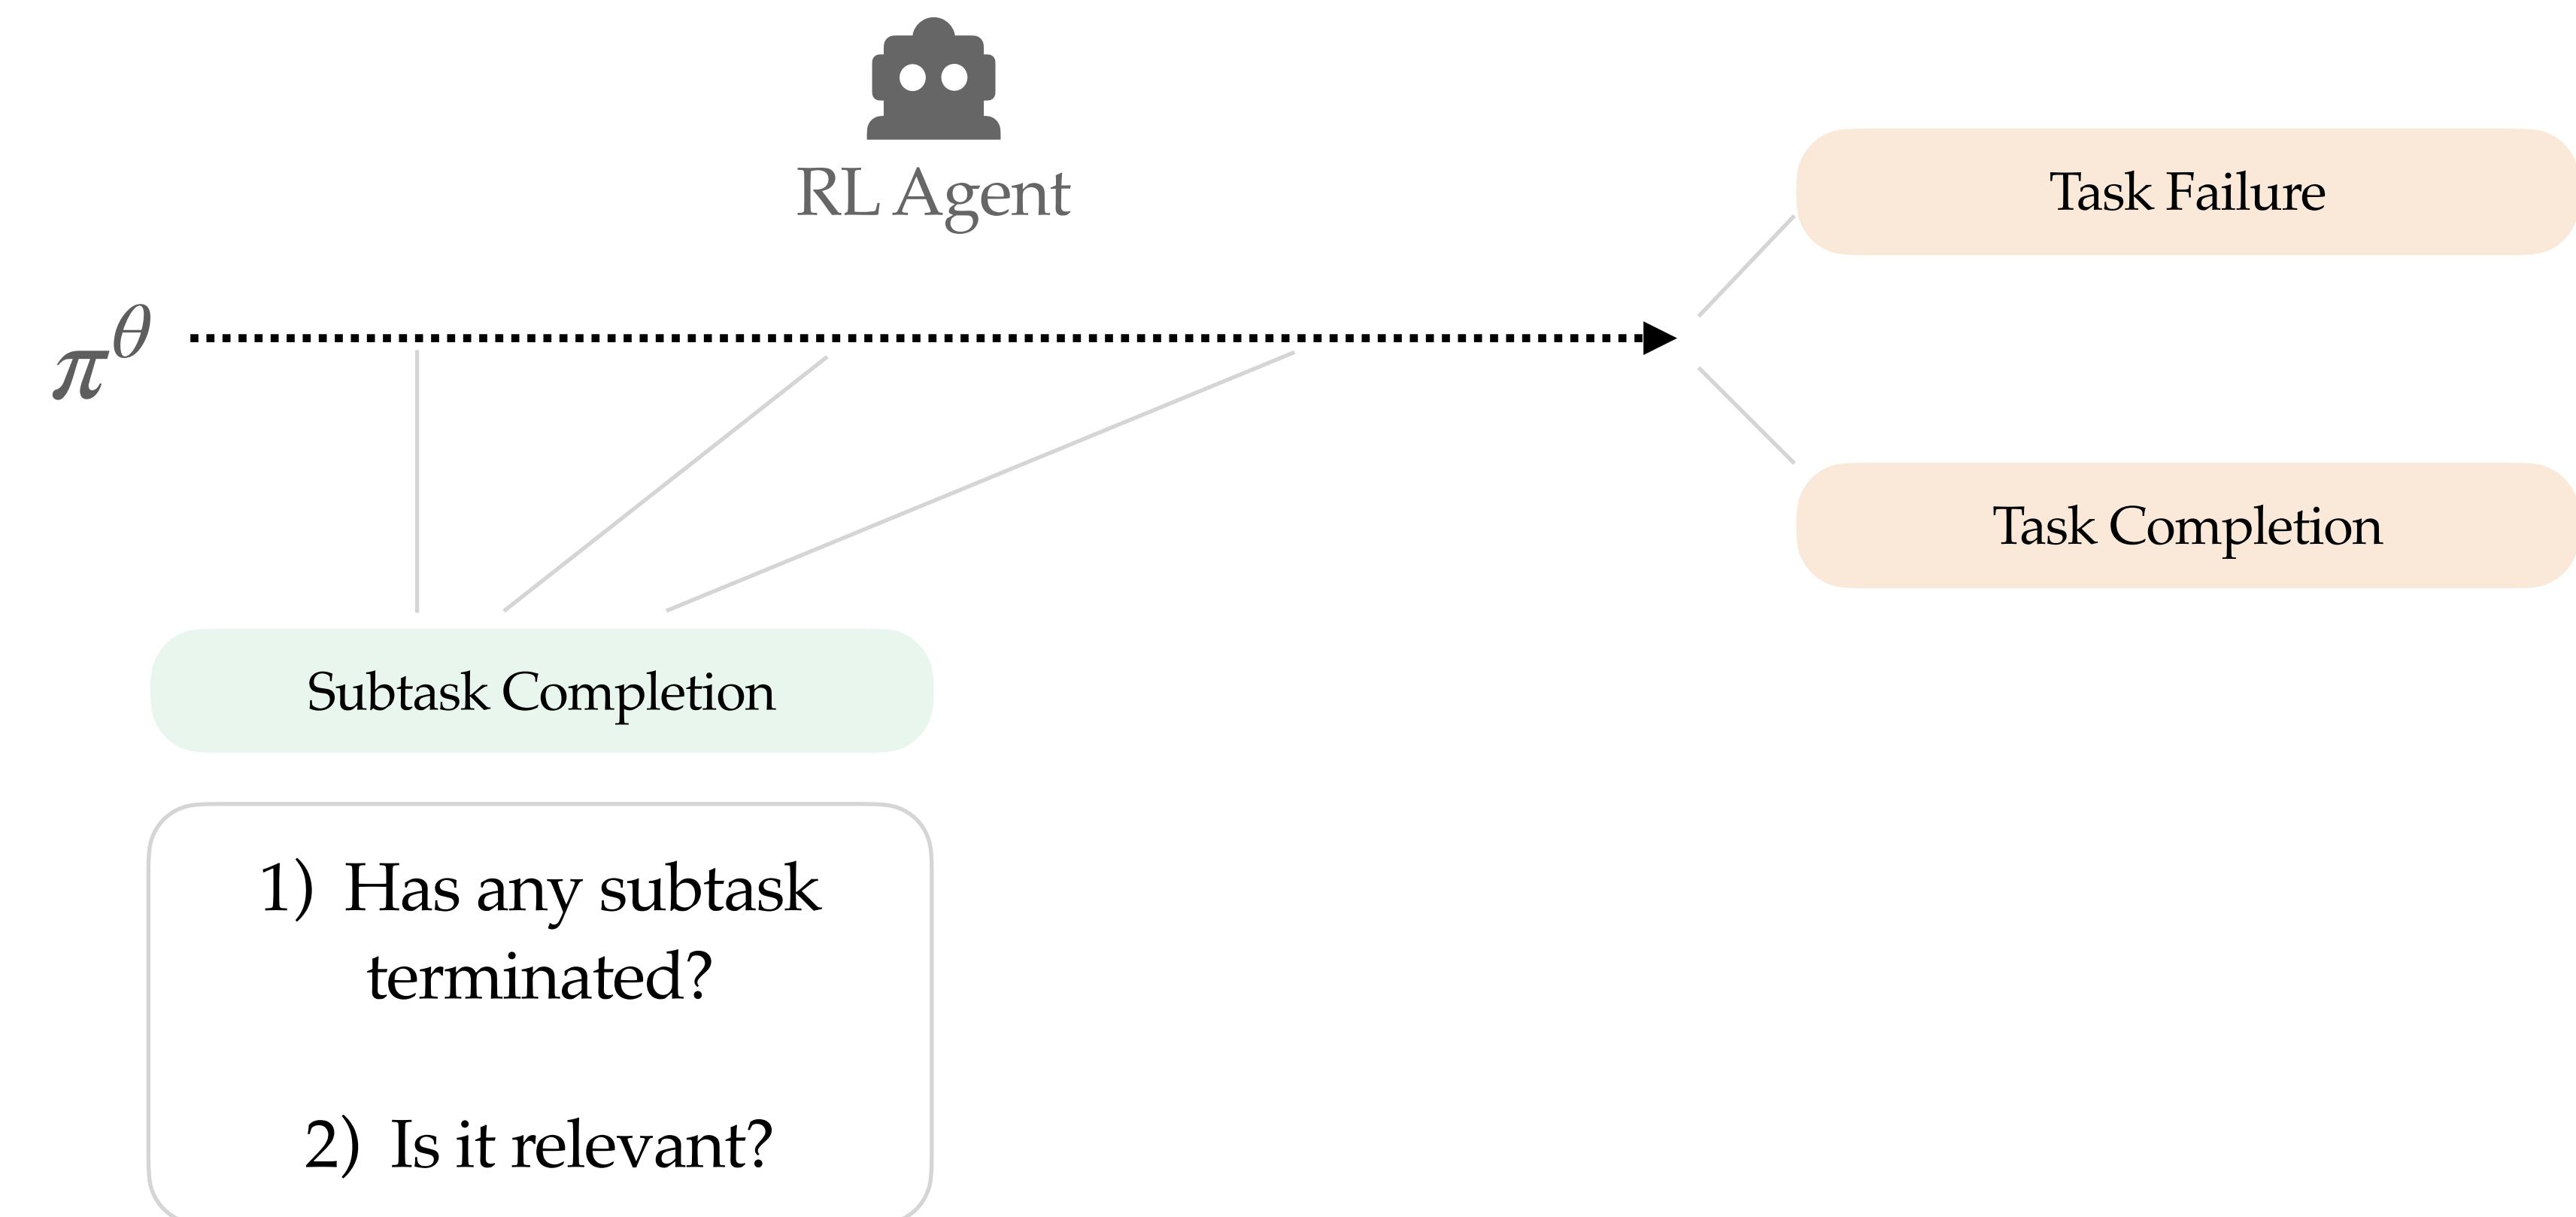
ELLA



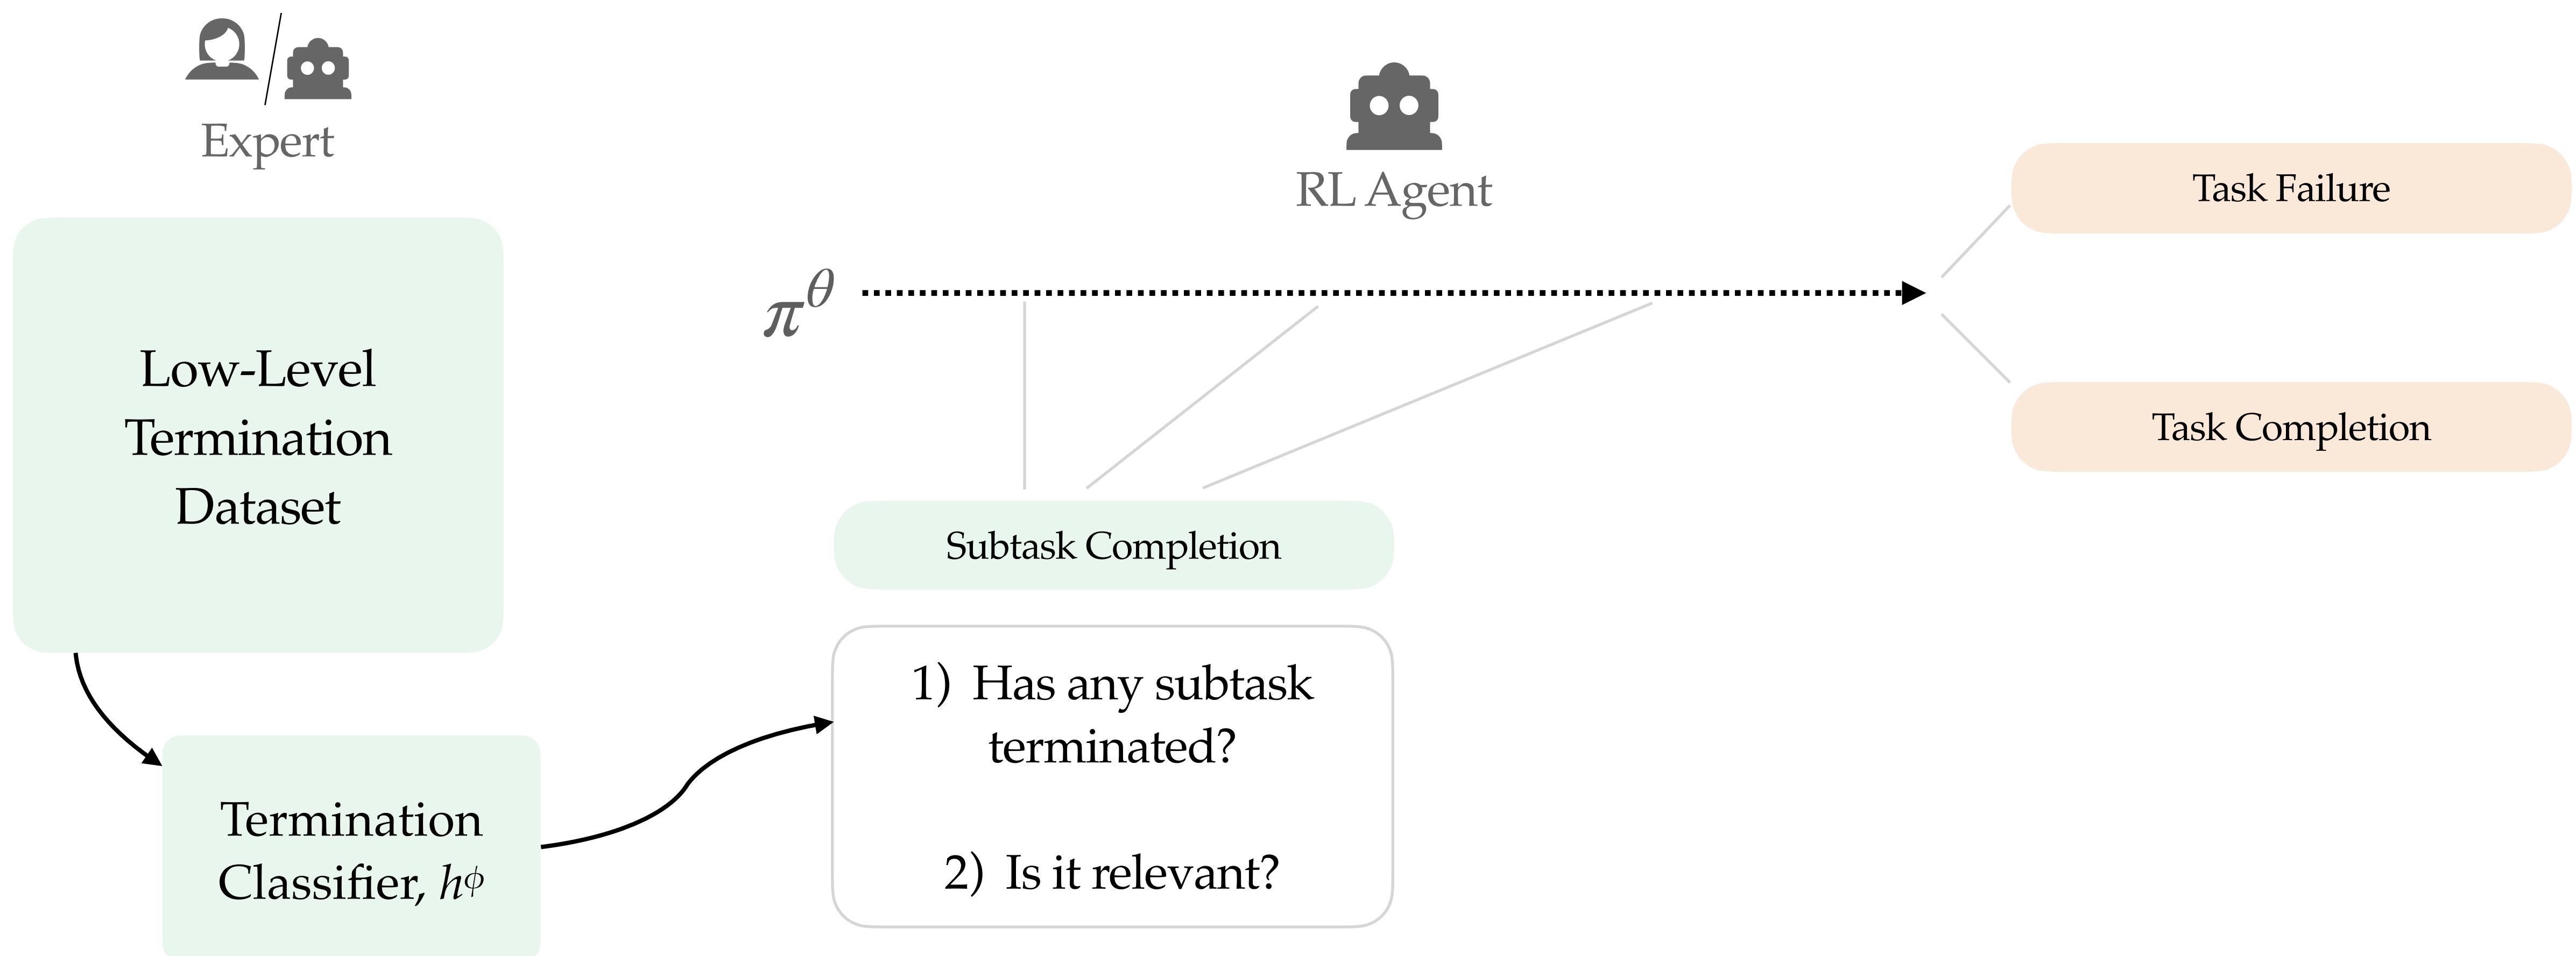
ELLA



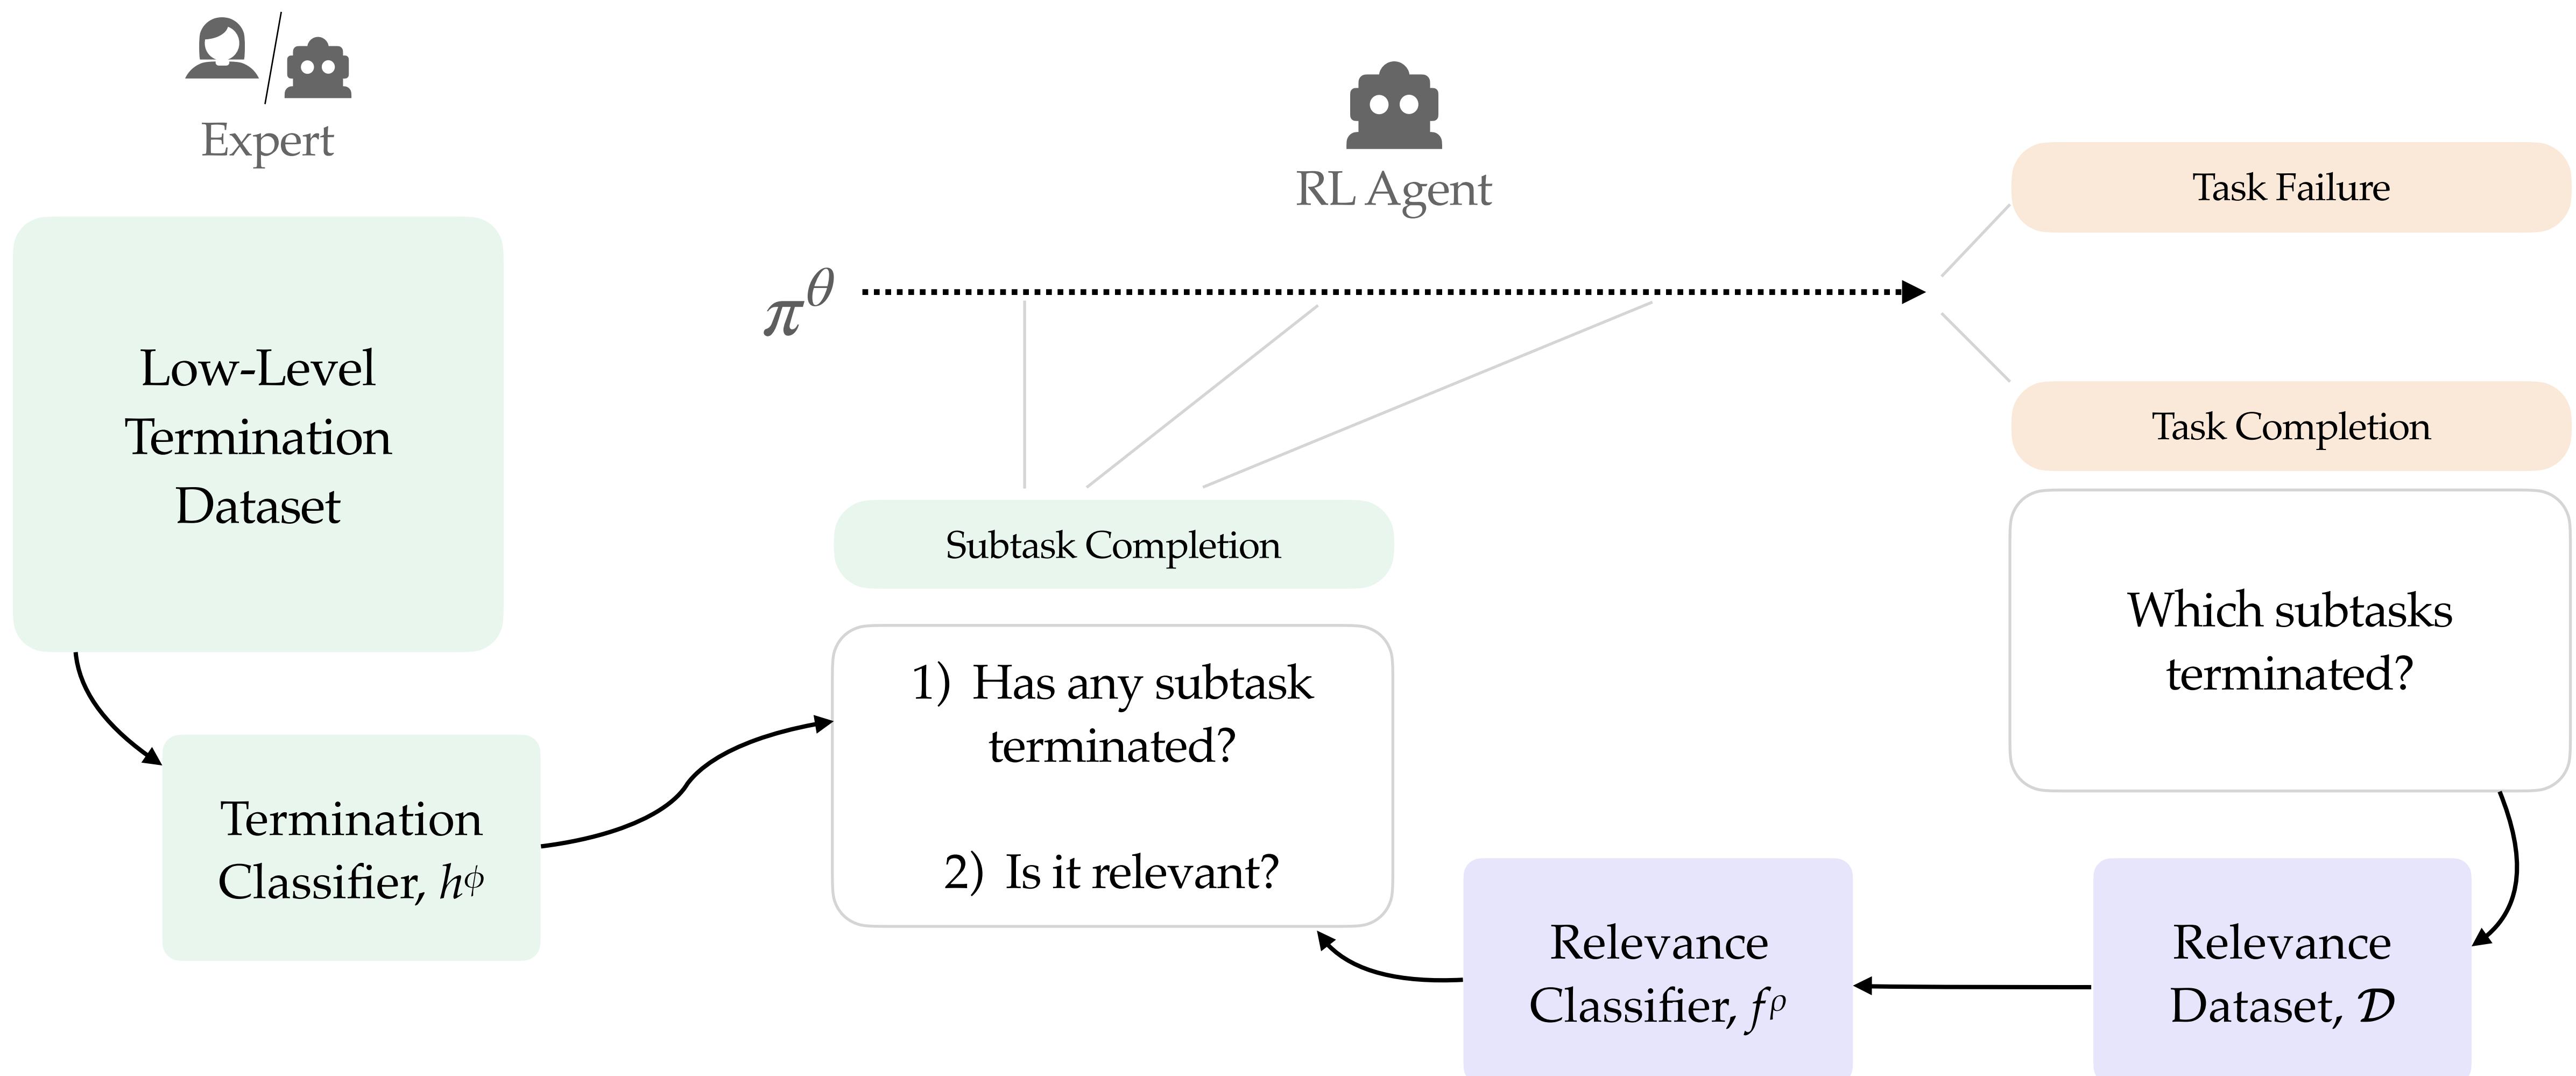
ELLA



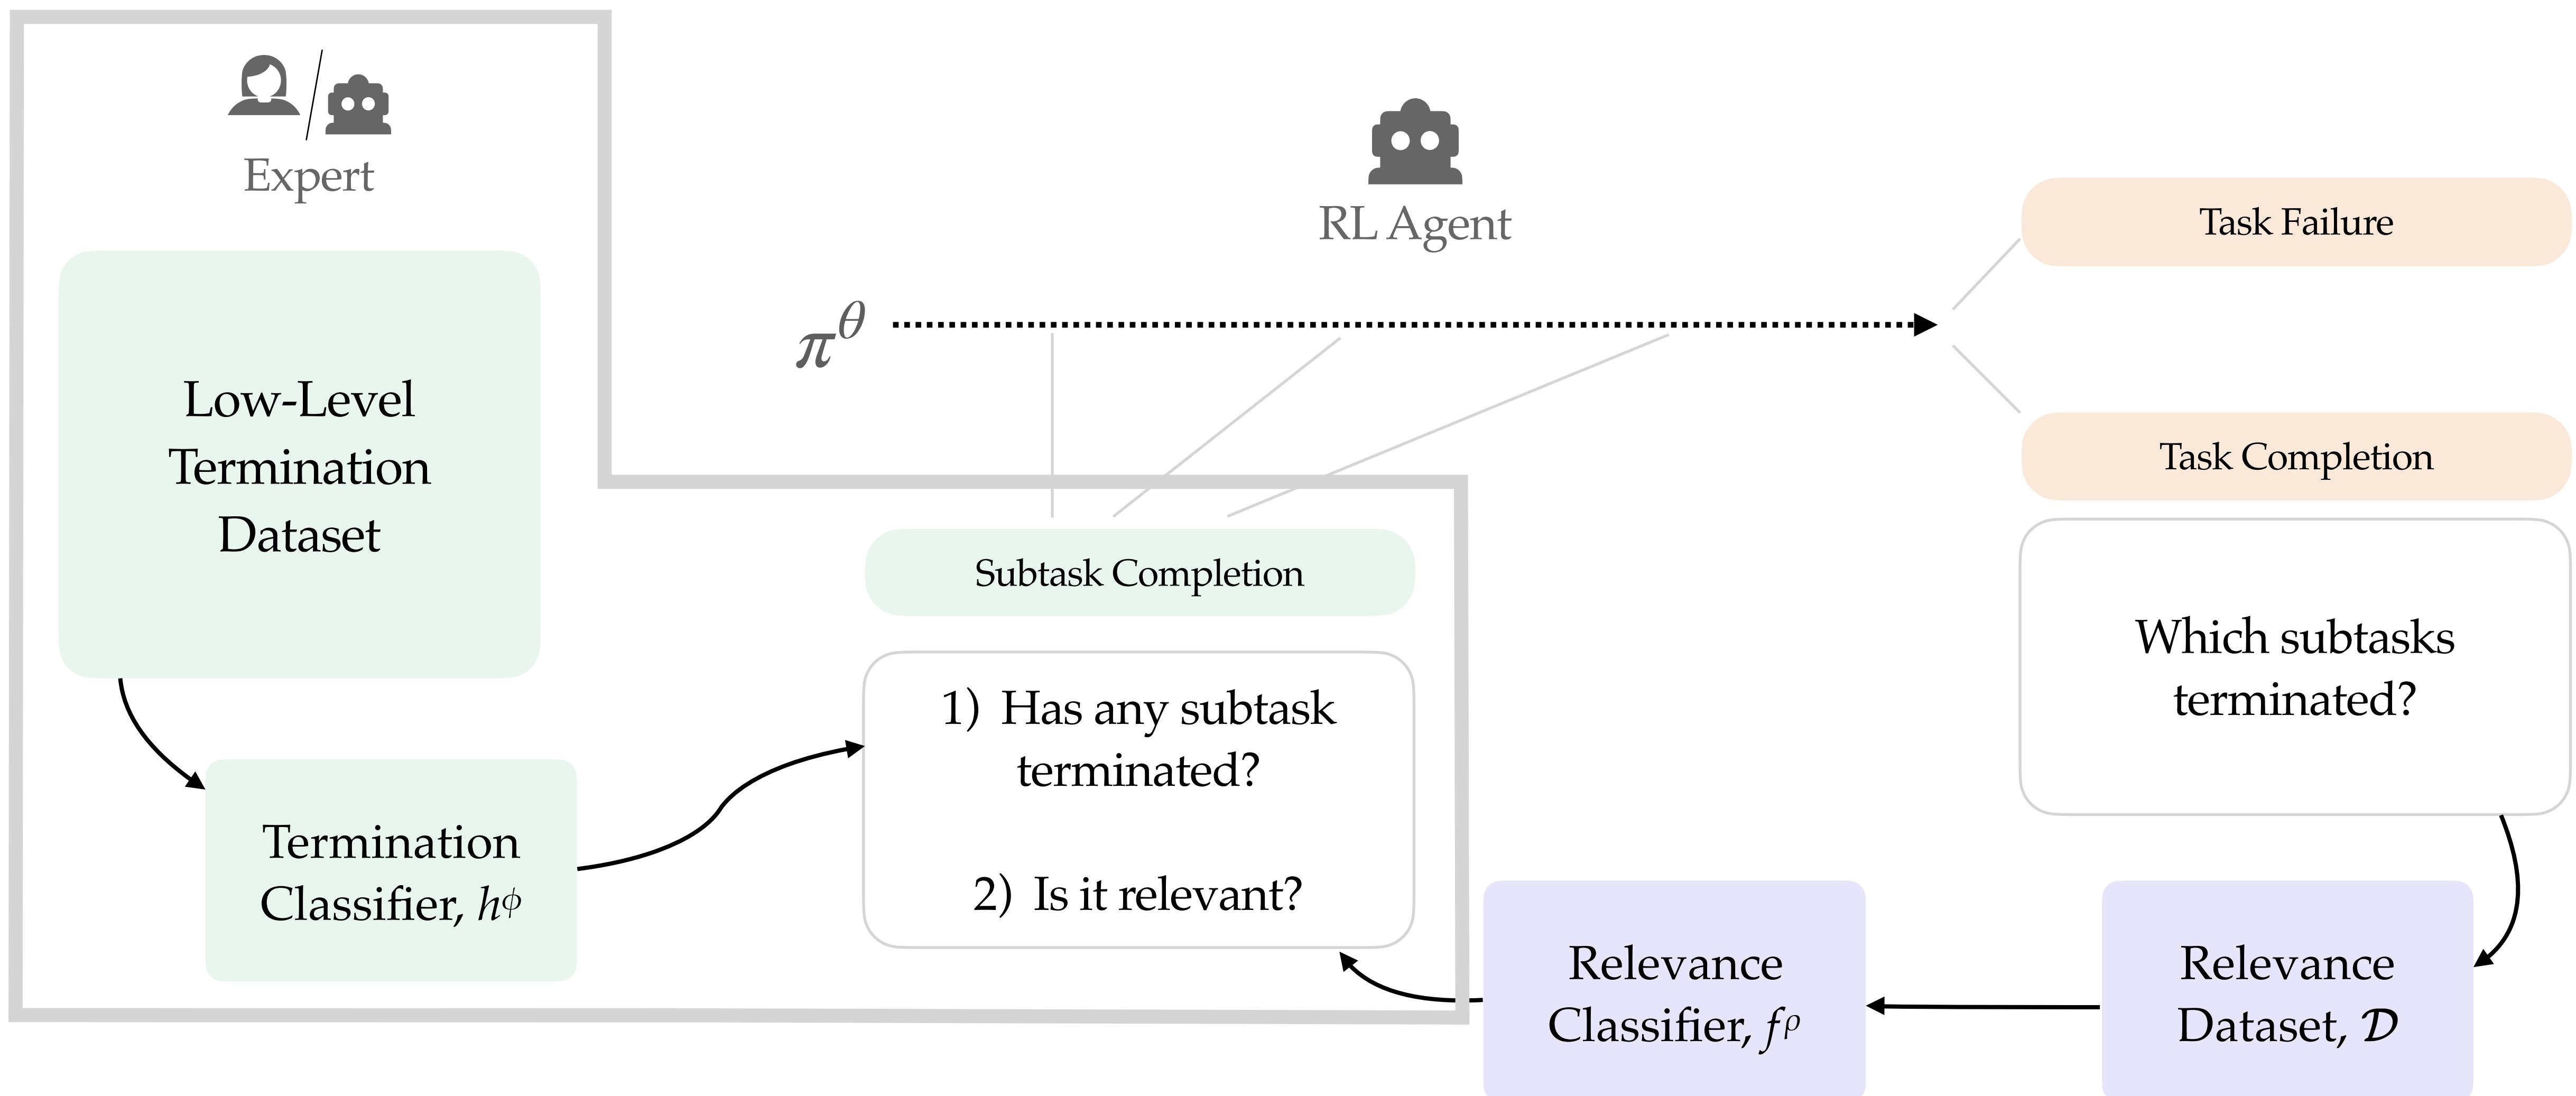
ELLA



ELLA



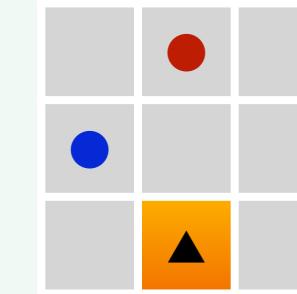
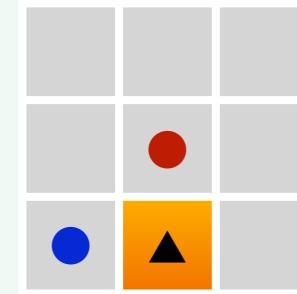
ELLA



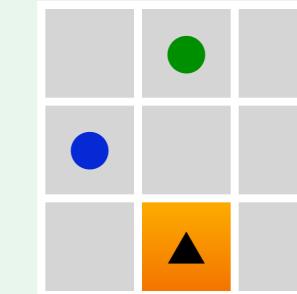
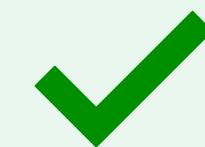
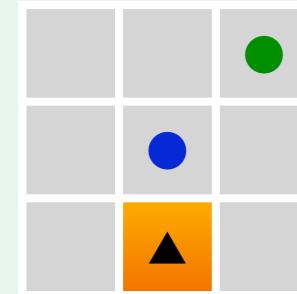
Low-Level Termination Classifier

- $h^\phi : S \times G_\ell \rightarrow \{0,1\}$

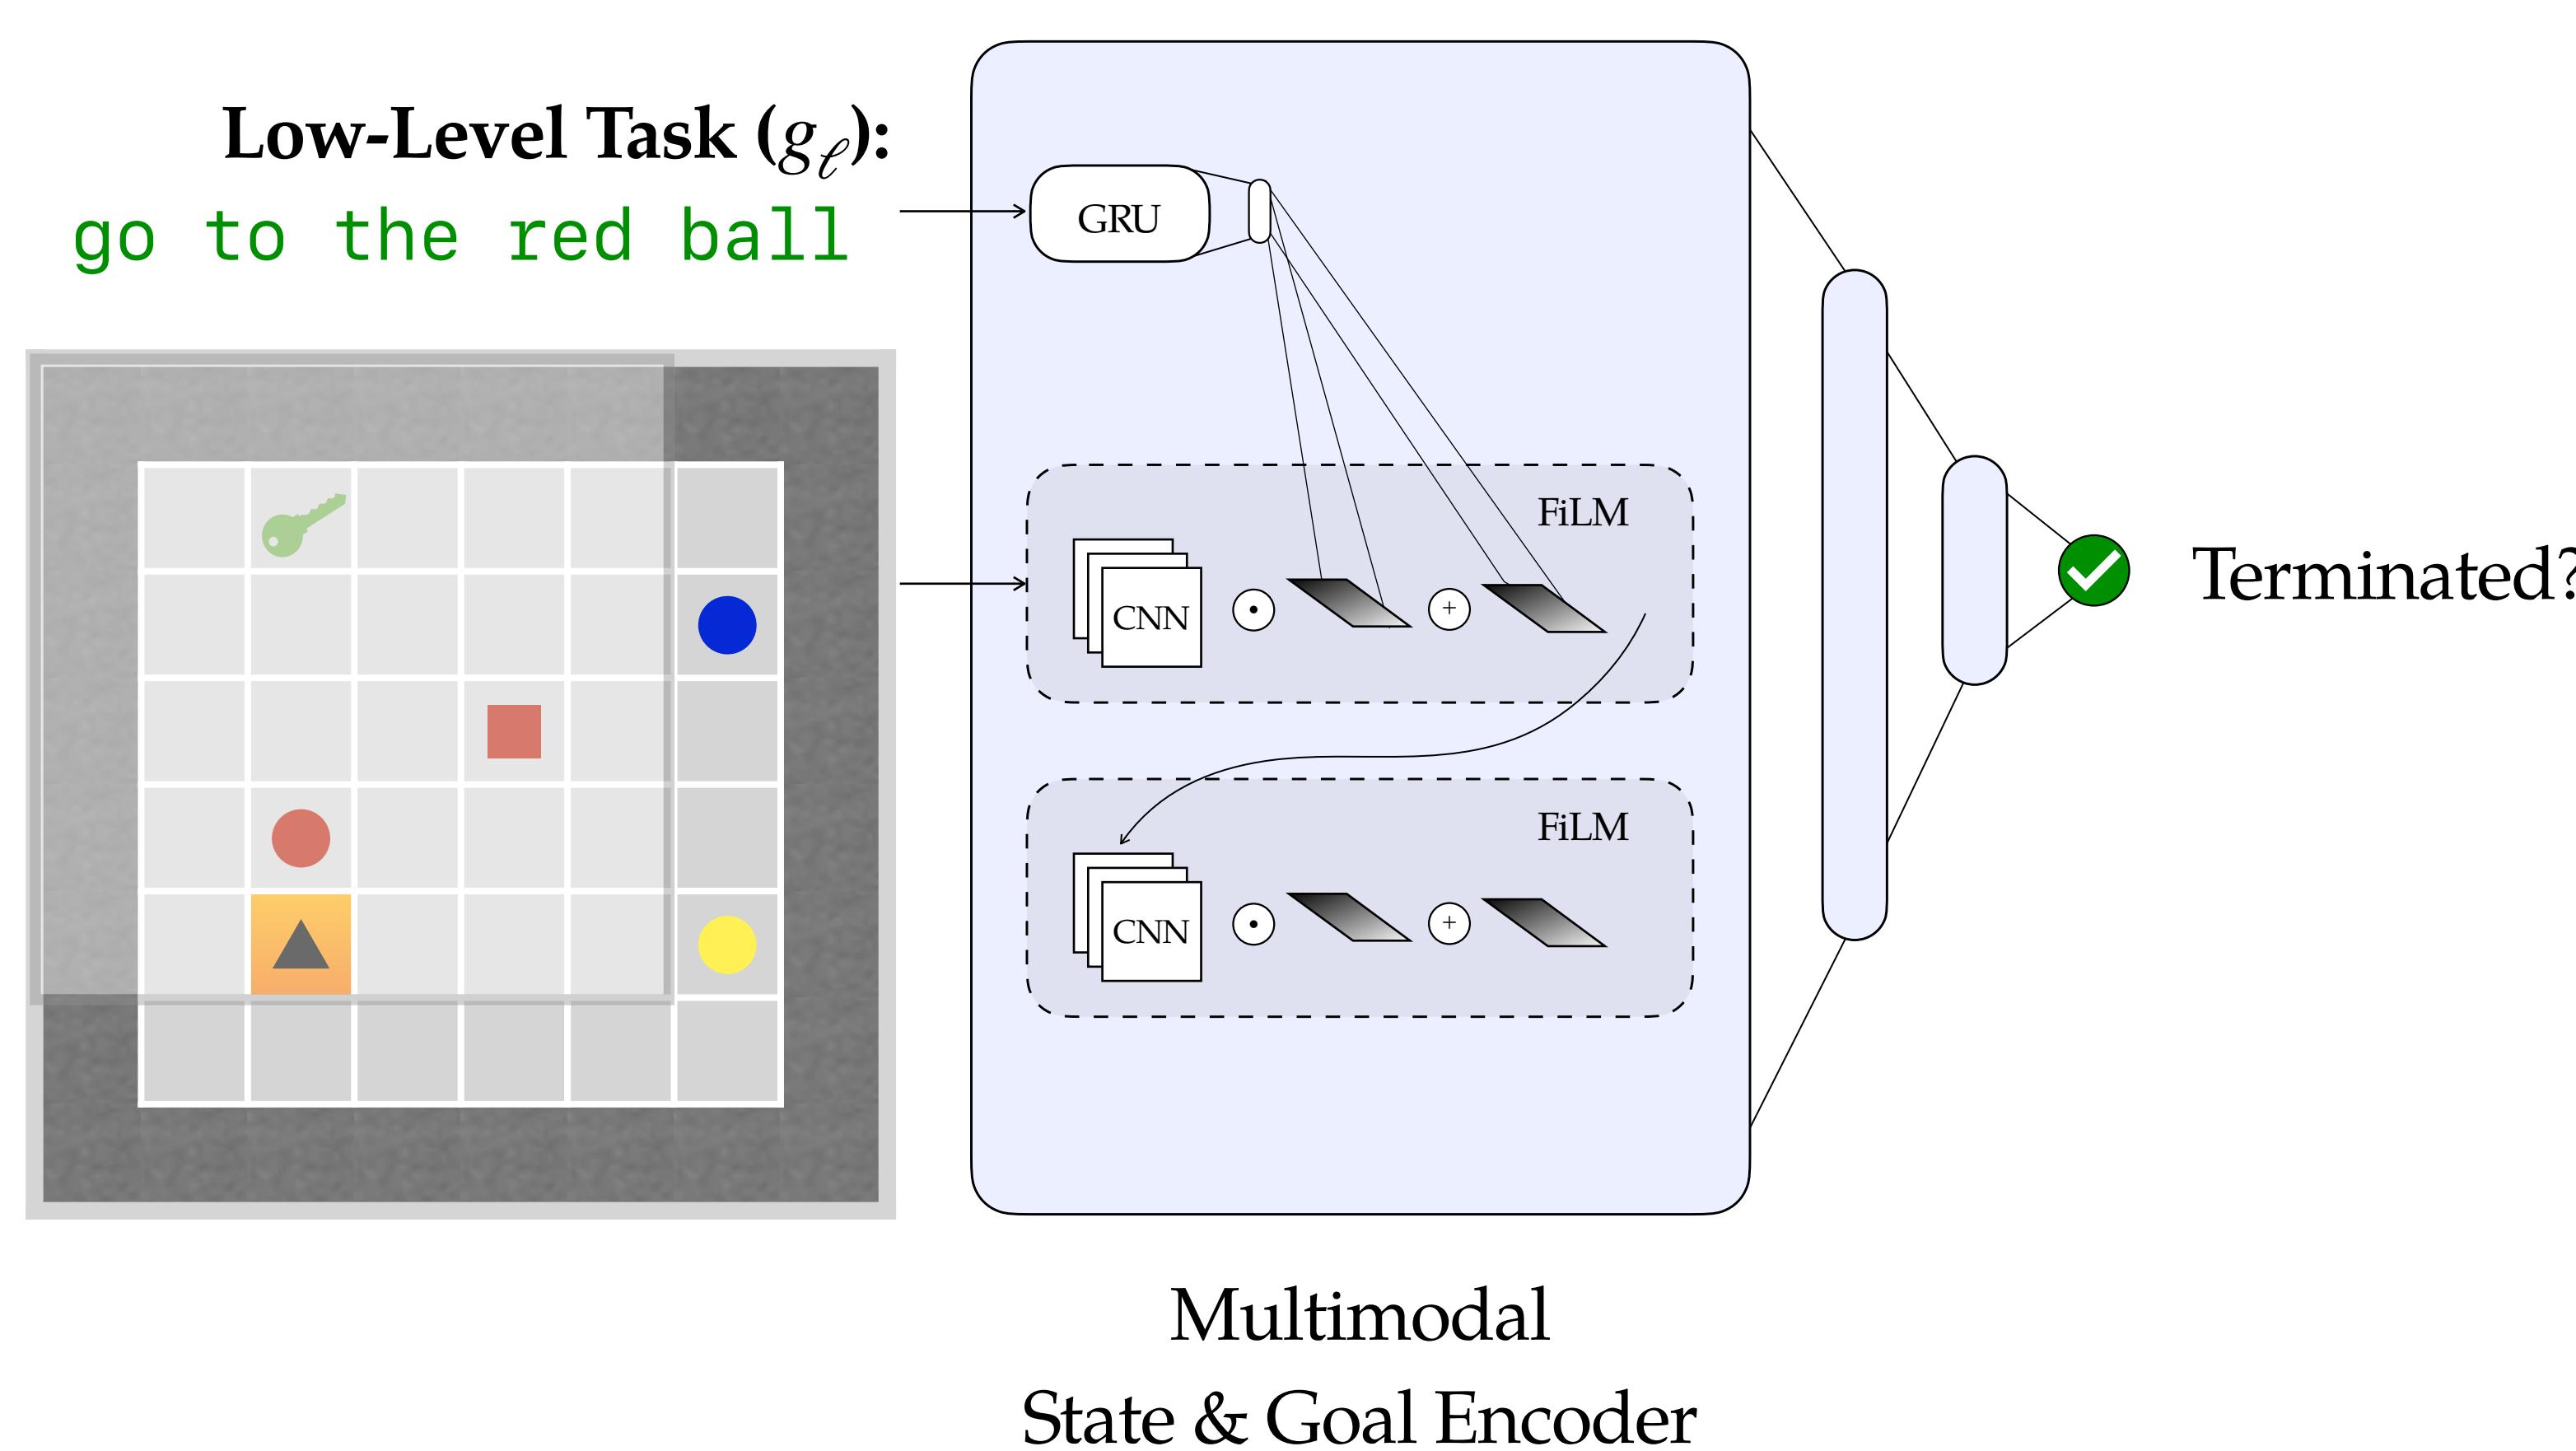
"go to the red ball"



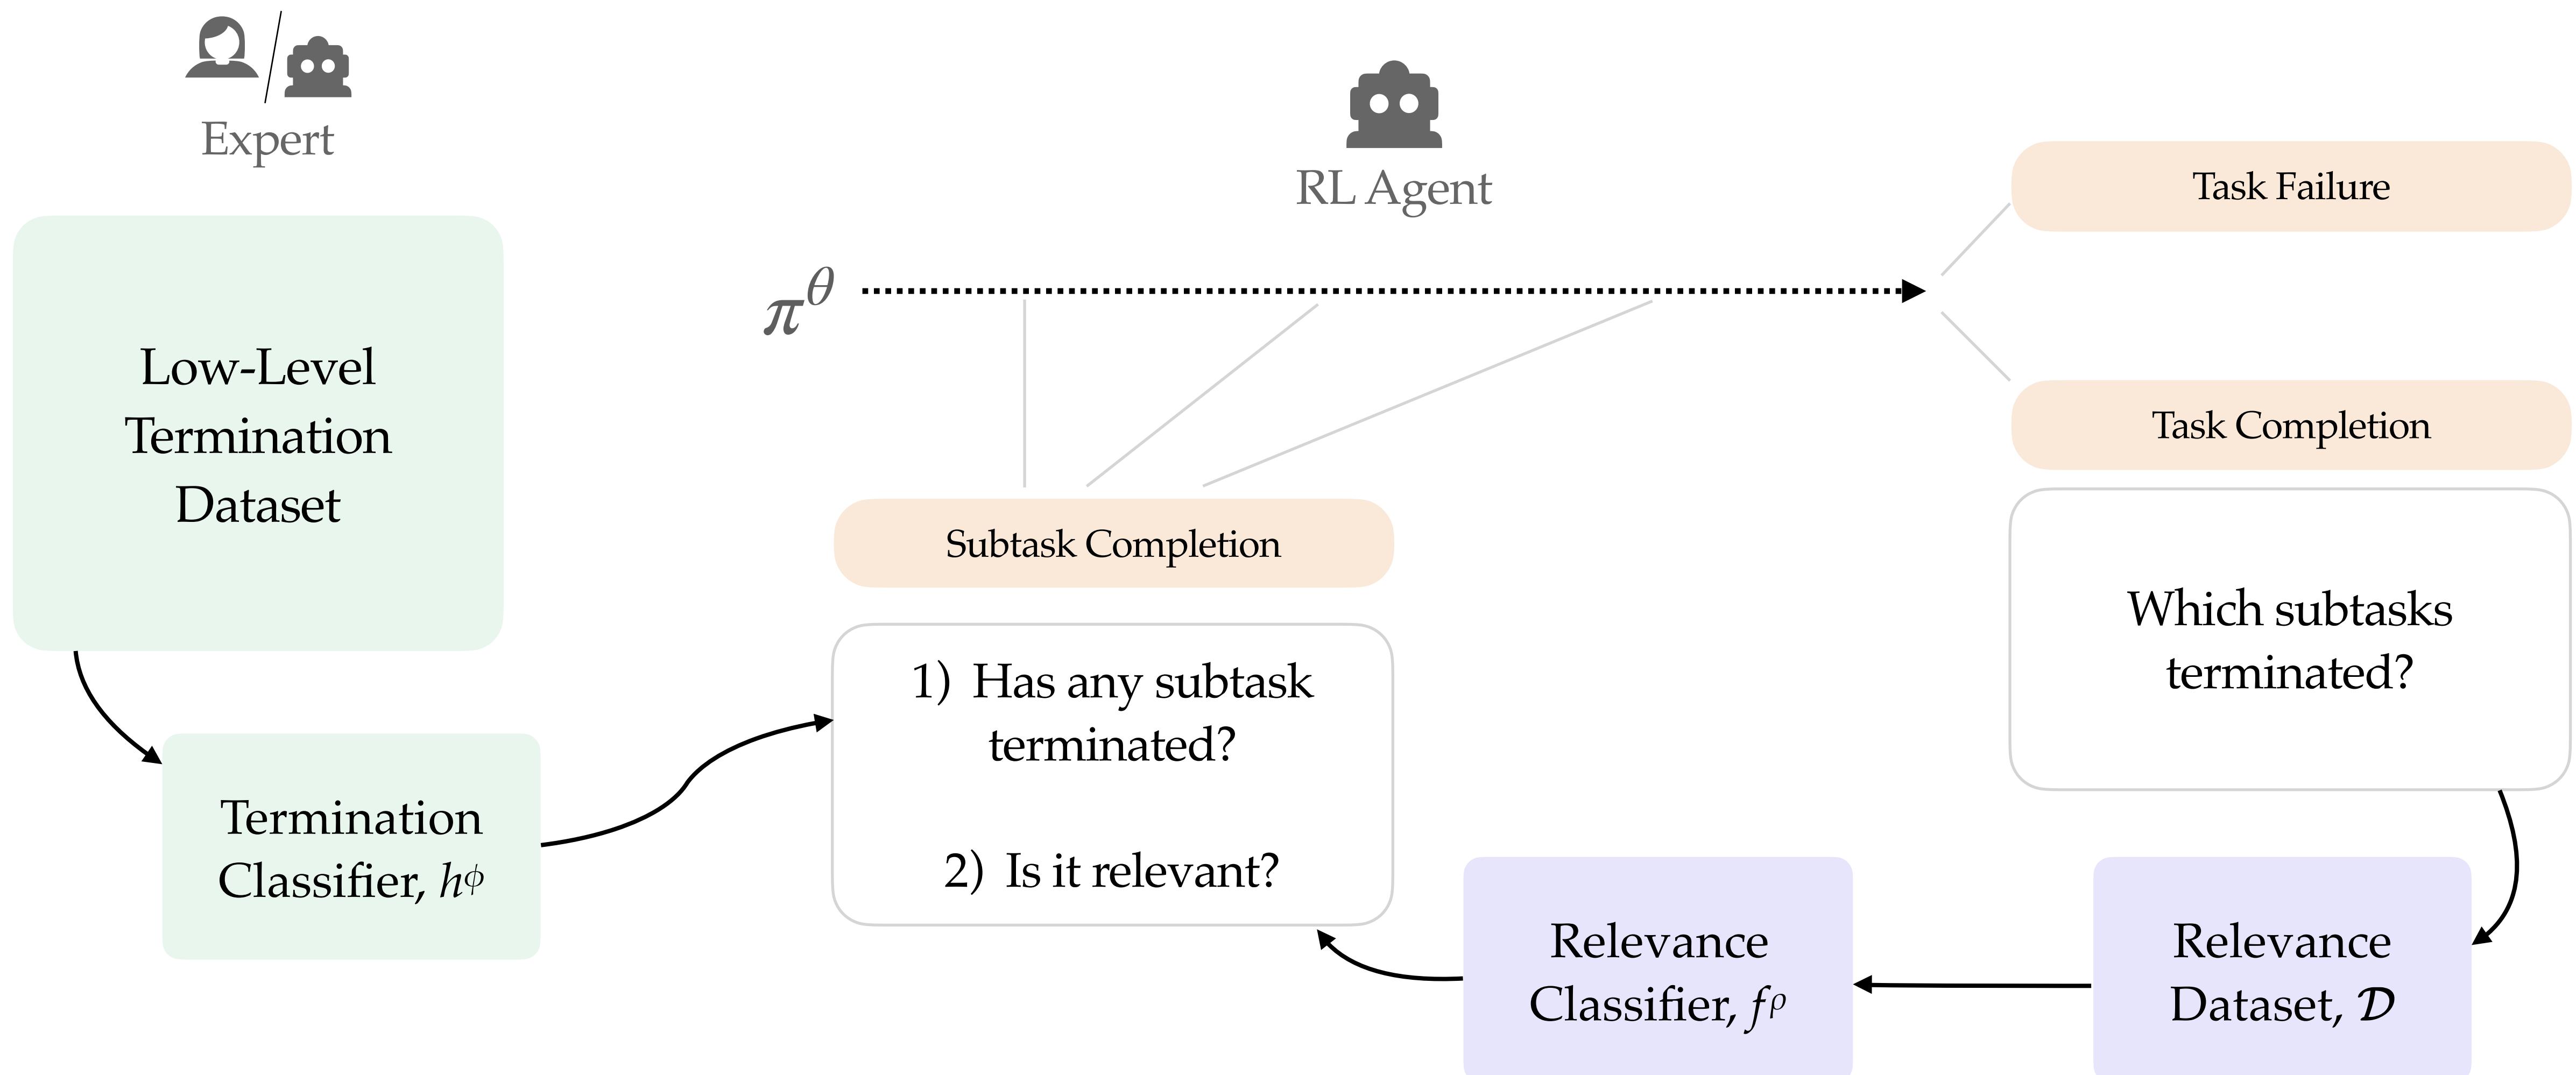
"go to the blue ball"



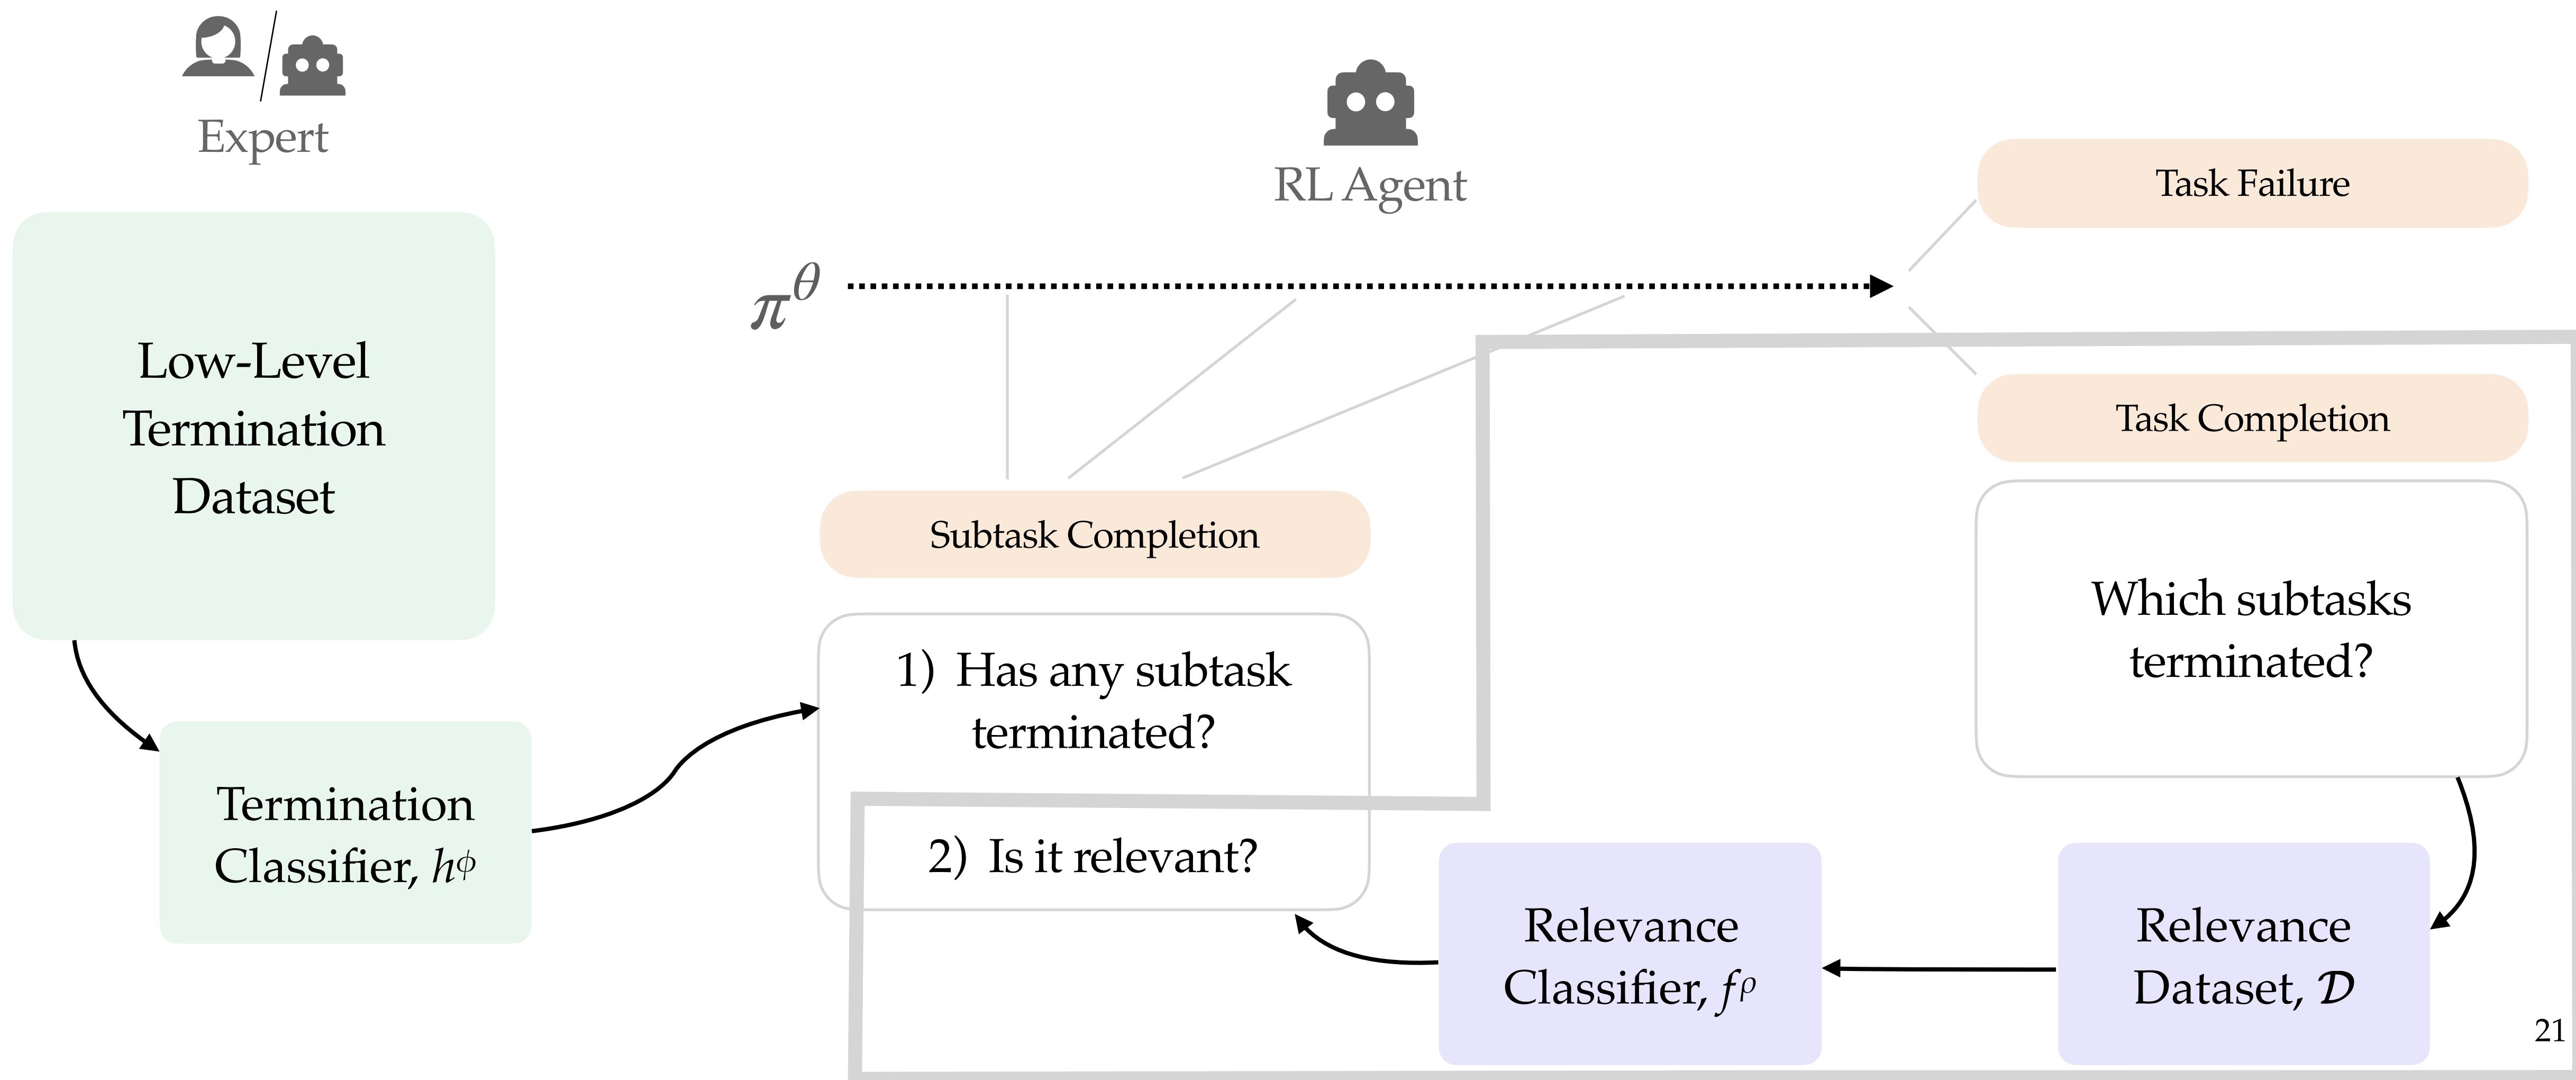
Low-Level Termination Classifier



Approach



Approach



Relevance Classifier

Relevance Classifier

- Mapping $G \rightarrow \mathcal{P}(G_\ell)$ is initially unknown

Relevance Classifier

- Mapping $G \rightarrow \mathcal{P}(G_\ell)$ is initially unknown
- Example: open the red door →
{ pick up the red key,
go to the red door }

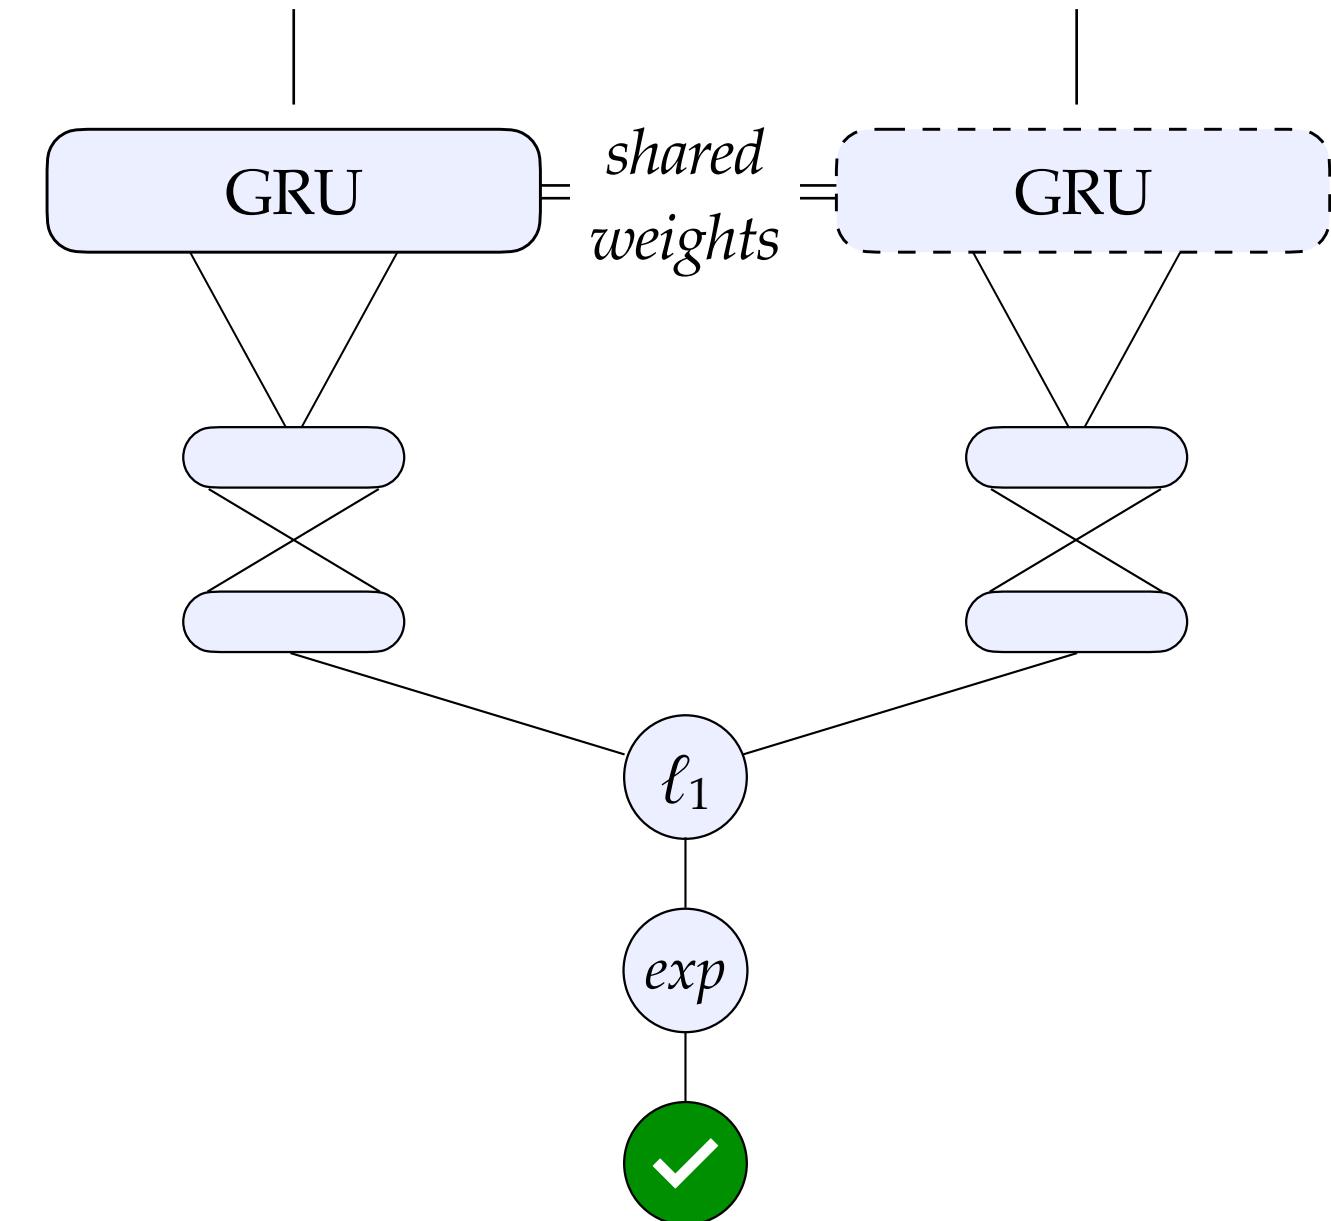
Relevance Classifier

- Mapping $G \rightarrow \mathcal{P}(G_\ell)$ is initially unknown
 - Example: open the red door → { pick up the red key, go to the red door }
- $f^\rho : G \times G_\ell \rightarrow \{0,1\}$

Relevance Classifier

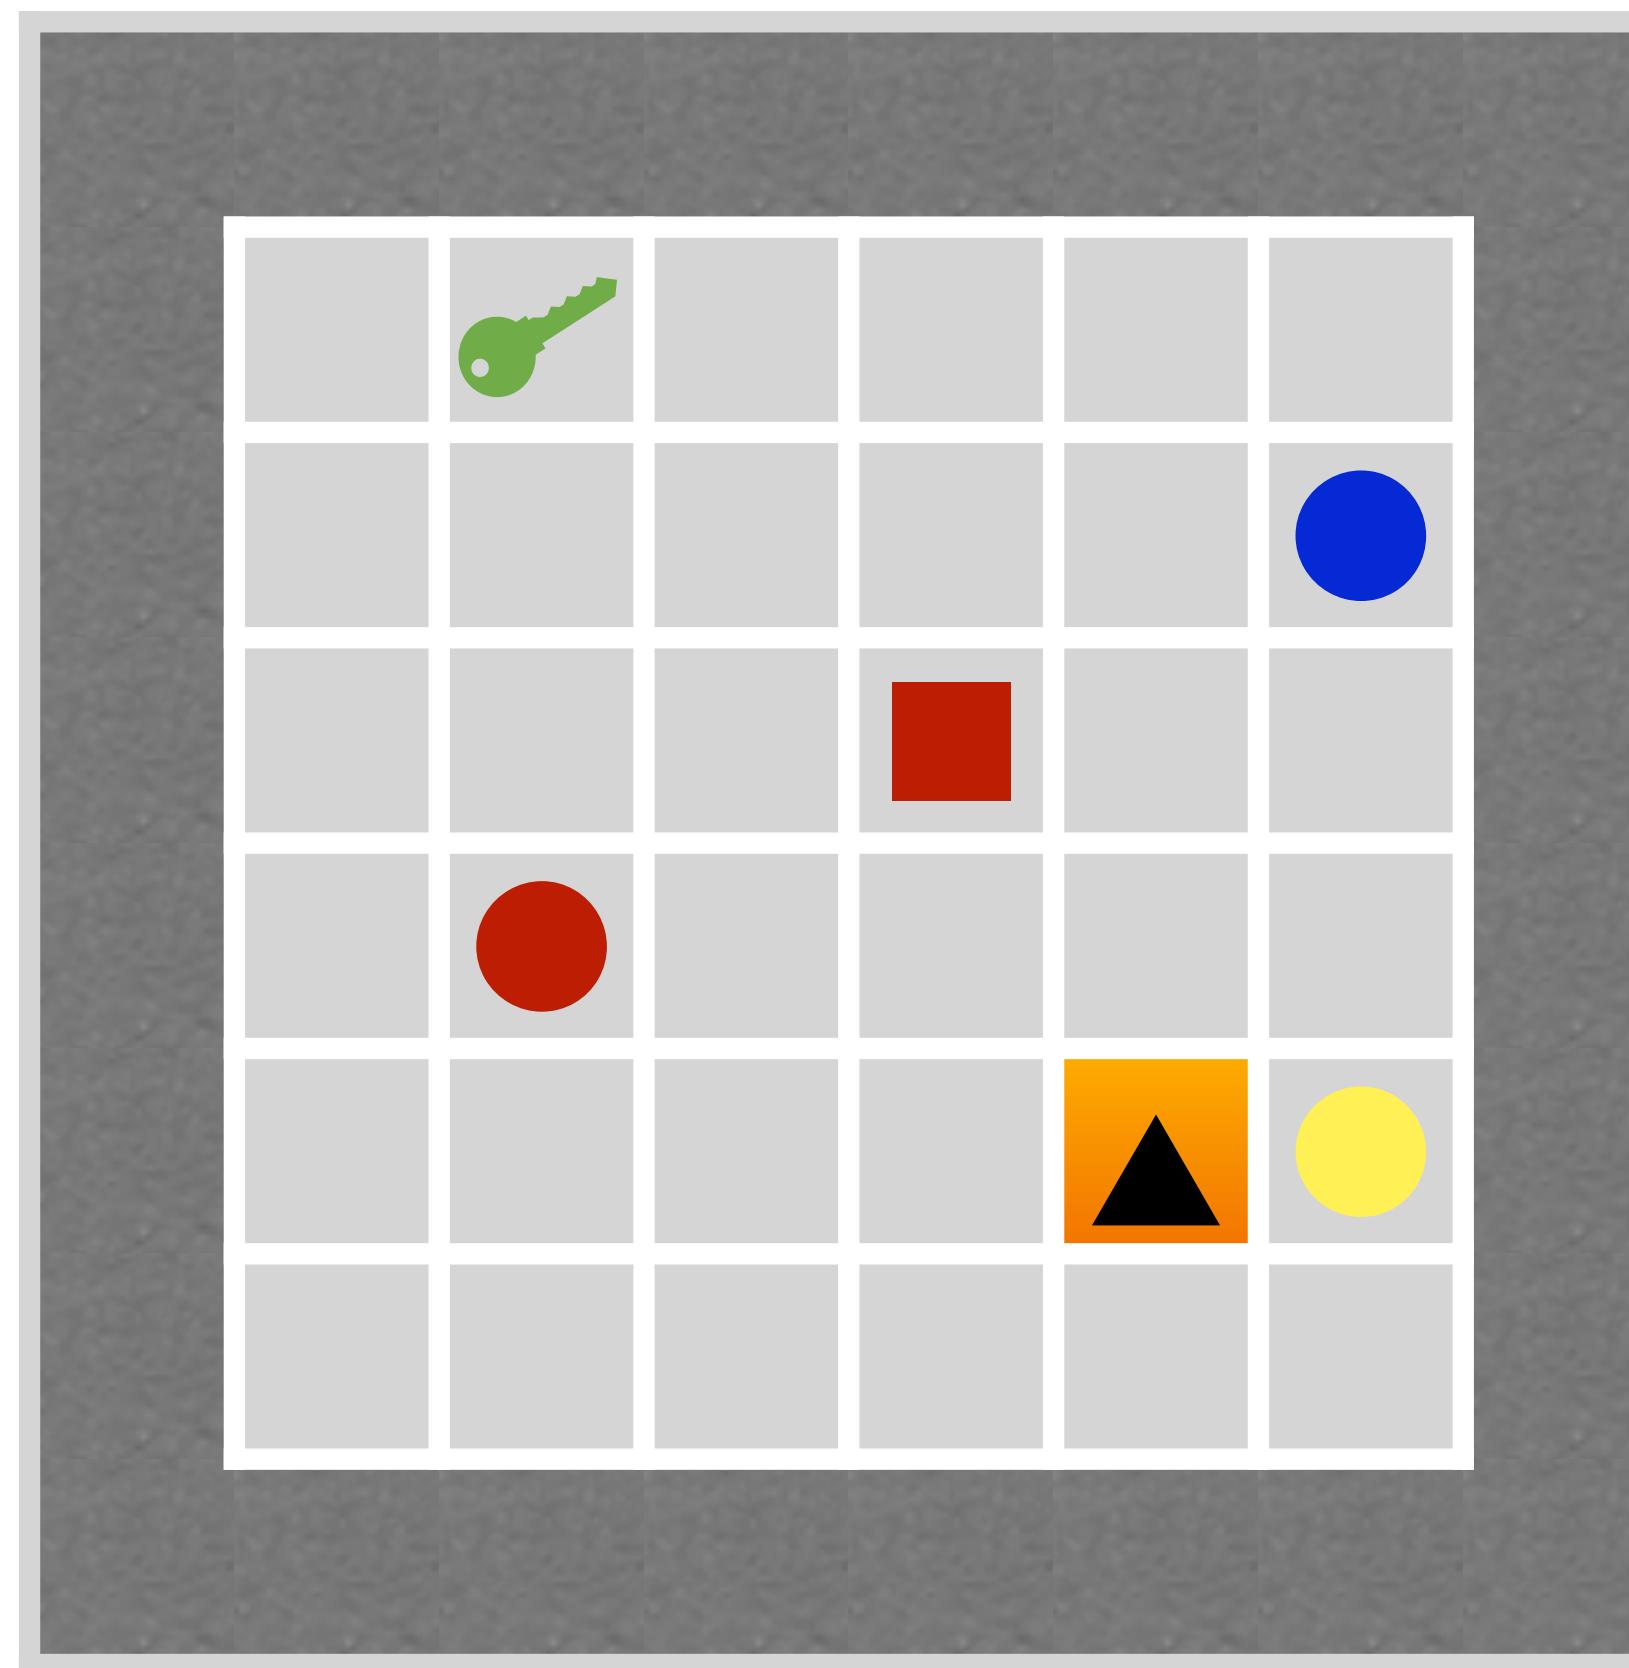
- Mapping $G \rightarrow \mathcal{P}(G_\ell)$ is initially unknown
- Example: open the red door → { pick up the red key, go to the red door }
- $f^\rho : G \times G_\ell \rightarrow \{0,1\}$

High-Level Task (g): Low-Level Task (g_ℓ):
open the red door pick up the red key



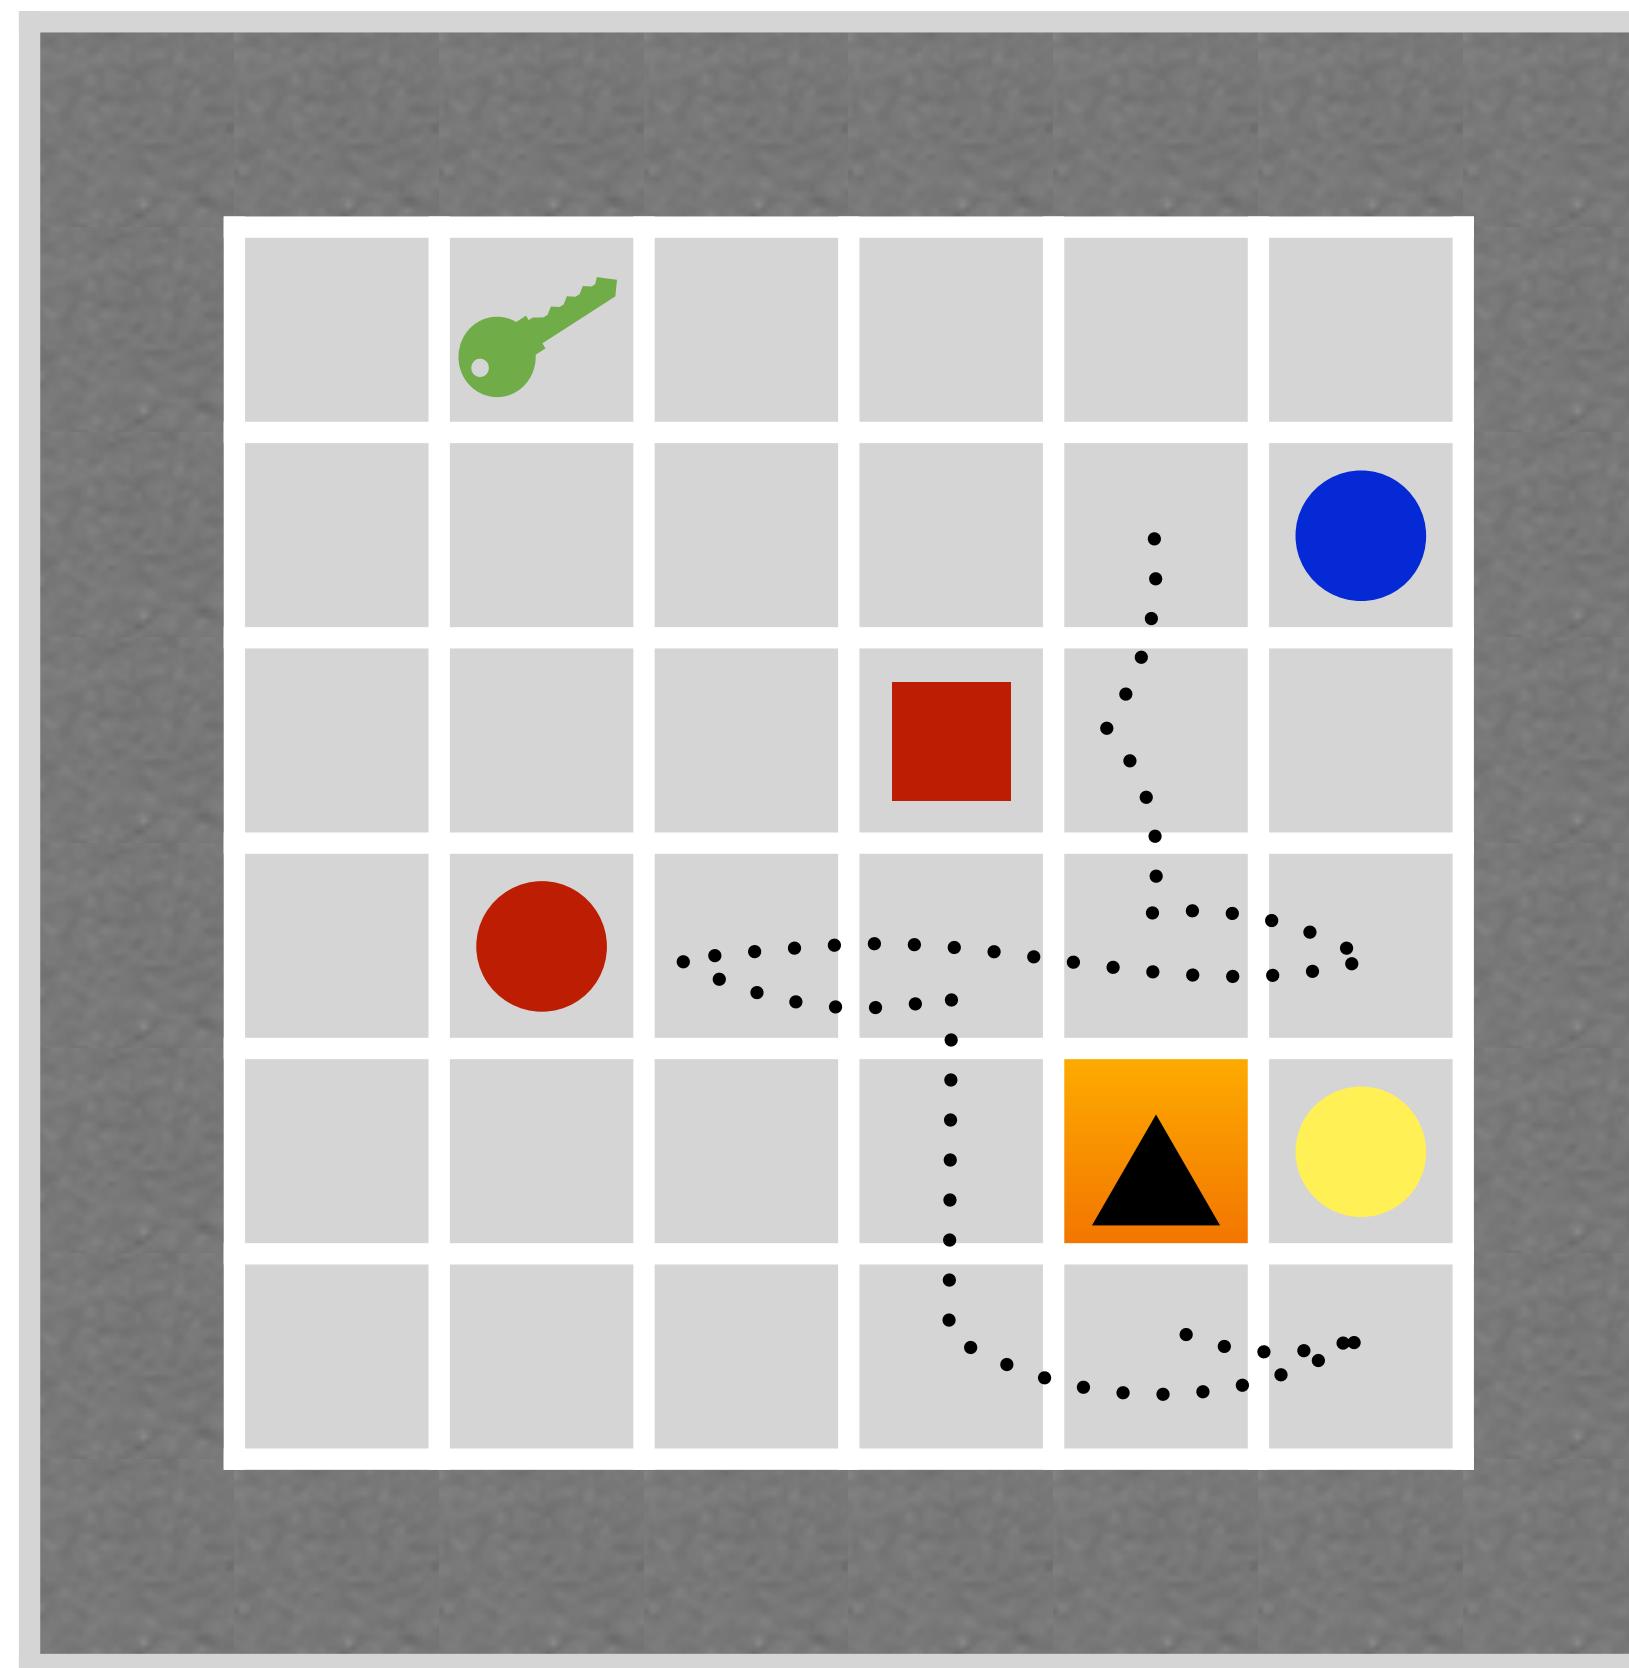
Relevant?

Collecting Relevance Data Online



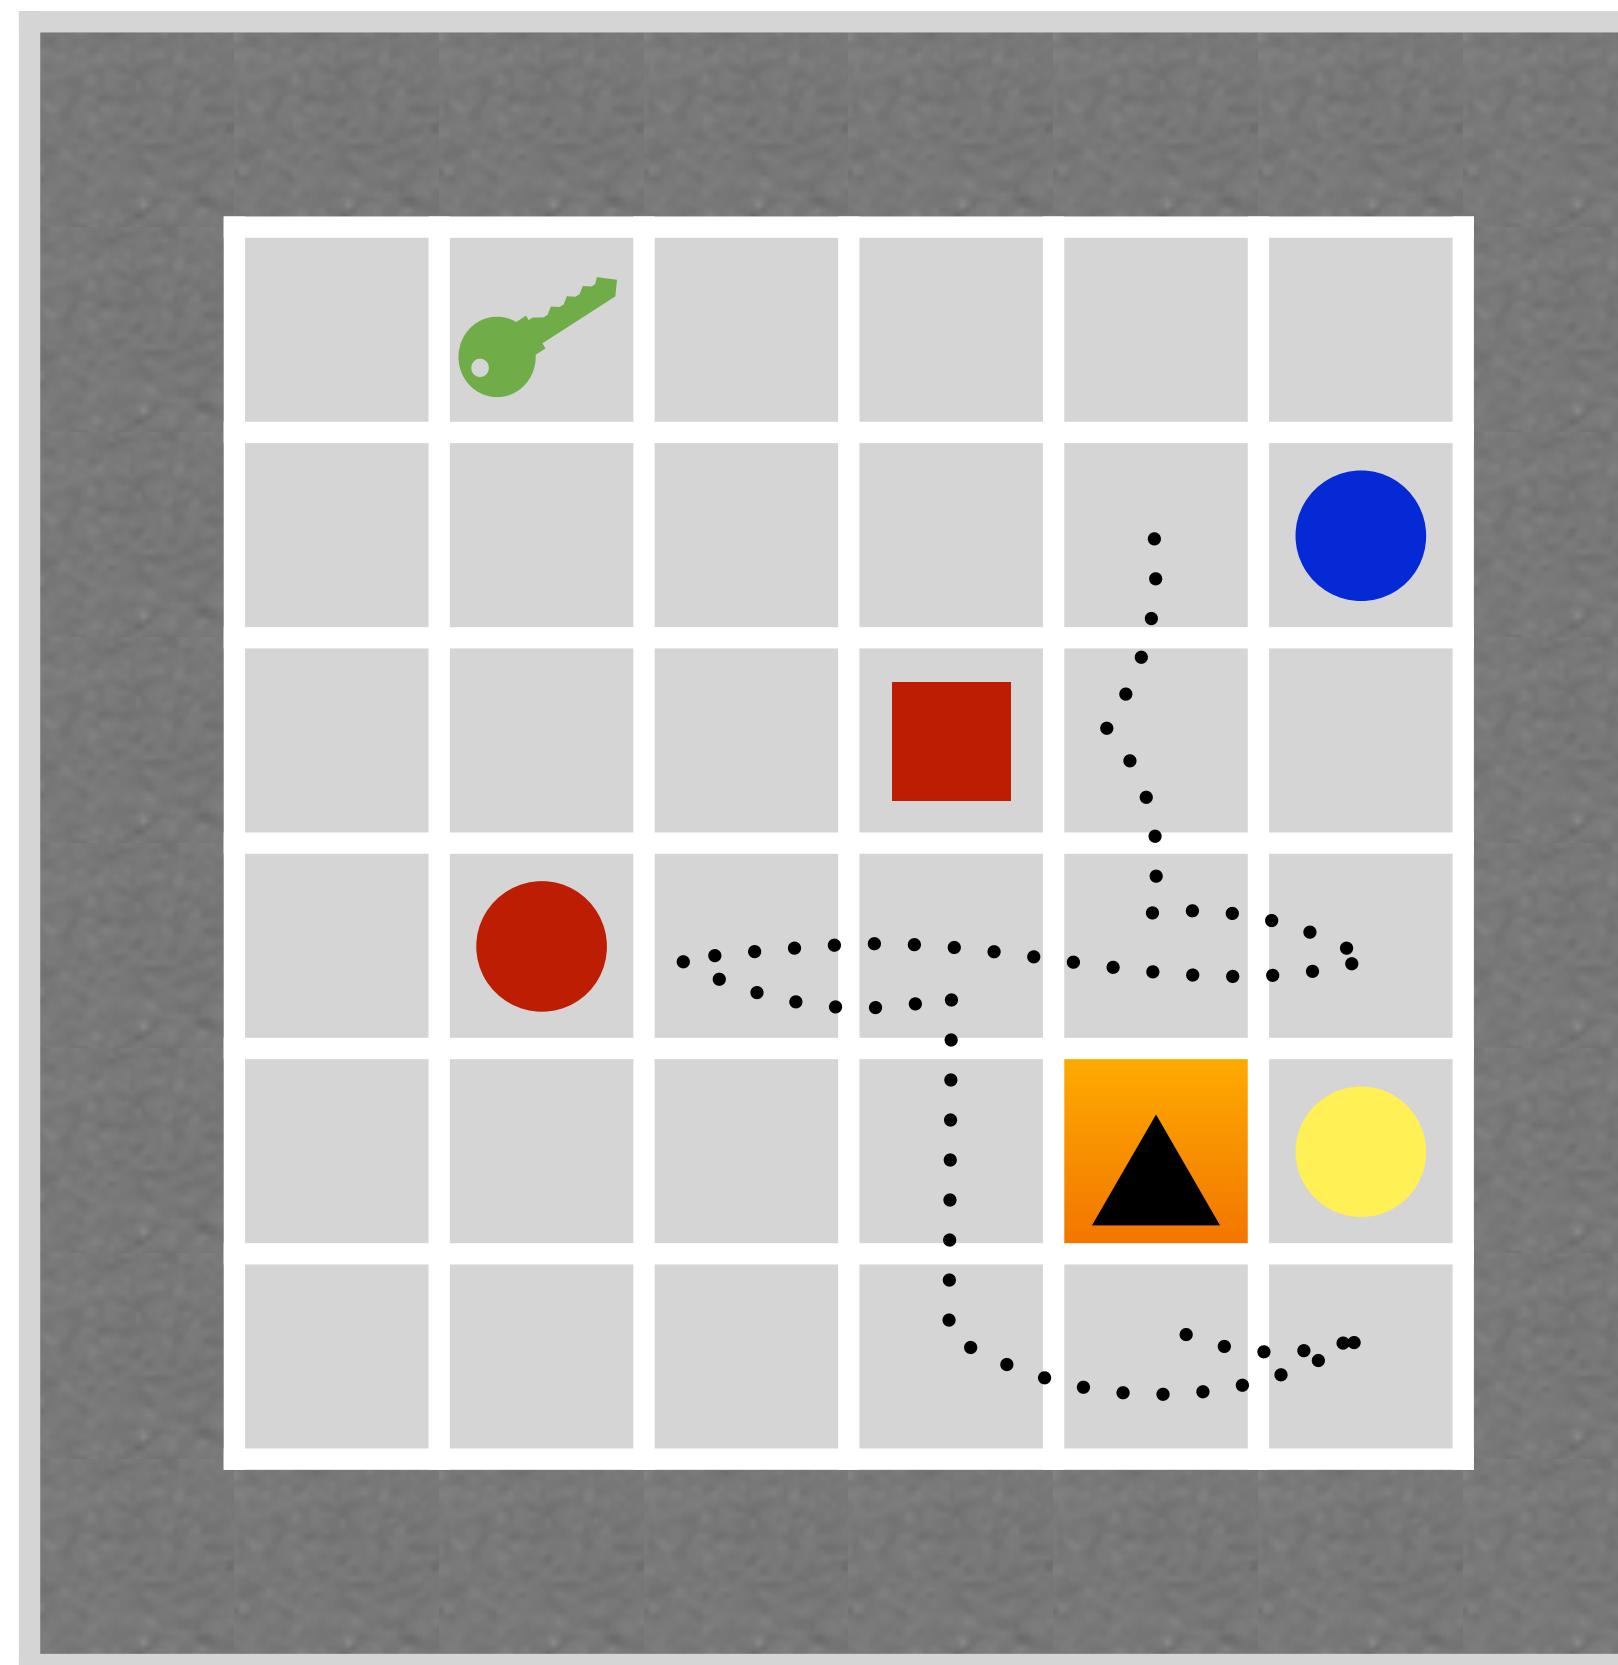
put the red ball next
to the blue ball

Collecting Relevance Data Online



put the red ball next
to the blue ball

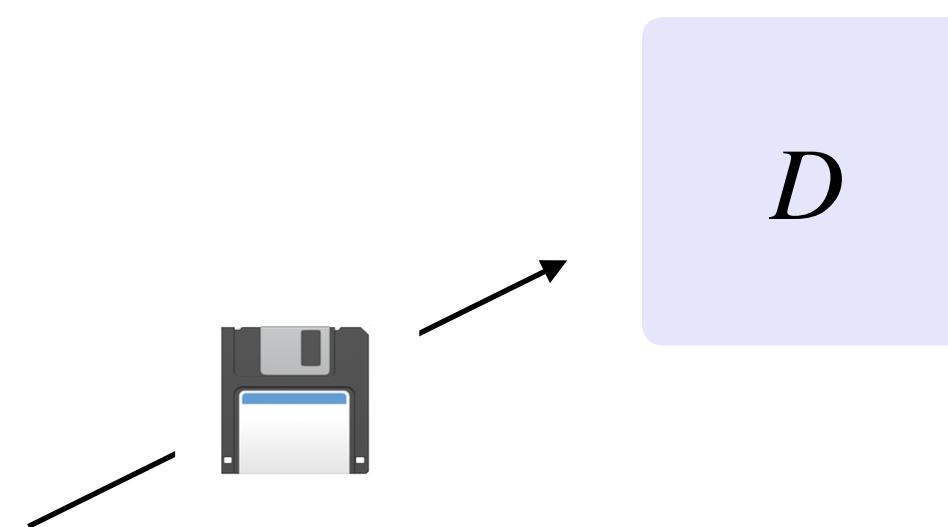
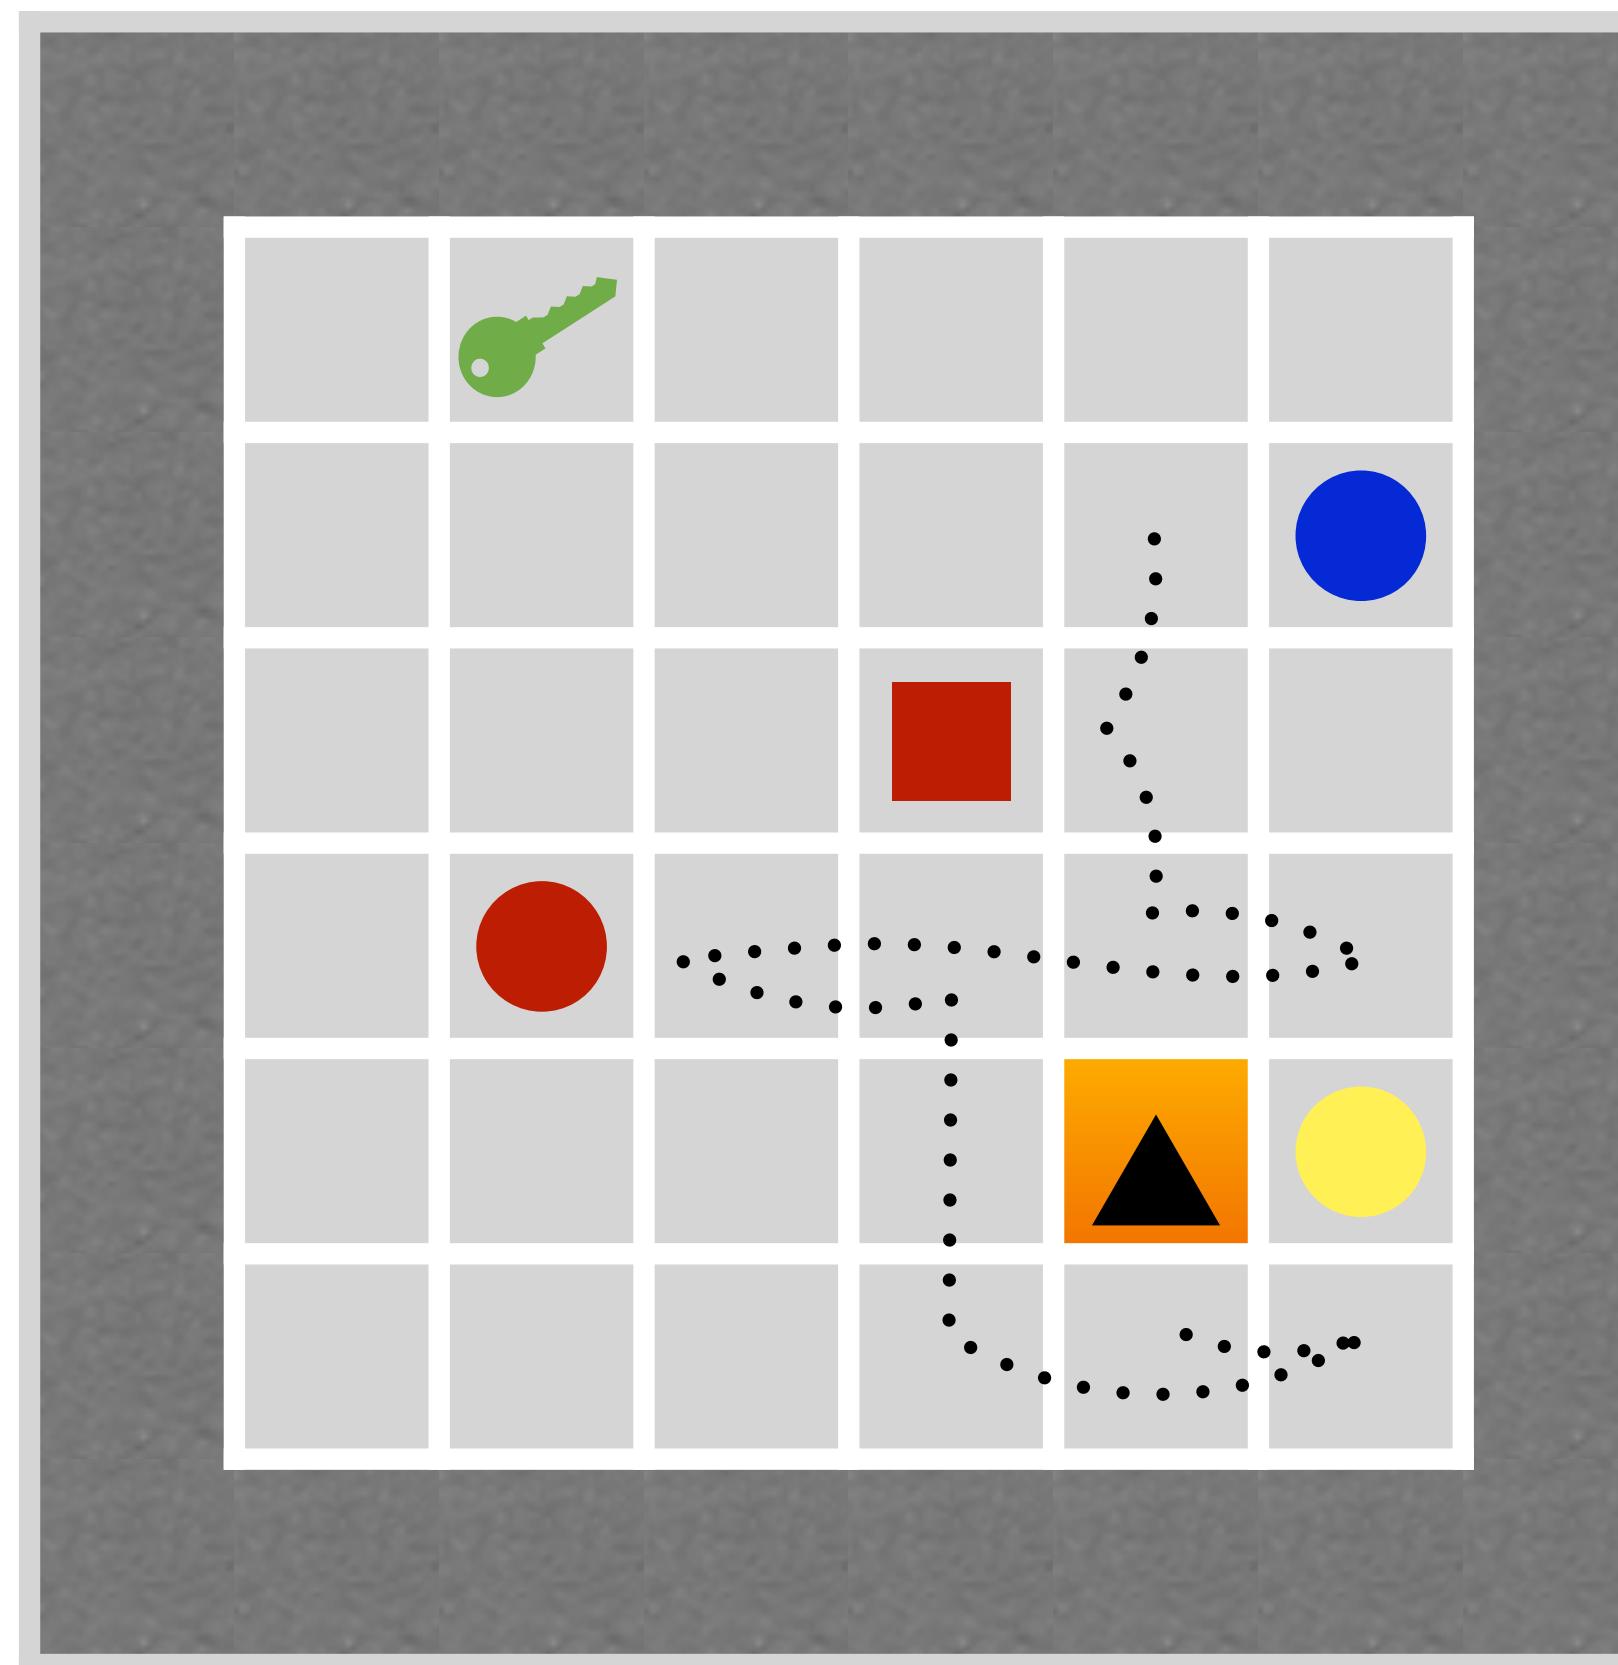
Collecting Relevance Data Online



put the red ball next
to the blue ball

- go to the yellow ball
- go to the red ball
- go to the red square
- go to the blue ball

Collecting Relevance Data Online

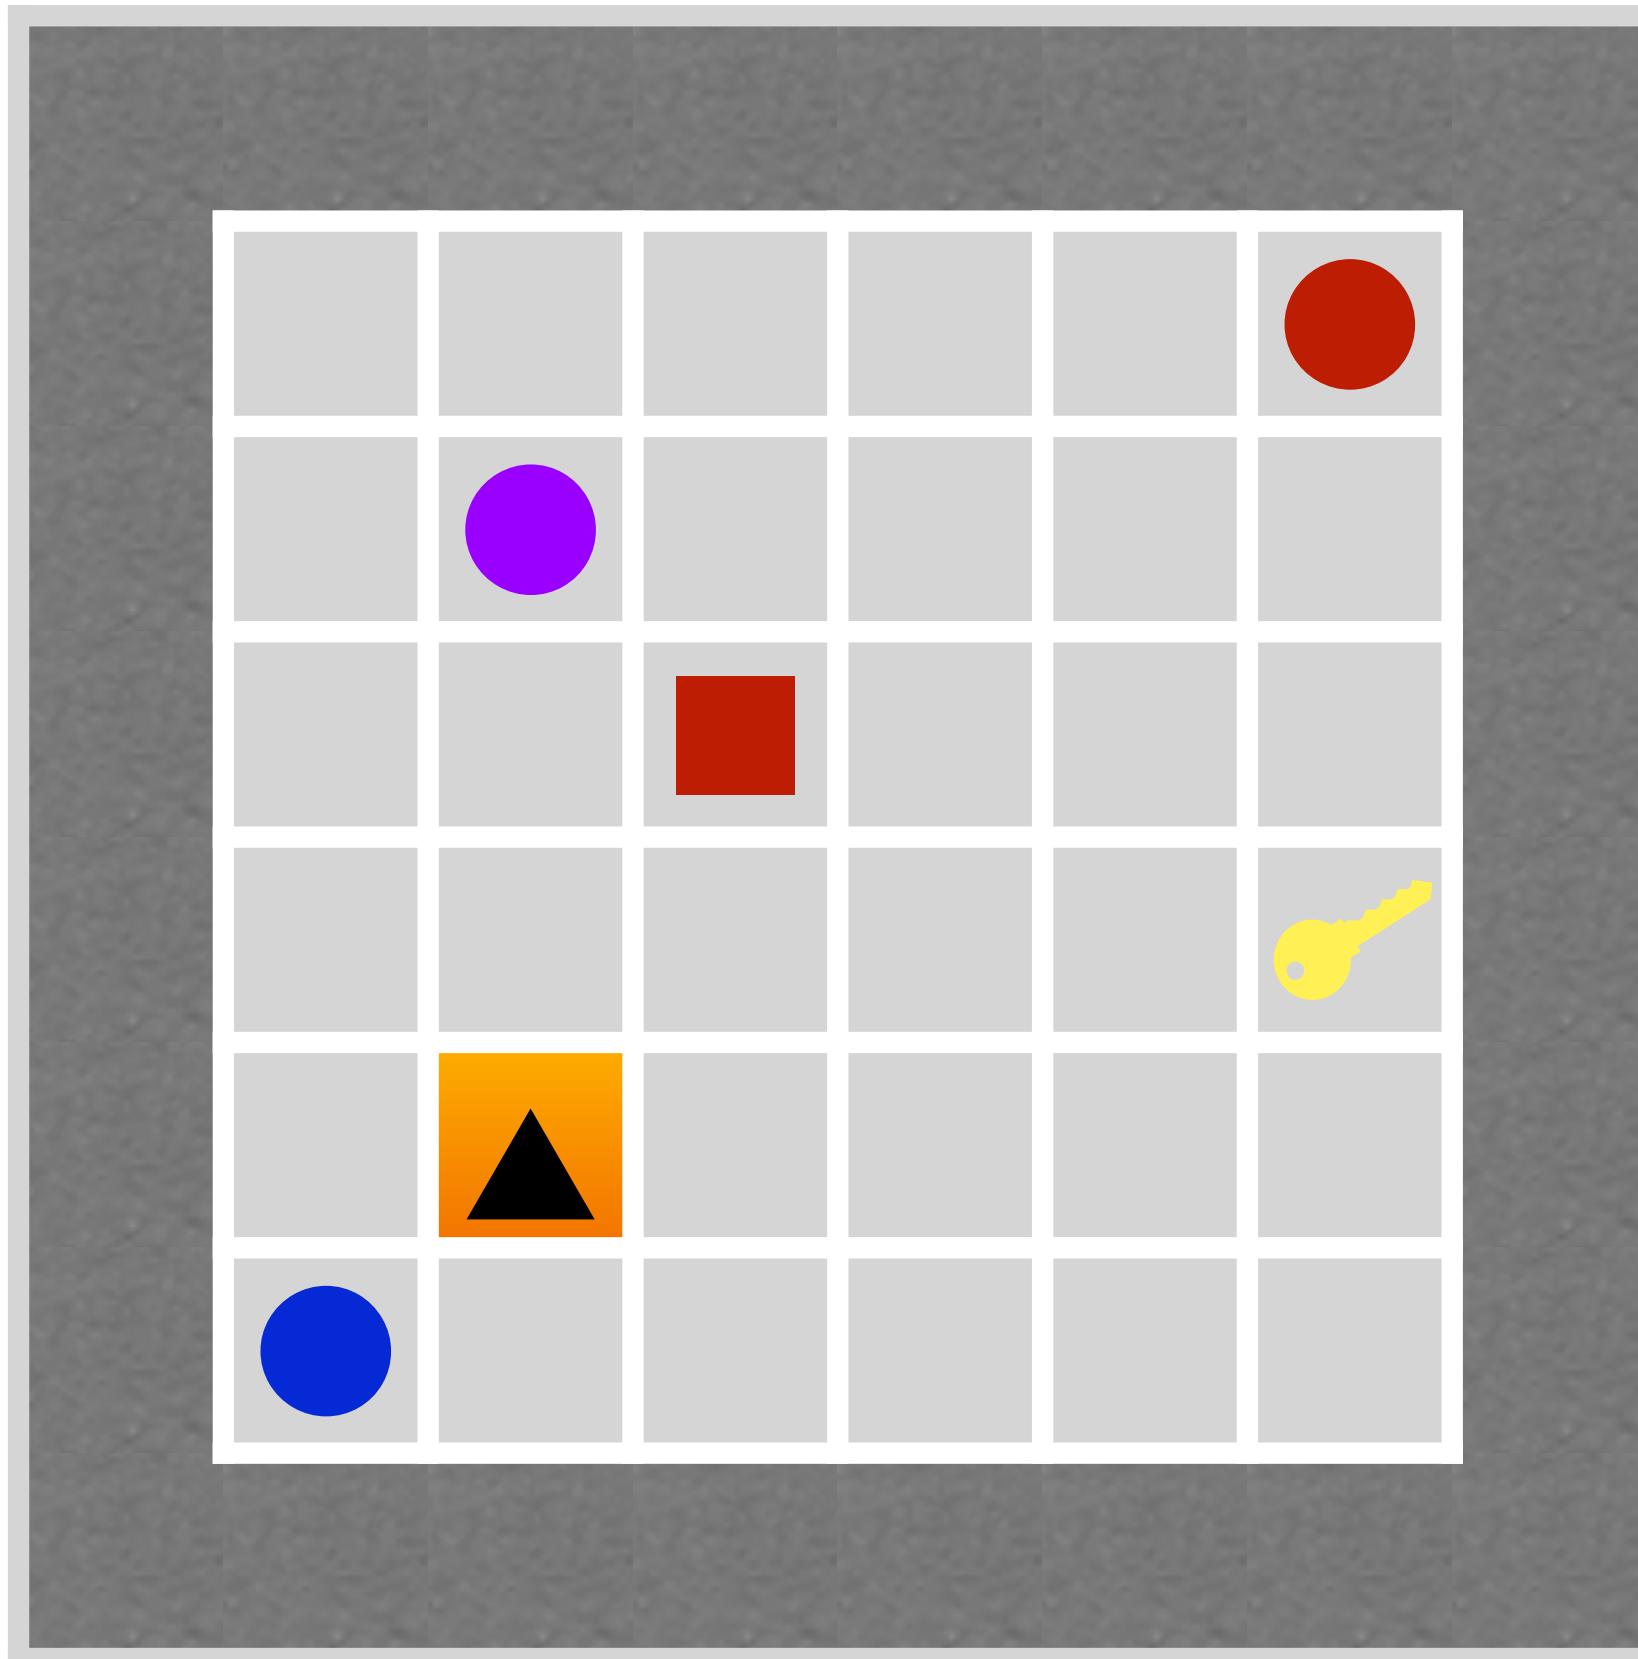


put the red ball next
to the blue ball

- go to the yellow ball
- go to the red ball
- go to the red square
- go to the blue ball

Deduplication

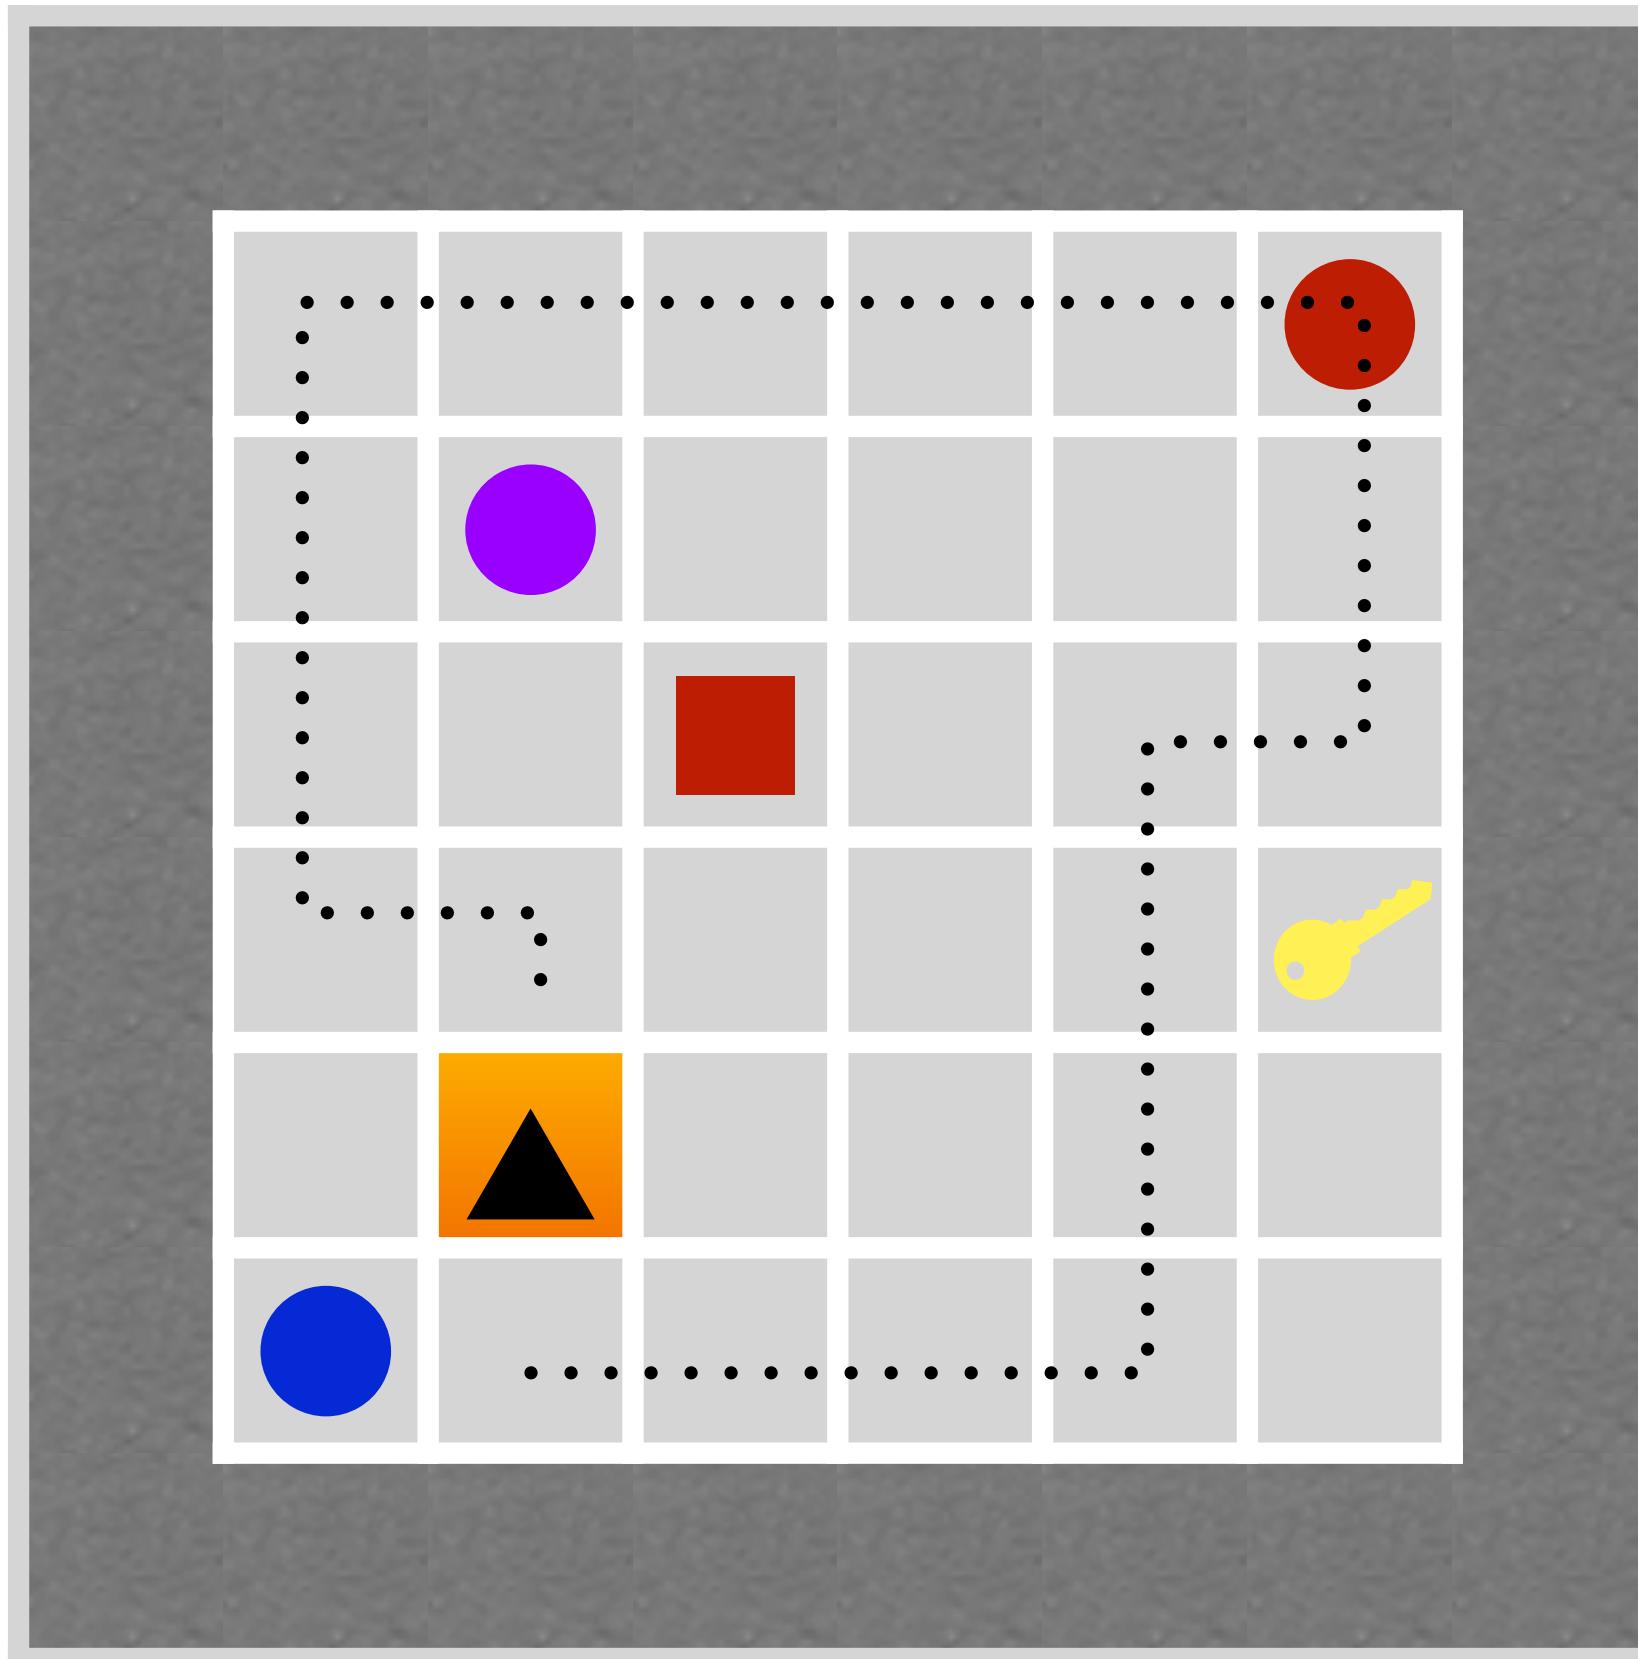
D



put the red ball next
to the blue ball

Deduplication

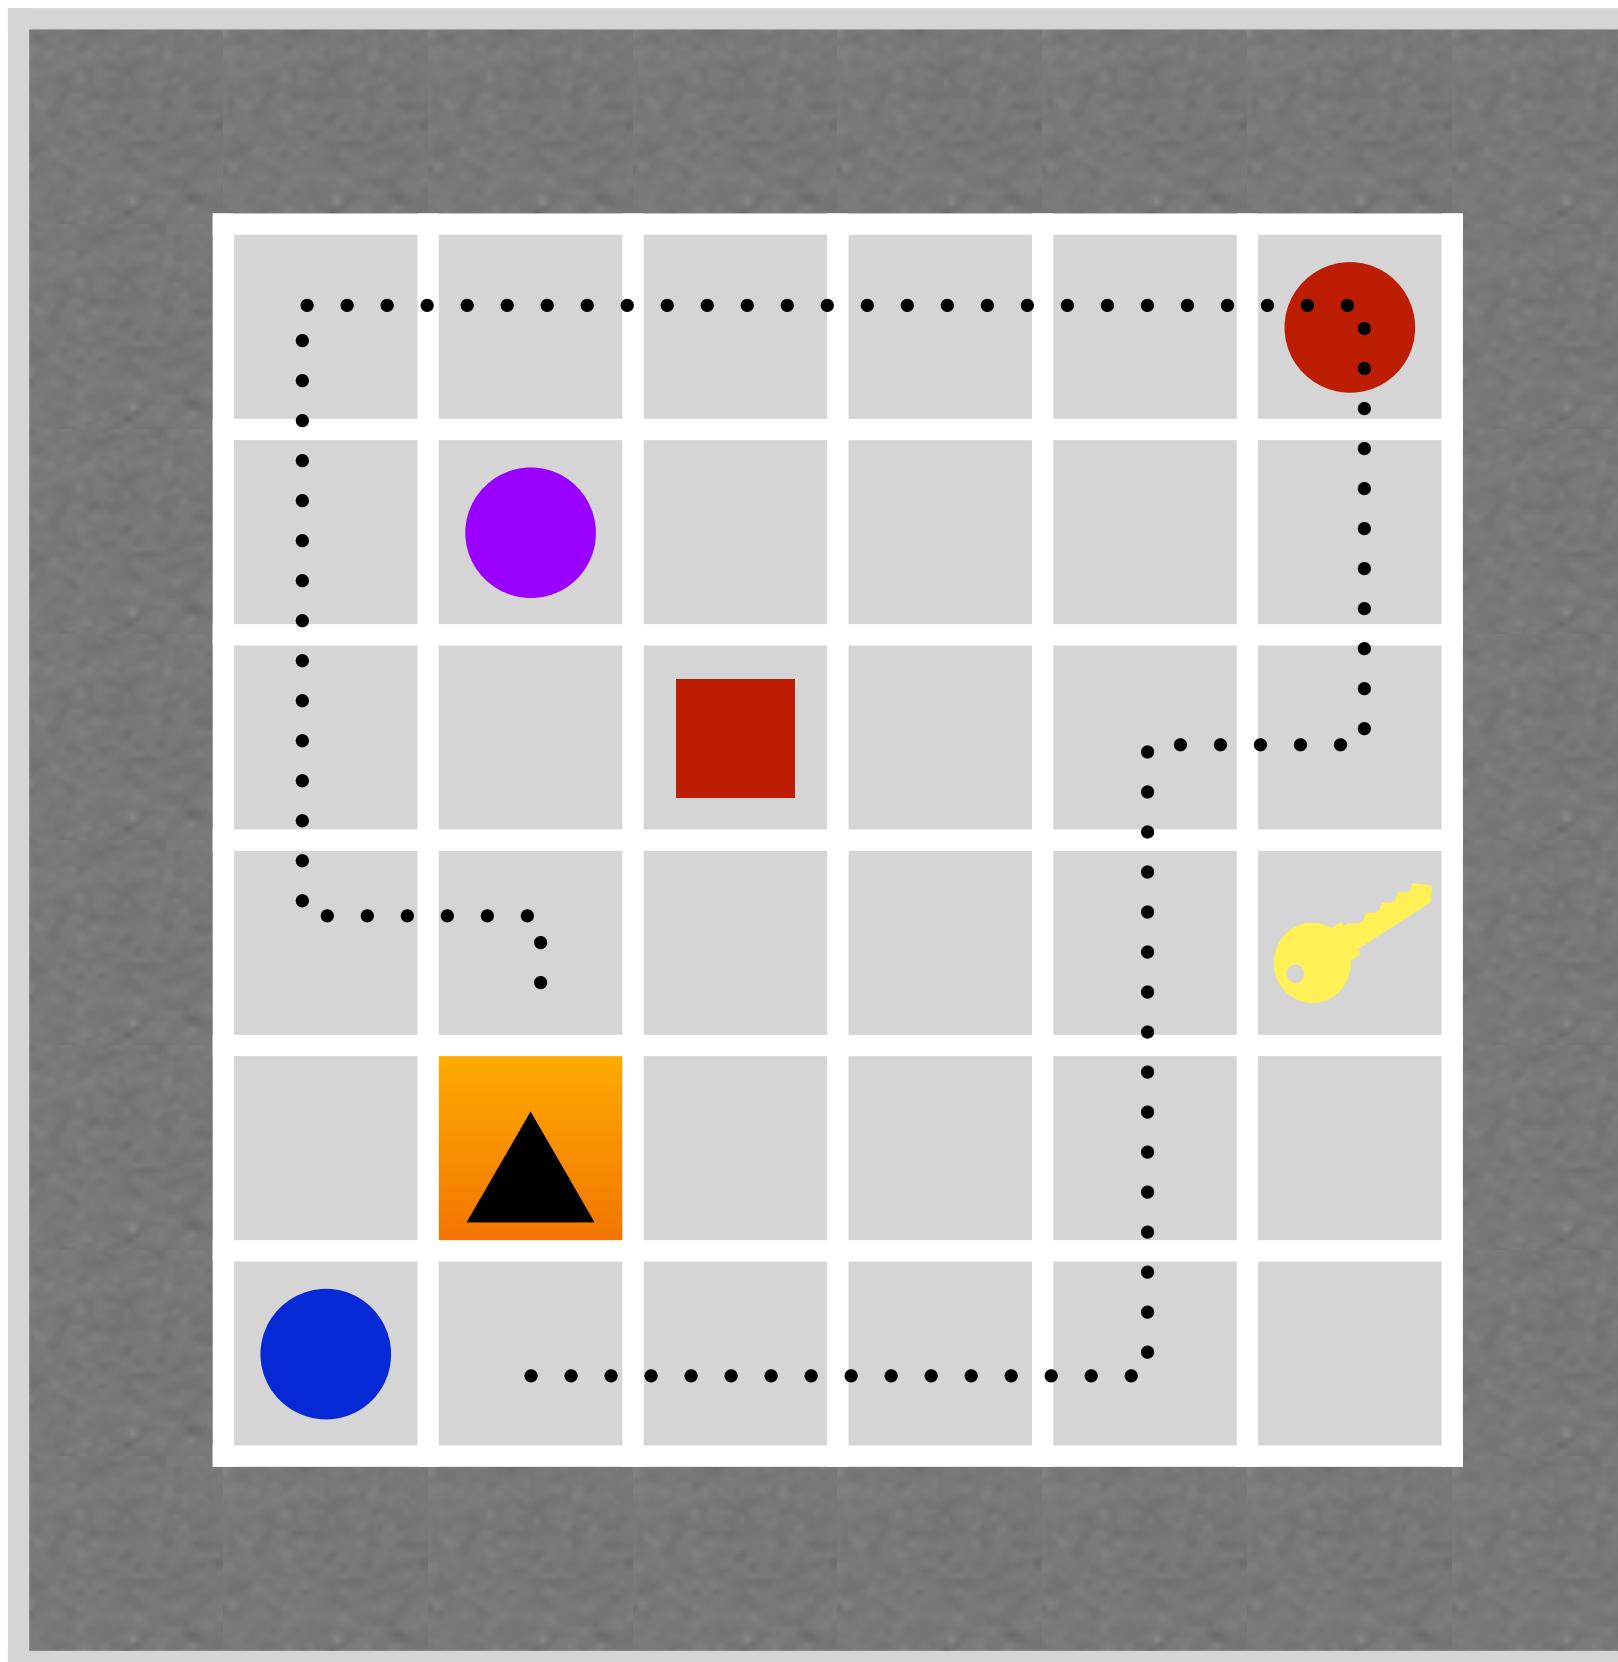
D



put the red ball next
to the blue ball

Deduplication

D



put the red ball next
to the blue ball

- go to the red ball
- go to the yellow key
- go to the blue ball

Deduplication

D

put the red ball next
to the blue ball

- go to the yellow ball
- go to the red ball
- go to the red square
- go to the blue ball

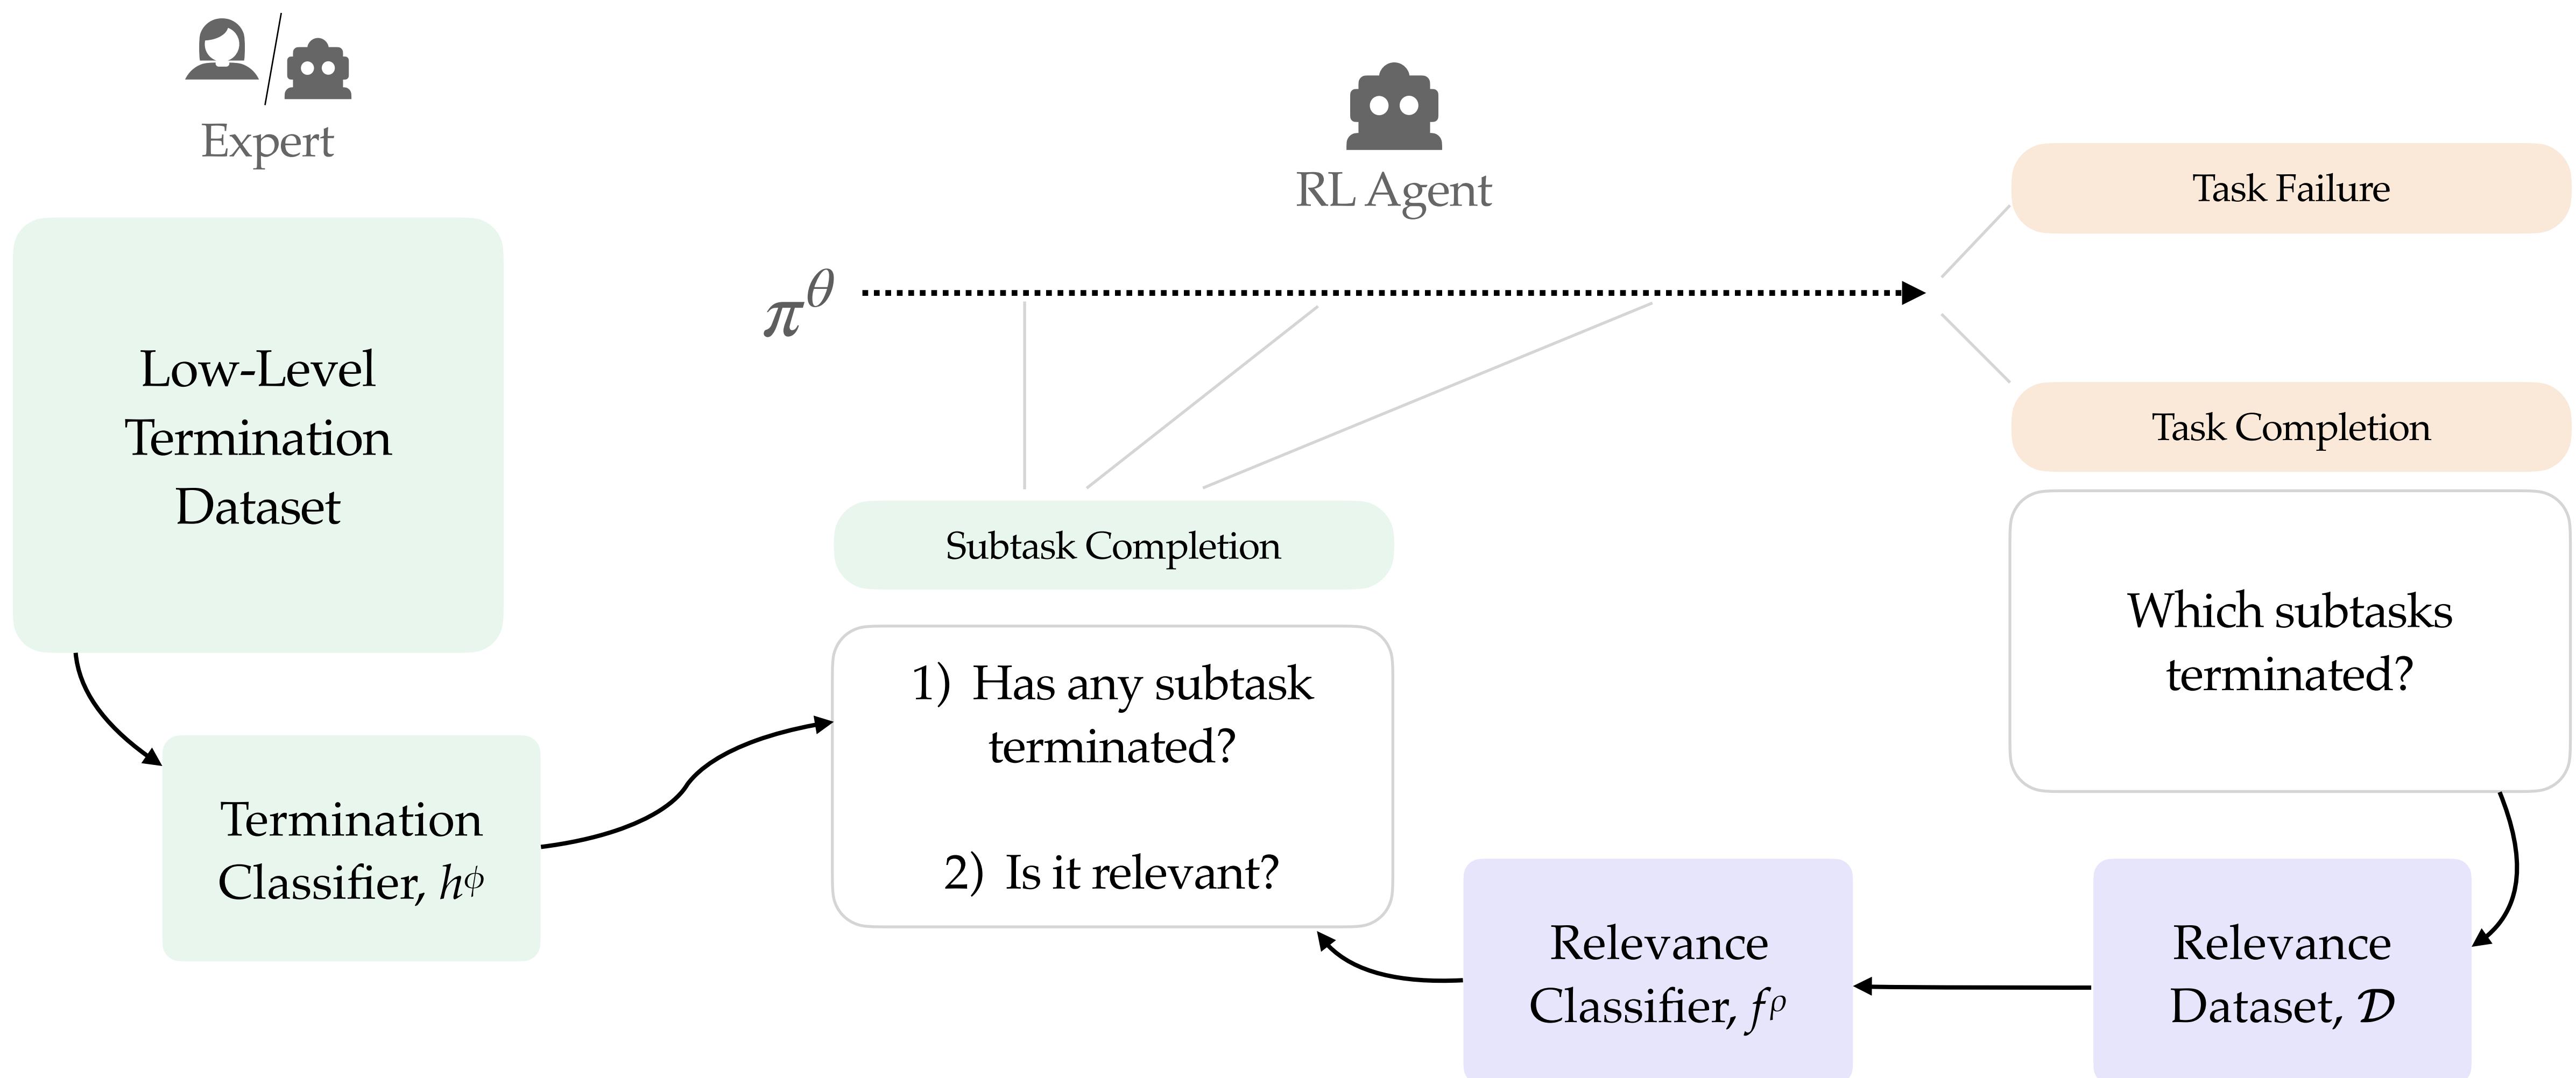
∩

put the red ball next
to the blue ball

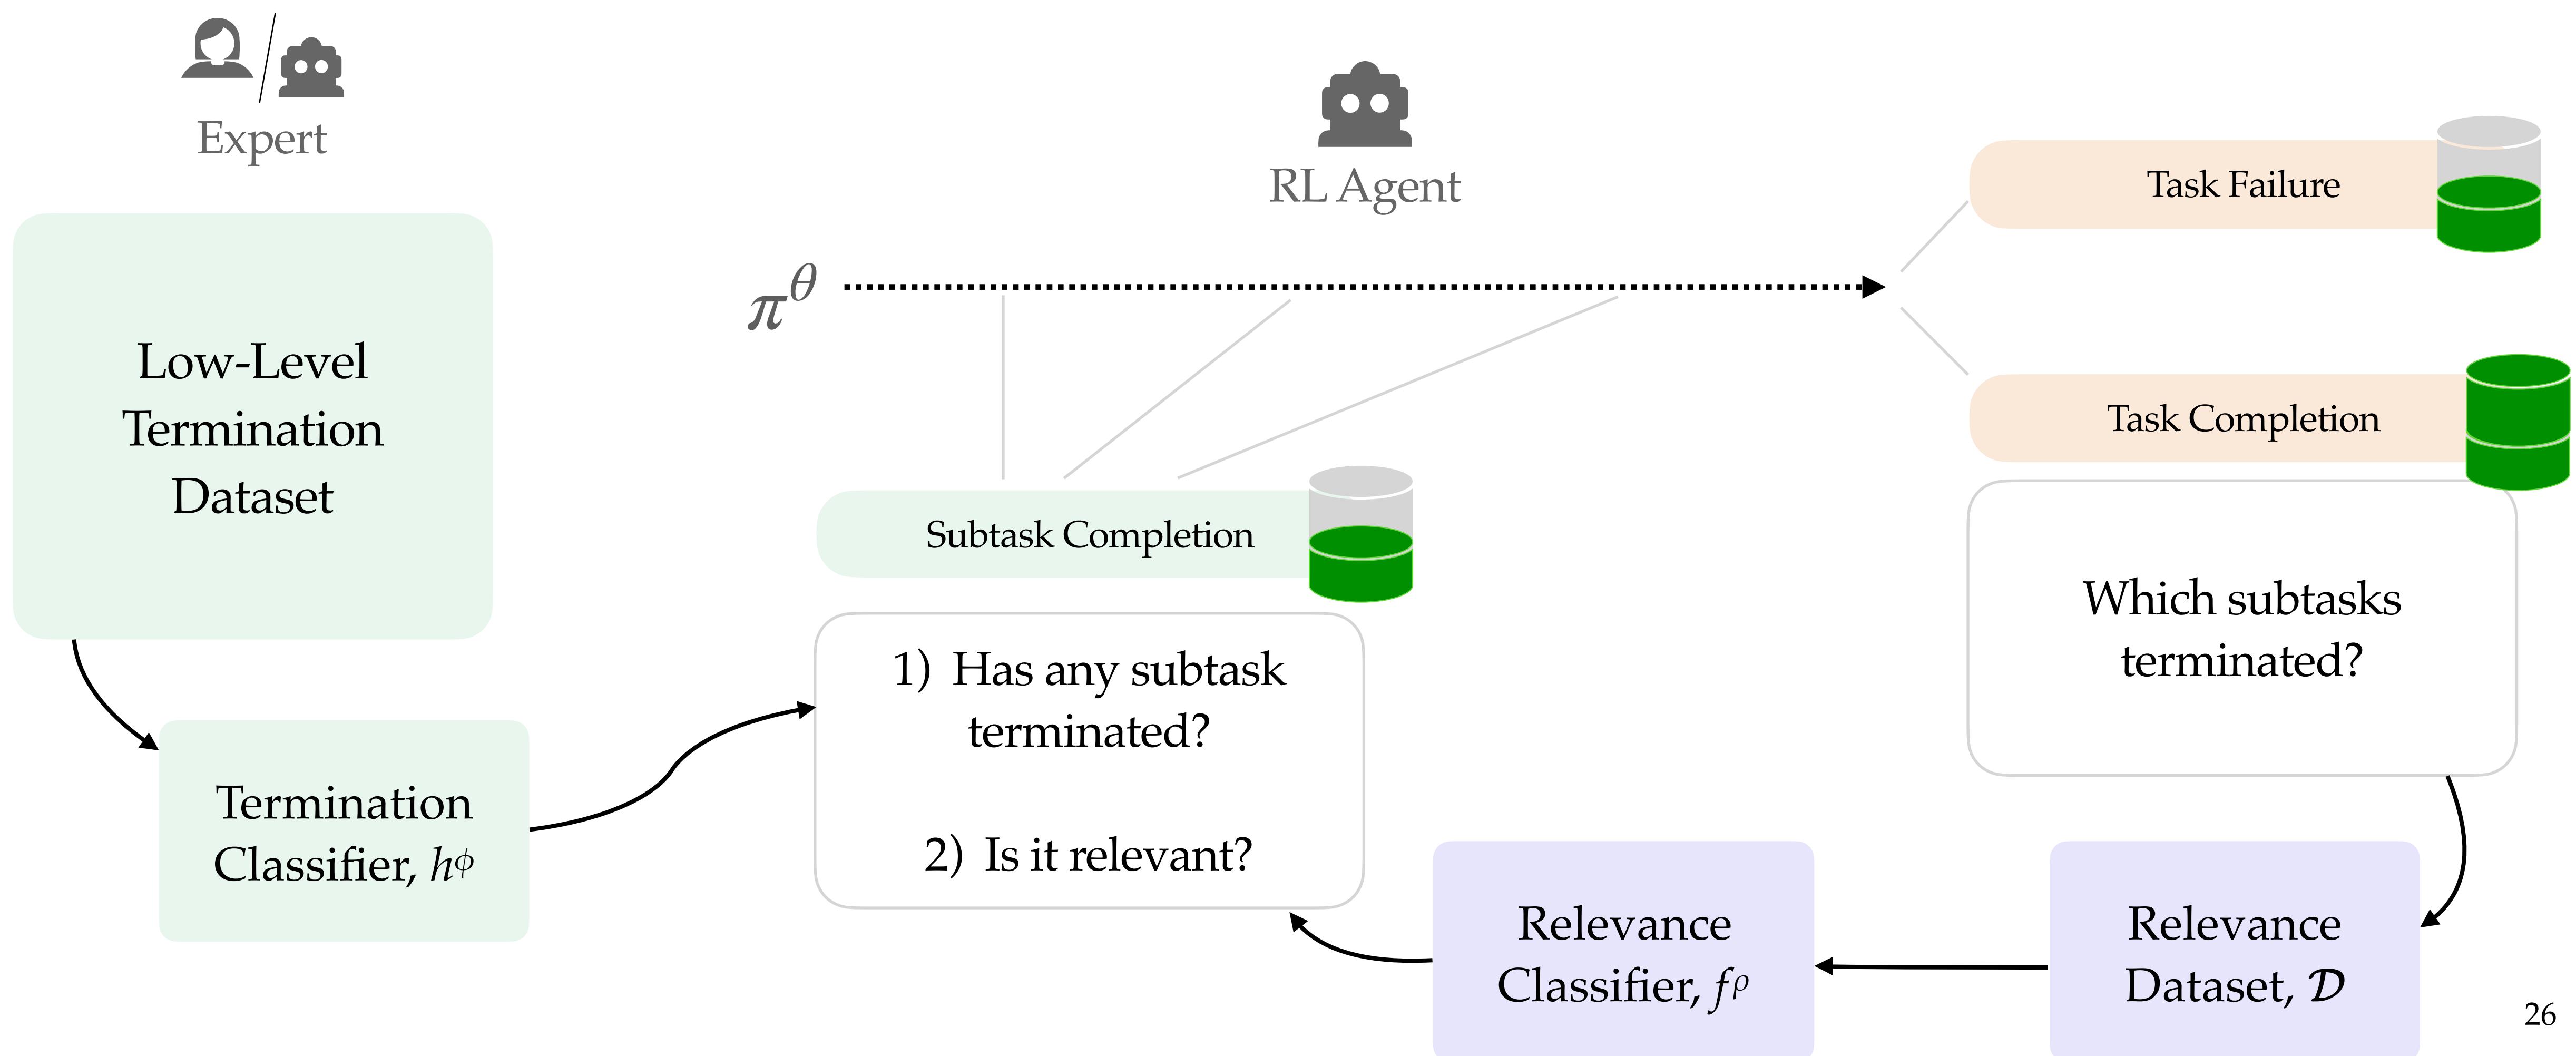
- go to the red ball
- go to the yellow key
- go to the blue ball



ELLA



ELLA



Reward Shaping

Reward Shaping

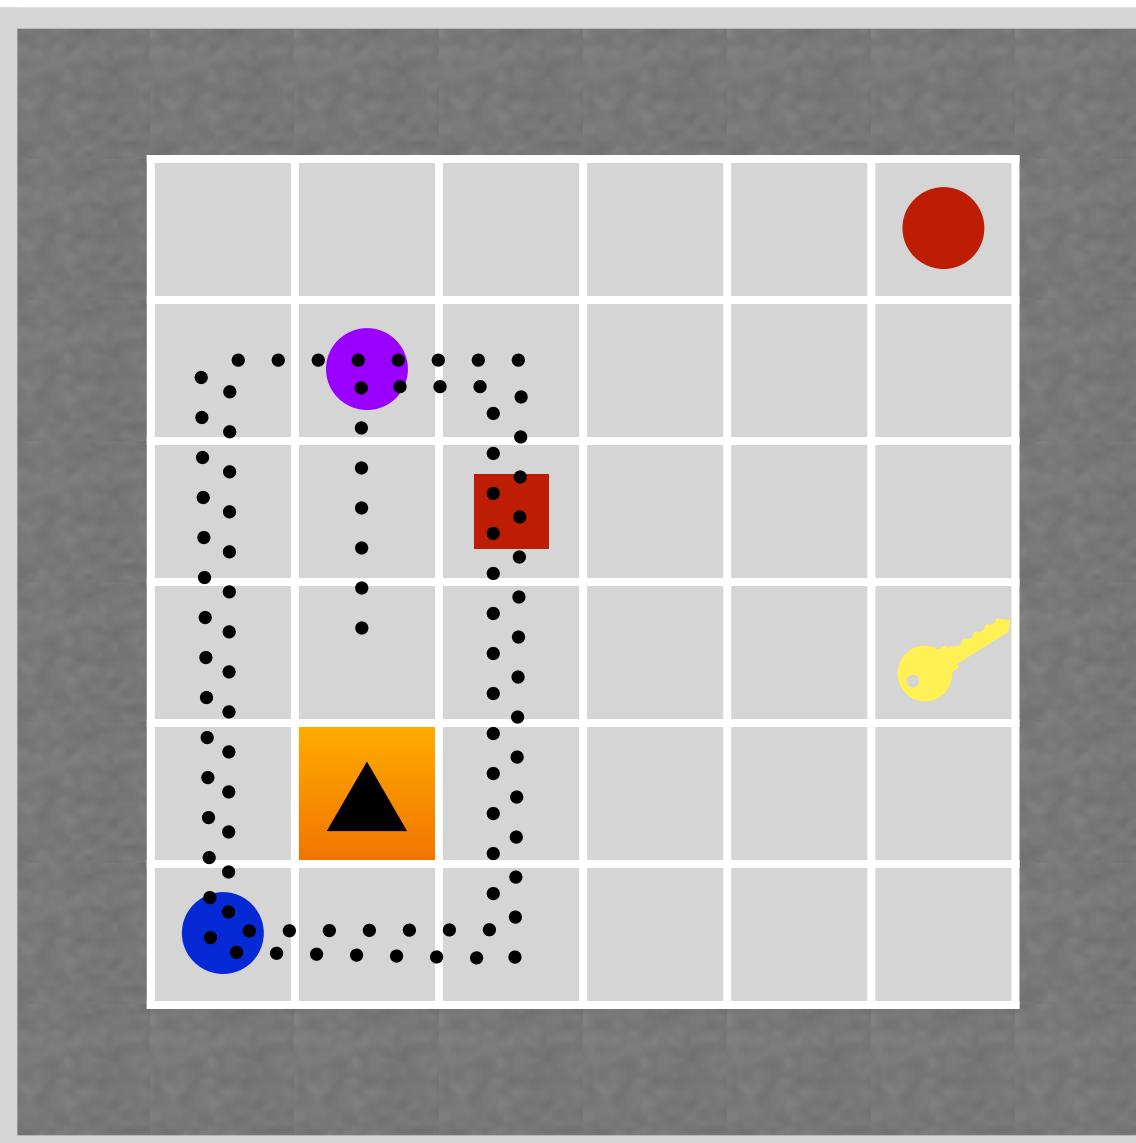
- Bonus λ for relevant low-level tasks

Reward Shaping

- Bonus λ for relevant low-level tasks
- However, we do not want “distraction” by subtasks

Reward Shaping

- Bonus λ for relevant low-level tasks
- However, we do not want “distraction” by subtasks



Reward Shaping

Reward Shaping

- Neutralize intermediate rewards in successful trajectories

Reward Shaping

- Neutralize intermediate rewards in successful trajectories
 - Set $r'_N = r_N - \sum_{t \in T_S} \gamma^{t-N} \lambda$ where T_S is the set of time steps when bonuses were applied

Reward Shaping

- Neutralize intermediate rewards in successful trajectories
 - Set $r'_N = r_N - \sum_{t \in T_S} \gamma^{t-N} \lambda$ where T_S is the set of time steps when bonuses were applied
- Limit return in unsuccessful trajectories by tuning λ

Reward Shaping

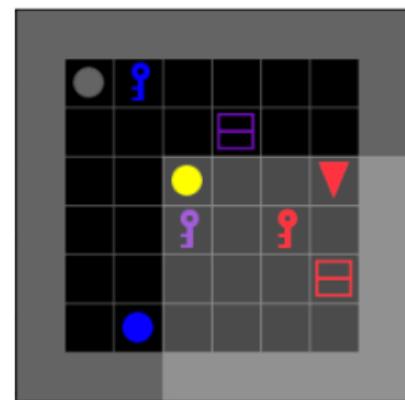
- Neutralize intermediate rewards in successful trajectories
 - Set $r'_N = r_N - \sum_{t \in T_S} \gamma^{t-N} \lambda$ where T_S is the set of time steps when bonuses were applied
- Limit return in unsuccessful trajectories by tuning λ
 - $\lambda < \frac{\gamma^H r_H}{|G_\ell|}$

Experiments

Experiments

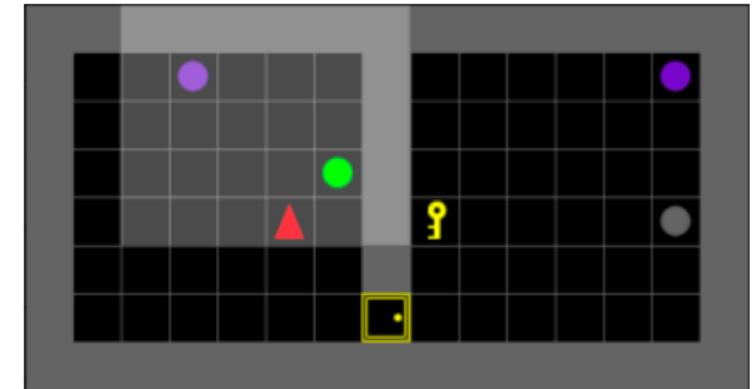
High-Level Tasks

PUTNEXT-ROOM



put the yellow ball
next to a purple key

OPEN&PICK-MAZE



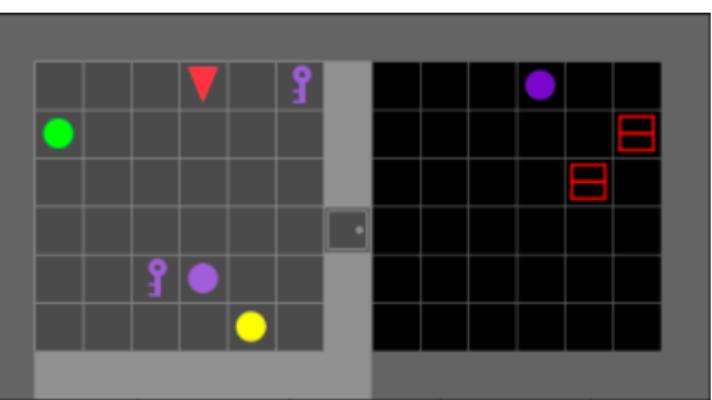
open the yellow door and
pick up the grey ball

COMBO-MAZE



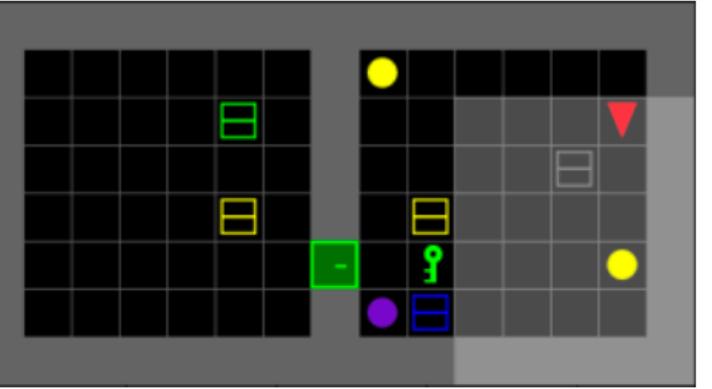
pick up the green ball

PUTNEXT-MAZE



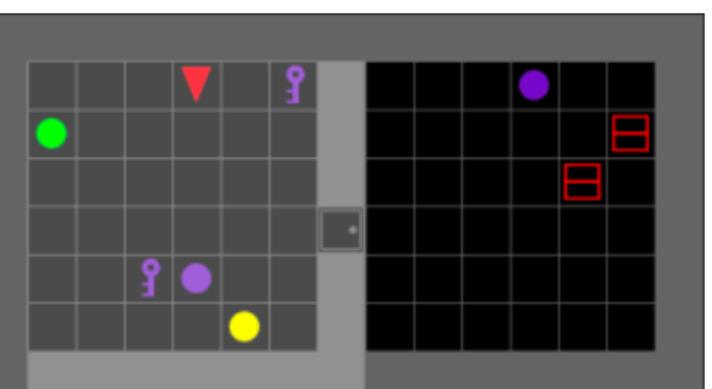
put the blue key next
to the yellow ball

UNLOCK-MAZE



open the green door

SEQUENCE-MAZE

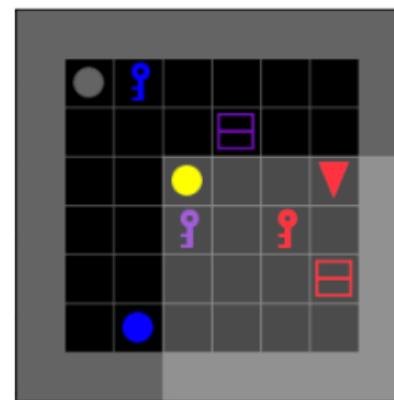


open the grey door after
you put the yellow ball
next to a purple key

Experiments

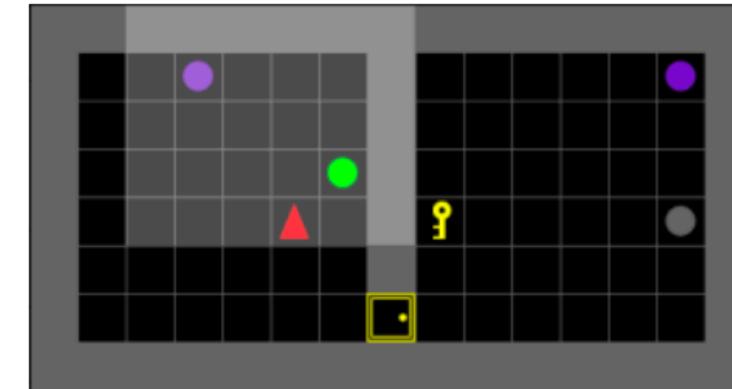
High-Level Tasks

PUTNEXT-ROOM



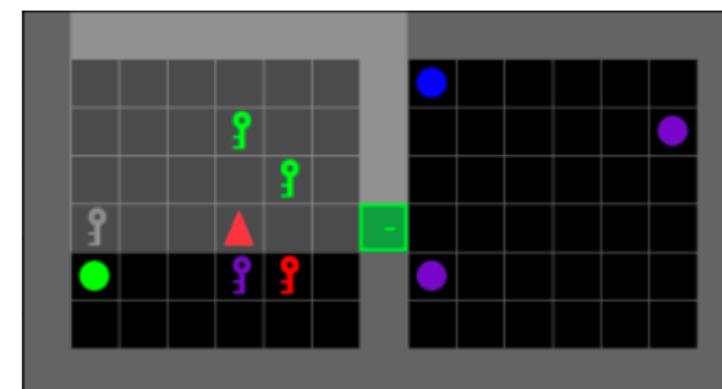
put the yellow ball
next to a purple key

OPEN&PICK-MAZE



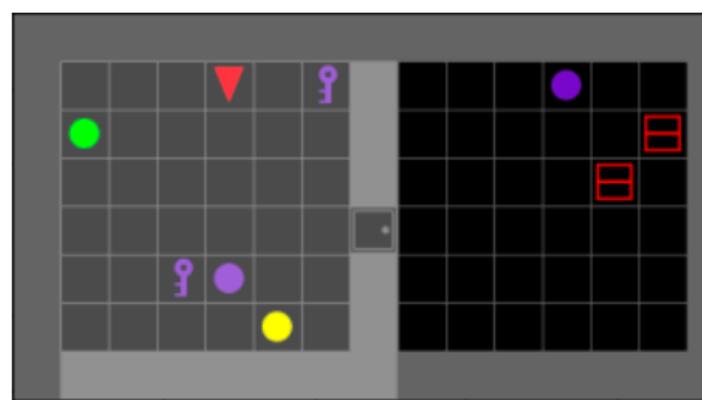
open the yellow door and
pick up the grey ball

COMBO-MAZE



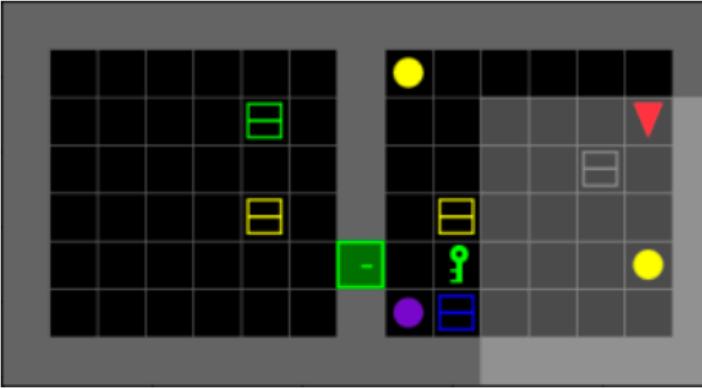
pick up the green ball

PUTNEXT-MAZE



put the blue key next
to the yellow ball

UNLOCK-MAZE



open the green door

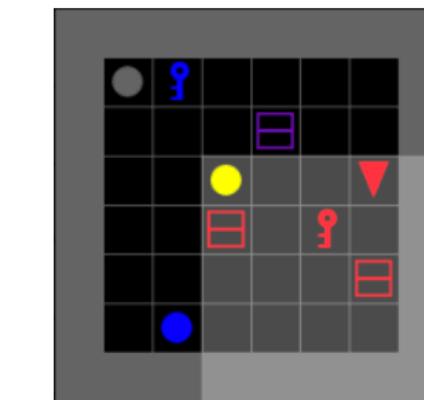
SEQUENCE-MAZE



open the grey door after
you put the yellow ball
next to a purple key

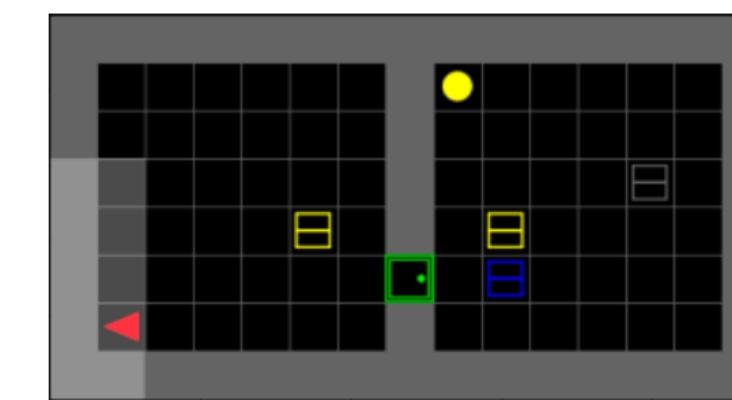
Low-Level Tasks

GoTo-ROOM



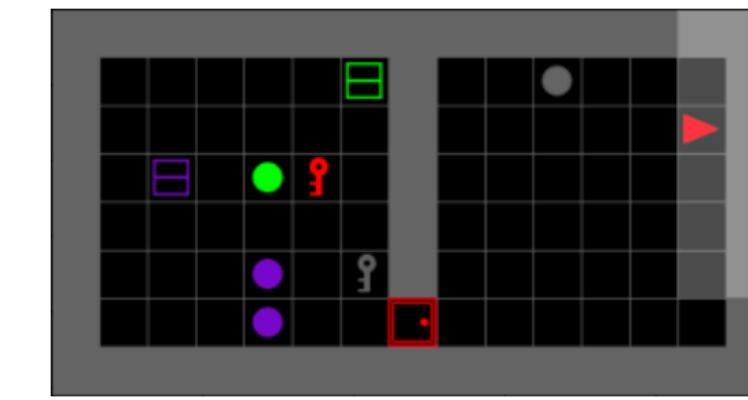
go to a yellow ball

OPEN-MAZE



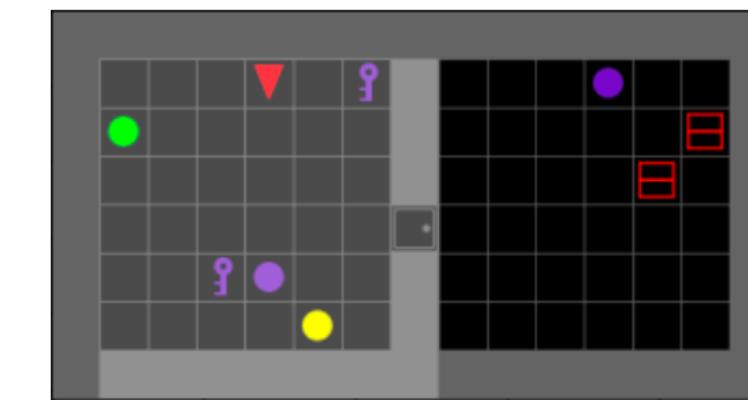
open the green door

GoTo-MAZE



go to a red key

PICK-MAZE

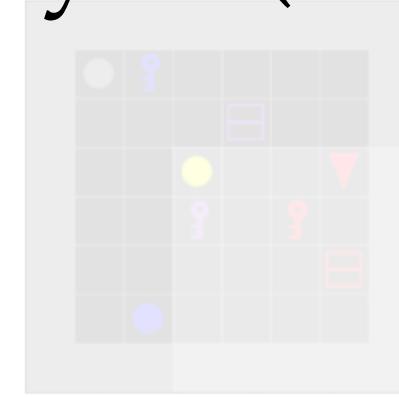


pick up a red box

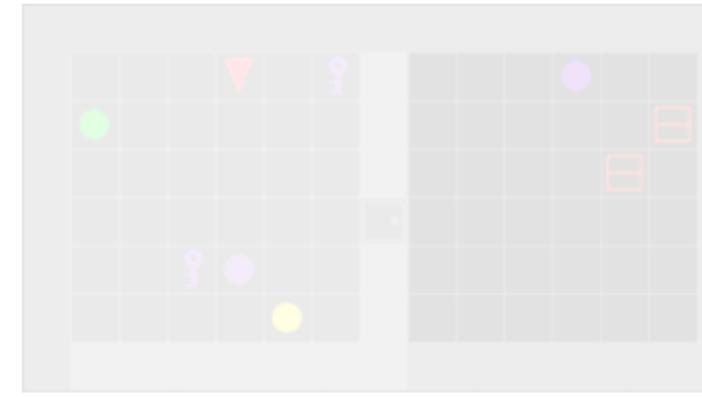
Experiments

High-Level Tasks

- BabyAI (Chevalier-Boisvert et al. 2018)

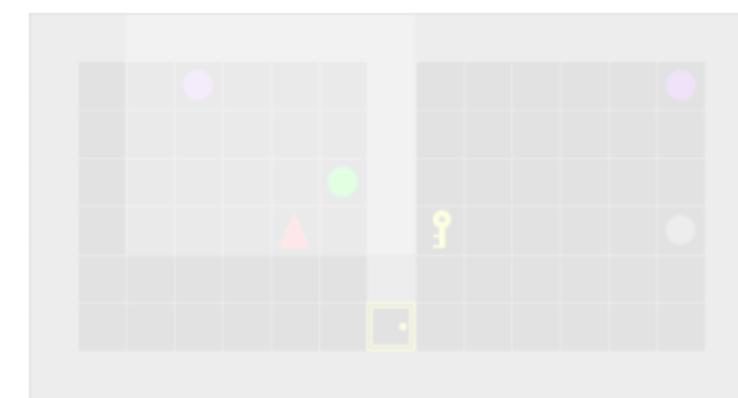


put the yellow ball
next to a purple key



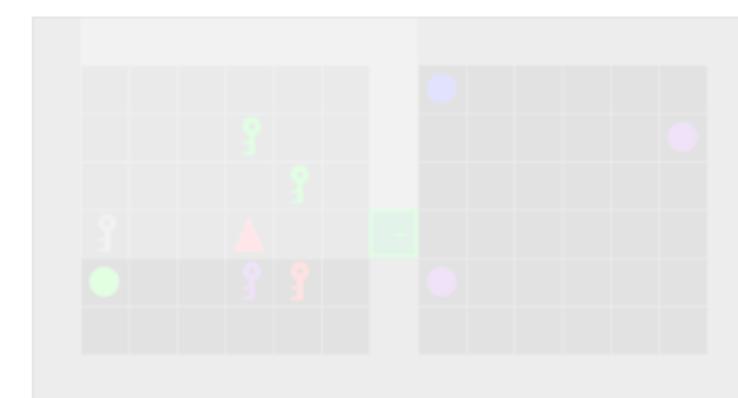
put the blue key next
to the yellow ball

OPEN&PICK-MAZE



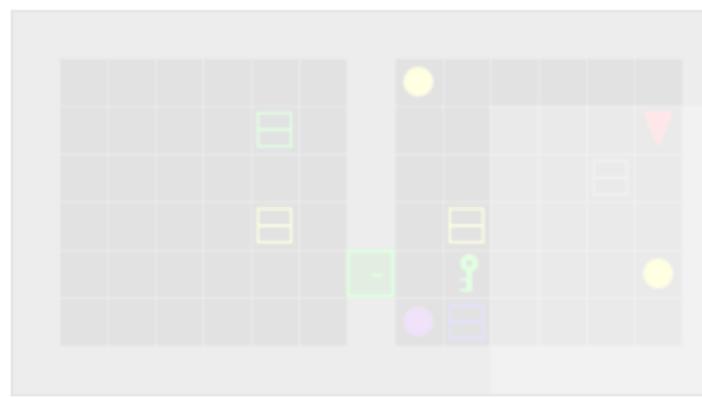
open the yellow door and
pick up the grey ball

COMBO-MAZE



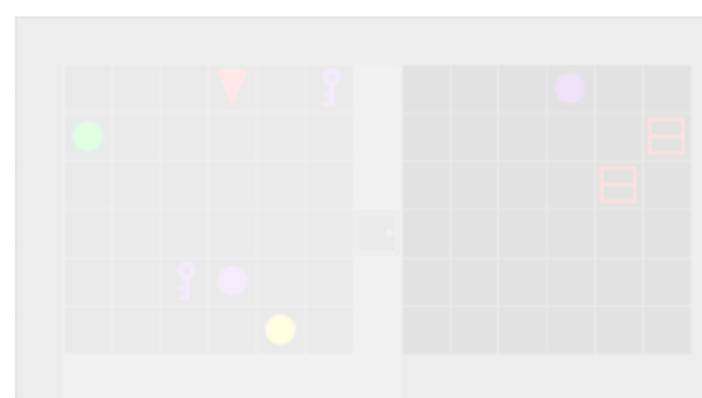
pick up the green ball

UNLOCK-MAZE



open the green door

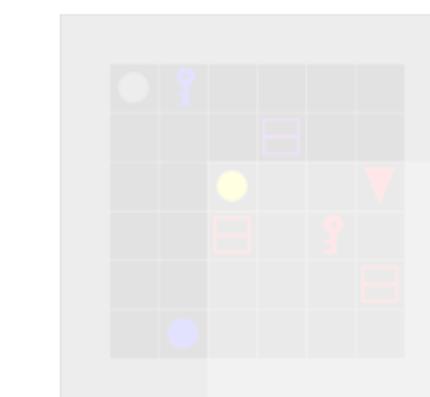
SEQUENCE-MAZE



open the grey door after
you put the yellow ball
next to a purple key

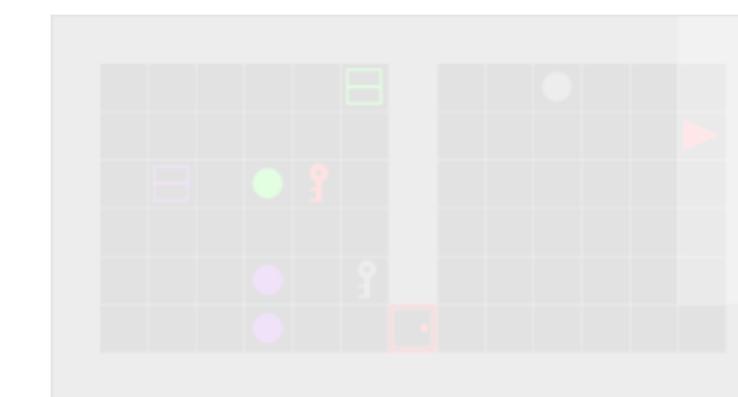
Low-Level Tasks

GoTo-Room



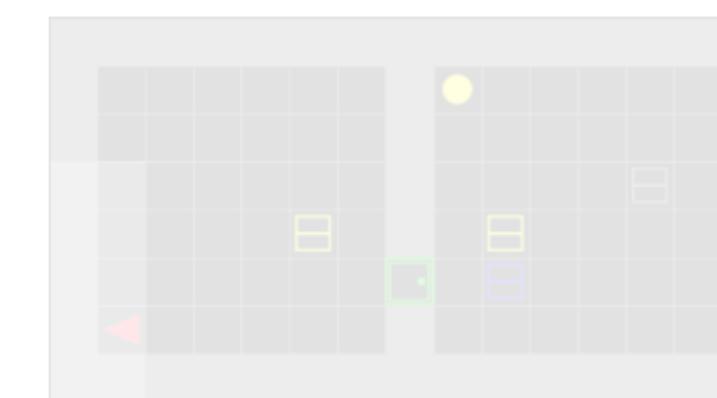
go to a yellow ball

GoTo-MAZE



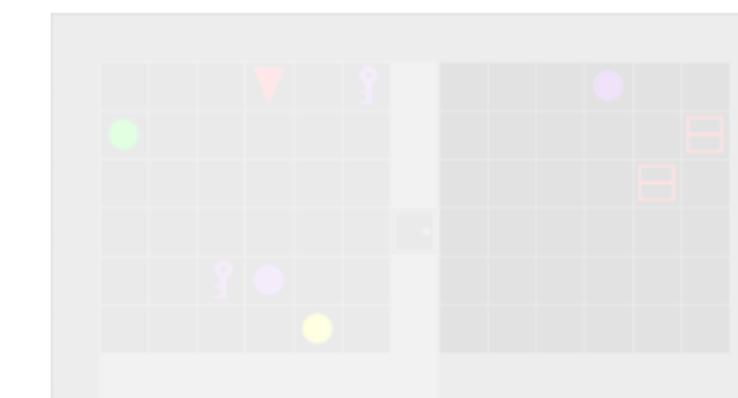
go to a red key

OPEN-MAZE



open the green door

PICK-MAZE



pick up a red box

Experiments

High-Level Tasks

- BabyAI (Chevalier-Boisvert et al. 2018)

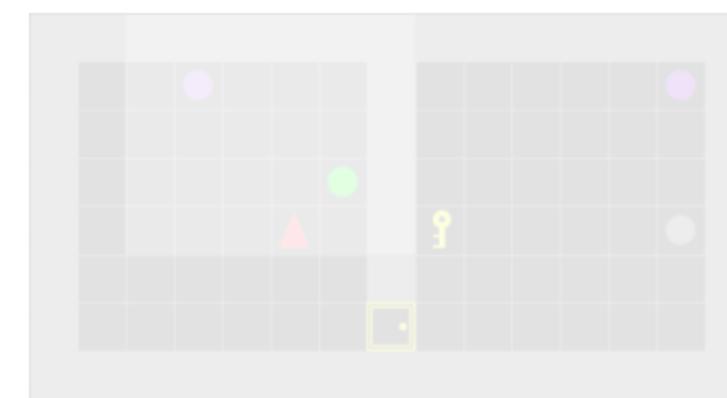


put the yellow ball
next to a purple key



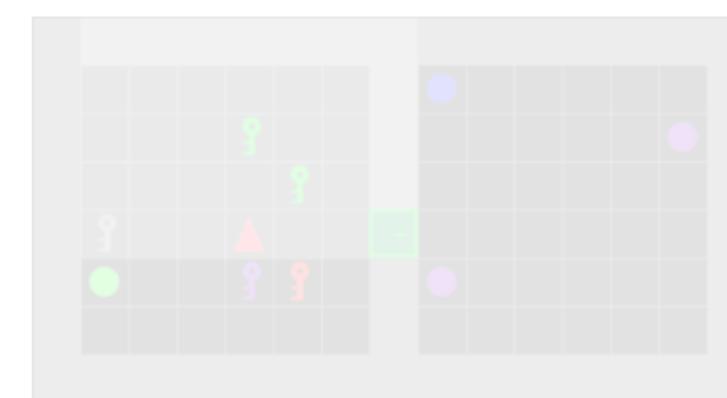
put the blue key next
to the yellow ball

OPEN&PICK-MAZE



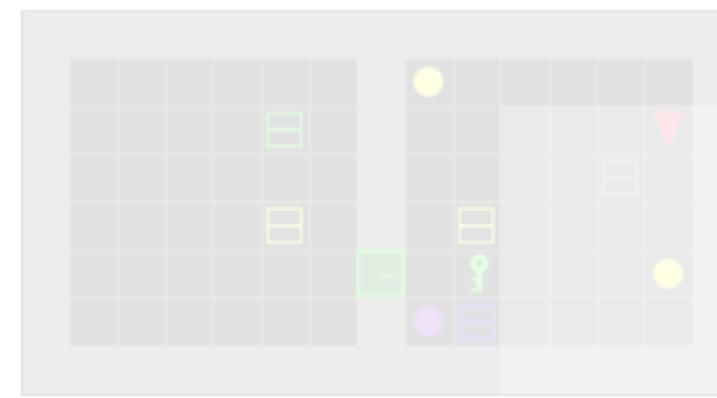
open the yellow door and
pick up the grey ball

COMBO-MAZE



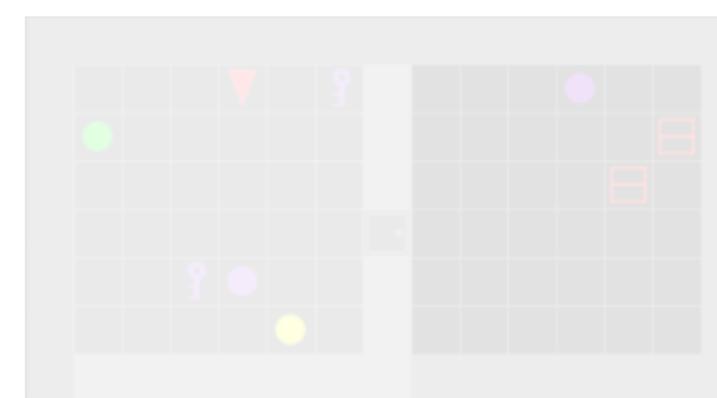
pick up the green ball

UNLOCK-MAZE



open the green door

SEQUENCE-MAZE



open the grey door after
you put the yellow ball
next to a purple key

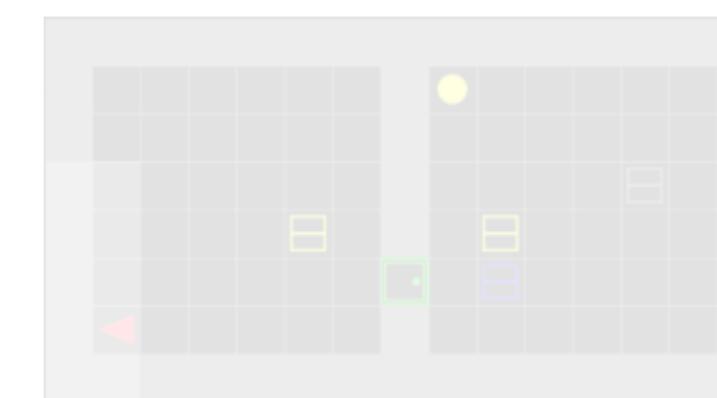
Low-Level Tasks

GoTo-Room



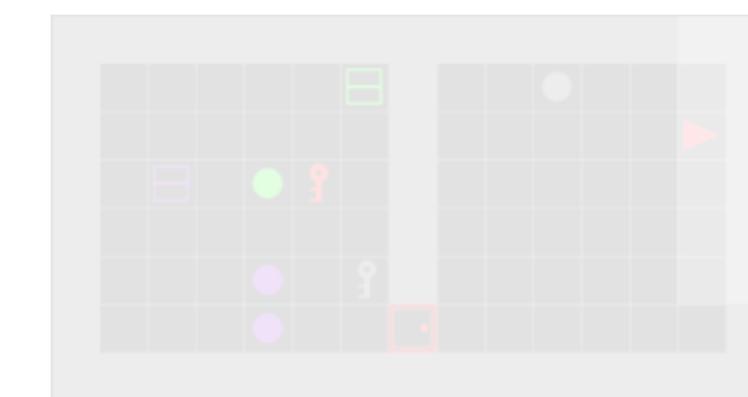
go to a yellow ball

OPEN-MAZE



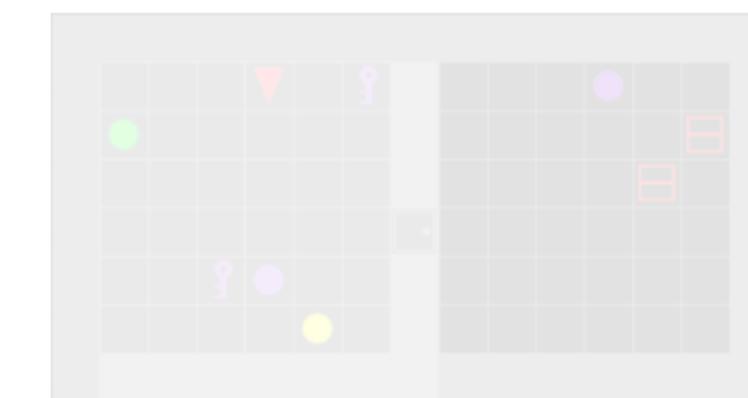
open the green door

GoTo-MAZE



go to a red key

PICK-MAZE



pick up a red box

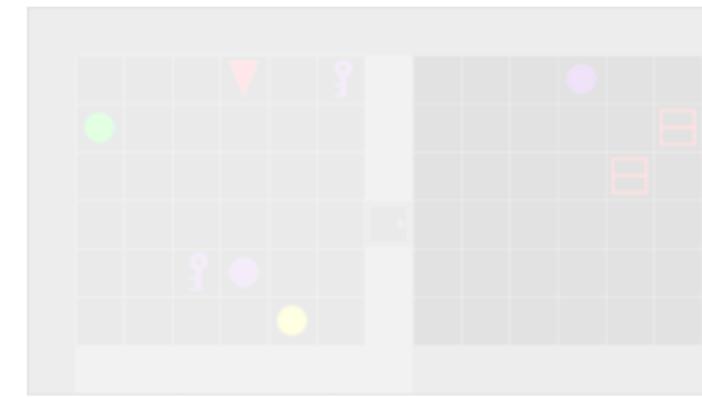
Experiments

High-Level Tasks

- BabyAI (Chevalier-Boisvert et al. 2018)



- Partial observability

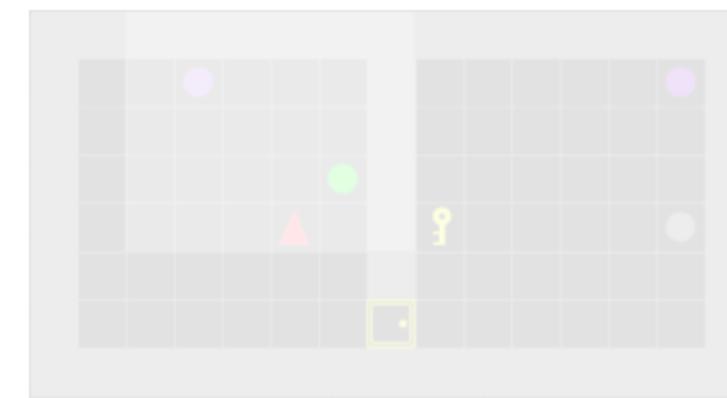


- Distractor objects

next to a purple key

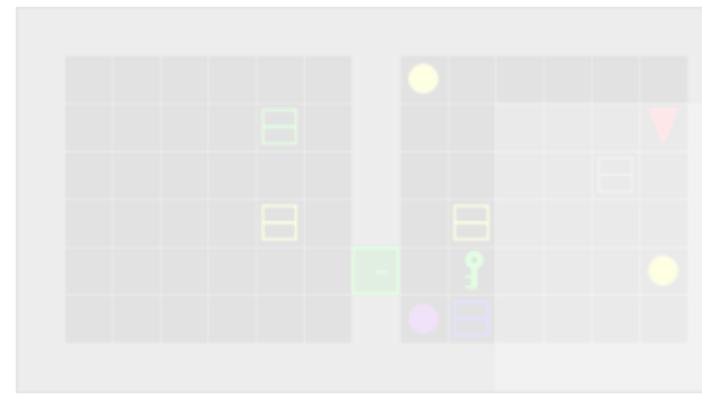
put the blue key next
to the yellow ball

OPEN&PICK-MAZE



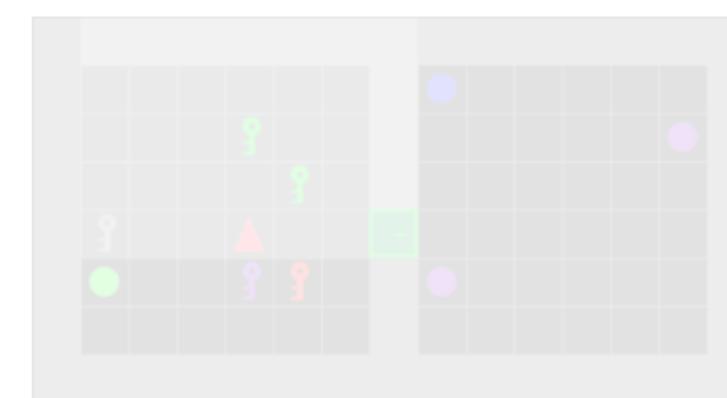
open the yellow door and
pick up the grey ball

UNLOCK-MAZE



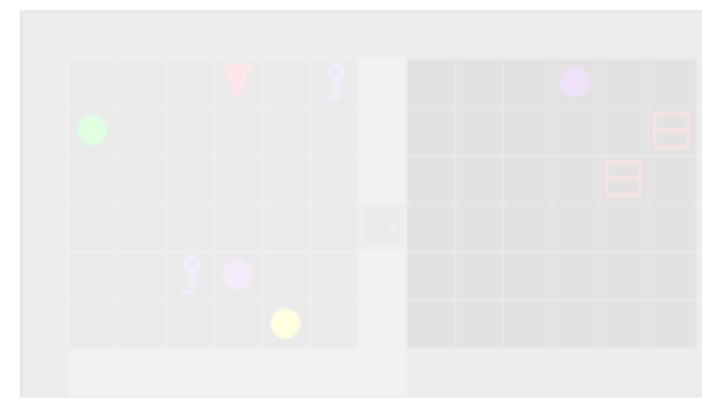
open the green door

COMBO-MAZE



pick up the green ball

SEQUENCE-MAZE



open the grey door after
you put the yellow ball
next to a purple key

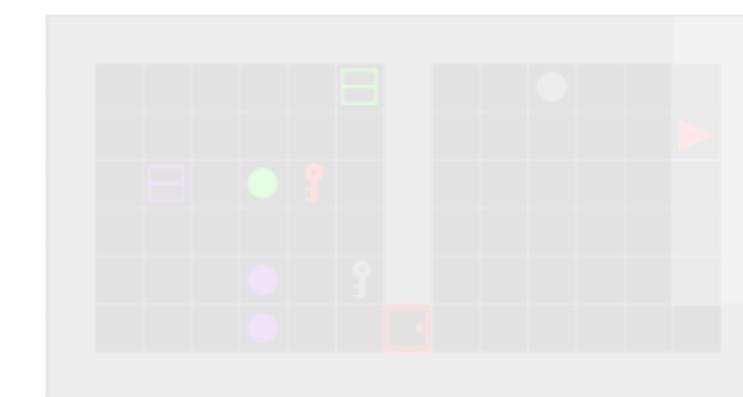
Low-Level Tasks

GoTo-Room



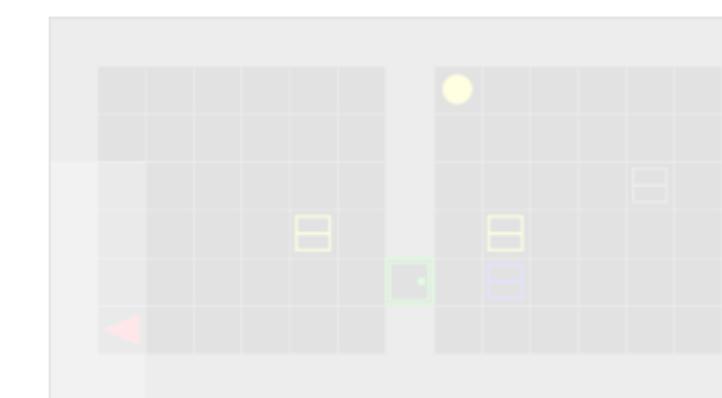
go to a yellow ball

GoTo-MAZE



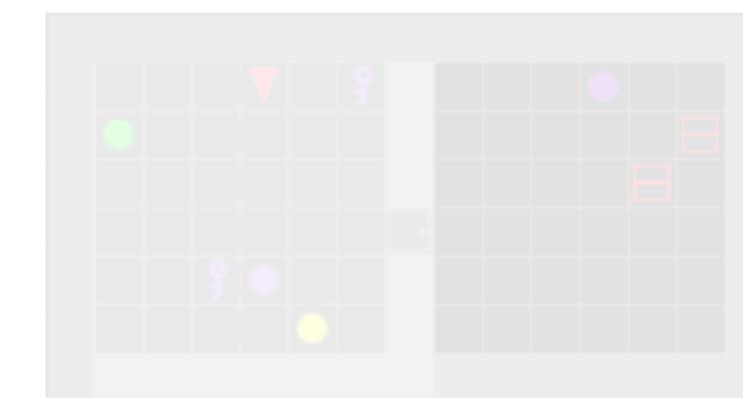
go to a red key

OPEN-MAZE



open the green door

PICK-MAZE

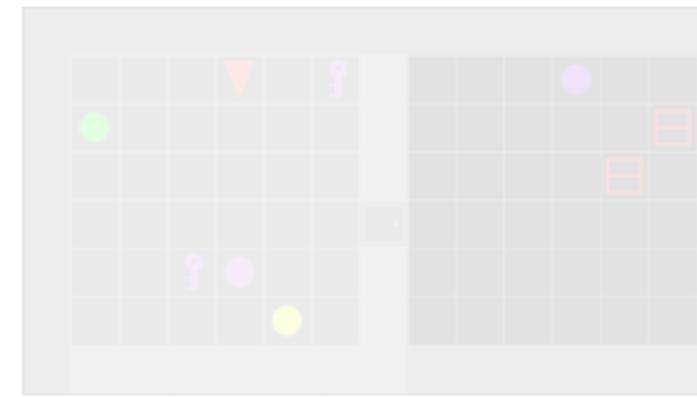


pick up a red box

Experiments

High-Level Tasks

- BabyAI (Chevalier-Boisvert et al. 2018)



- Partial observability

put the blue key next
to the yellow ball

- Distractor objects
next to a purple key

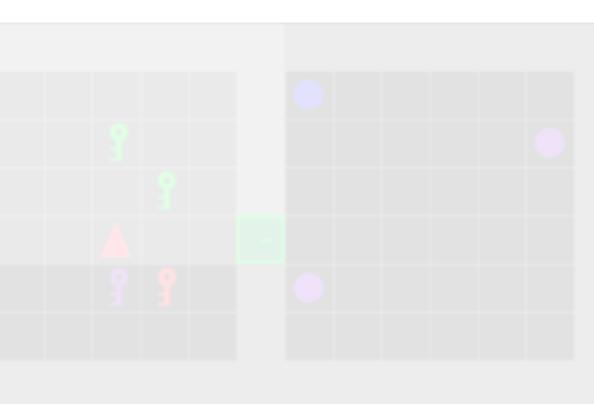


- ROOM levels: single 7×7 grid

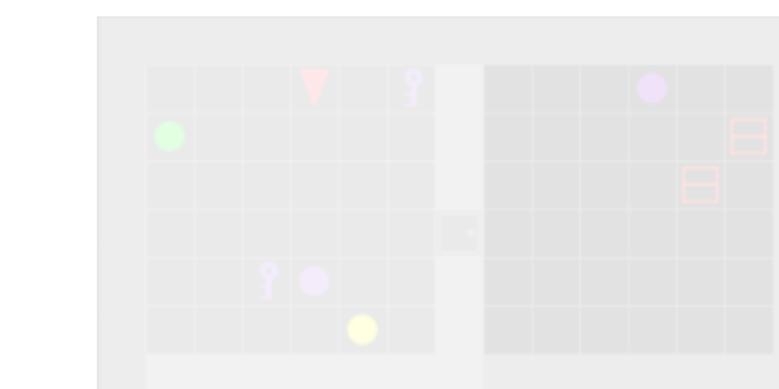
open the yellow door and
pick up the grey ball



open the green door



COMBO-MAZE

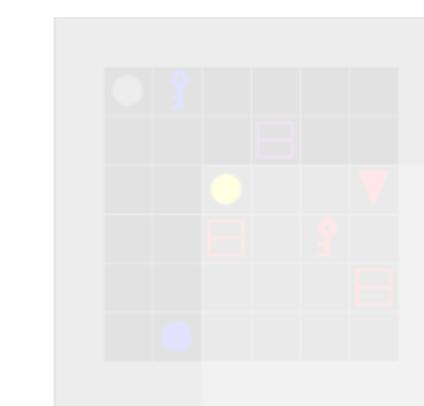


pick up the green ball

open the grey door after
you put the yellow ball
next to a purple key

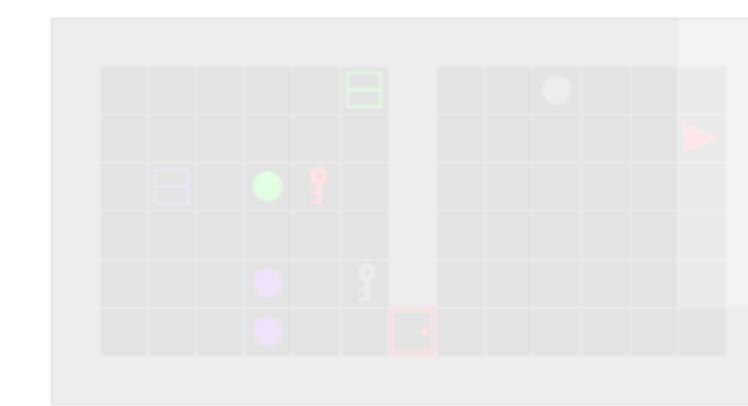
Low-Level Tasks

GoTo-Room



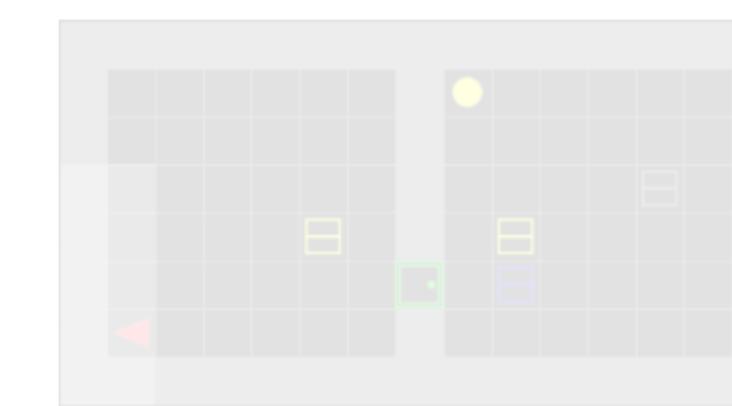
go to a yellow ball

GoTo-MAZE



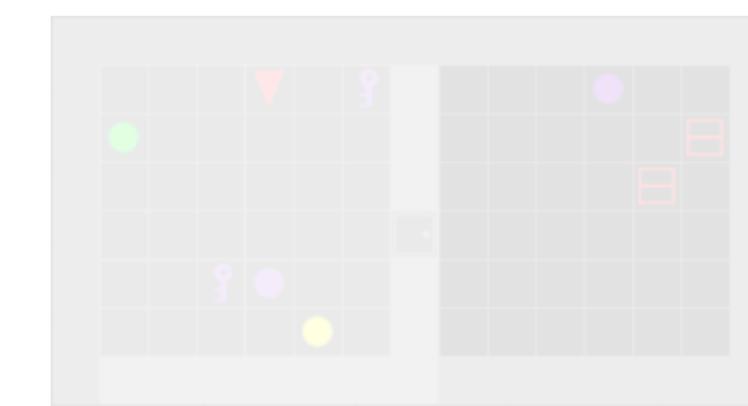
go to a red key

OPEN-MAZE



open the green door

PICK-MAZE



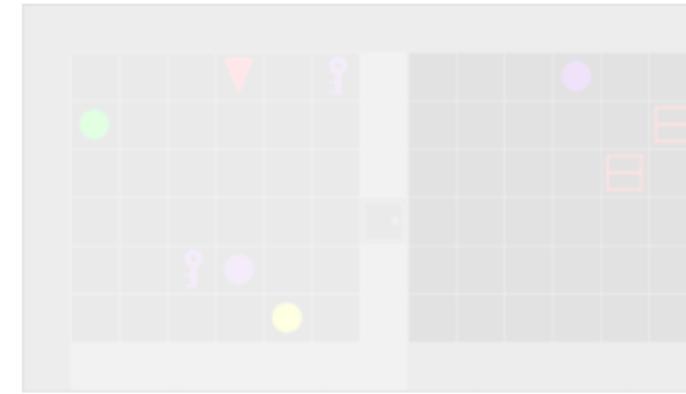
pick up a red box

SEQUENCE-MAZE

Experiments

High-Level Tasks

- BabyAI (Chevalier-Boisvert et al. 2018)



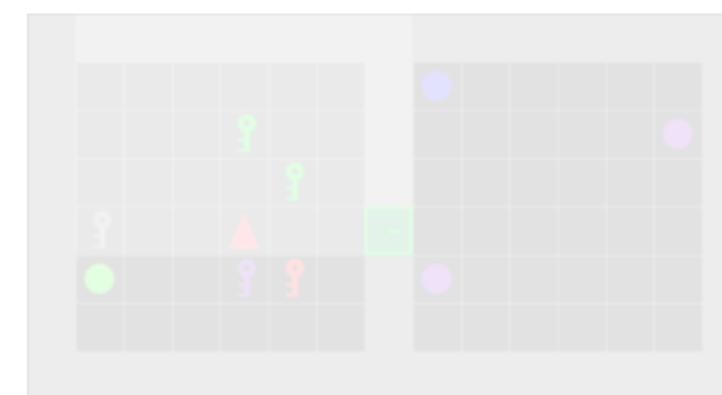
- Partial observability

put the blue key next
to the yellow ball

- Distractor objects
next to a purple key

open the yellow door and
pick up the grey ball

COMBO-MAZE



pick up the green ball

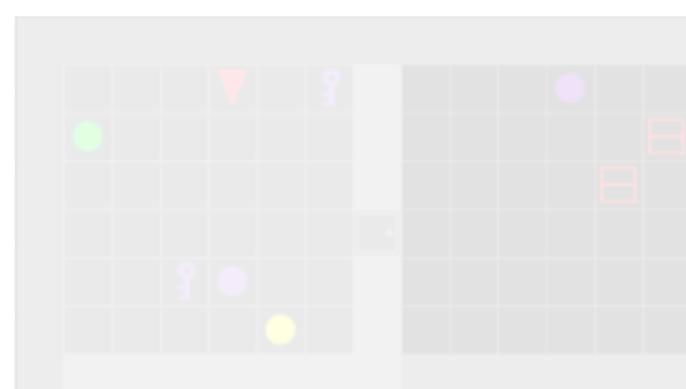
- ROOM levels: single 7×7 grid

UNLOCK-MAZE



open the green door

SEQUENCE-MAZE



open the grey door after
you put the yellow ball
next to a purple key

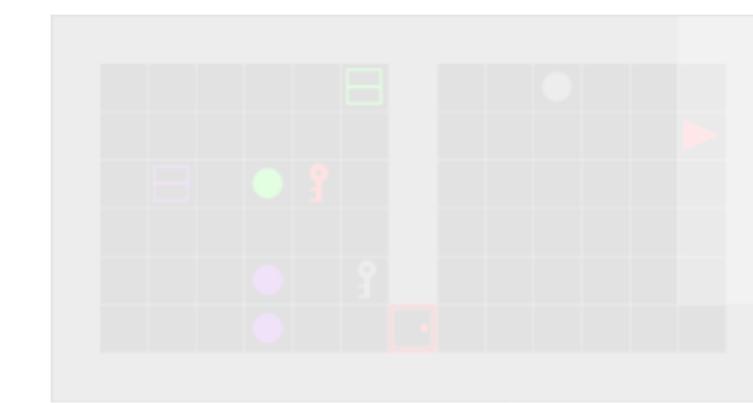
Low-Level Tasks

GoTo-Room



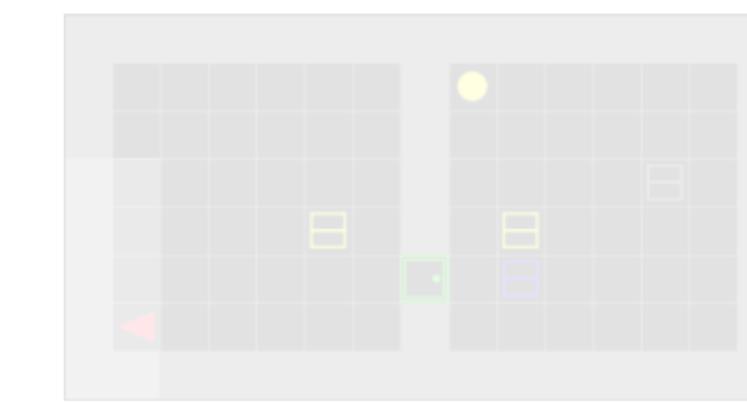
go to a yellow ball

GoTo-MAZE



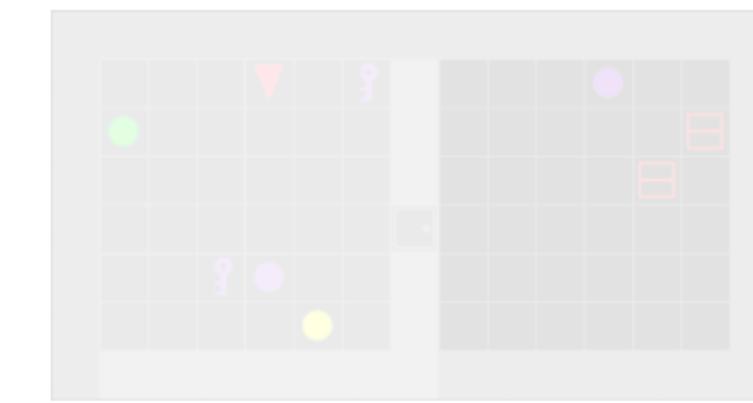
go to a red key

OPEN-MAZE



open the green door

PICK-MAZE



pick up a red box

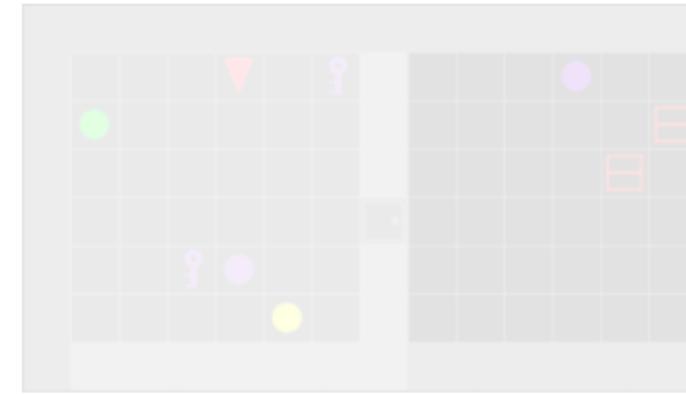
Experiments

High-Level Tasks

- BabyAI (Chevalier-Boisvert et al. 2018)



- Partial observability



- Distractor objects

put the blue key next
next to a purple key

to the yellow ball

- ROOM levels: single 7×7 grid

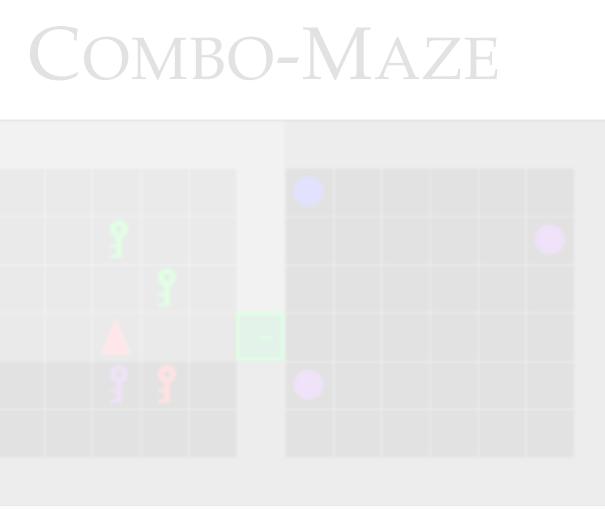
open the yellow door and
pick up the grey ball

UNLOCK-MAZE



- MAZE levels: two rooms connected by
a closed/locked door

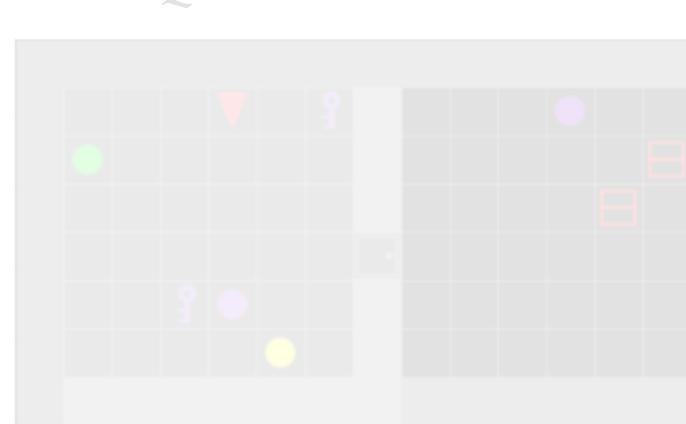
open the green door



COMBO-MAZE

pick up the green ball

SEQUENCE-MAZE

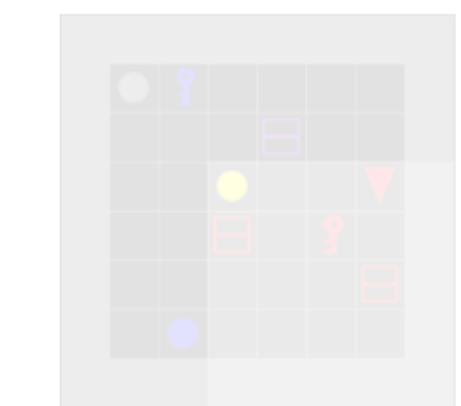


open the grey door after
you put the yellow ball
next to a purple key

Low-Level Tasks

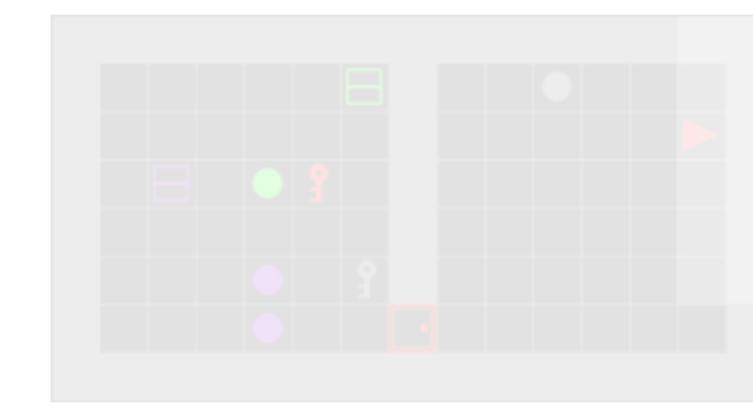
- Tasks differ on several axes

GoTo-Room



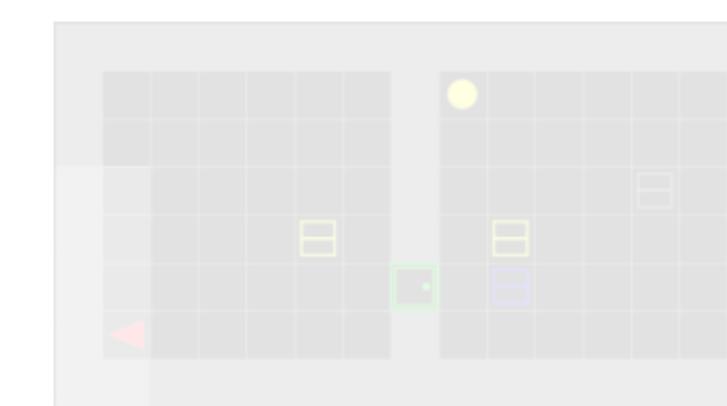
go to a yellow ball

GoTo-MAZE



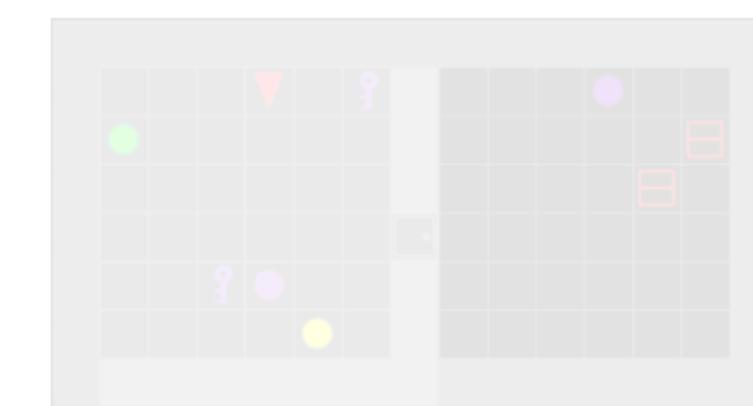
go to a red key

OPEN-MAZE



open the green door

PICK-MAZE



pick up a red box

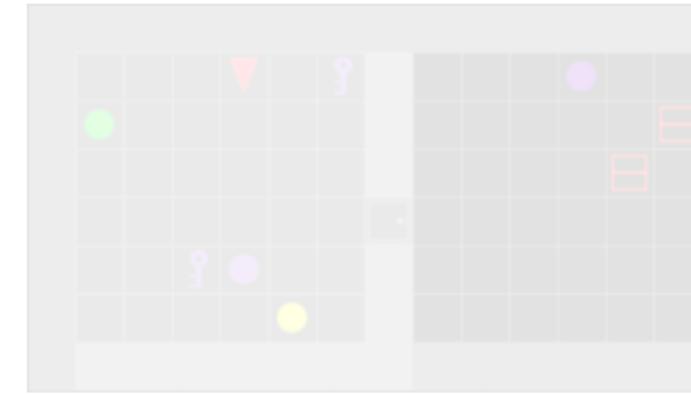
Experiments

High-Level Tasks

- BabyAI (Chevalier-Boisvert et al. 2018)



- Partial observability



- Distractor objects

next to a purple key

put the blue key next
to the yellow ball

- ROOM levels: single 7×7 grid

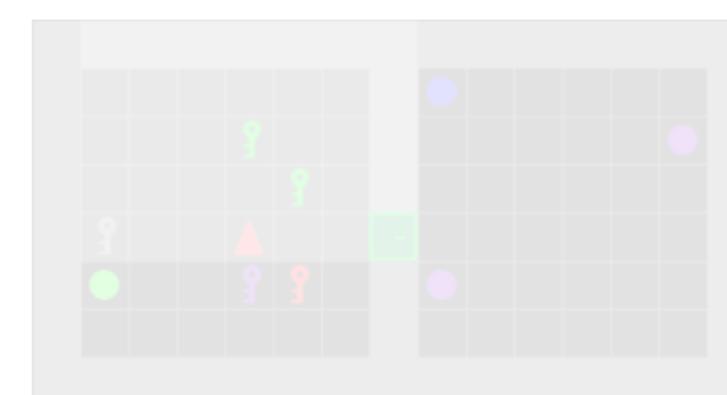
open the yellow door and
pick up the grey ball

UNLOCK-MAZE



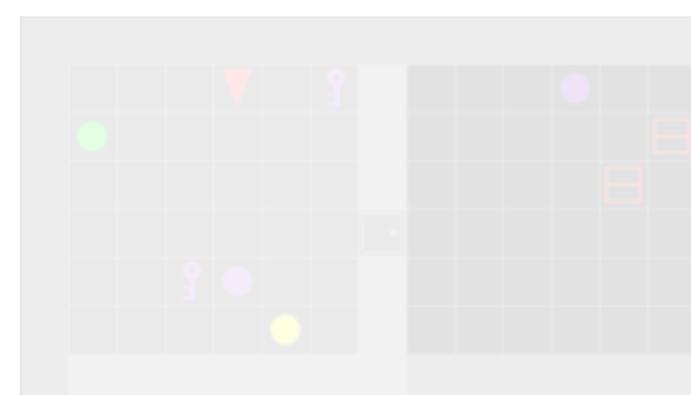
open the green door

COMBO-MAZE



pick up the green ball

SEQUENCE-MAZE



open the grey door after
you put the yellow ball
next to a purple key

Low-Level Tasks

- Tasks differ on several axes

GoTo-Room



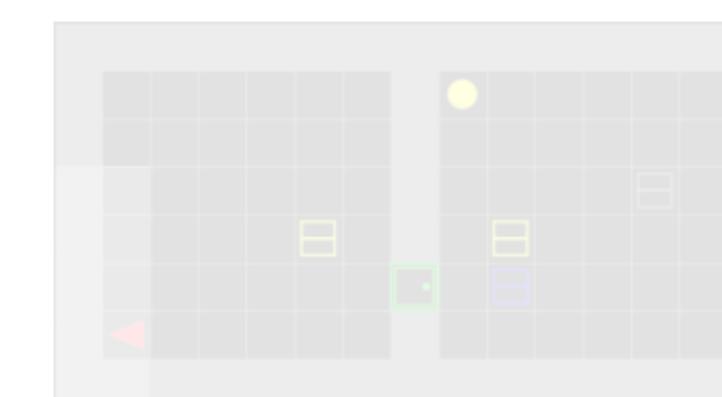
go to a yellow ball

GoTo-MAZE



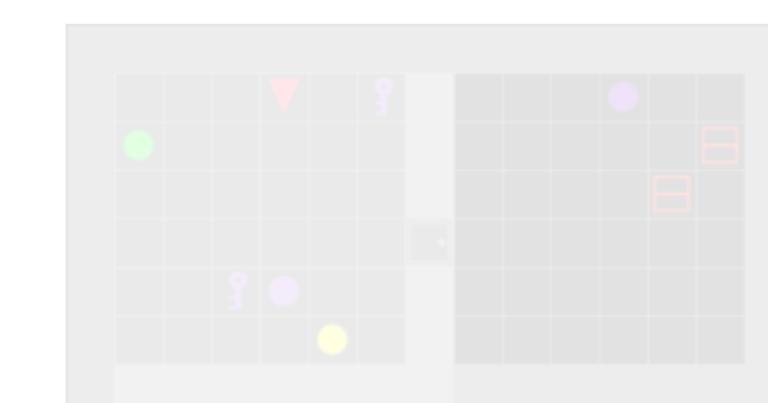
go to a red key

OPEN-MAZE



open the green door

PICK-MAZE



pick up a red box

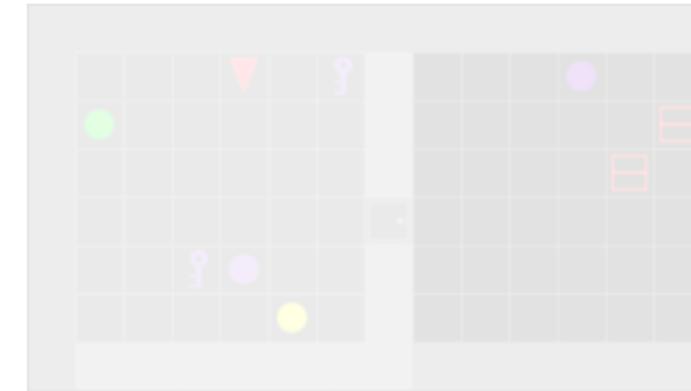
Experiments

High-Level Tasks

- BabyAI (Chevalier-Boisvert et al. 2018)



- Partial observability

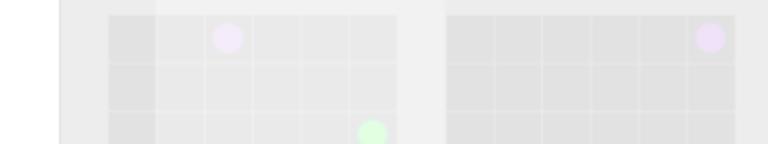


- Distractor objects

next to a purple key

put the blue key next
to the yellow ball

- ROOM levels: single 7×7 grid



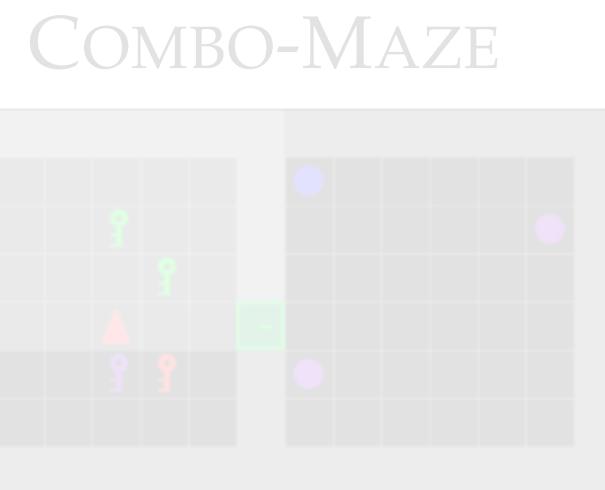
UNLOCK-MAZE



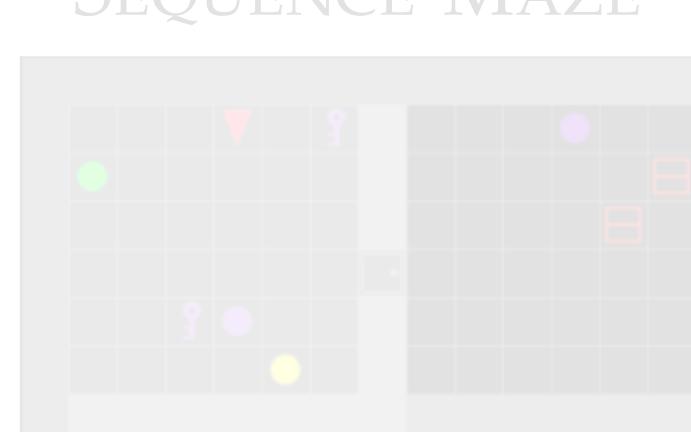
- MAZE levels: two rooms connected by a closed/locked door

open the yellow door and
pick up the grey ball

open the green door



COMBO-MAZE



SEQUENCE-MAZE

open the grey door after
you put the yellow ball
next to a purple key

pick up the green ball

Low-Level Tasks

- Tasks differ on several axes

GoTo-Room

- *Sparsity* of the high-level task

GoTo-MAZE

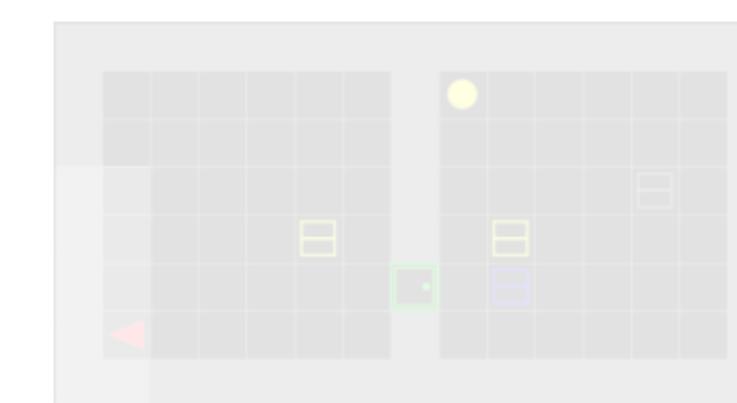


- *Similarity* of the low- and high-level tasks

go to a yellow ball

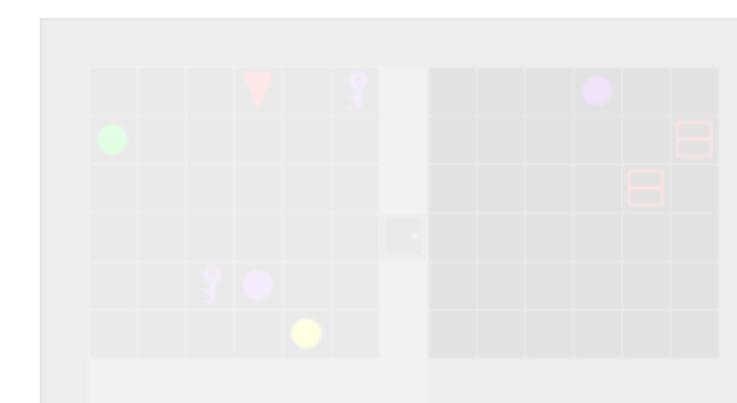
go to a red key

OPEN-MAZE



open the green door

PICK-MAZE



pick up a red box

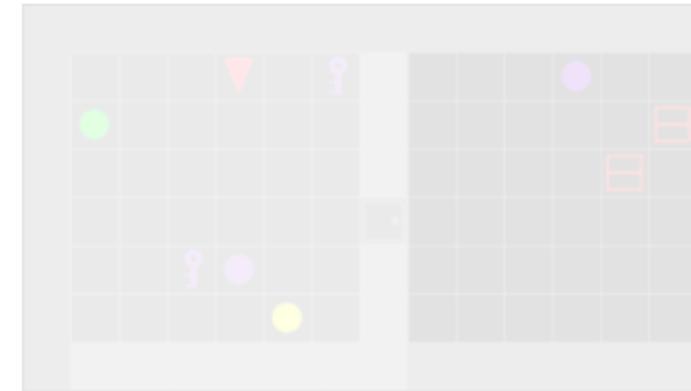
Experiments

High-Level Tasks

- BabyAI (Chevalier-Boisvert et al. 2018)



- Partial observability



- Distractor objects

next to a purple key

put the blue key next
to the yellow ball

- ROOM levels: single 7×7 grid

open the yellow door and
pick up the grey ball

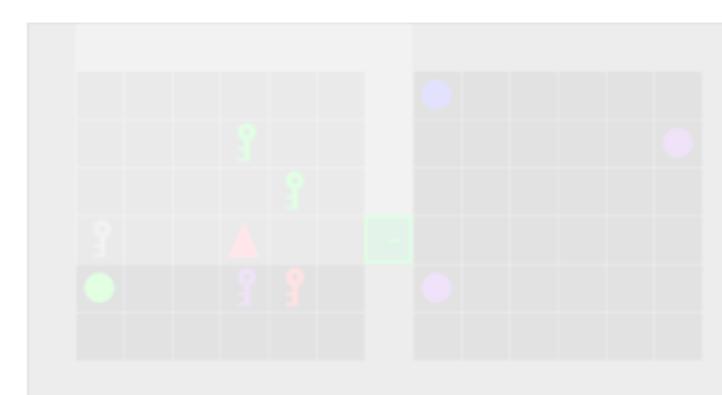
UNLOCK-MAZE



- MAZE levels: two rooms connected by
a closed/locked door

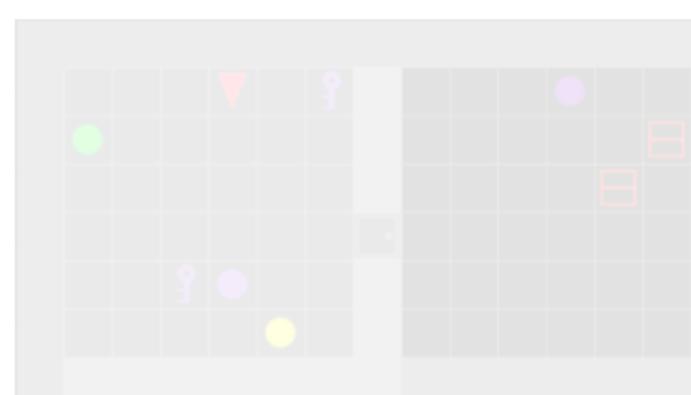
open the green door

COMBO-MAZE



pick up the green ball

SEQUENCE-MAZE



open the grey door after
you put the yellow ball
next to a purple key

Low-Level Tasks

- Tasks differ on several axes

GoTo-Room

- *Sparsity* of the high-level task

GoTo-MAZE

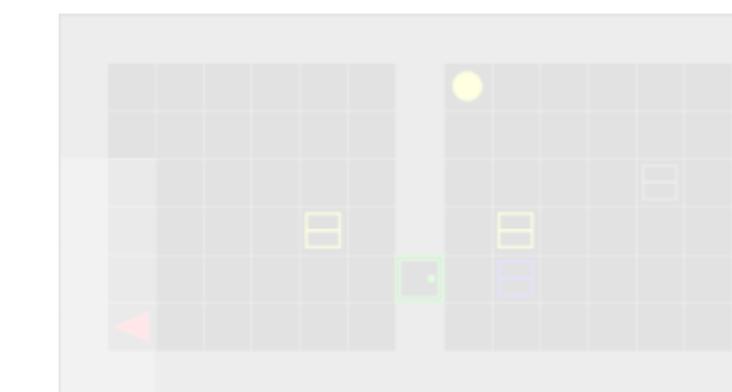


- *Similarity* of the low- and high-level tasks

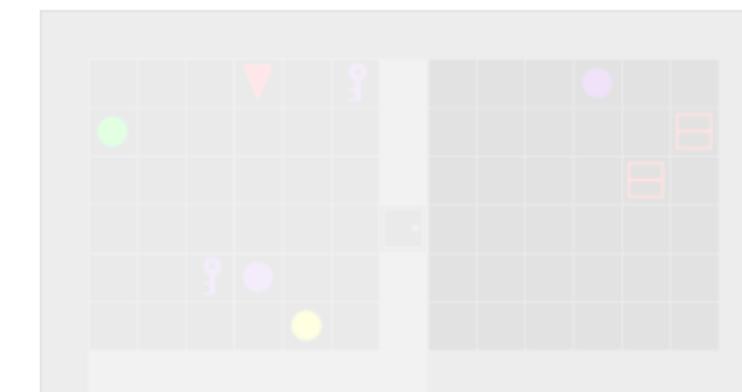
go to a yellow ball

go to a red key

- *Compositionality* of the tasks in G



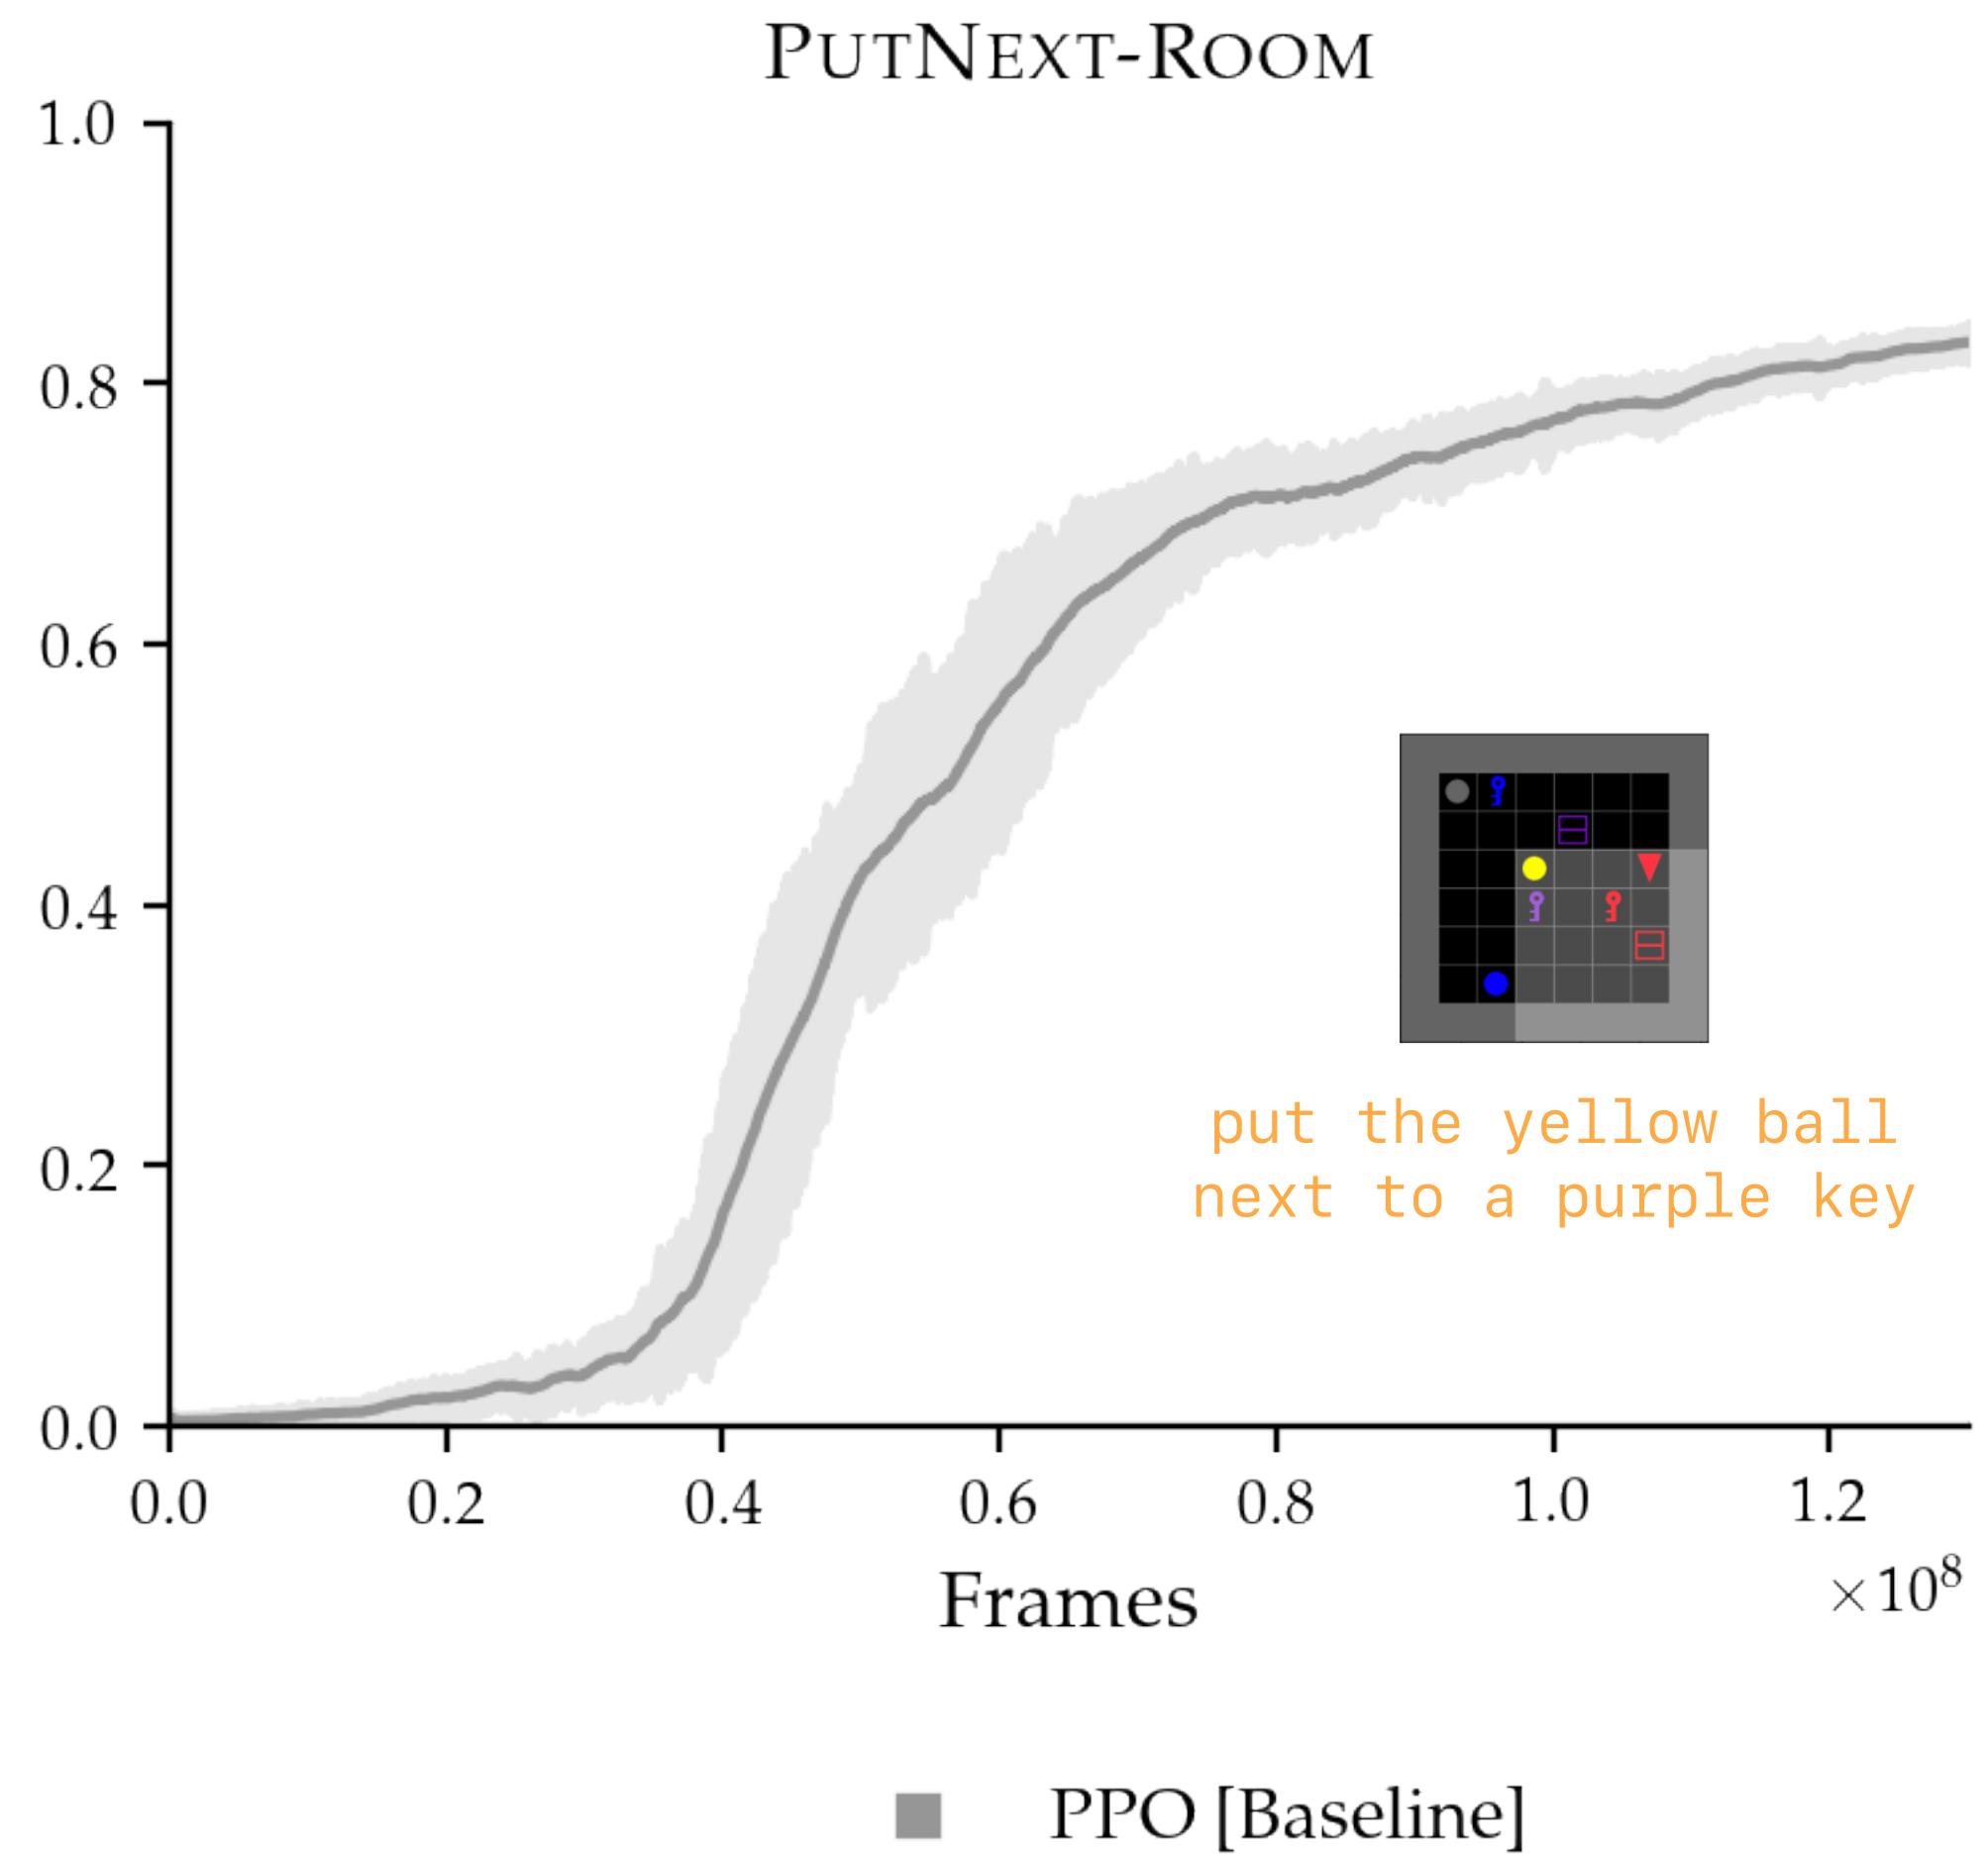
open the green door



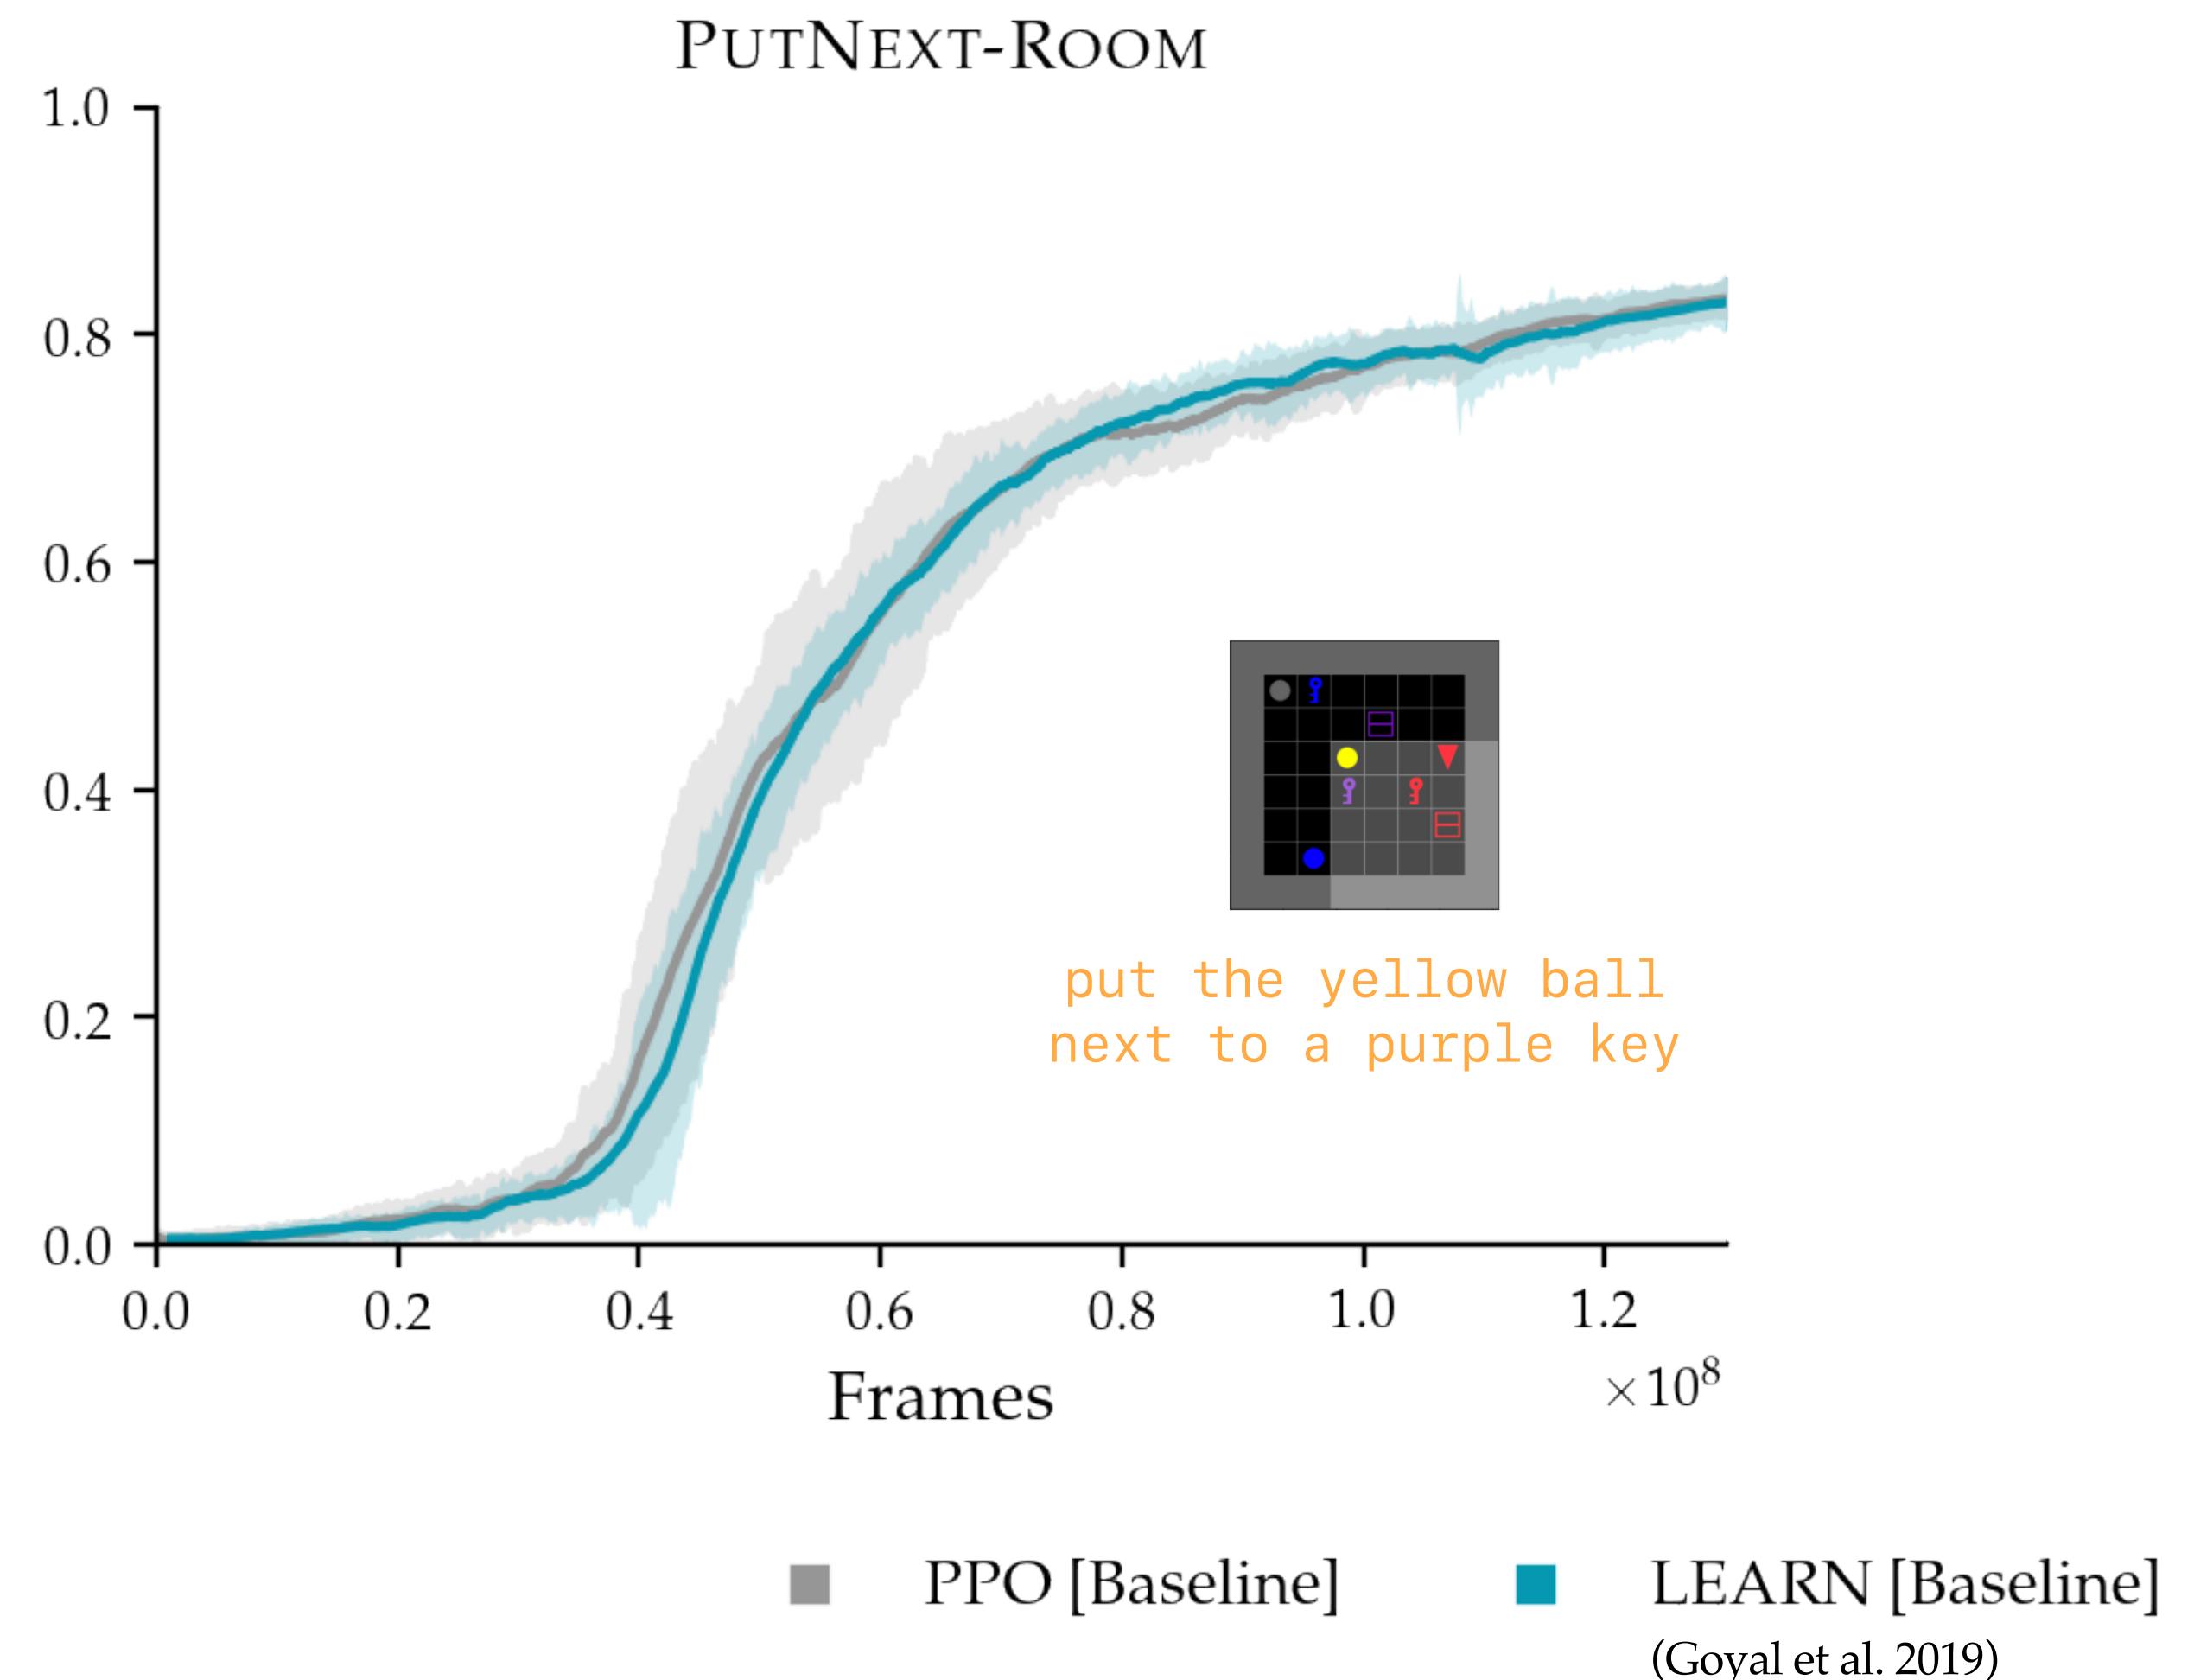
pick up a red box

Results: Sparsity

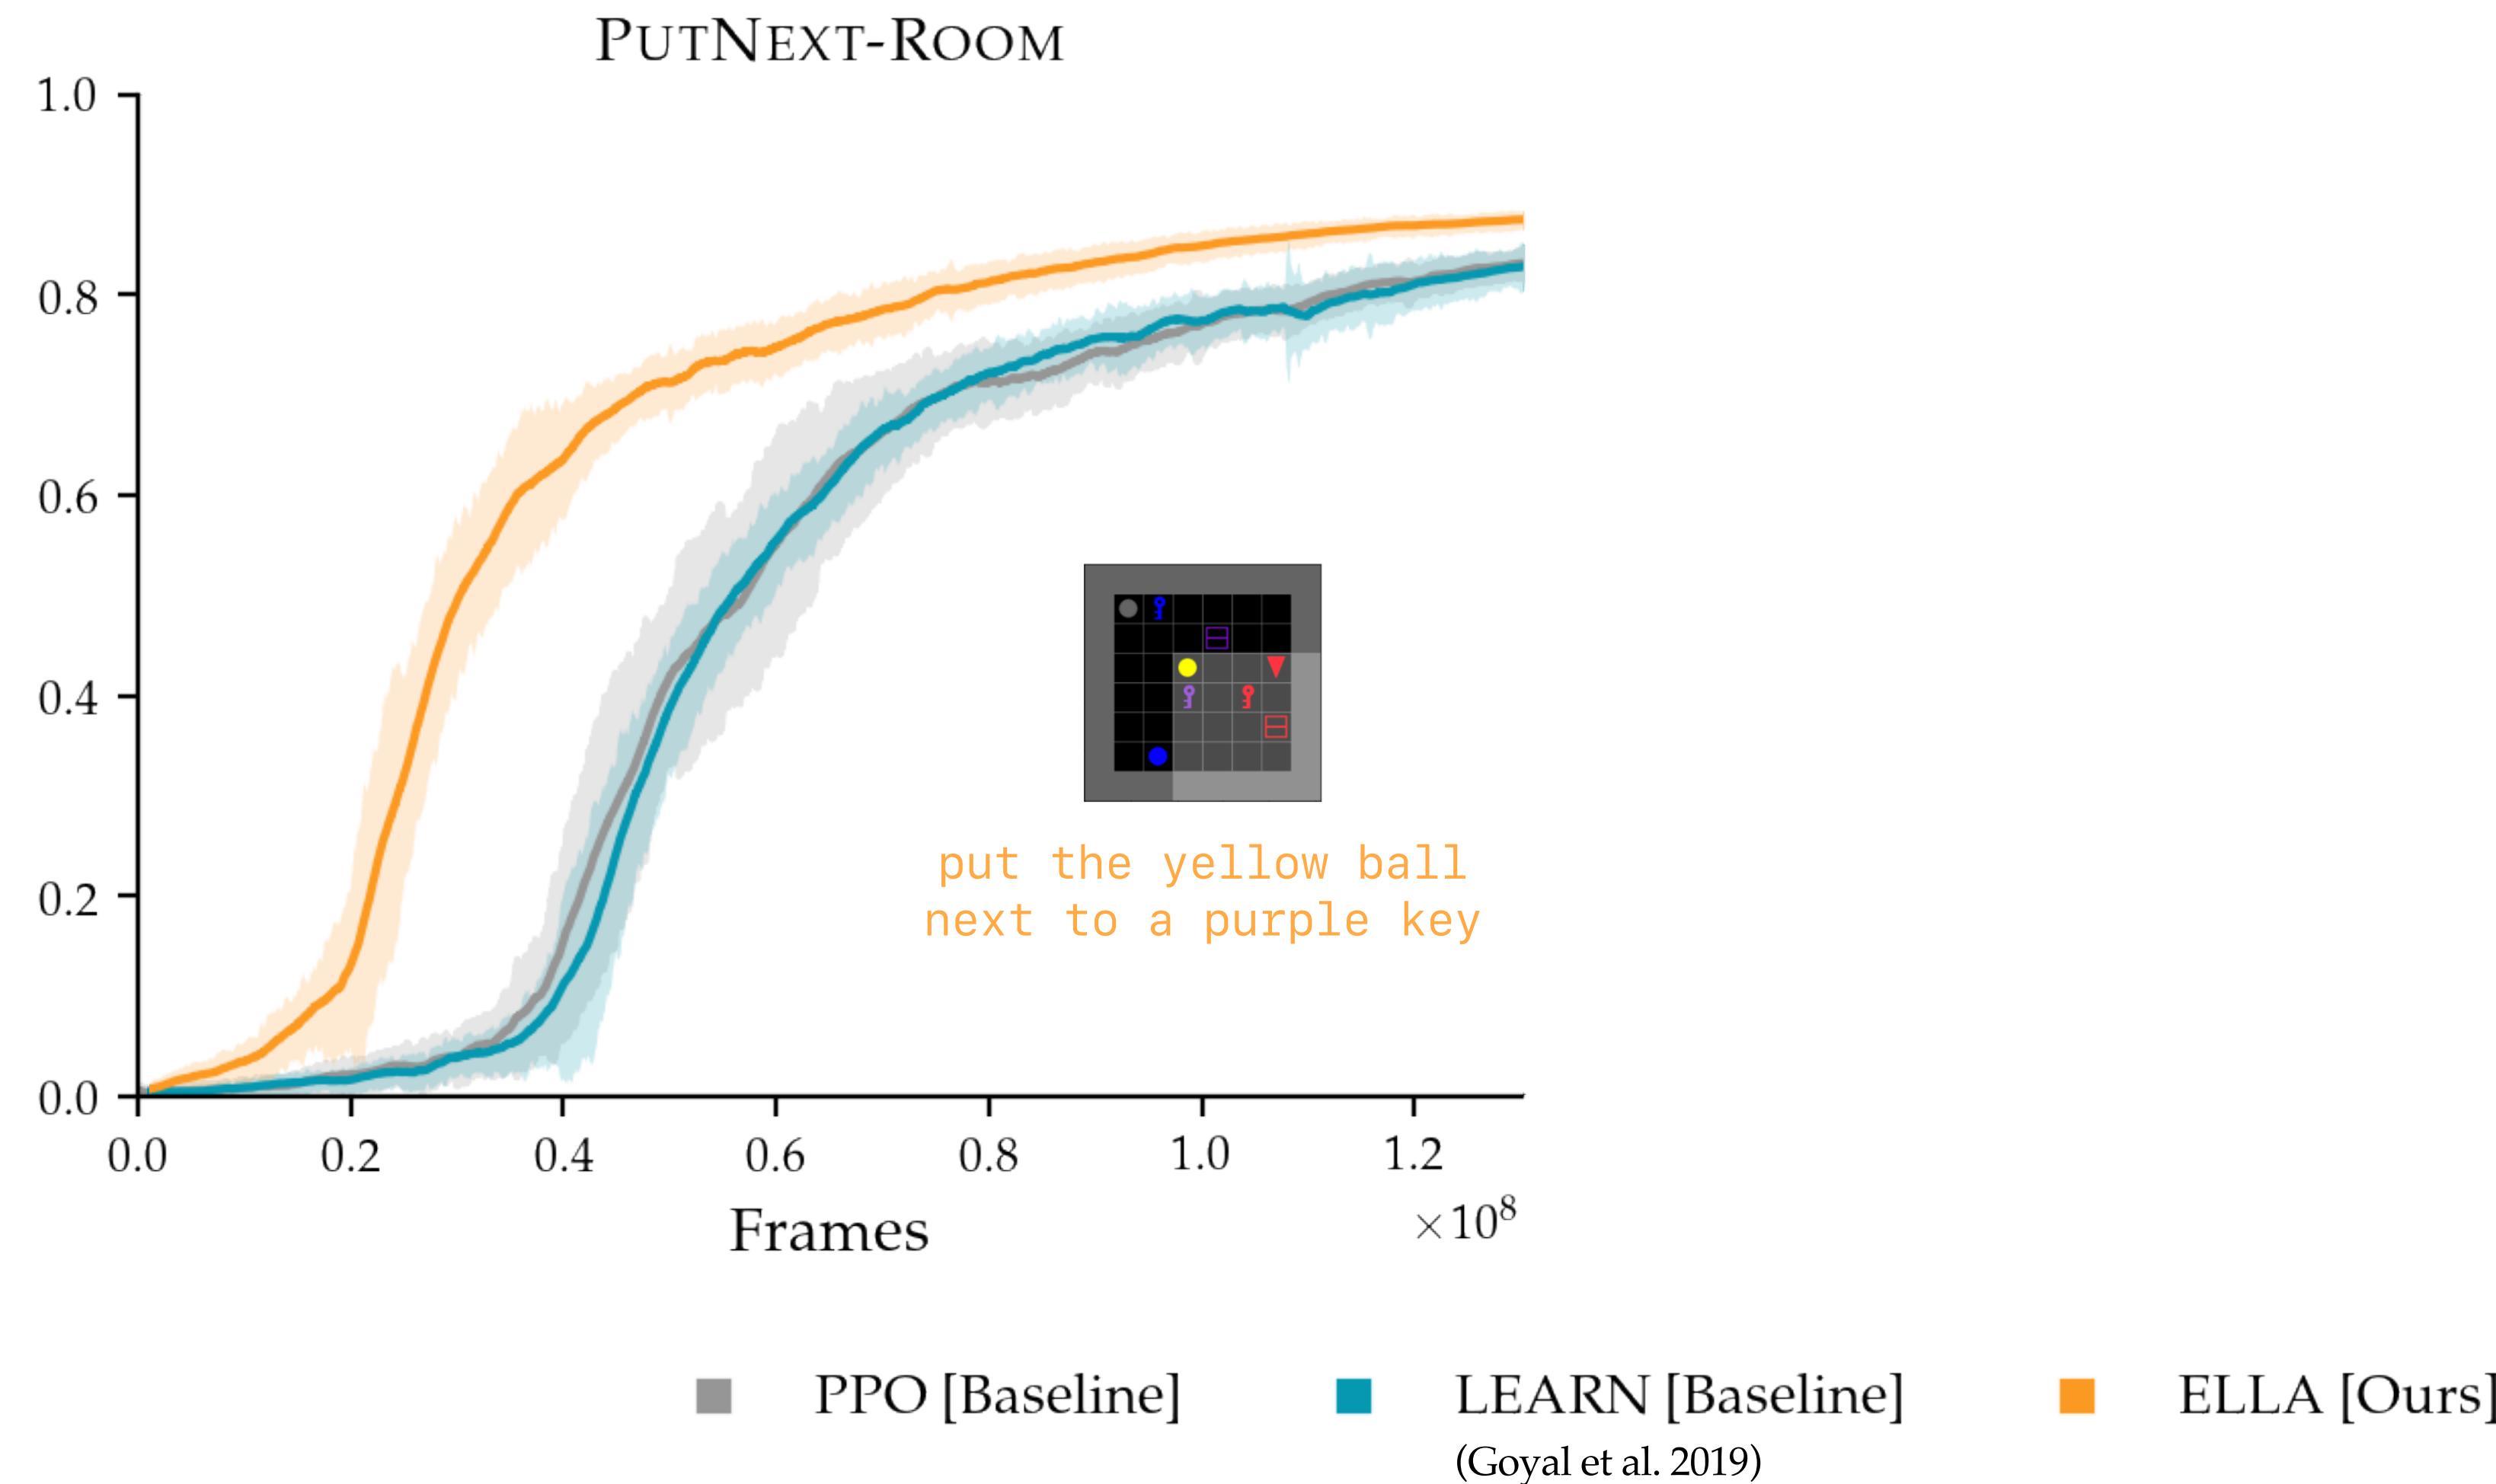
Results: Sparsity



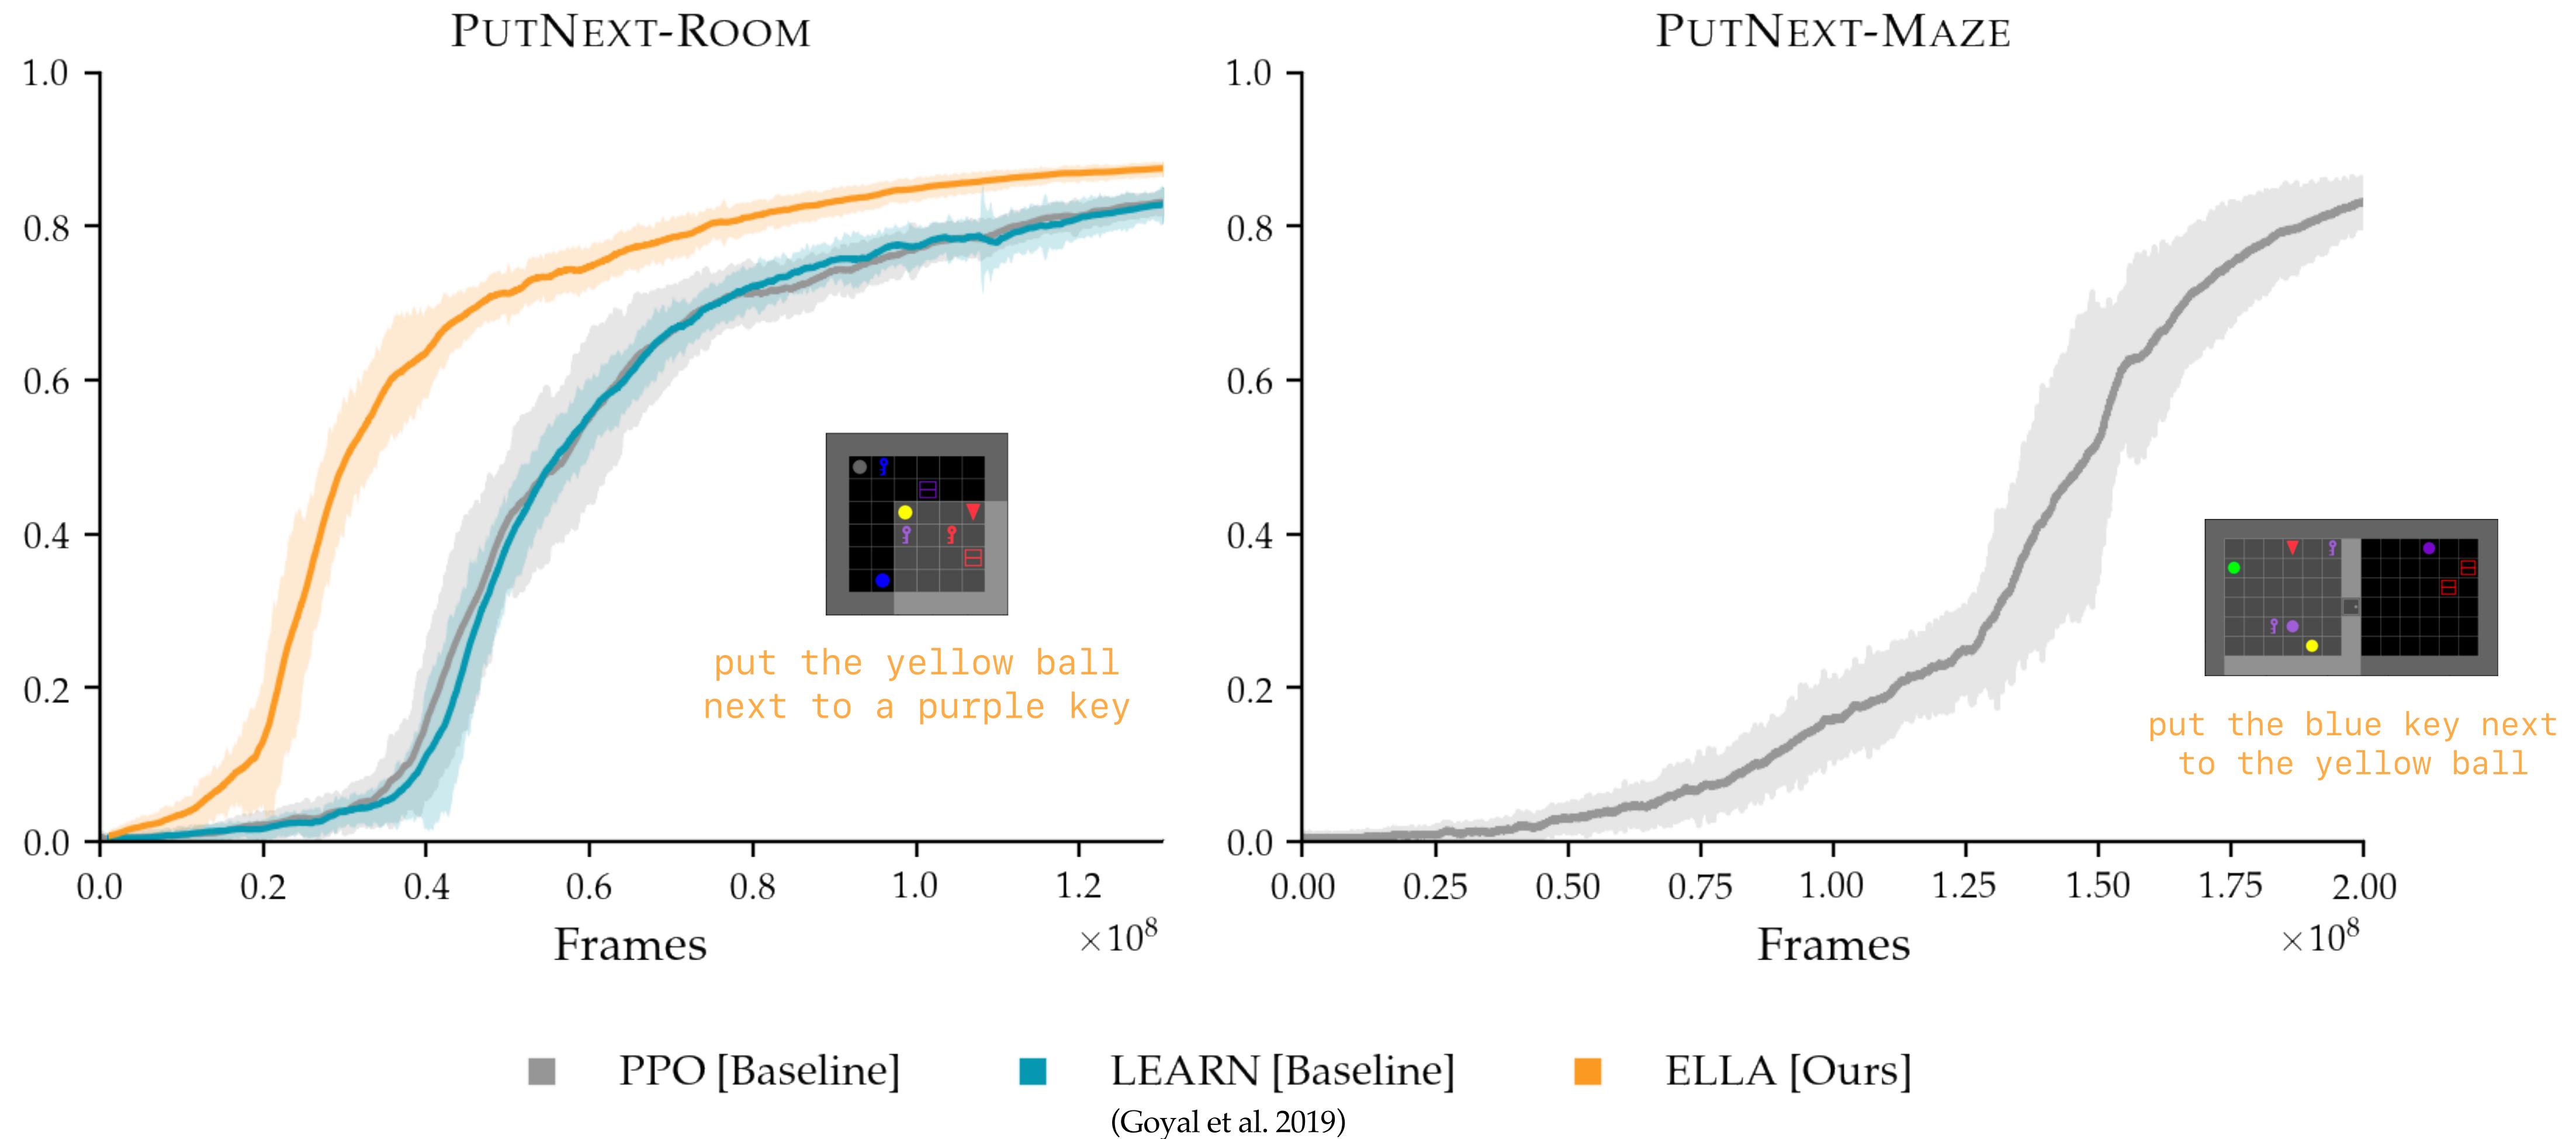
Results: Sparsity



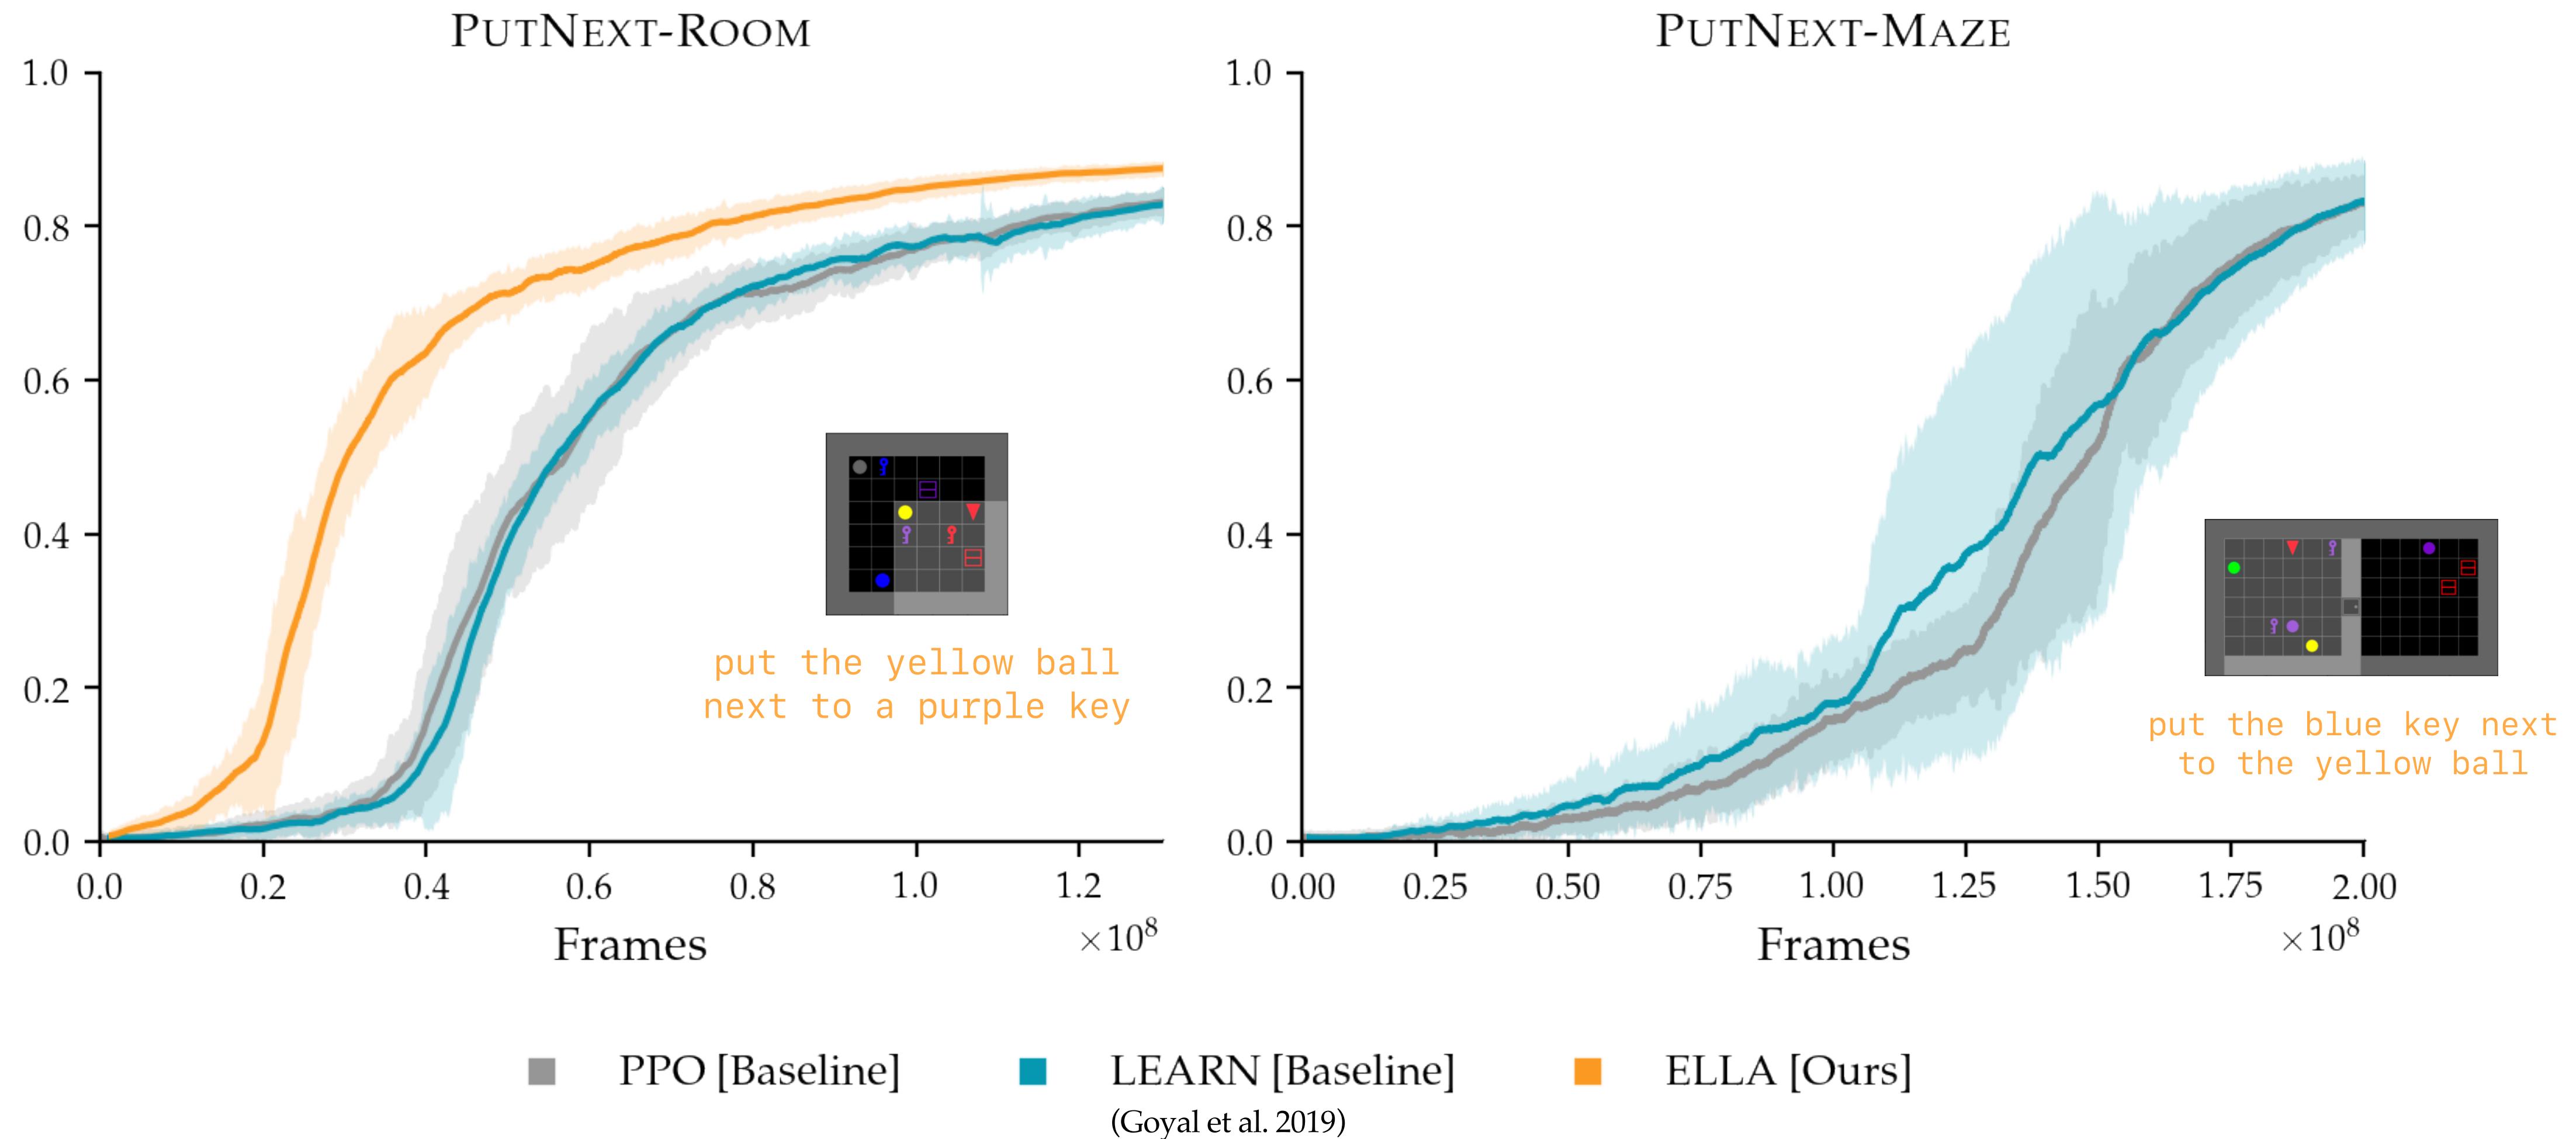
Results: Sparsity



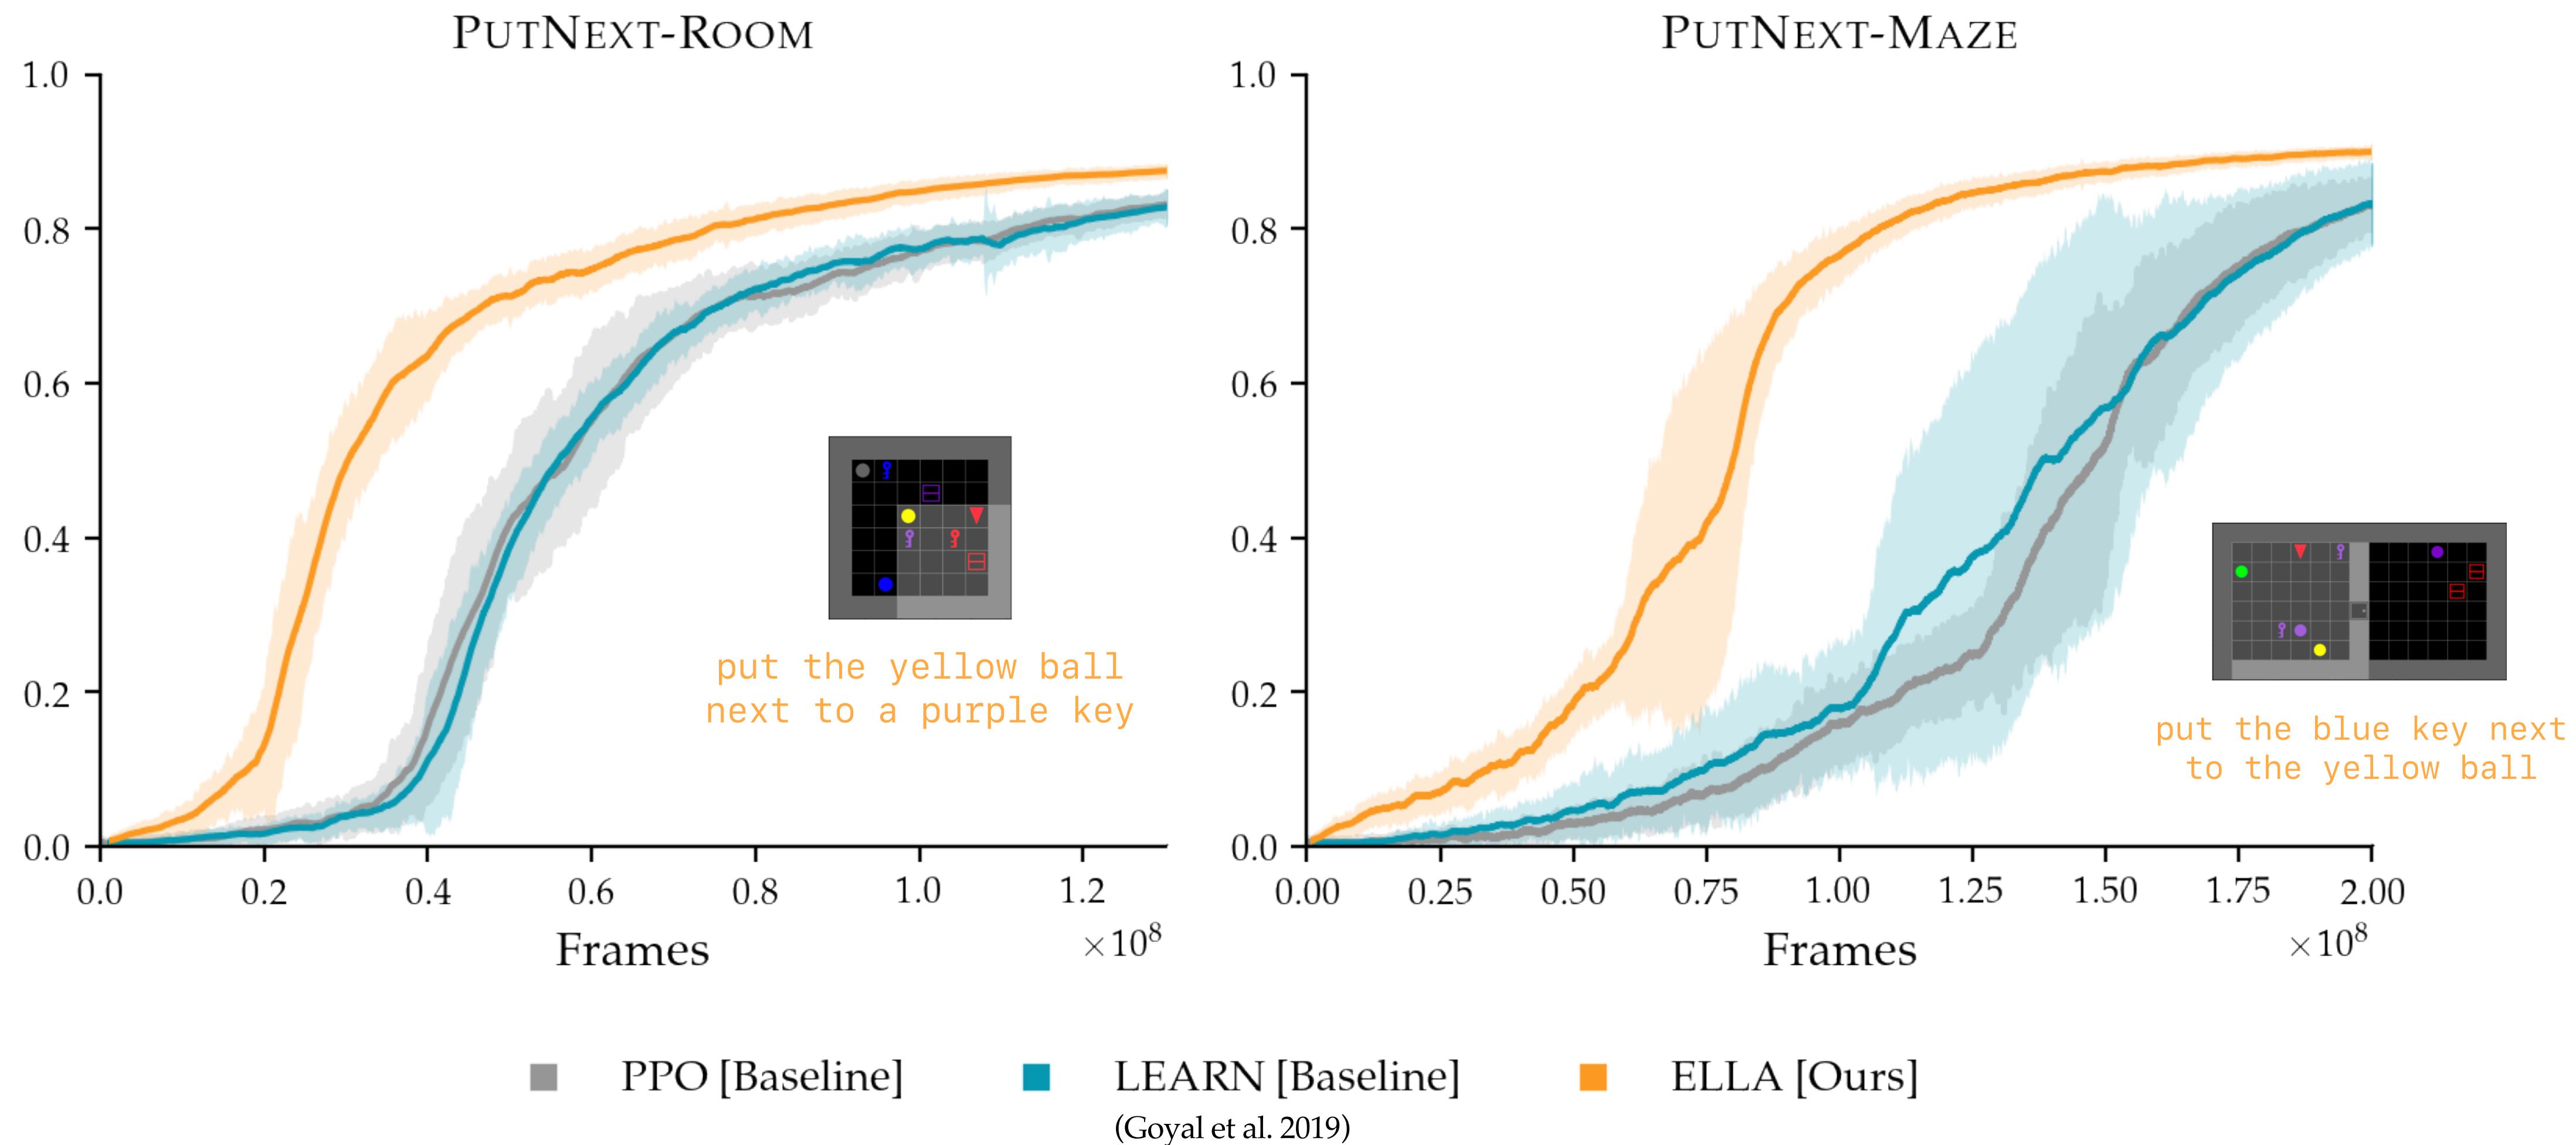
Results: Sparsity



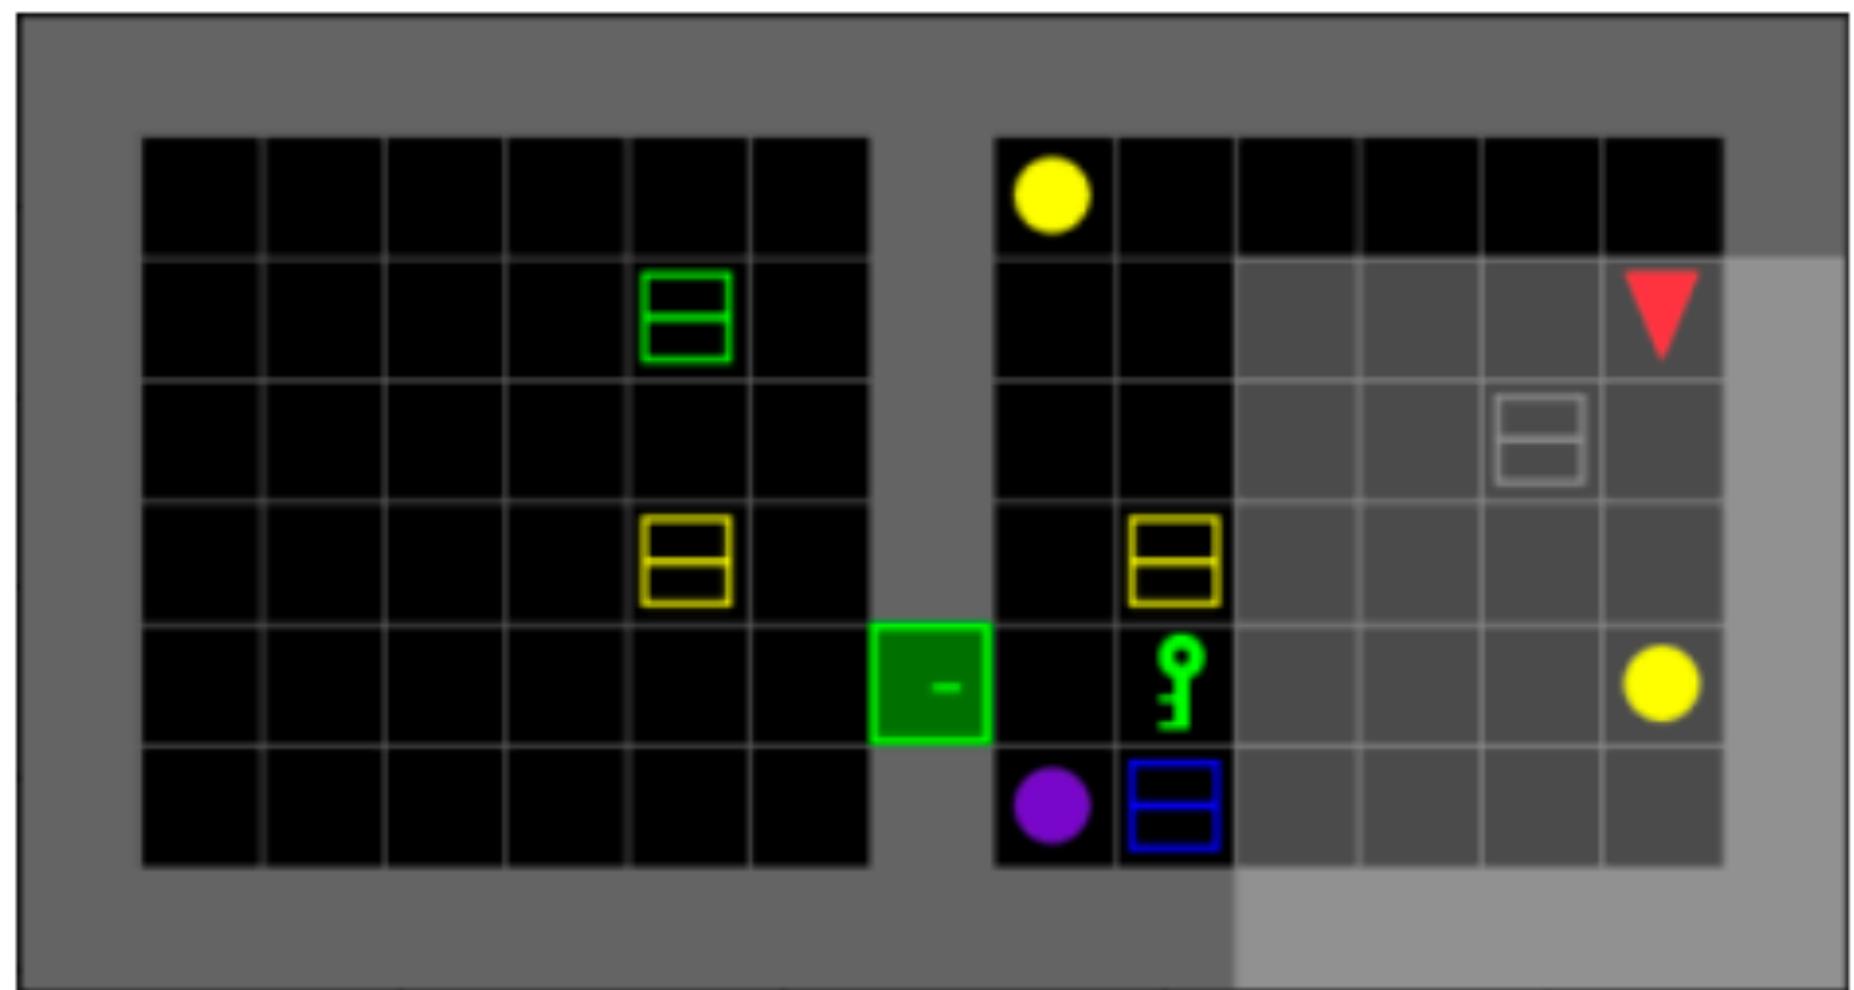
Results: Sparsity



Results: Sparsity

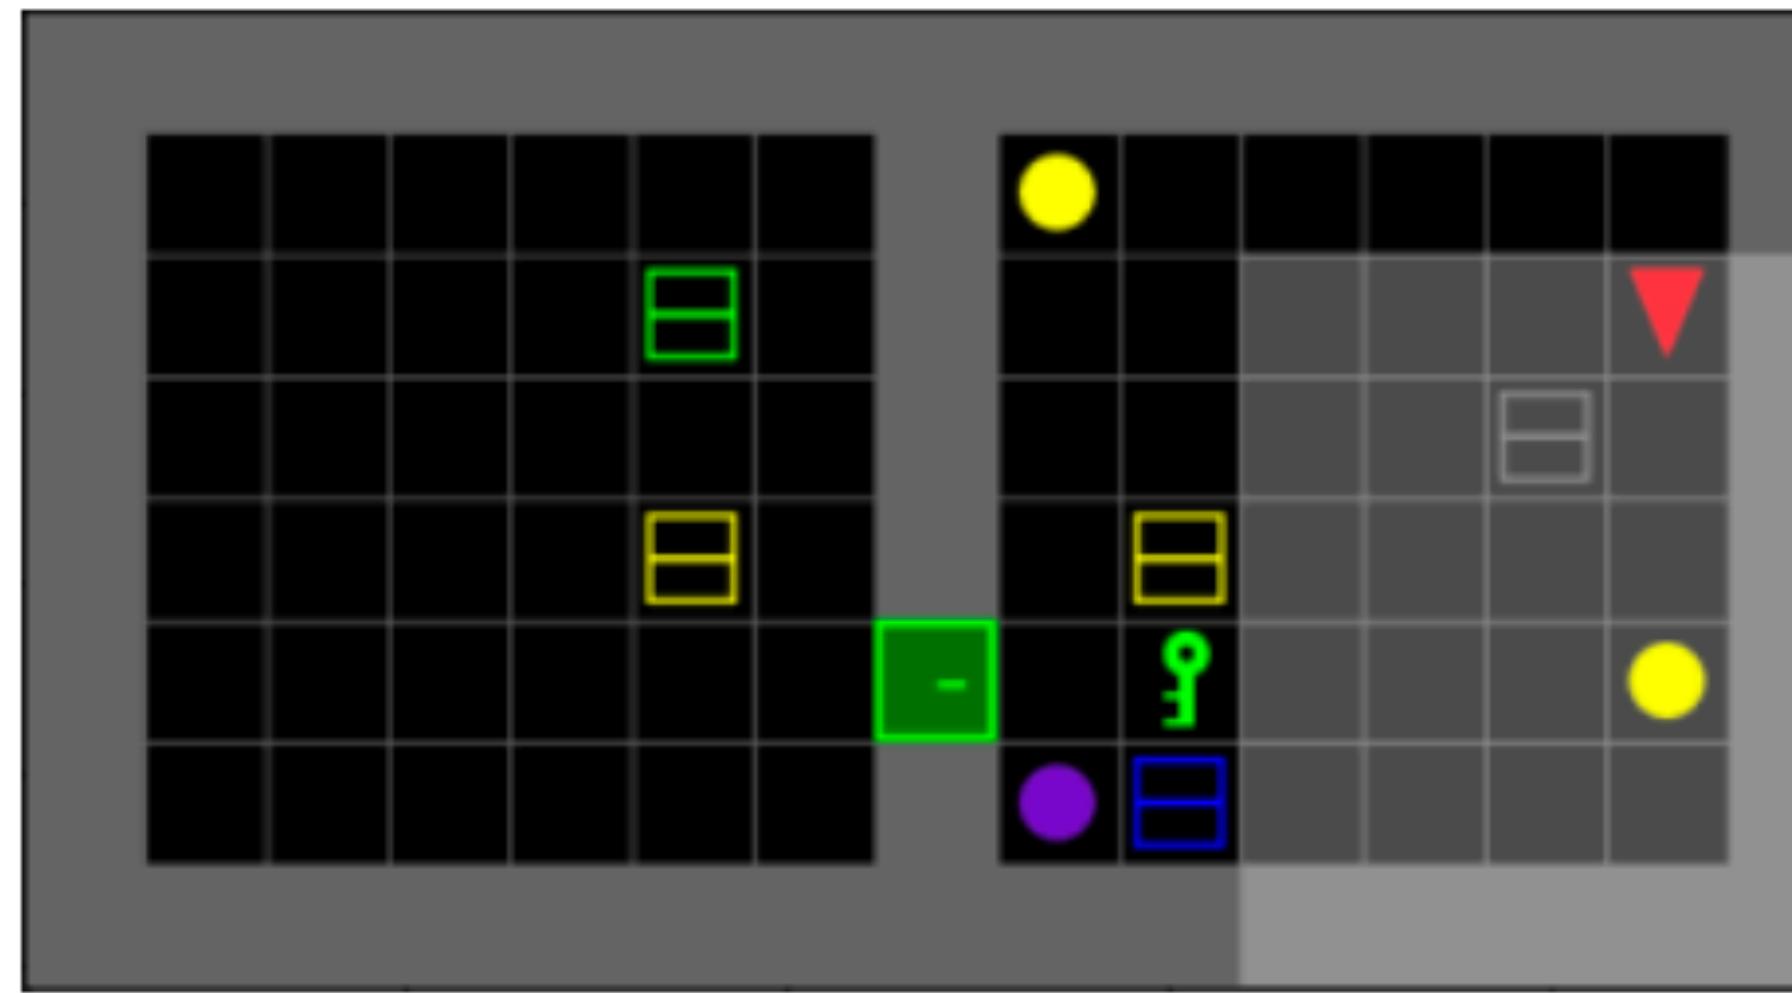


Results: Sparsity

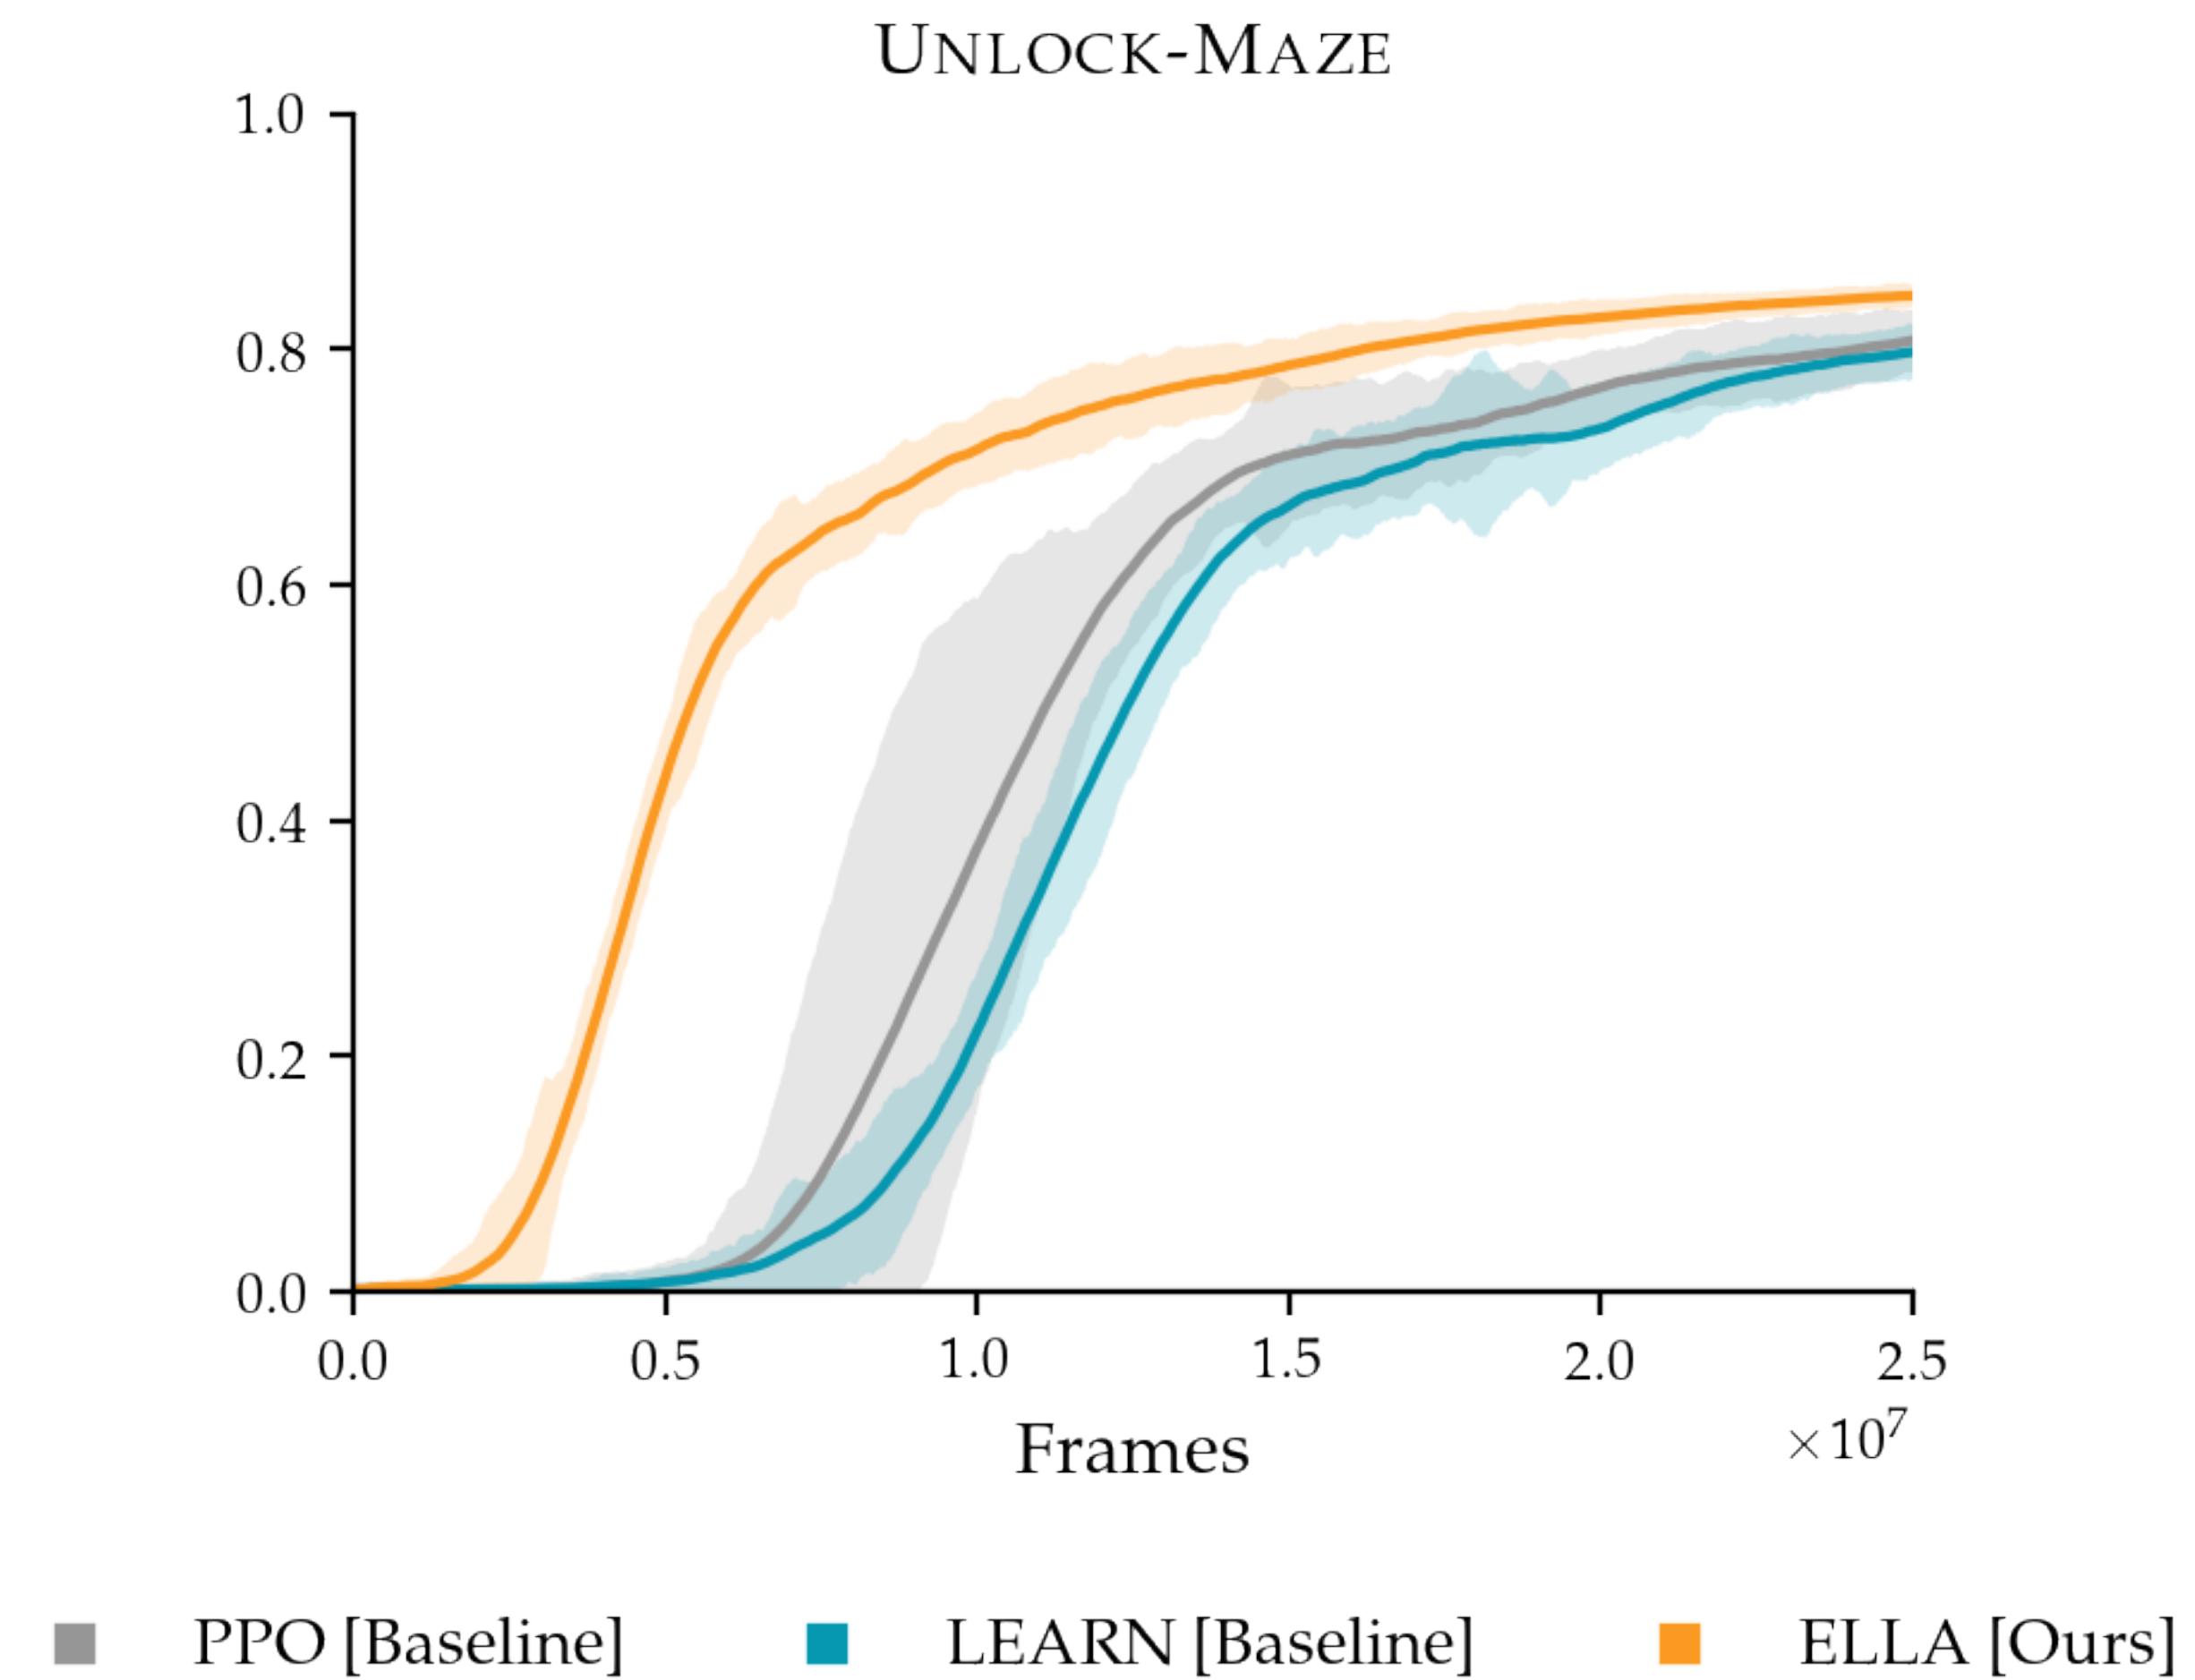


open the green door

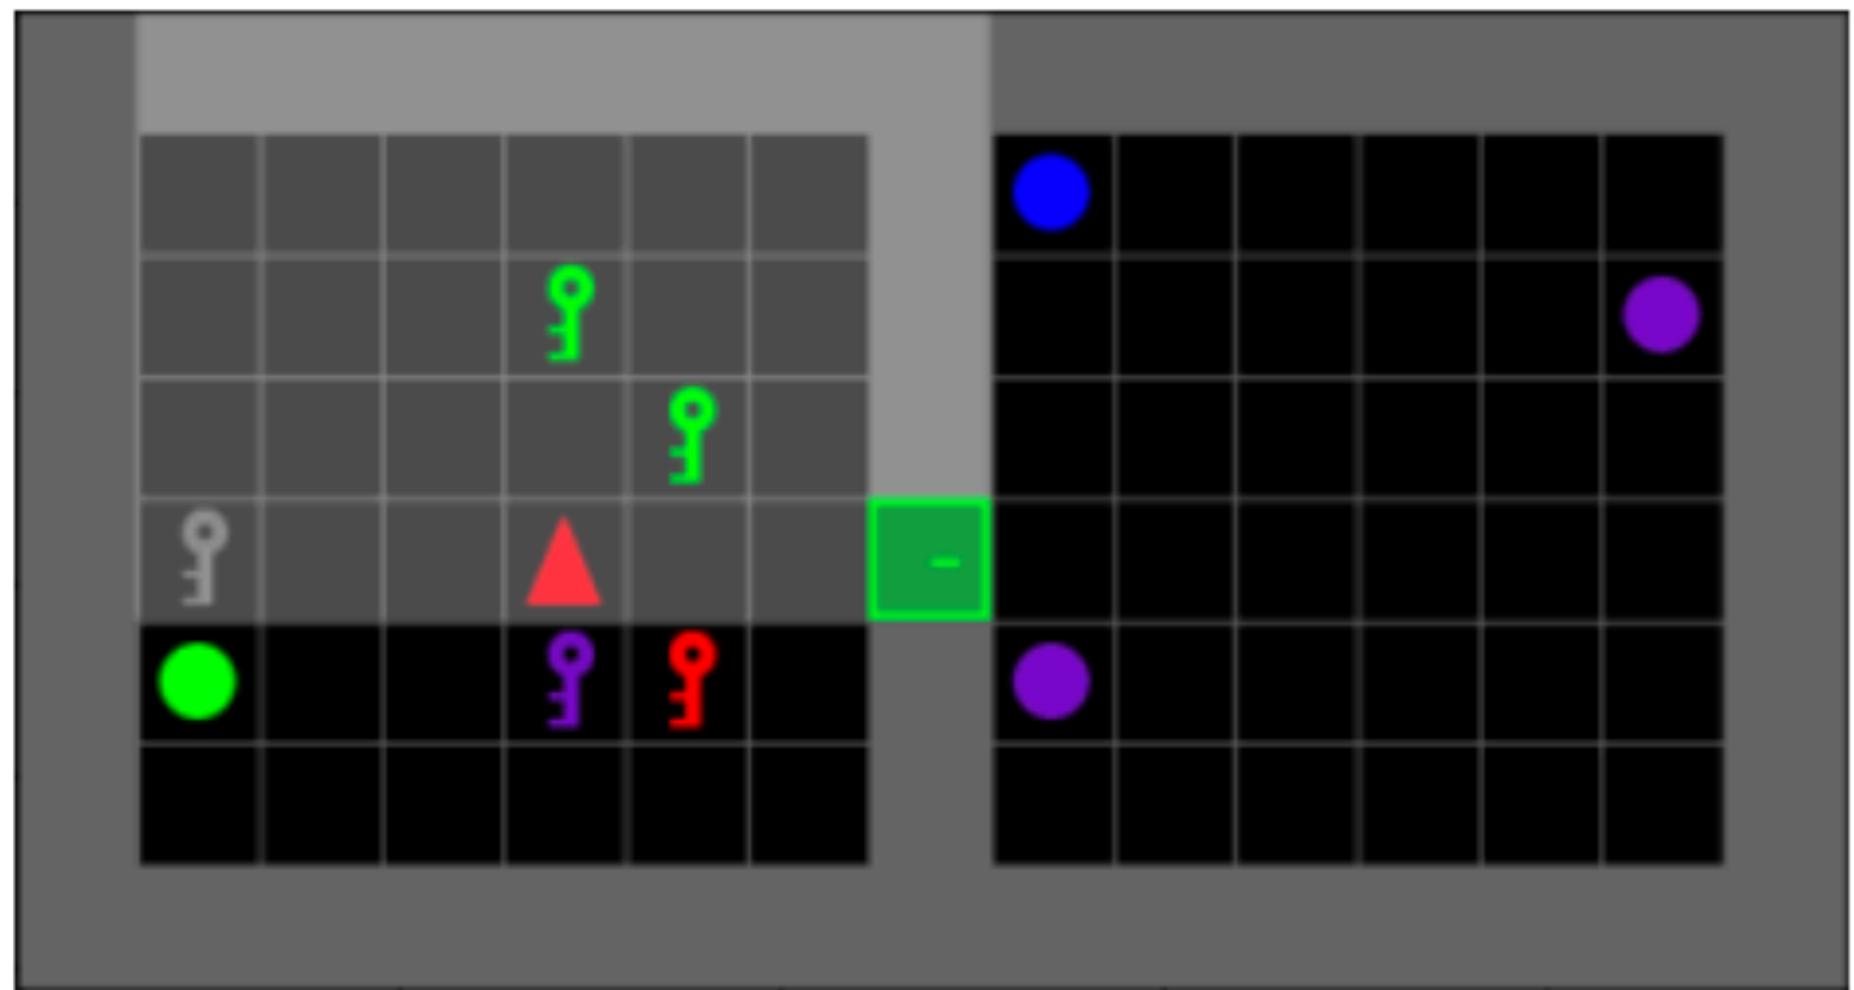
Results: Sparsity



open the green door



Results: Similarity

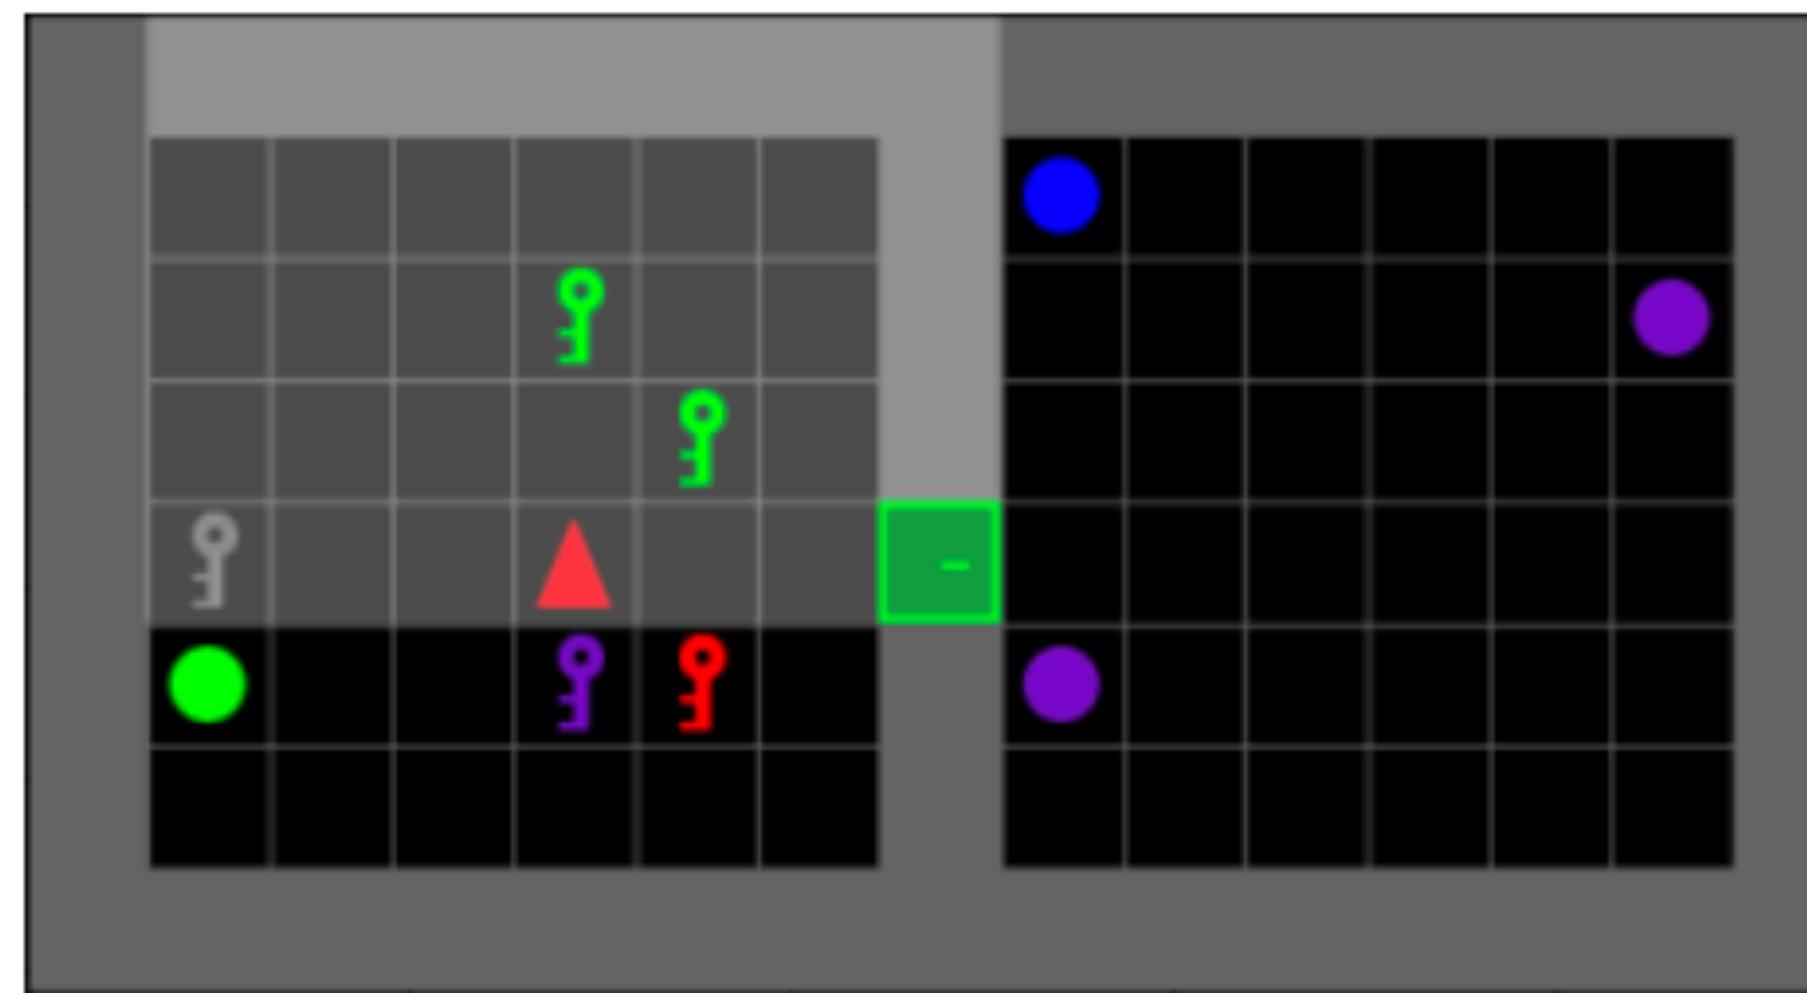


open the green door

pickup the red key

put the green ball next
to the purple ball

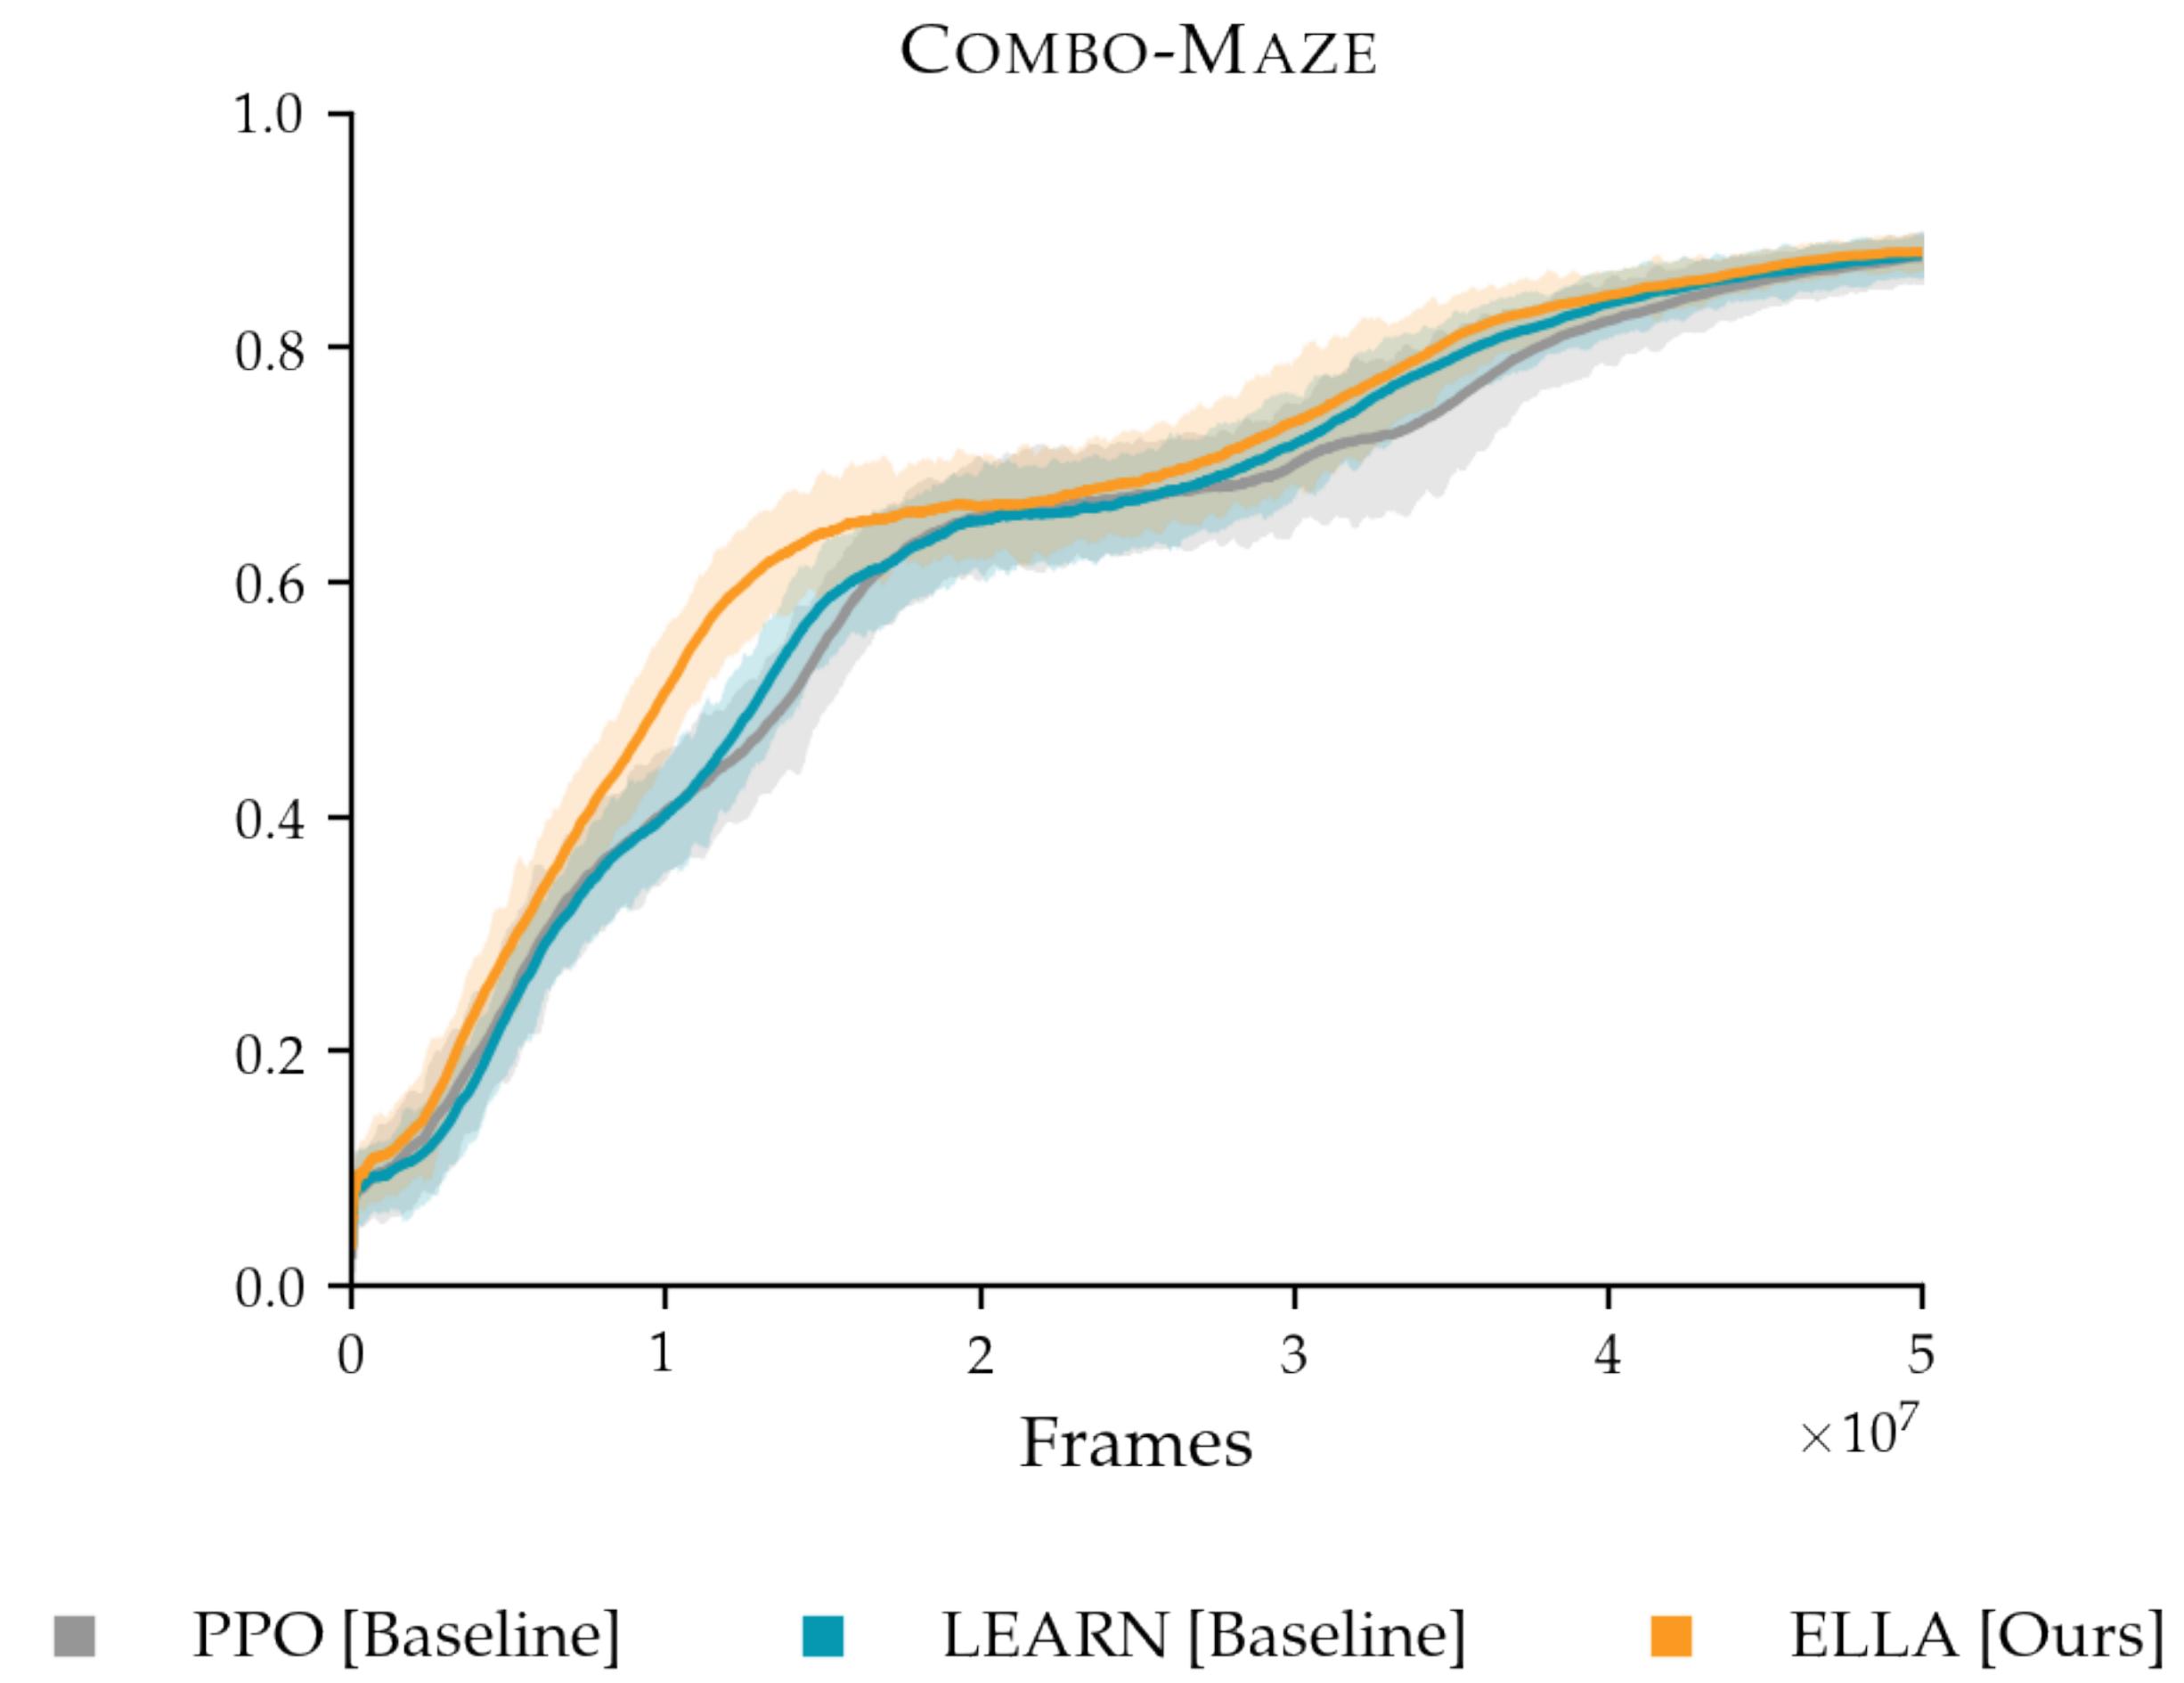
Results: Similarity



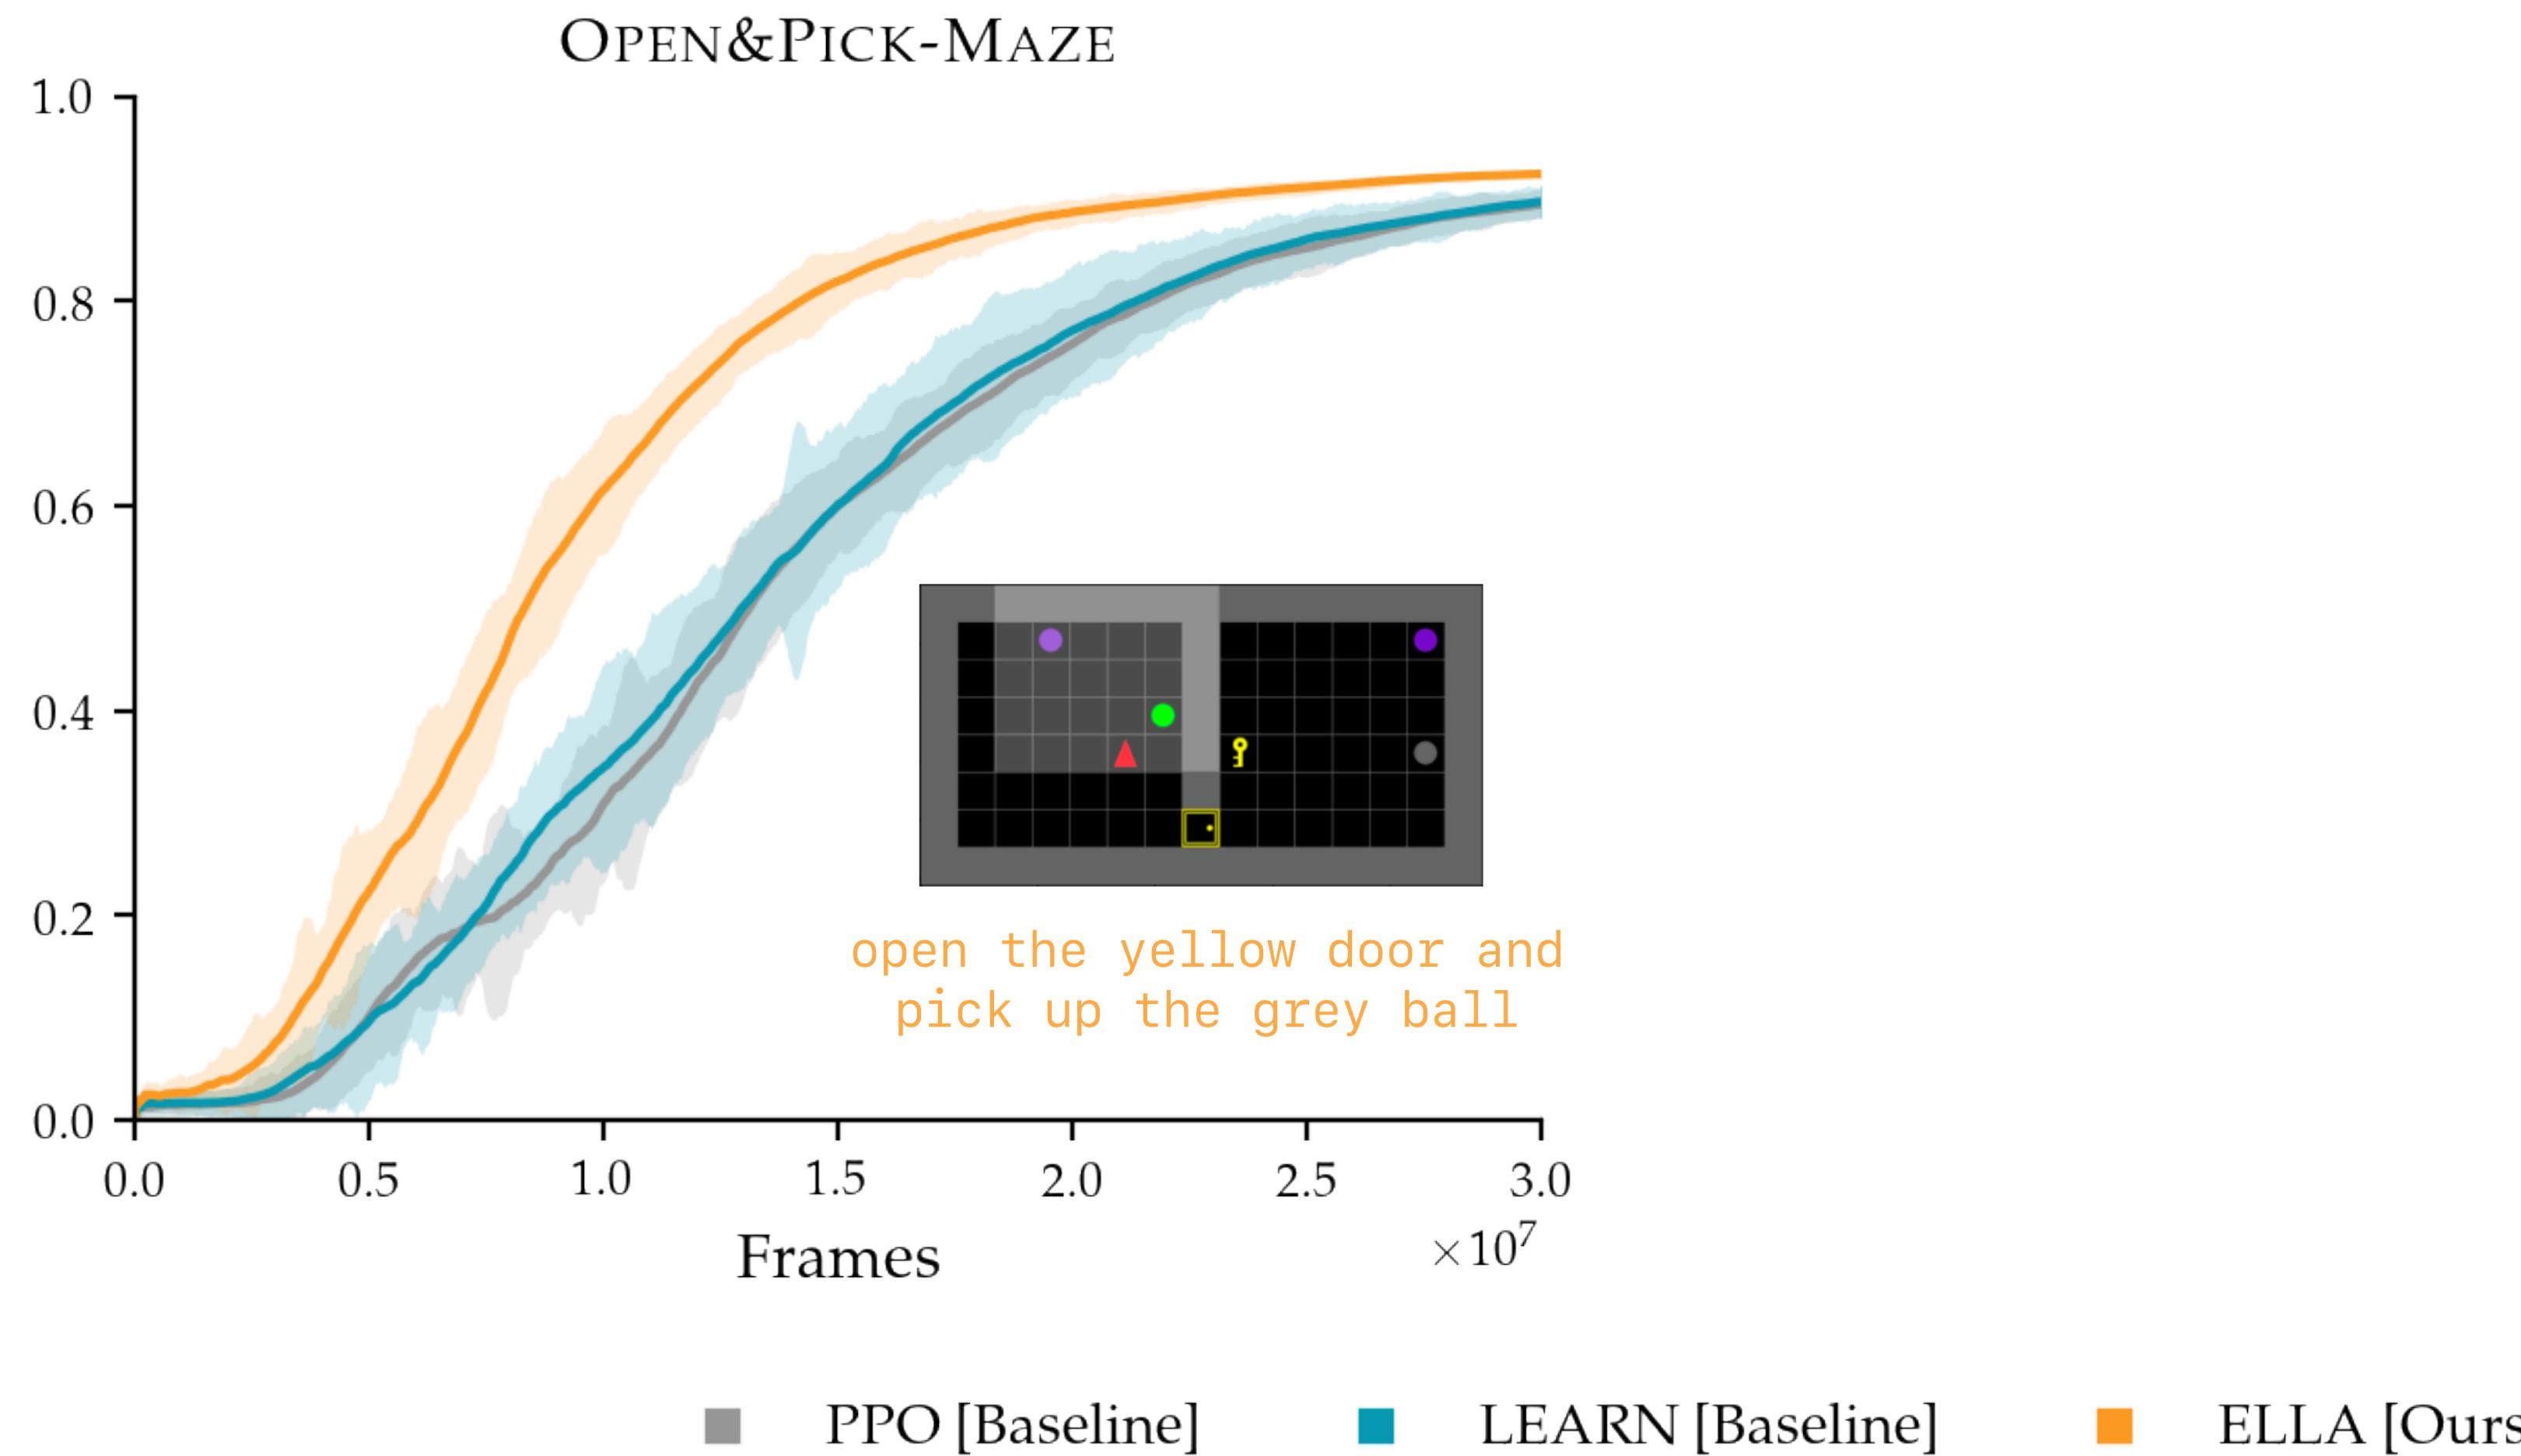
open the green door

pickup the red key

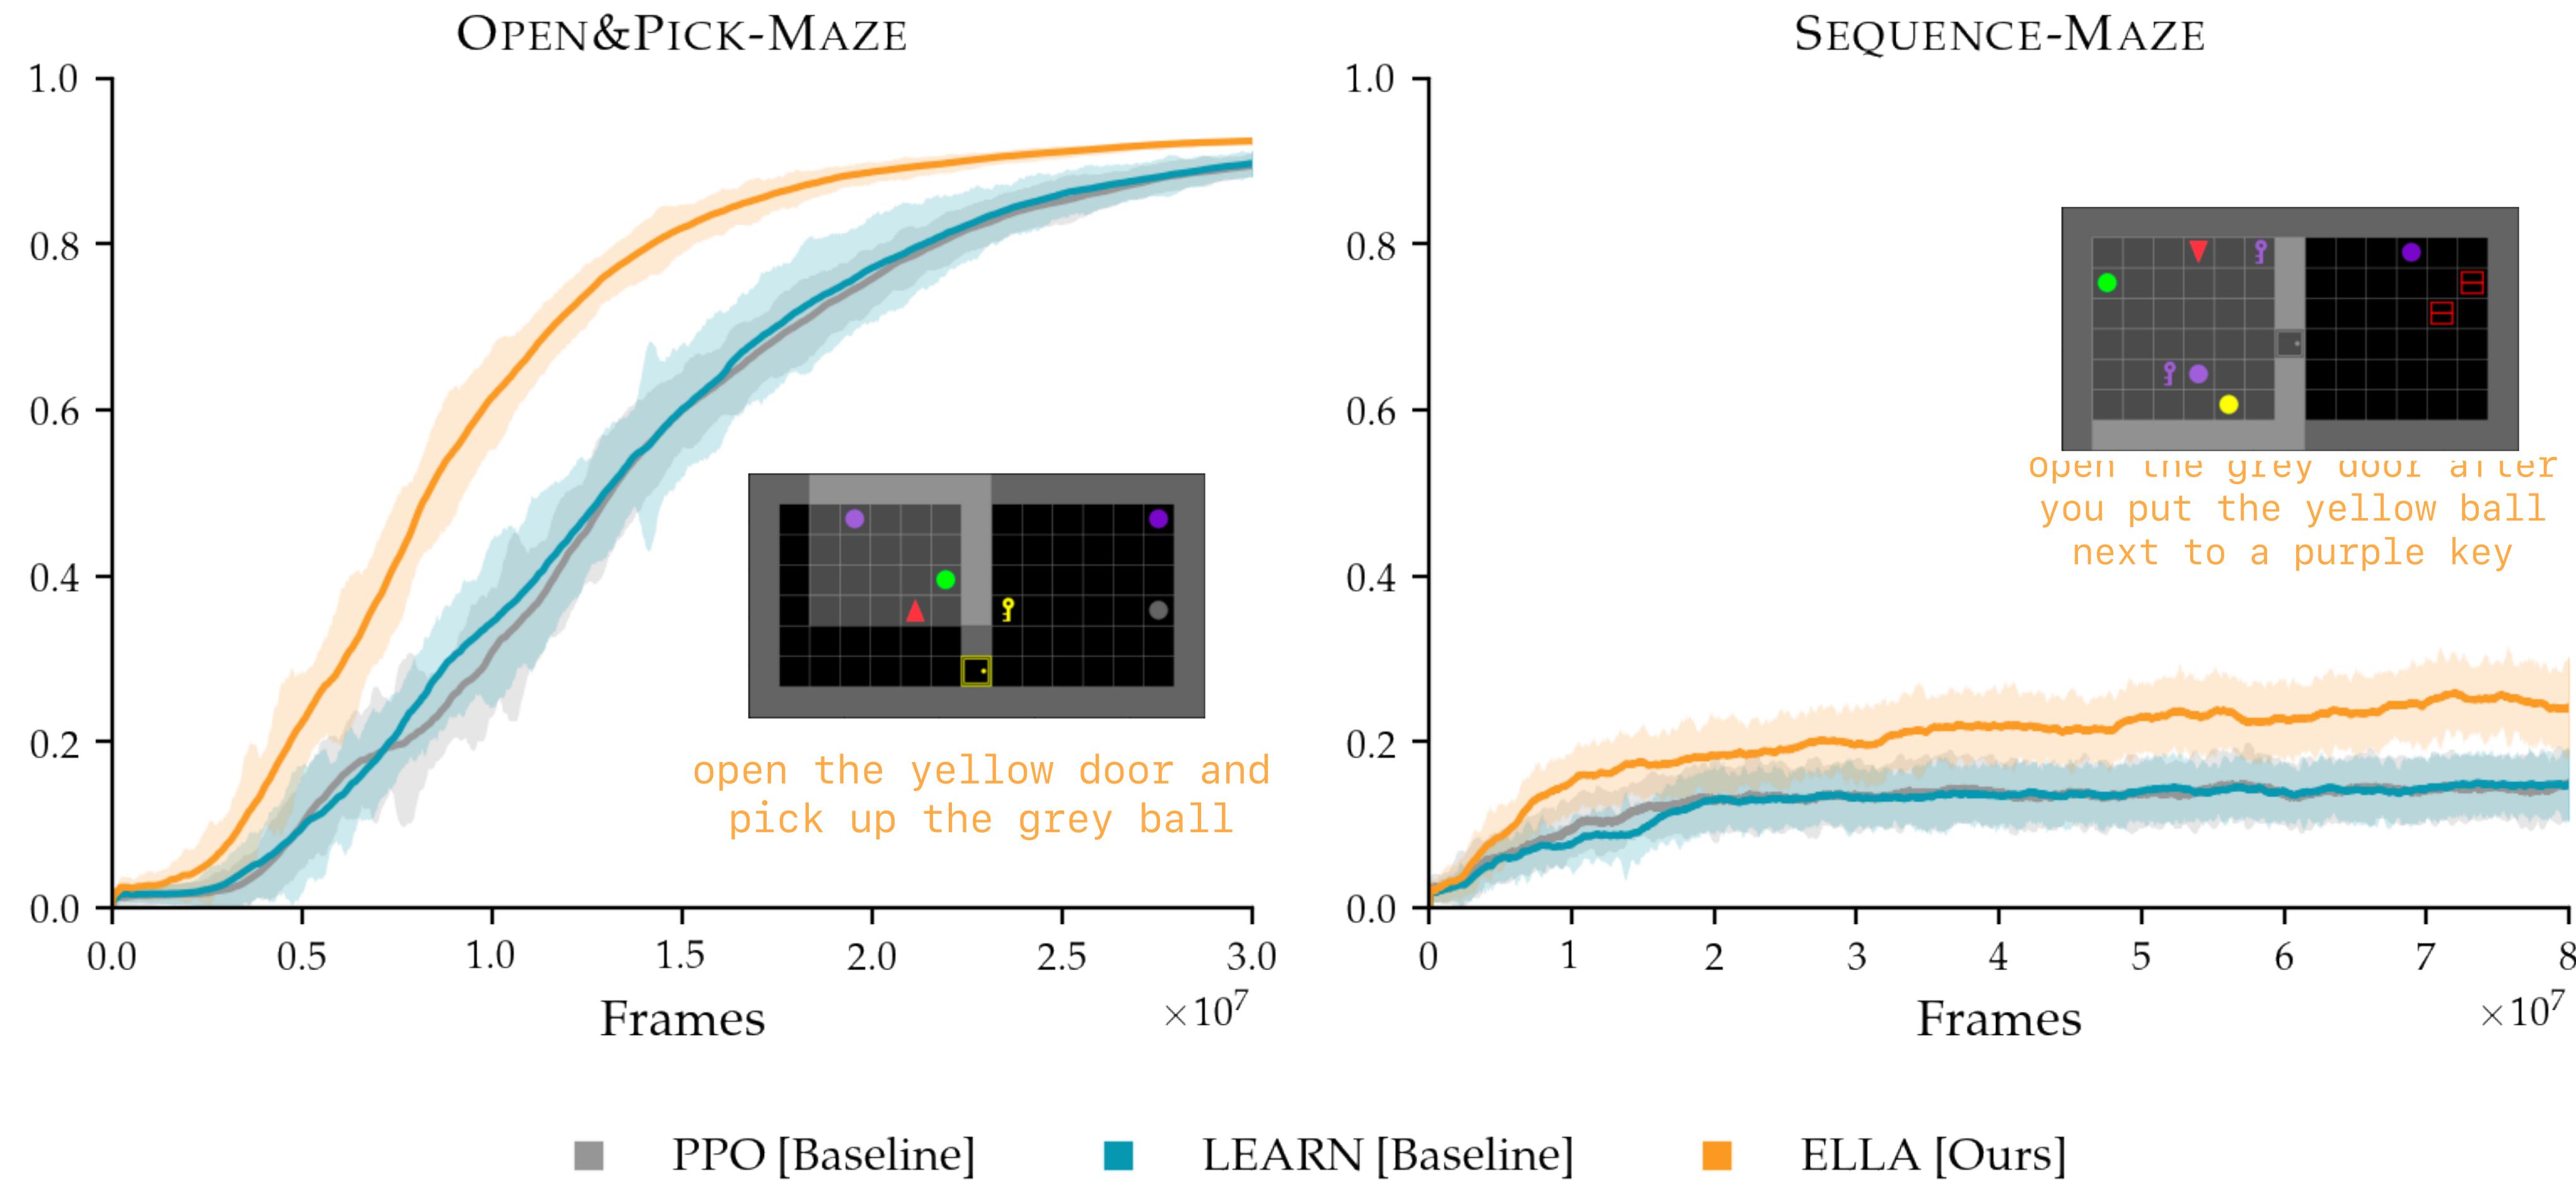
put the green ball next
to the purple ball



Results: Compositionality



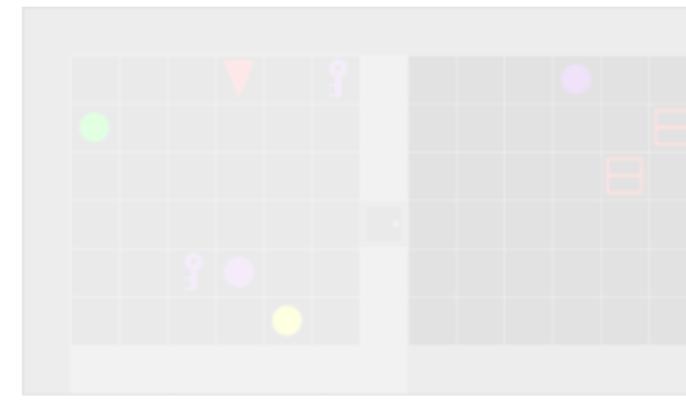
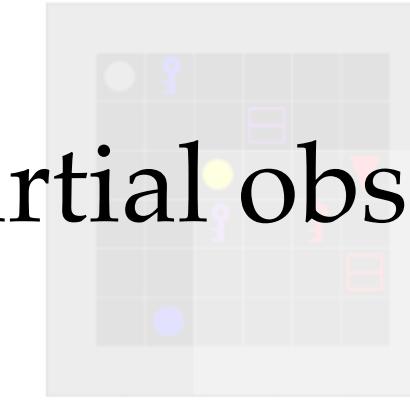
Results: Compositionality



Experiments

High-Level Tasks

- BabyAI (Chevalier-Boisvert et al. 2018)



- Partial observability

put the yellow ball
next to a purple key

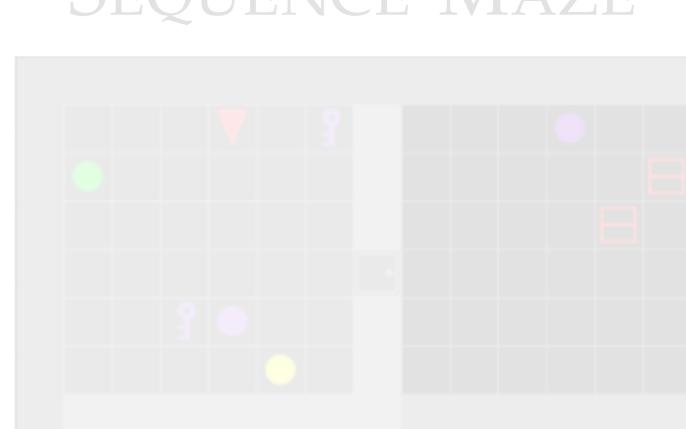
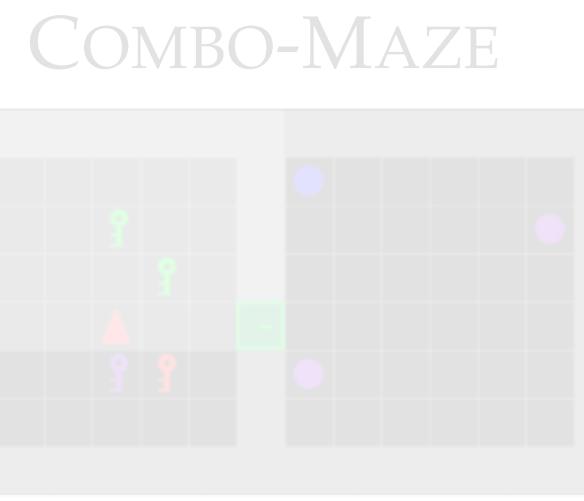
- ROOM levels: single 7×7 grid



- MAZE levels: two rooms connected by a closed/locked door

open the yellow door and
pick up the grey ball

open the green door



pick up the green ball

open the grey door after
you put the yellow ball
next to a purple key

Low-Level Tasks

- Tasks differ on several axes

- *Sparsity* of the high-level task

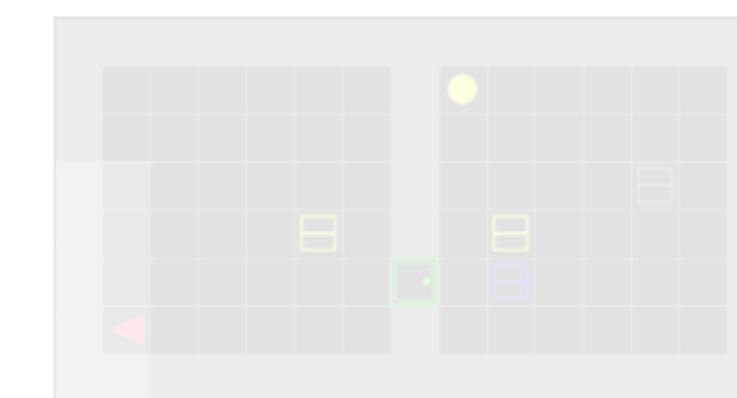


- *Similarity* of the low- and high-level tasks

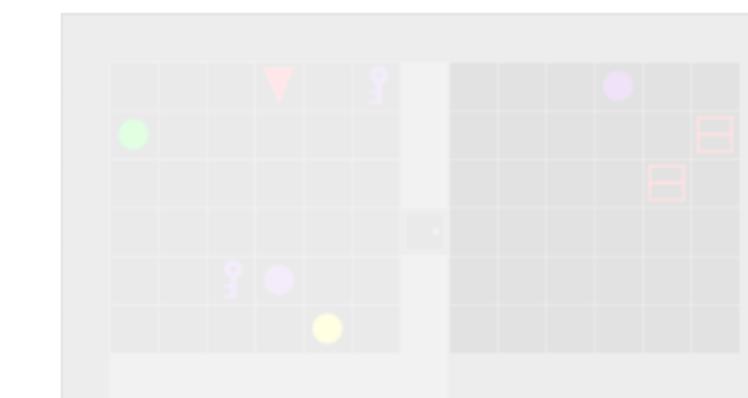
go to a yellow ball

go to a red key

- *Compositionality* of the tasks in G

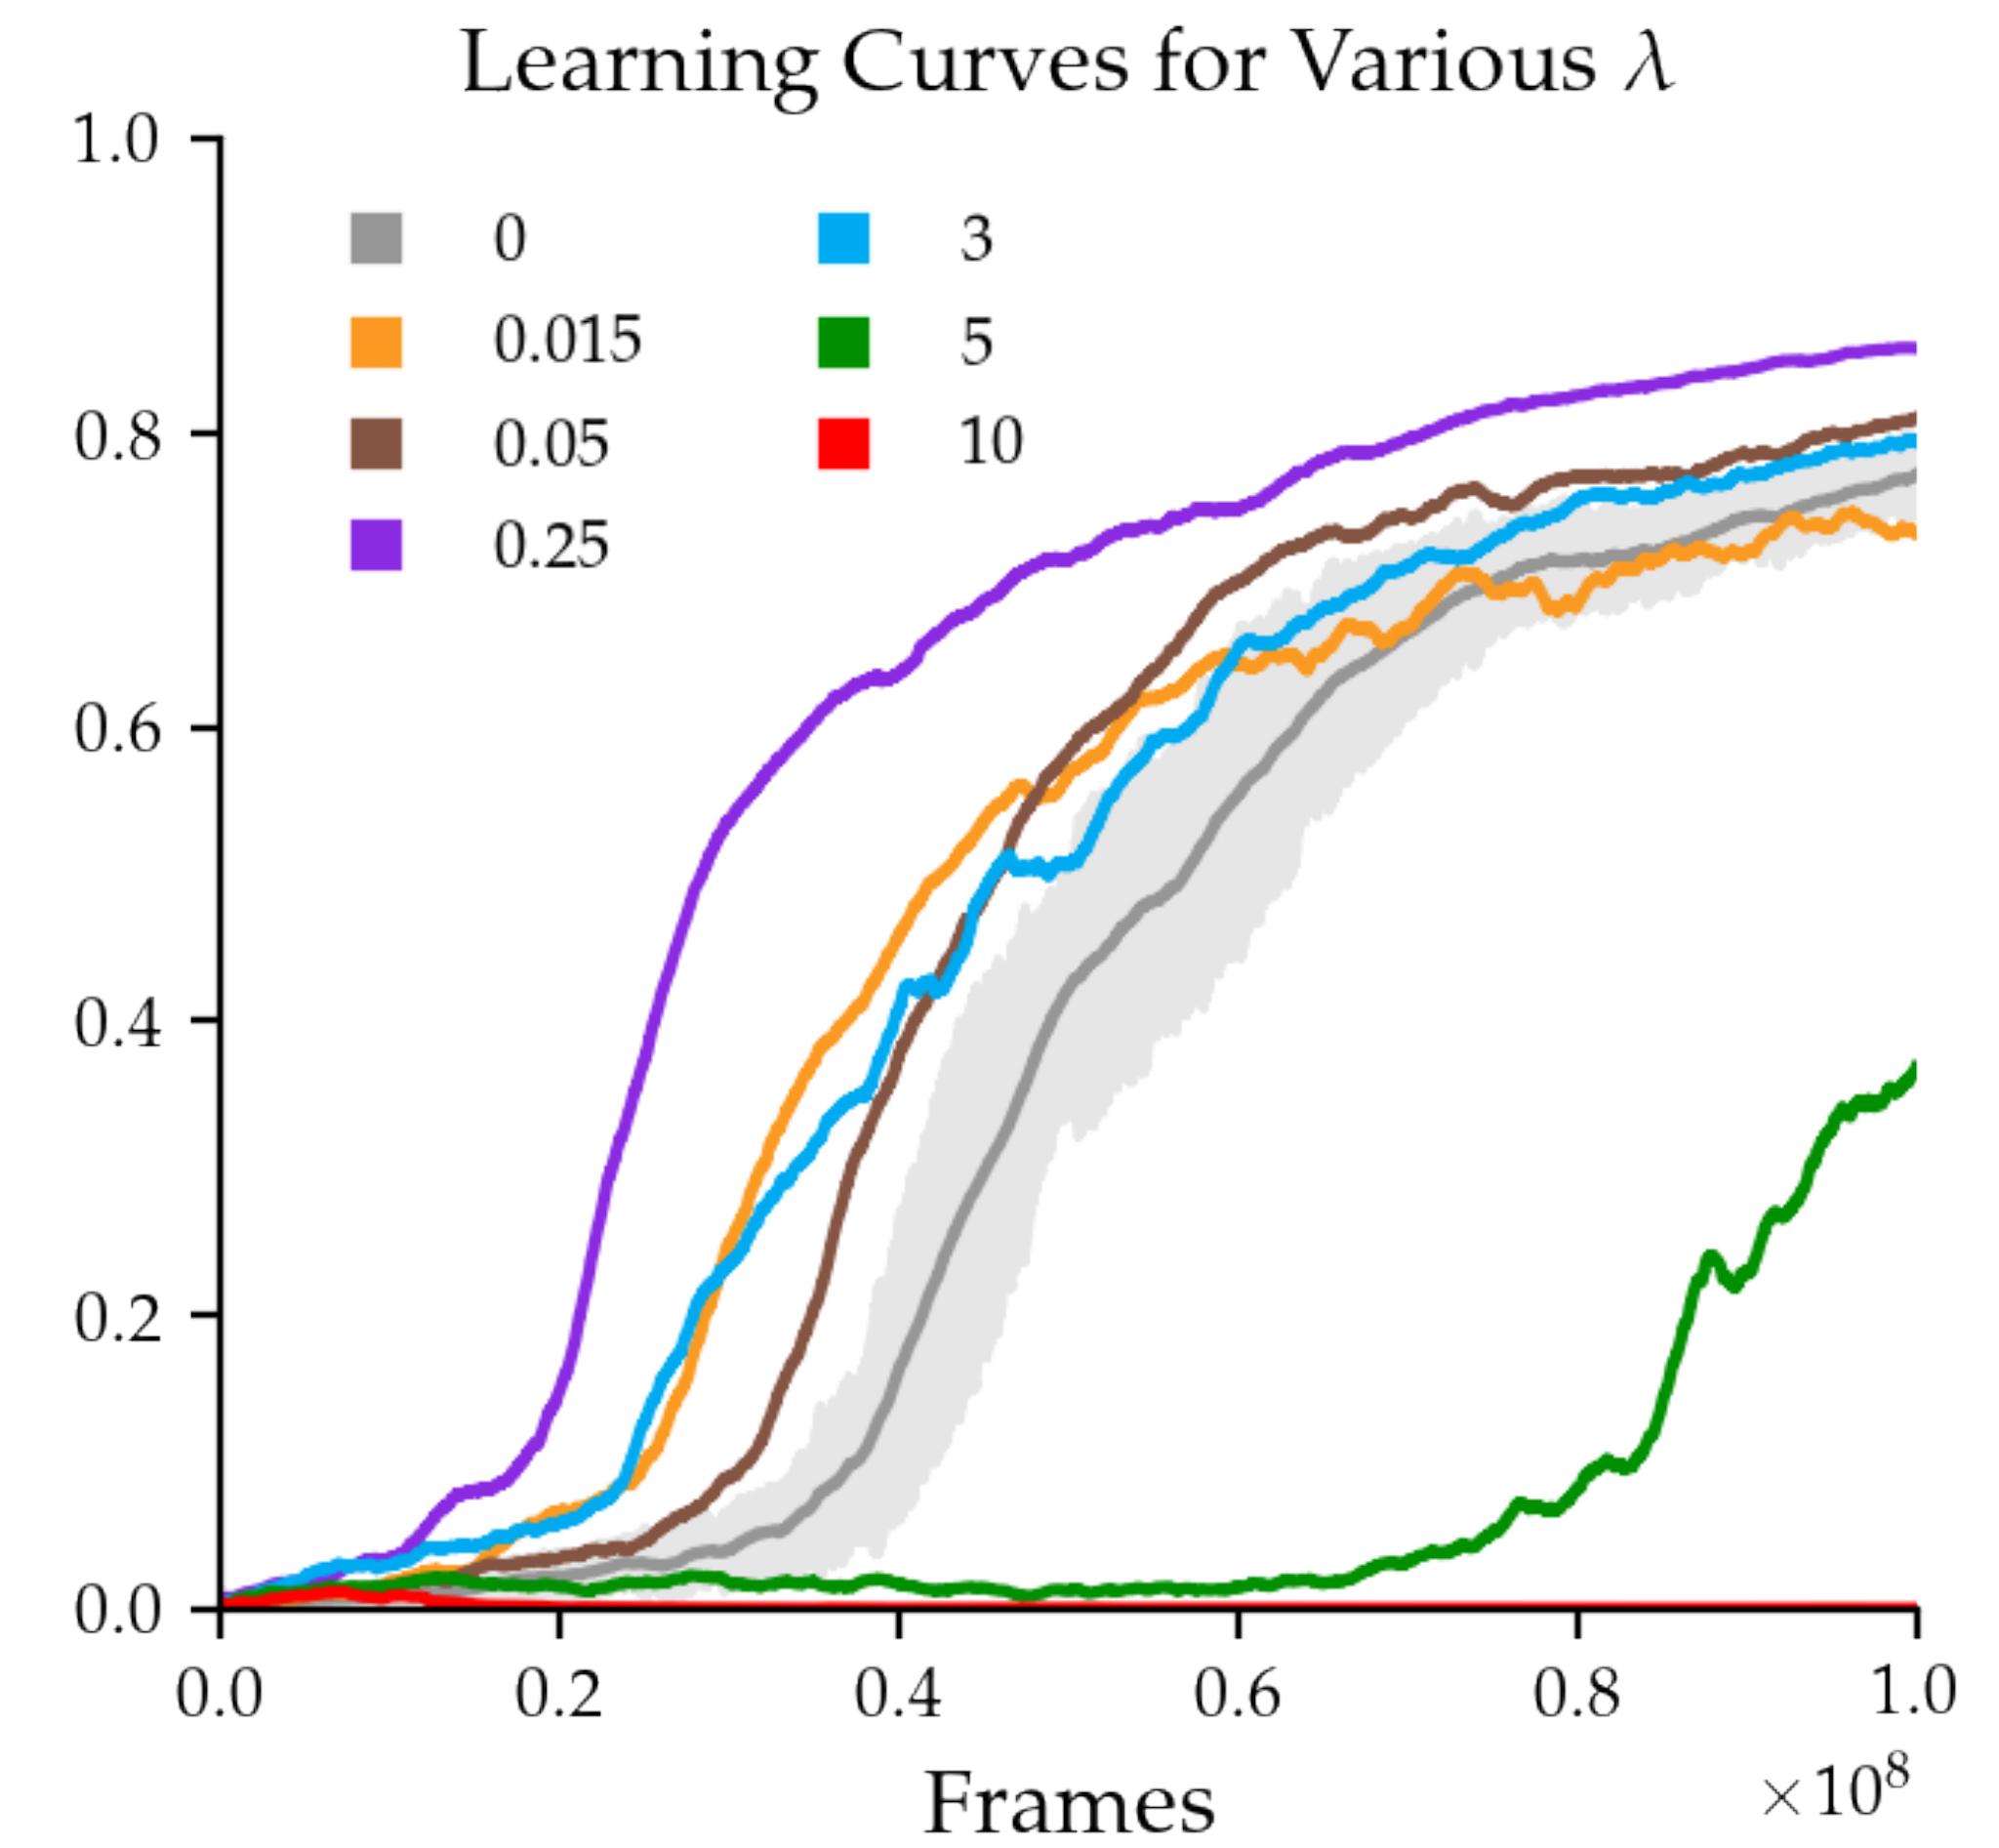


open the green door



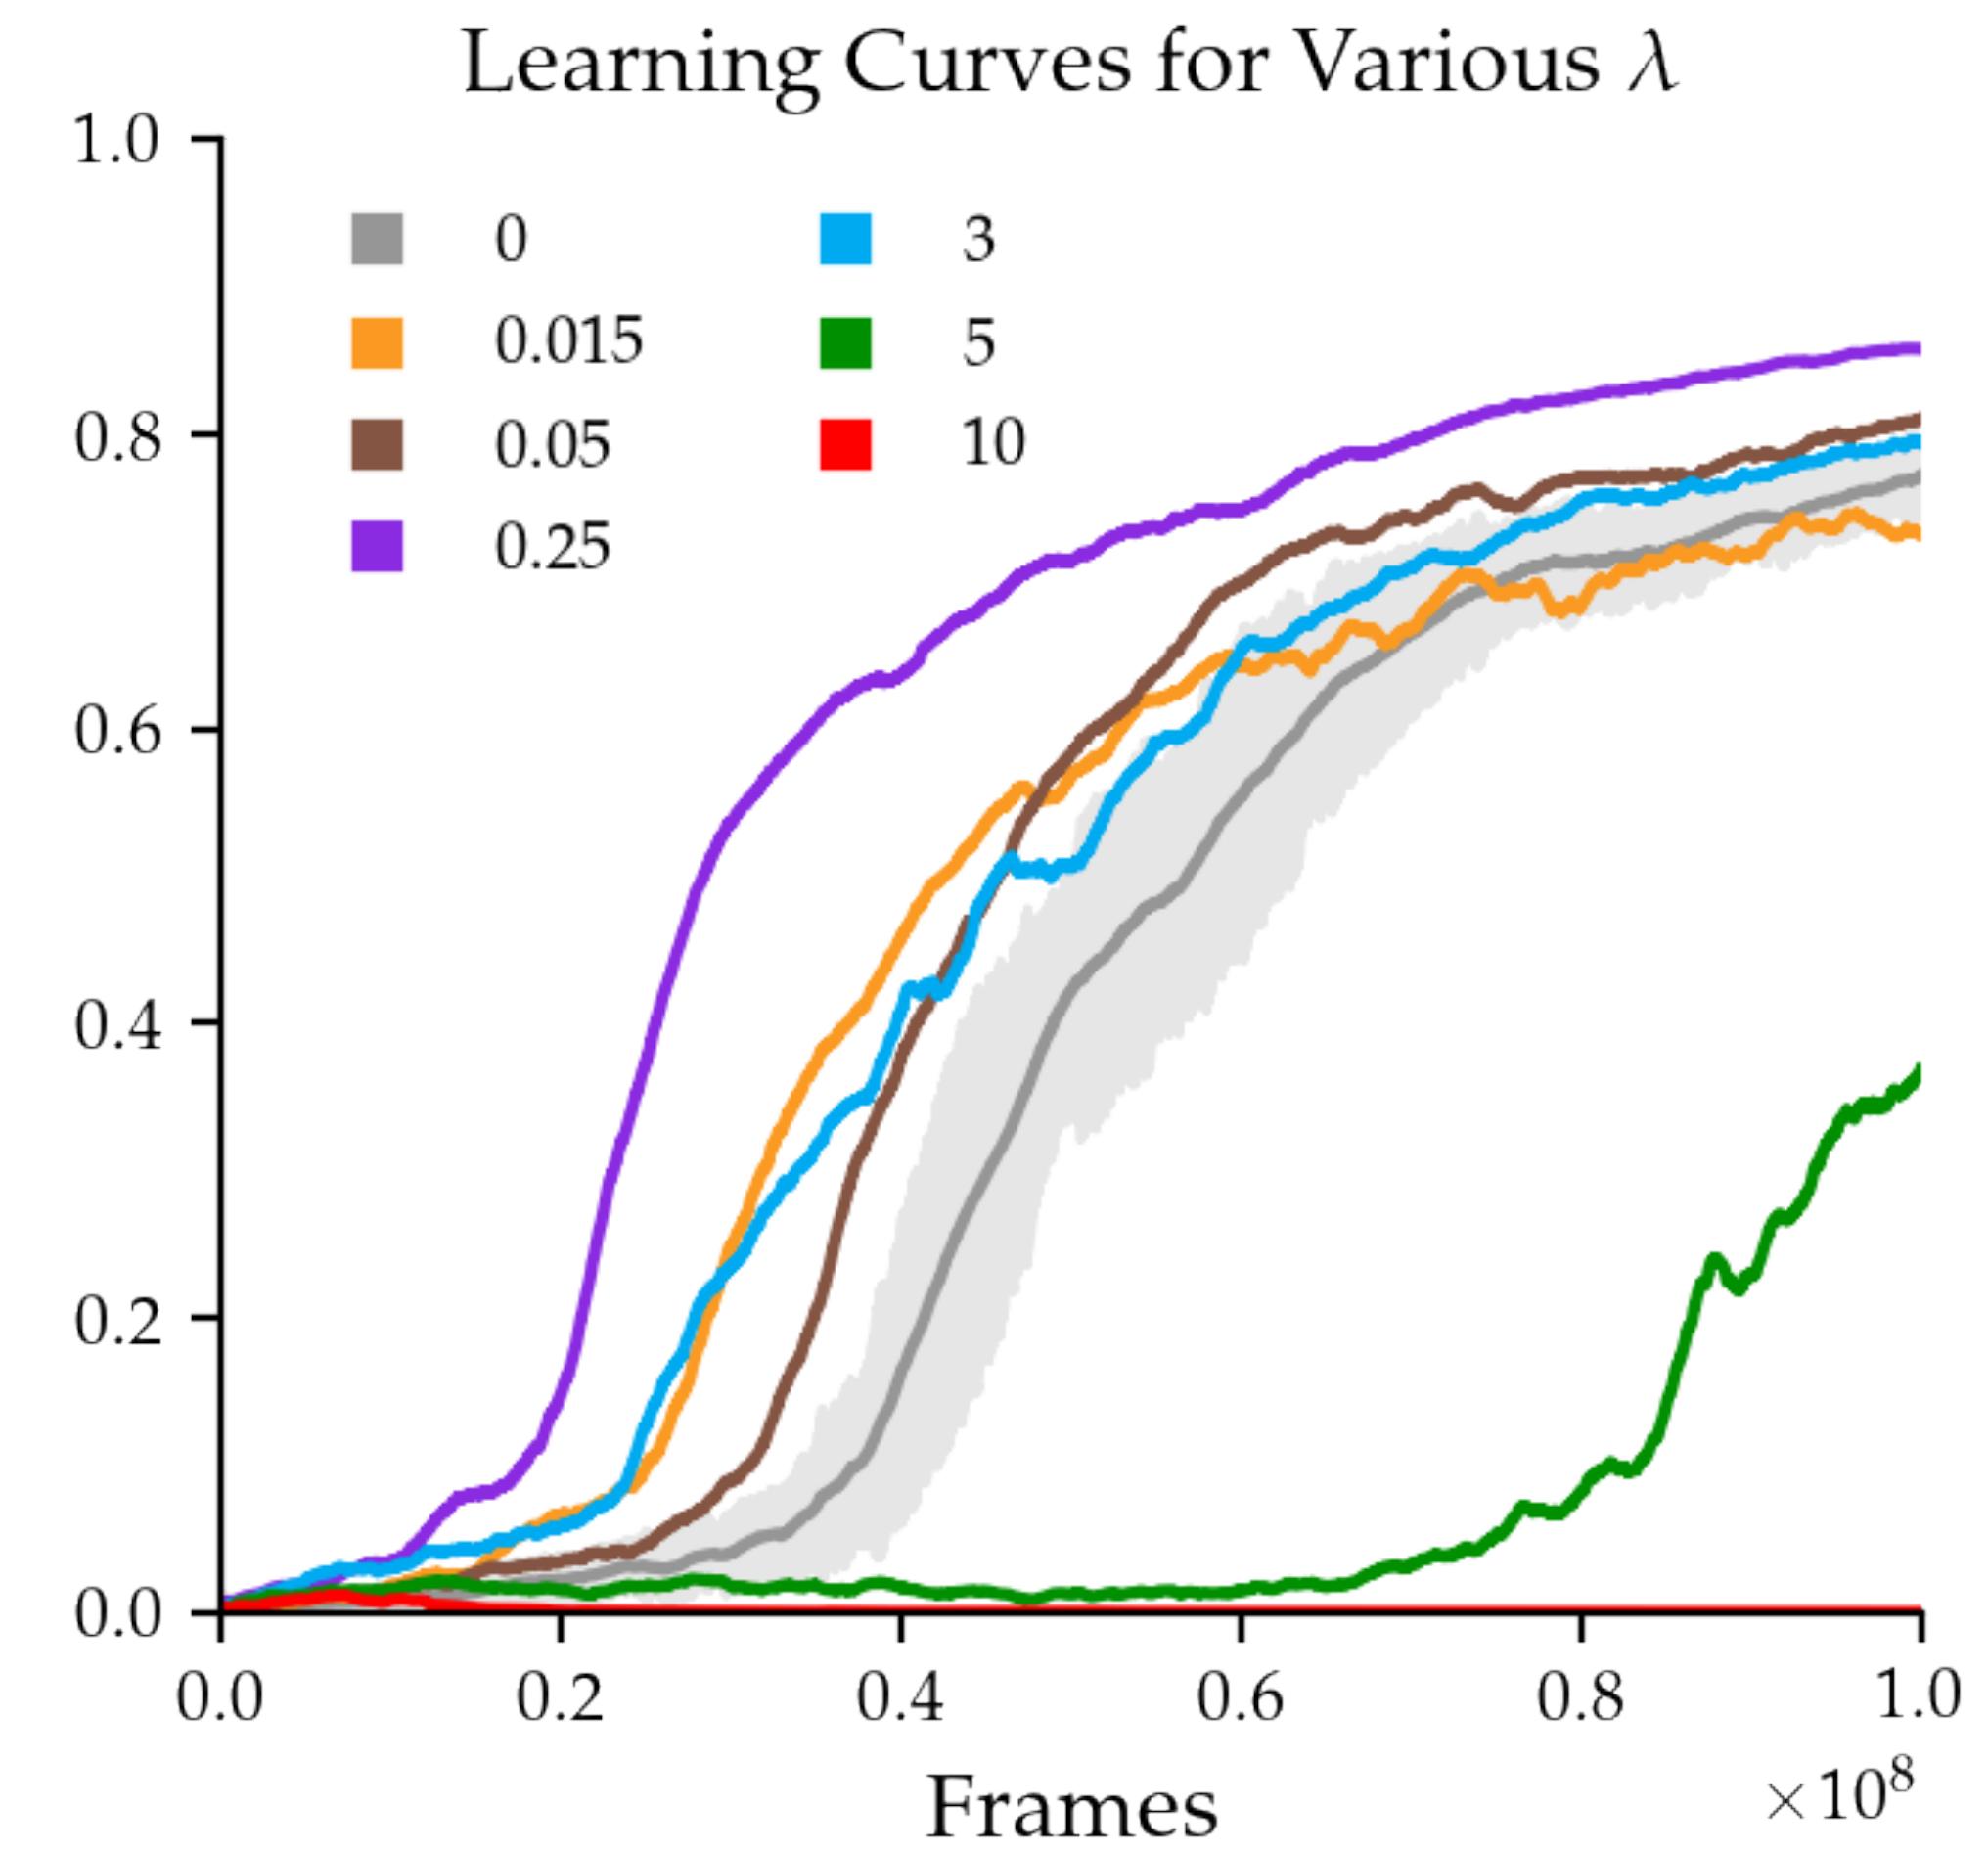
pick up a red box

Effect of λ



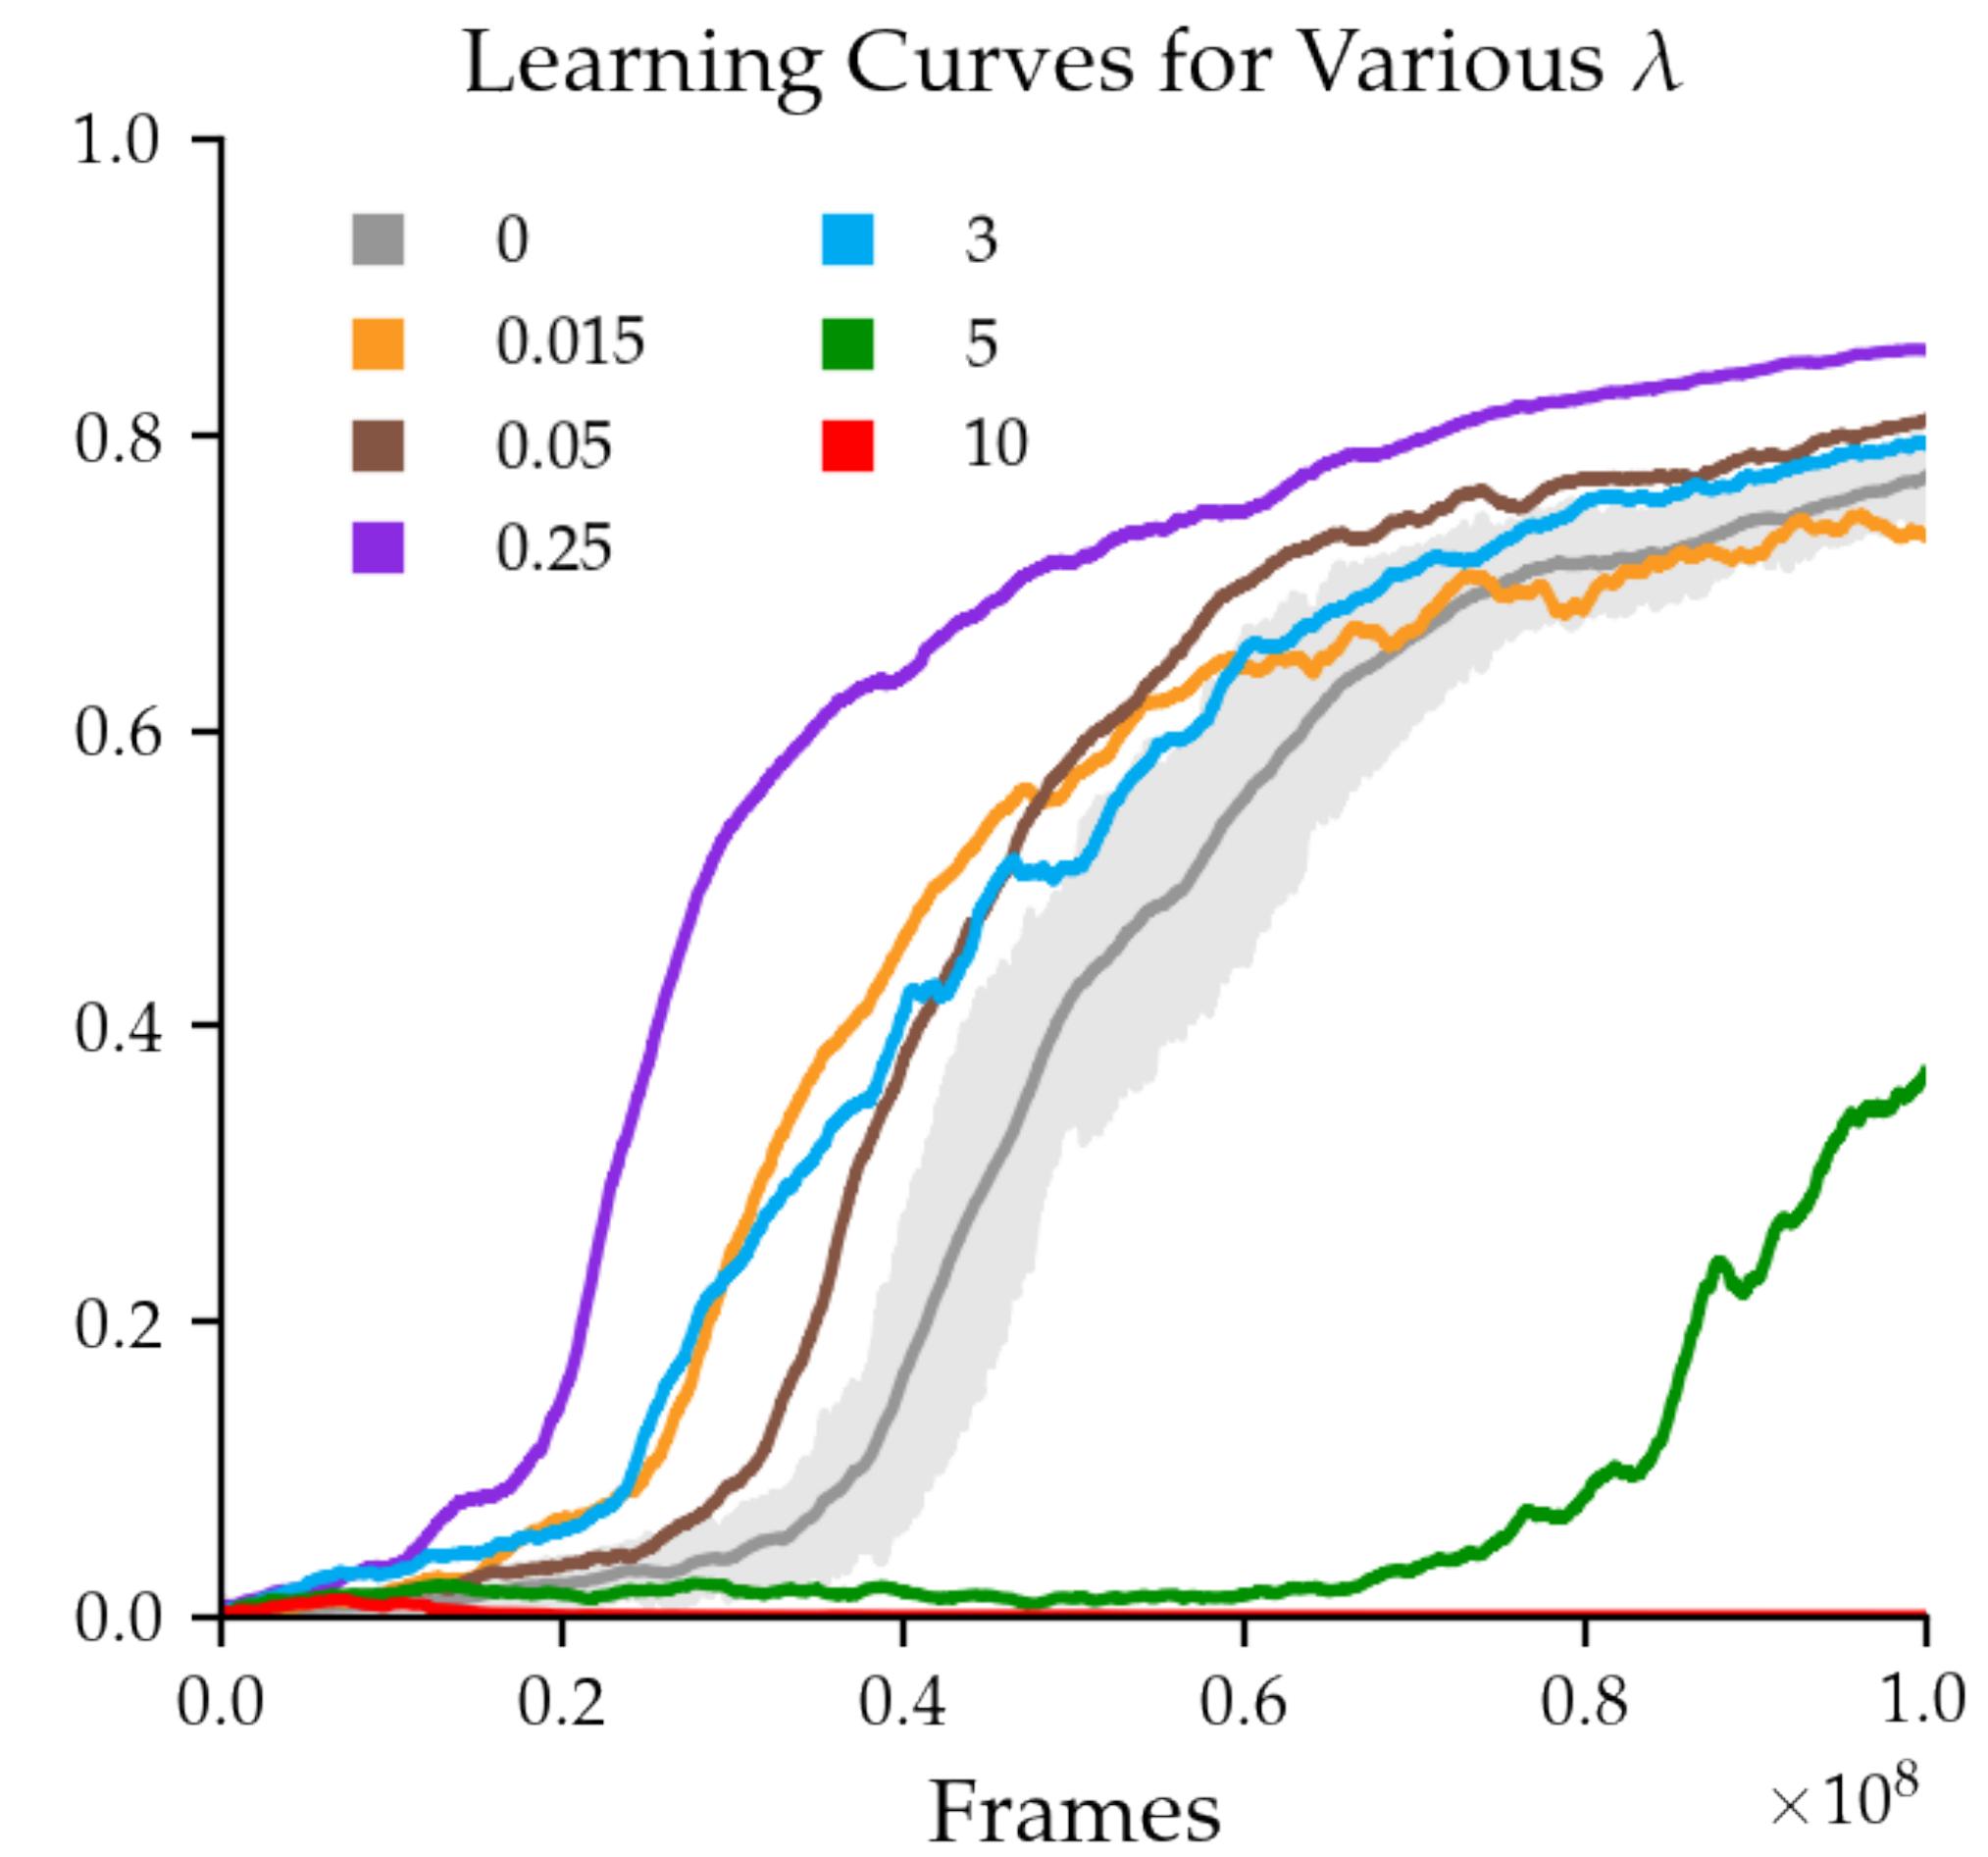
Effect of λ

- Set λ intelligently to increase sample efficiency

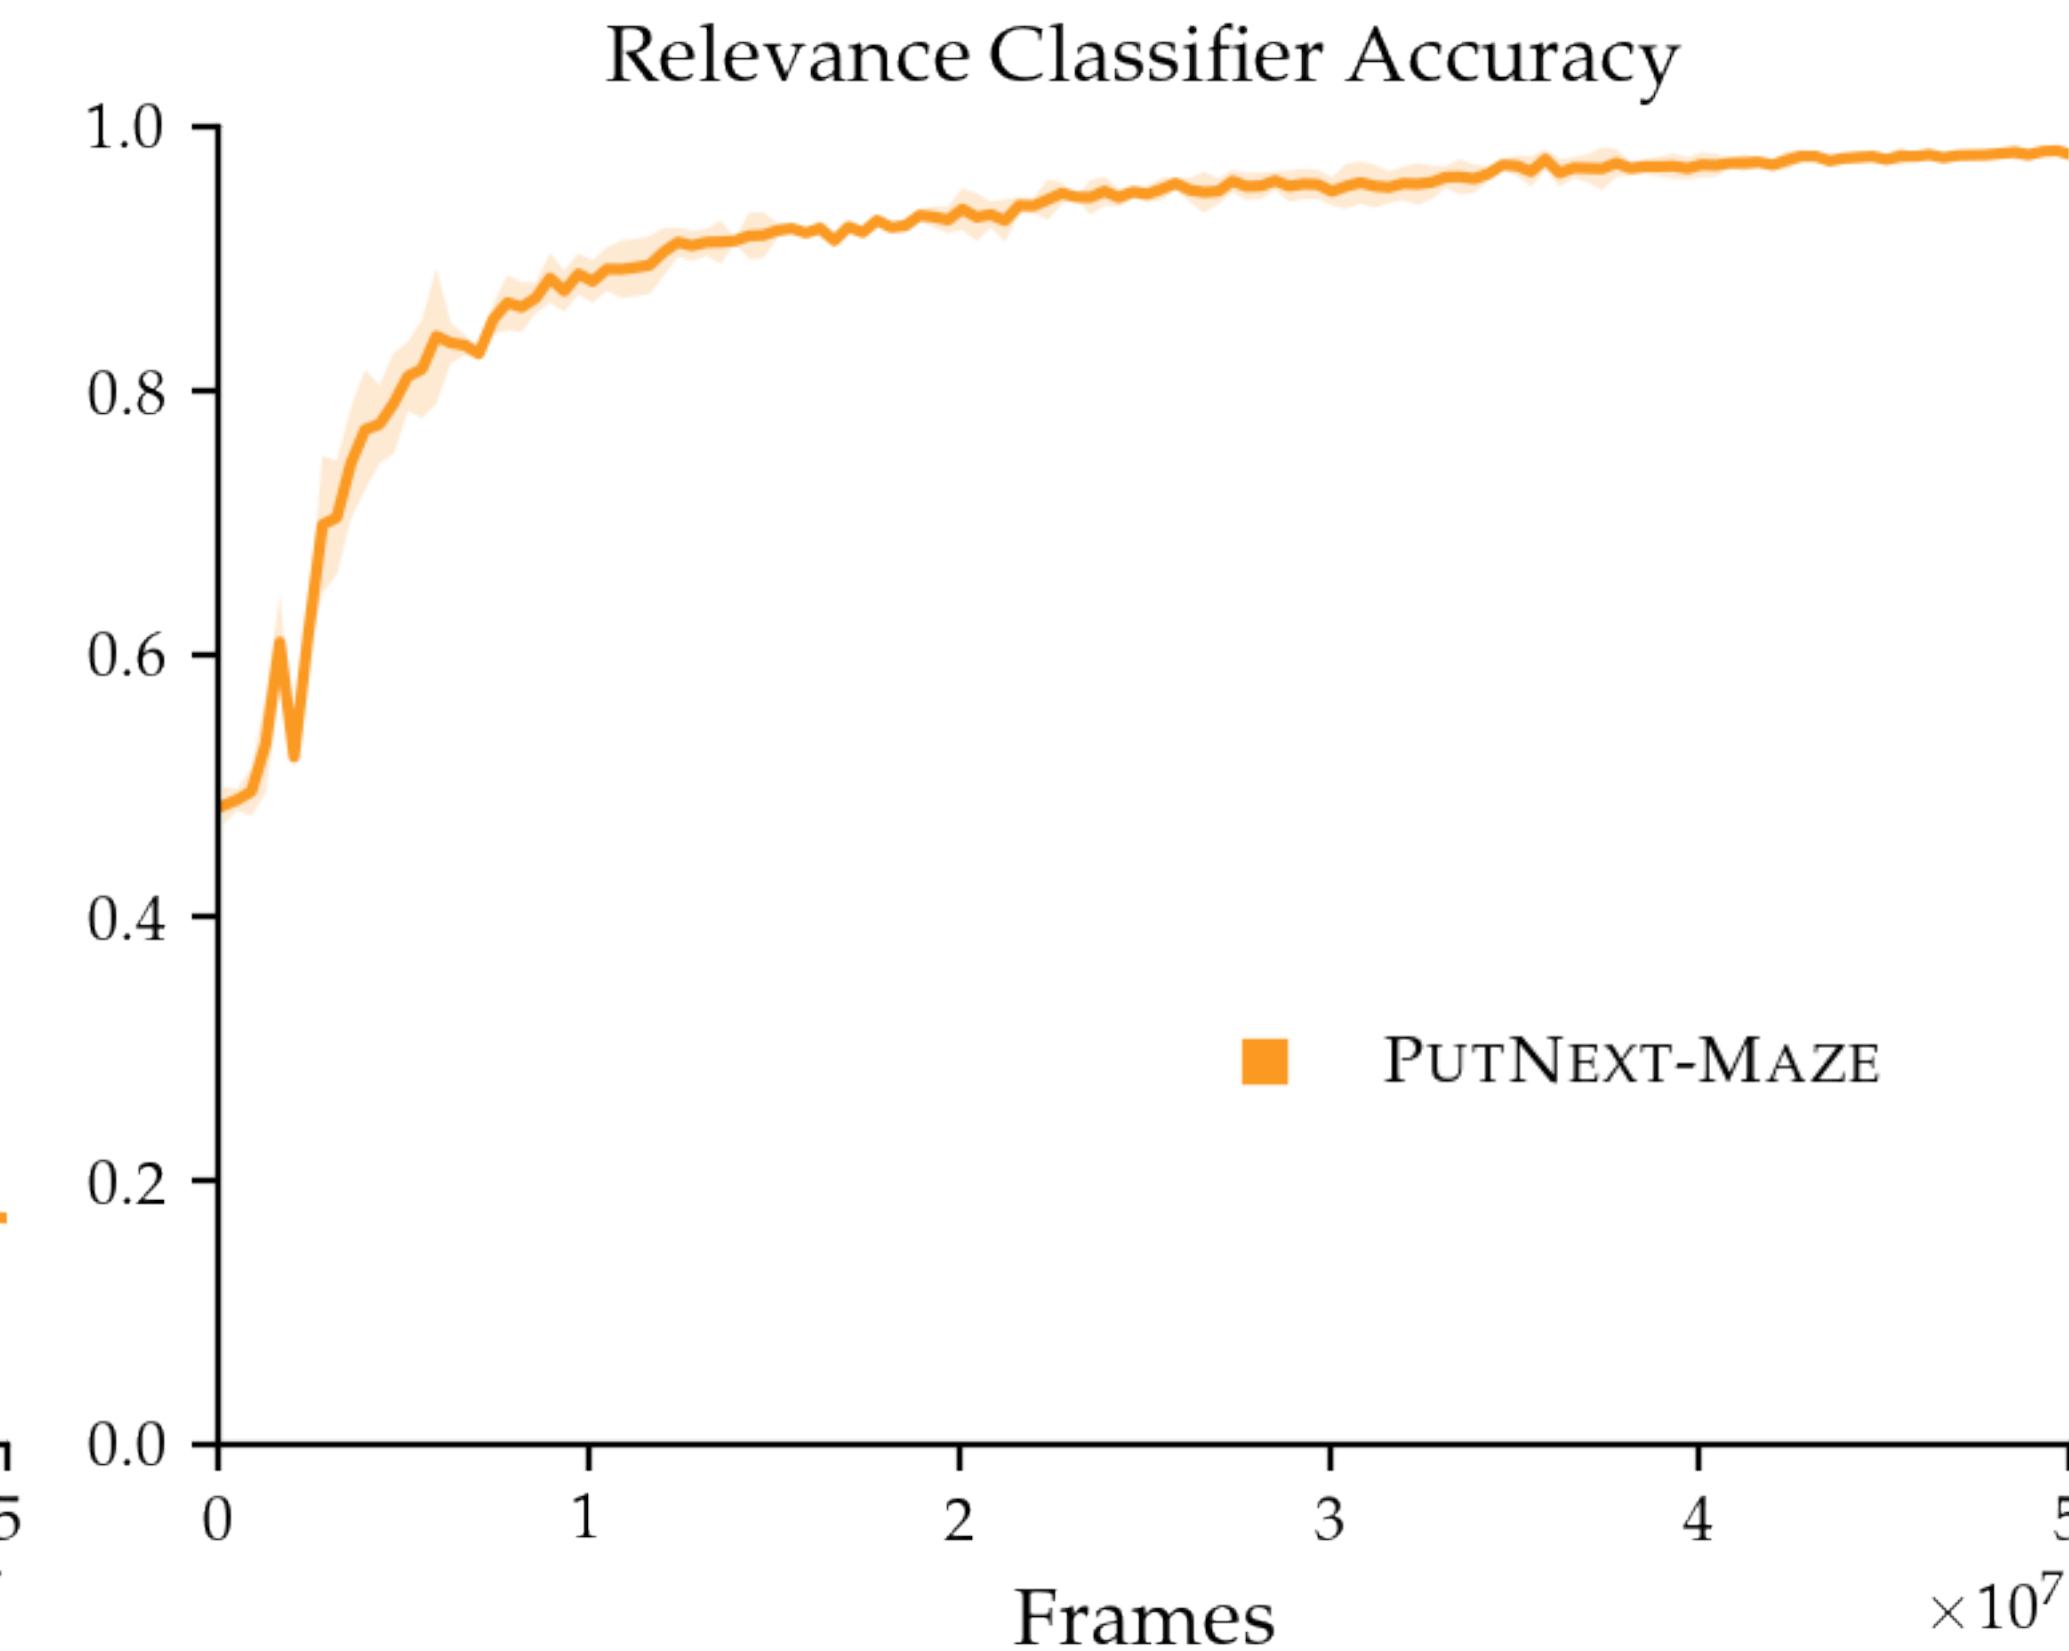
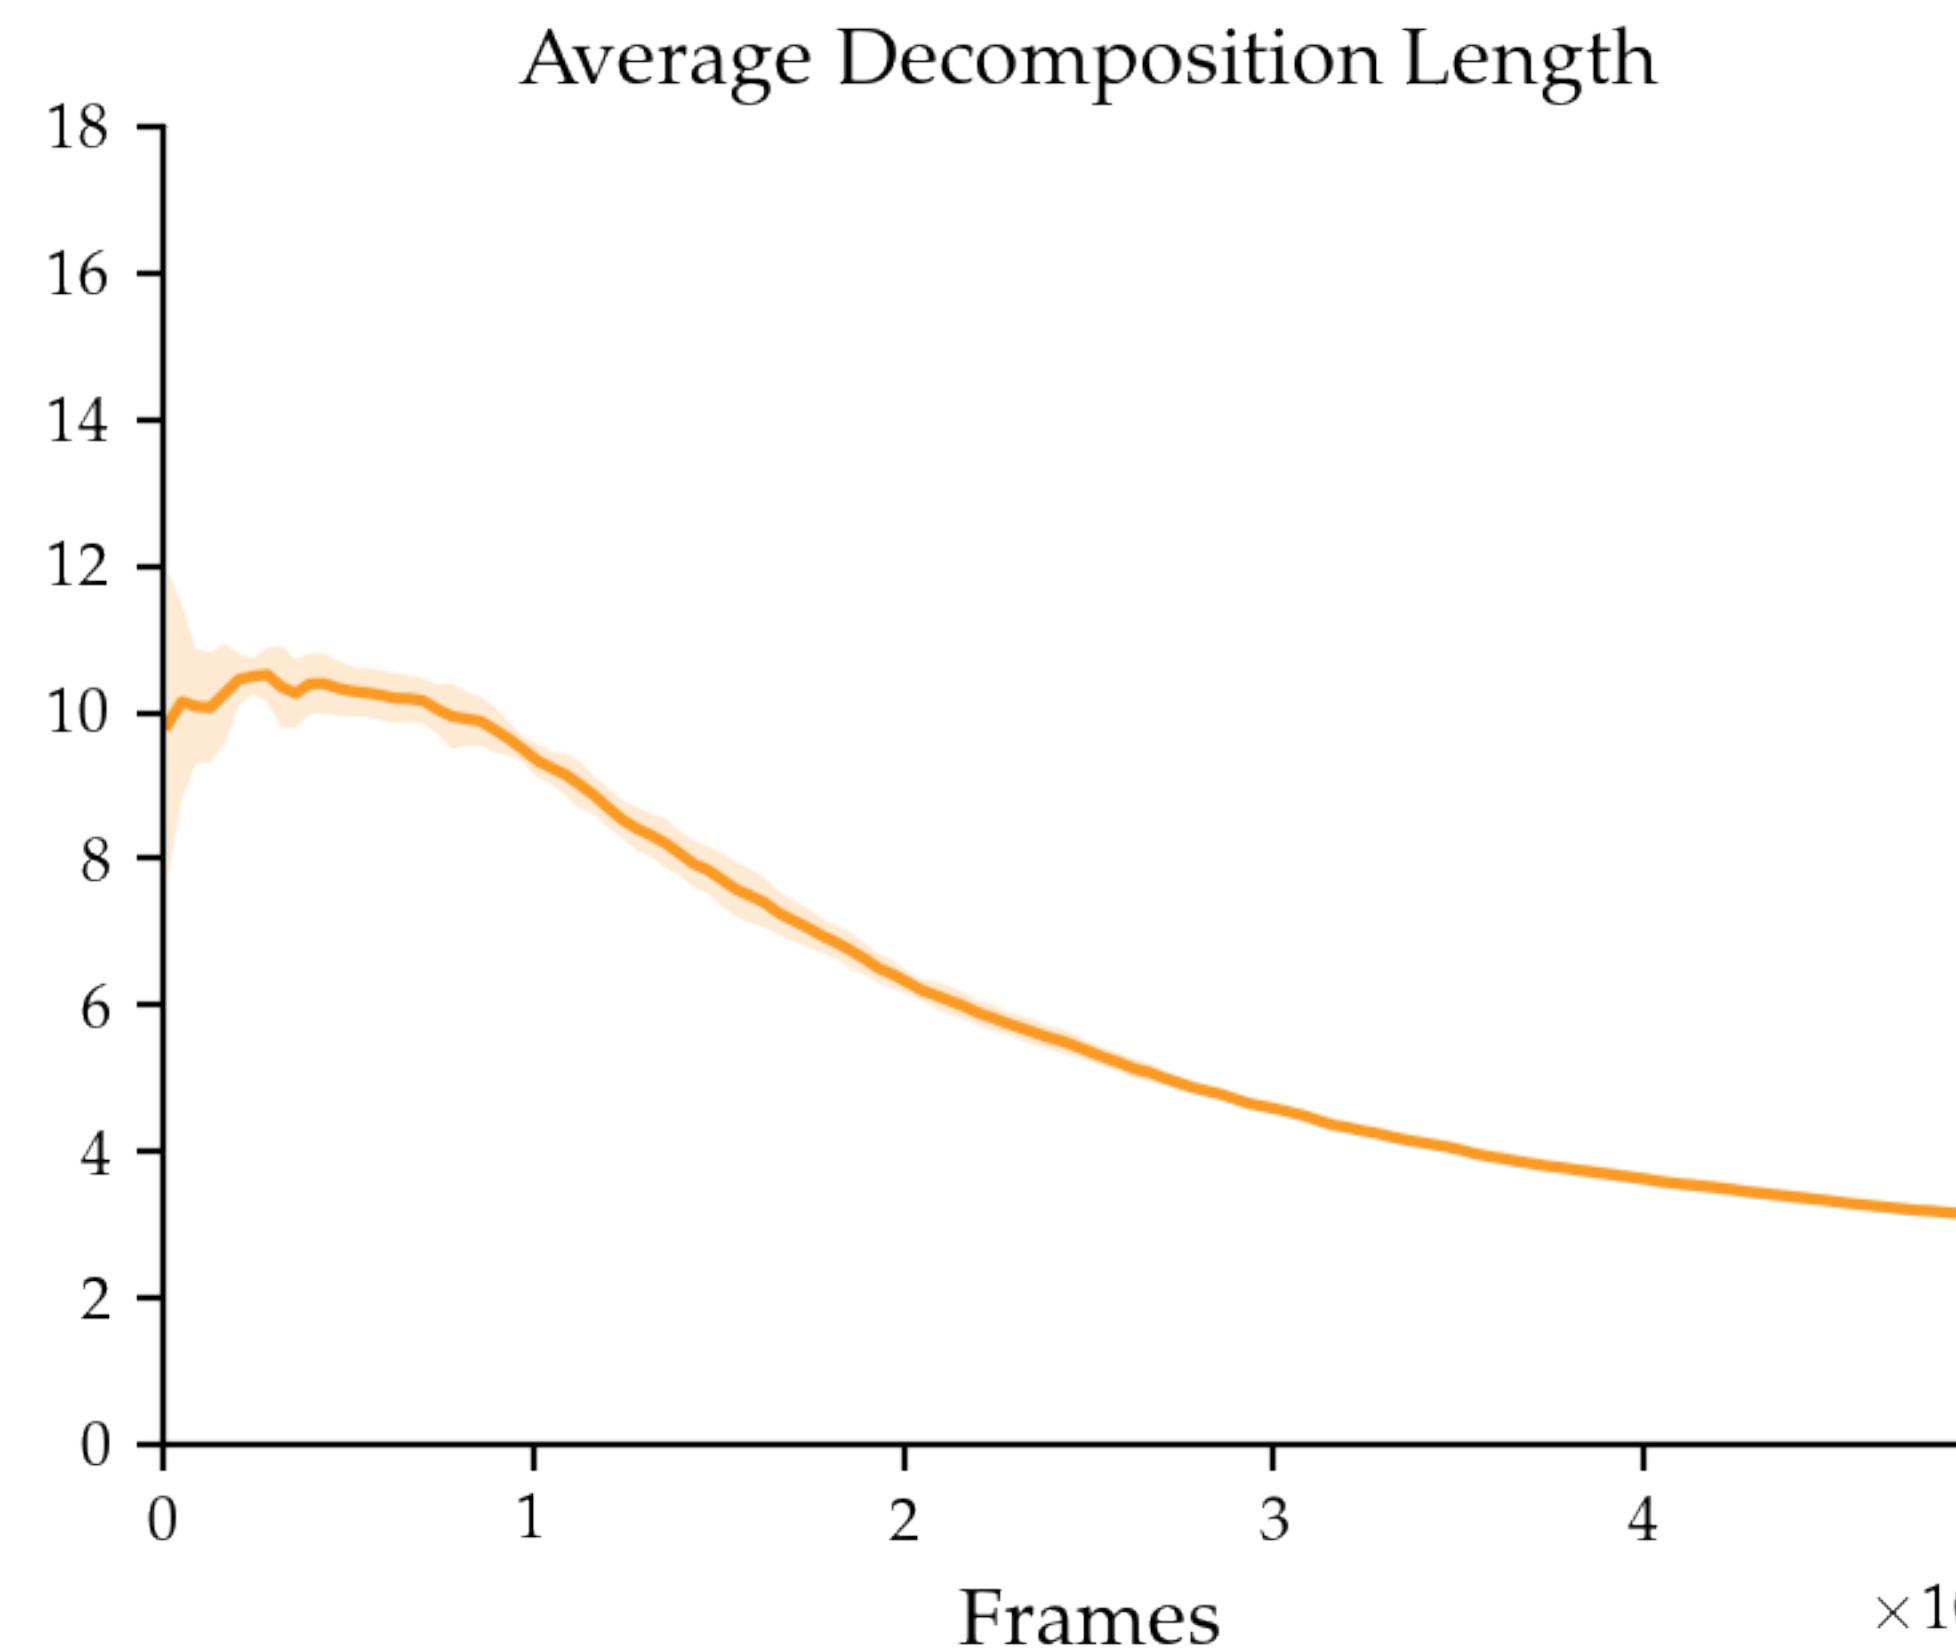


Effect of λ

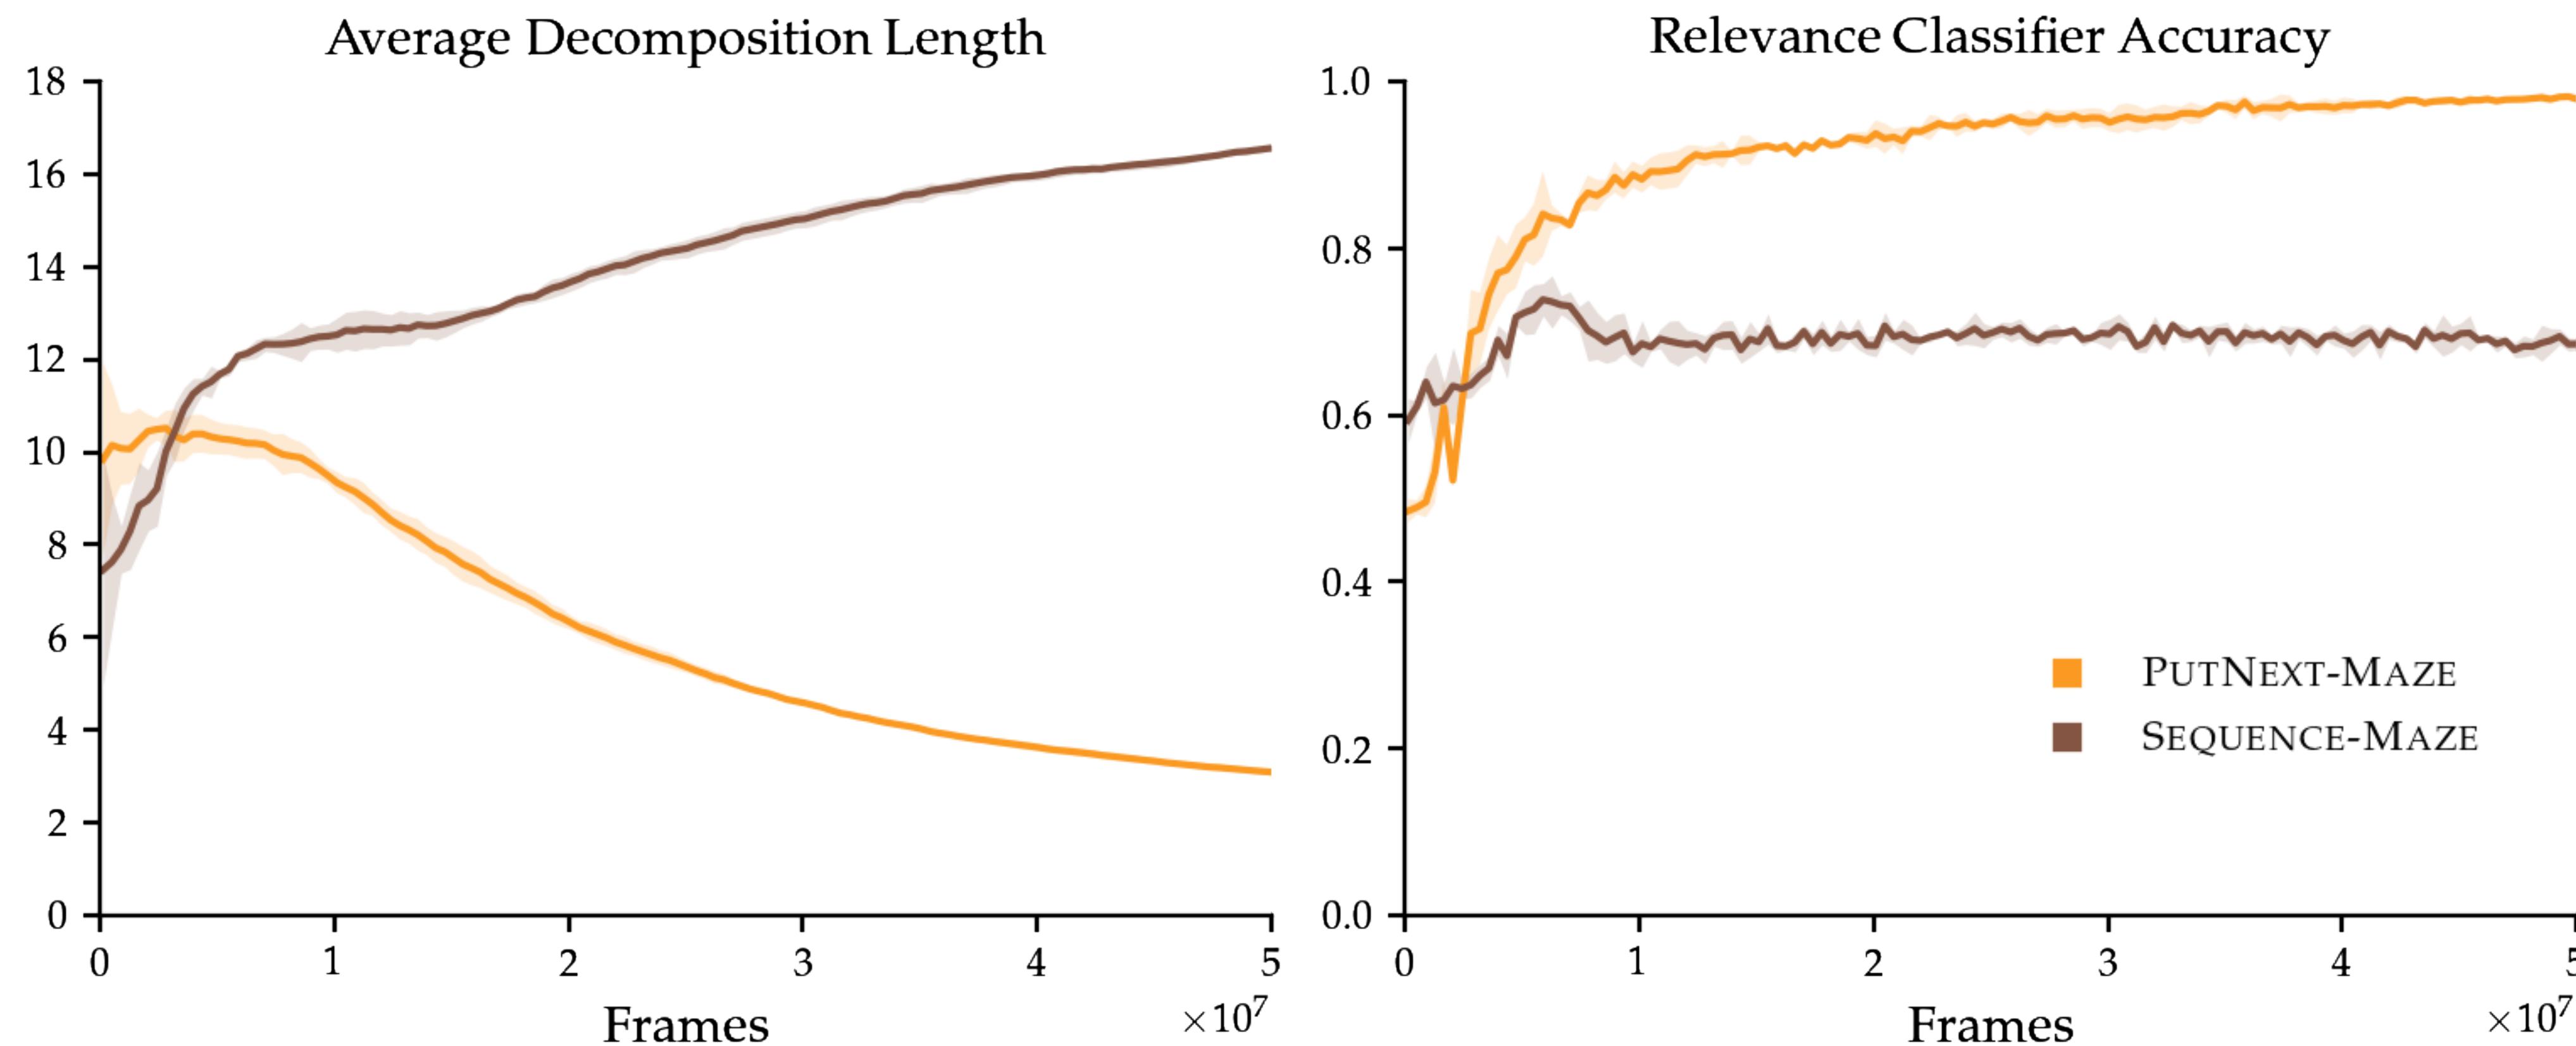
- Set λ intelligently to increase sample efficiency
- Large values of λ lead to unstable learning



Relevance Classifier Performance

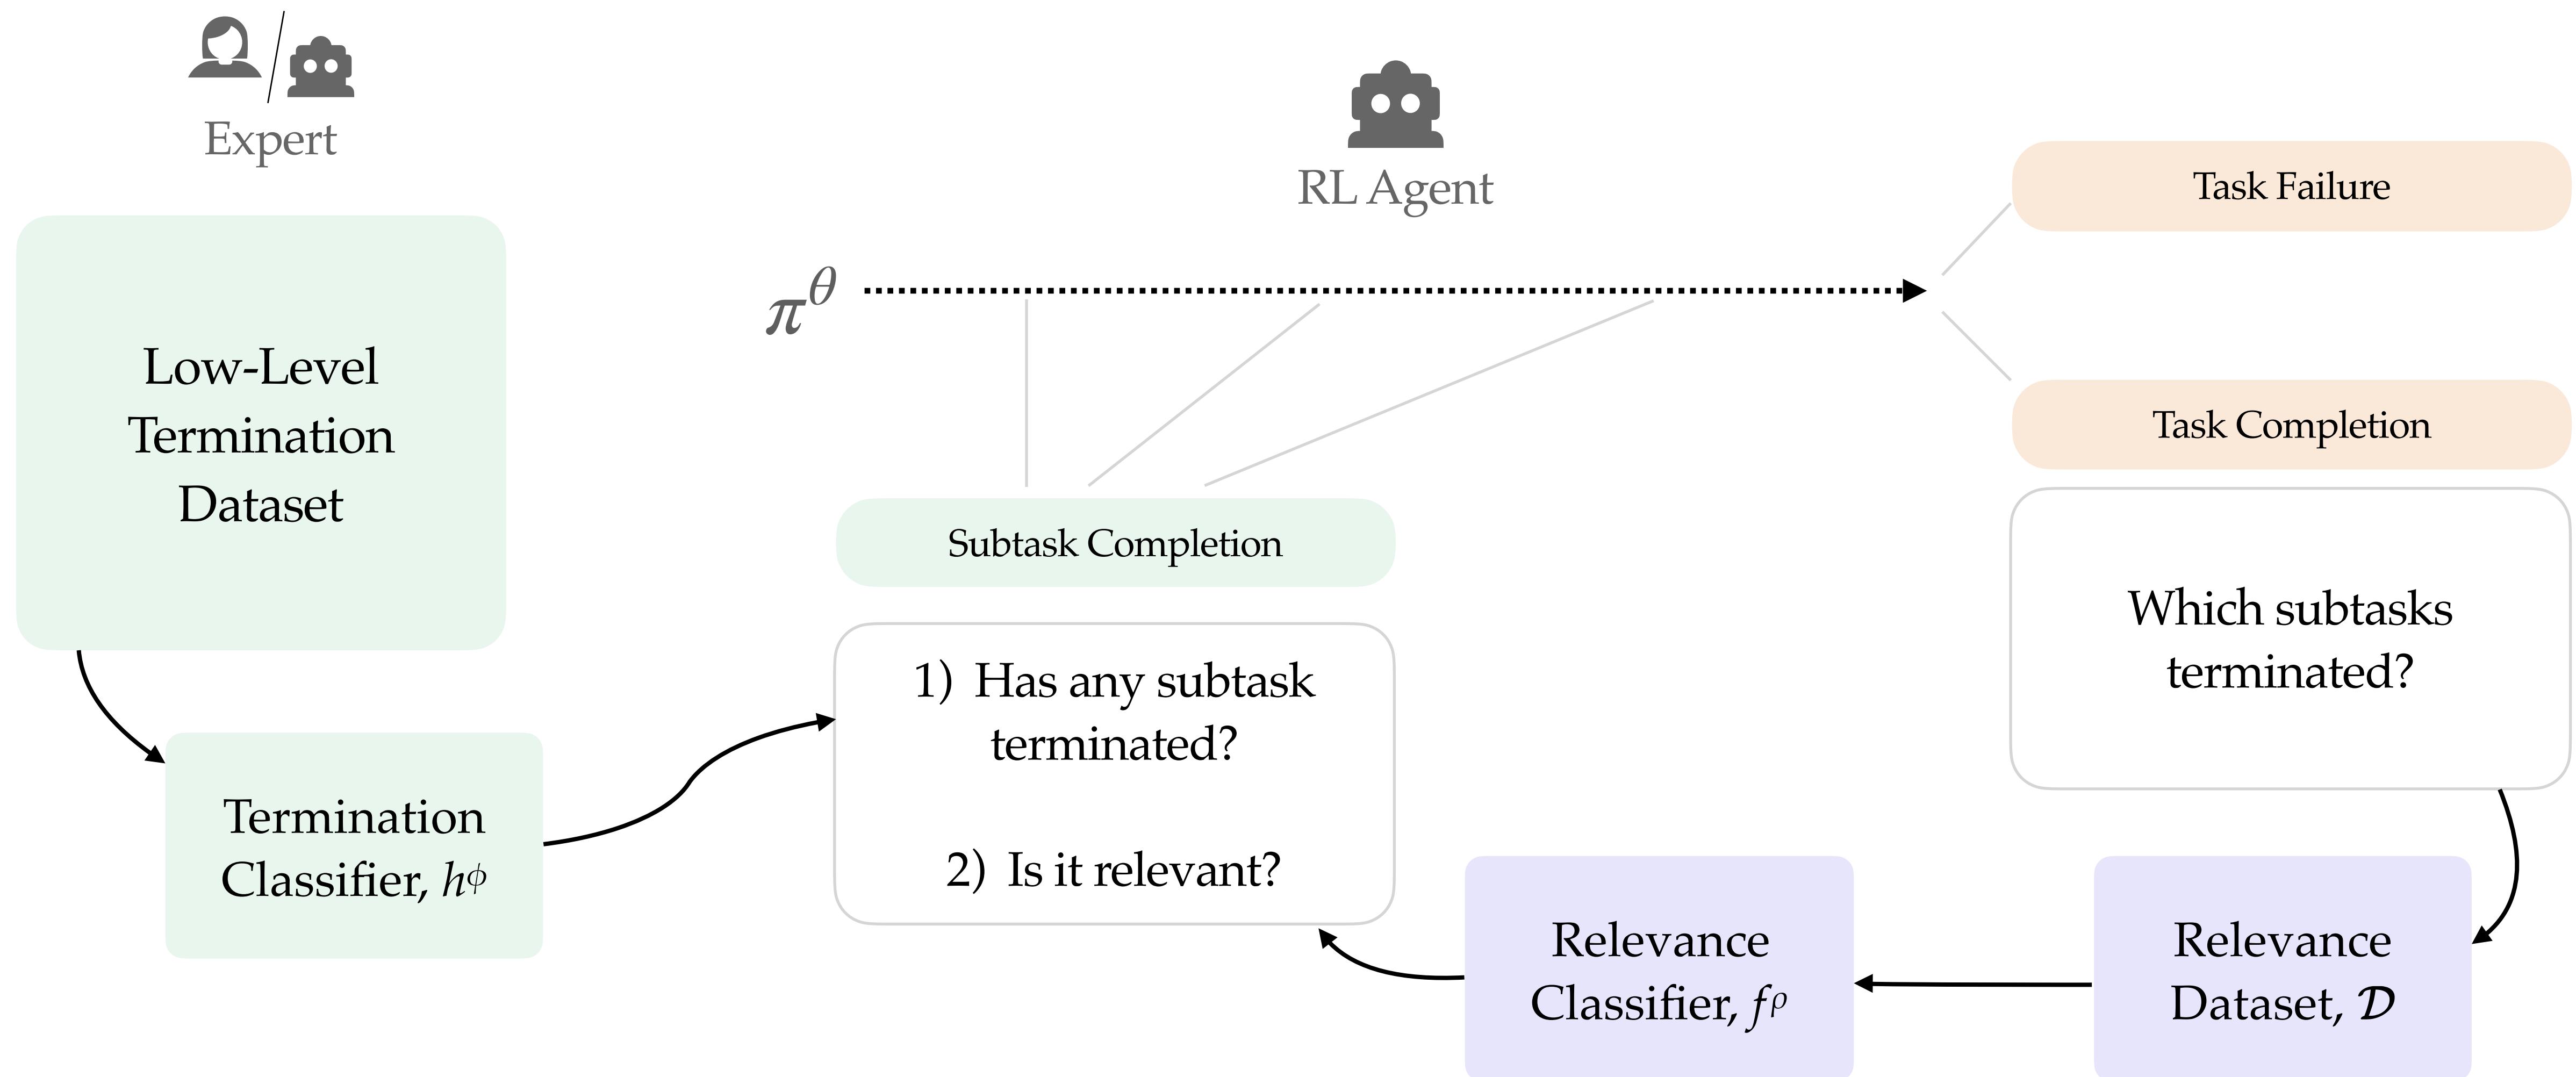


Relevance Classifier Performance





ELLA



ELLA

