



ADAPTIVE GREY-BOX FUZZ-TESTING WITH THOMPSON SAMPLING

*Sidd Karamcheti, David Rosenberg, Gideon Mann
Bloomberg - CTO Data Science*

FUZZING - A BACKGROUND

- Fuzzing is a technique for automated software testing.
 - Core Idea: Provide programs with unexpected inputs, with goal of finding bugs, maximizing coverage, etc.
- Different types of fuzzers and testing tools:
 - **Black-Box:** Assume no transparency into program.
 - Random generators (good for testing parsers)
 - **White-Box:** Assume lots of transparency (e.g. KLEE)
 - Suffer from path explosion as programs get big

GREY-BOX FUZZING & AFL

- Grey-Box Mutational Fuzzers

- Lightweight instrumentation, to check code path
- Mutate existing inputs to generate new test inputs
- Fast, efficient, and proven - can find lots of bugs

```
1: // Core Algorithm for American Fuzzy Lop (AFL)
2: // time: Fixed time window to fuzz (e.g. 24 hours)
3: // queue: Queue of inputs that exercise new code paths.
4: while time has not elapsed do
5:   parent, energy ← pick_input(queue)
6:   for i ∈ range(energy) do
7:     child ← parent
8:     for j ∈ 1 to sample_num_mutations() do
9:       mutation ← sample_mutation()
10:      site ← sample_mutation_site()
11:      child ← apply_mutation(mutation, child, site)
12:    end for
13:    path ← execute_path(child, code)
14:    if (path is new) then queue ← child
15:  end for
16: end while
```

- AFL is best of the bunch!

GREY-BOX FUZZING & AFL - HEURISTICS

```
1: // Core Algorithm for American Fuzzy Lop (AFL)
2: // time: Fixed time window to fuzz (e.g. 24 hours)
3: // queue: Queue of inputs that exercise new code paths.
4: while time has not elapsed do
5:   parent, energy ← pick_input(queue) → Weighted by scores (execution
6:   for i ∈ range(energy) do time, program depth)
7:     child ← parent
8:     for j ∈ 1 to sample_num_mutations() do → RANDOM!
9:       mutation ← sample_mutation()
10:      site ← sample_mutation_site()
11:      child ← apply_mutation(mutation, child, site)
12:    end for
13:    path ← execute_path(child, code)
14:    if (path is new) then queue ← child
15:  end for
16: end while
```

Q: How can we do better?

A: Data-Driven, Adaptive Control!

GREY-BOX FUZZING & AFL - RELATED WORK

```
1: // Core Algorithm for American Fuzzy Lop (AFL)
2: // time: Fixed time window to fuzz (e.g. 24 hours)
3: // queue: Queue of inputs that exercise new code paths.
4: while time has not elapsed do
5:   parent, energy ← pick_input(queue)
6:   for i ∈ range(energy) do
7:     child ← parent
8:     for j ∈ 1 to sample_num_mutations() do
9:       mutation ← sample_mutation()
10:      site ← sample_mutation_site()
11:      child ← apply_mutation(mutation, child, site)
12:    end for
13:    path ← execute_path(child, code)
14:    if (path is new) then queue ← child
15:  end for
16: end while
```

AFLFast

Coverage-Based Grey-Box Fuzzing

Böhme et. al, 2016

AFLGo

Directed Greybox Fuzzing

Böhme et. al, 2017

GREY-BOX FUZZING & AFL - RELATED WORK

```
1: // Core Algorithm for American Fuzzy Lop (AFL)
2: // time: Fixed time window to fuzz (e.g. 24 hours)
3: // queue: Queue of inputs that exercise new code paths.
4: while time has not elapsed do
5:   parent, energy ← pick_input(queue)
6:   for i ∈ range(energy) do
7:     child ← parent
8:     for j ∈ 1 to sample_num_mutations() do
9:       mutation ← sample_mutation()
10:      site ← sample_mutation_site()
11:      child ← apply_mutation(mutation, child, site)
12:    end for
13:    path ← execute_path(child, code)
14:    if (path is new) then queue ← child
15:  end for
16: end while
```

Neural Byte Sieve for Fuzzing
Rajpal et. al., 2017

FairFuzz: Targeting Rare Branches
to Rapidly Increase Greybox
Fuzz Testing Coverage
Lemieux and Sen, 2017

GREY-BOX FUZZING & AFL - WHAT'S NEXT?

```
1: // Core Algorithm for American Fuzzy Lop (AFL)
2: // time: Fixed time window to fuzz (e.g. 24 hours)
3: // queue: Queue of inputs that exercise new code paths.
4: while time has not elapsed do
5:   parent, energy ← pick_input(queue)
6:   for i ∈ range(energy) do
7:     child ← parent
8:     for j ∈ 1 to sample_num_mutations() do
9:       mutation ← sample_mutation() → ???
10:      site ← sample_mutation_site()
11:      child ← apply_mutation(mutation, child, site)
12:    end for
13:    path ← execute_path(child, code)
14:    if (path is new) then queue ← child
15:  end for
16: end while
```

Mutation Operation	Notes
Bitflips	Flip single bit
Interesting Values	NULL, -1, 0, etc.
Addition	Add random value
Subtraction	Subtract random value
Random Value	Insert random value
Deletion	Delete from parent
Cloning	Clone/add from parent
Overwrite	Replace with random
Extra Overwrite	Extras: strings scraped
Extra Insertion	from binary

MOTIVATING EXAMPLE

Mutation Operation	Notes
Bitflips	Flip single bit
Interesting Values	NULL, -1, 0, etc.
Addition	Add random value
Subtraction	Subtract random value
Random Value	Insert random value
Deletion	Delete from parent
Cloning	Clone/add from parent
Overwrite	Replace with random
Extra Overwrite	Extras: strings scraped from binary
Extra Insertion	

```
// Runs the ASCII Content Server
int main(void) {
    // Initialize server
    InitializeTree();

    // Respond to commands
    Command command, int more = {}, 1;
    while (more) {
        ReceiveCommand(&command, &more);
        HandleCommand(&command);
    }
    return 0;
}

// Receives and parses an incoming command
int ReceiveCommand(Command *command, int *more) {
    char buffer[64], size_t bytes_received;
    read_until(buffer, ':', sizeof(buffer));

    switch (buffer) {
        case "REQUEST": command->c = REQUEST; break;
        case "QUERY": command->c = QUERY; break;
        case "SEND": command->c = SEND; break;
        case "VISUALIZE": command->c = VISUALIZE; break;
        case "INTERACT": command->c = INTERACT; break;
        default: more = 0; return -1;
    }
    parse_data(command, read_rest(), more);
    return 0;
}
```

Q: How do we identify the right mutators?

A: Learn it!

Best indicator of future success is past success

PICK_MUTATION AS A MULTI-ARMED BANDIT

- Multi-Armed Bandit Problem
 - k "arms" (mutators), each with different probability of paying out (discovering new code path)
 - Starts out unobserved
 - Need to discover the "best" arm (or best distribution over arms) to maximize payout
- Requires a balance between exploration and exploitation

THOMPSON SAMPLING FOR EXPLORATION/EXPLOITATION

- Exploration-Exploitation as Bayesian Posterior Estimation
 - Each mutator has prior $\pi(\theta)$ - draws parameterize Bernoulli distribution (0/1 reward)
 - $\pi(\theta) \propto \theta^{\alpha-1} (1 - \theta)^{\beta-1}$ (Beta-Bernoulli form)
 - Collect data D for fixed interval (we used 3 minutes)
 - Count how many times each mutator was used in generating a "successful" input (S) vs otherwise (F)
 - Compute Posterior with new info
 - $\pi(\theta | D) \propto \theta^{\alpha-1+S} (1 - \theta)^{\beta-1+F}$
 - Use posteriors for each arm to obtain mutator distribution

EXPERIMENTS - DARPA CGC & LAVA-M

- We test our approach on two datasets:
 - DARPA Cyber Grand Challenge Binaries
 - Set of 200 binaries released by DARPA
 - Each binary has a real "bug" added by a human user
 - We utilize the 150 binaries that read from STDIN
 - LAVA-M Binaries
 - 4 Binaries from Coreutils
 - Injected with 100s of synthetic bugs
 - If $\text{MAGIC_1} < \text{INPUT}[10:18] < \text{MAGIC_2}$: Crash

EXPERIMENTS - BASELINES

- We utilize 3 baselines in comparison to our Thompson Sampling approach:
 - AFL (Vanilla): Same as original algorithm, with extra deterministic step.
 - FidgetyAFL (Havoc): Original algorithm, implements Böhme et. al Power Schedule for input selection!
 - Empirical: Test "adaptive" in Thompson Sampling
 - On 75 of CGC binaries, estimate "empirical" mutator distribution

RESULTS - CGC BINARIES

- CGC Binaries (on 75 test programs):

	6 hr	12 hr	18 hr	24 hr
AFL	0.64 ± 0.03	0.63 ± 0.03	0.63 ± 0.03	0.63 ± 0.03
FidgetyAFL	0.84 ± 0.02	0.84 ± 0.02	0.85 ± 0.02	0.84 ± 0.02
Empirical	0.85 ± 0.02	0.86 ± 0.02	0.86 ± 0.02	0.87 ± 0.02
Thompson	0.91 ± 0.02	0.92 ± 0.02	0.92 ± 0.02	0.93 ± 0.02

Relative Coverage Statistics (# paths discovered / max)

	Crashes	Wins / FidgetyAFL	Wins / All
AFL	554	18	4
FidgetyAFL	780	—	14
Empirical	766	41	5
Thompson	1336	52	47

Crash Statistics (Unique Paths triggering Crash)

RESULTS - LAVA-M BINARIES

➤ LAVA-M Binaries:

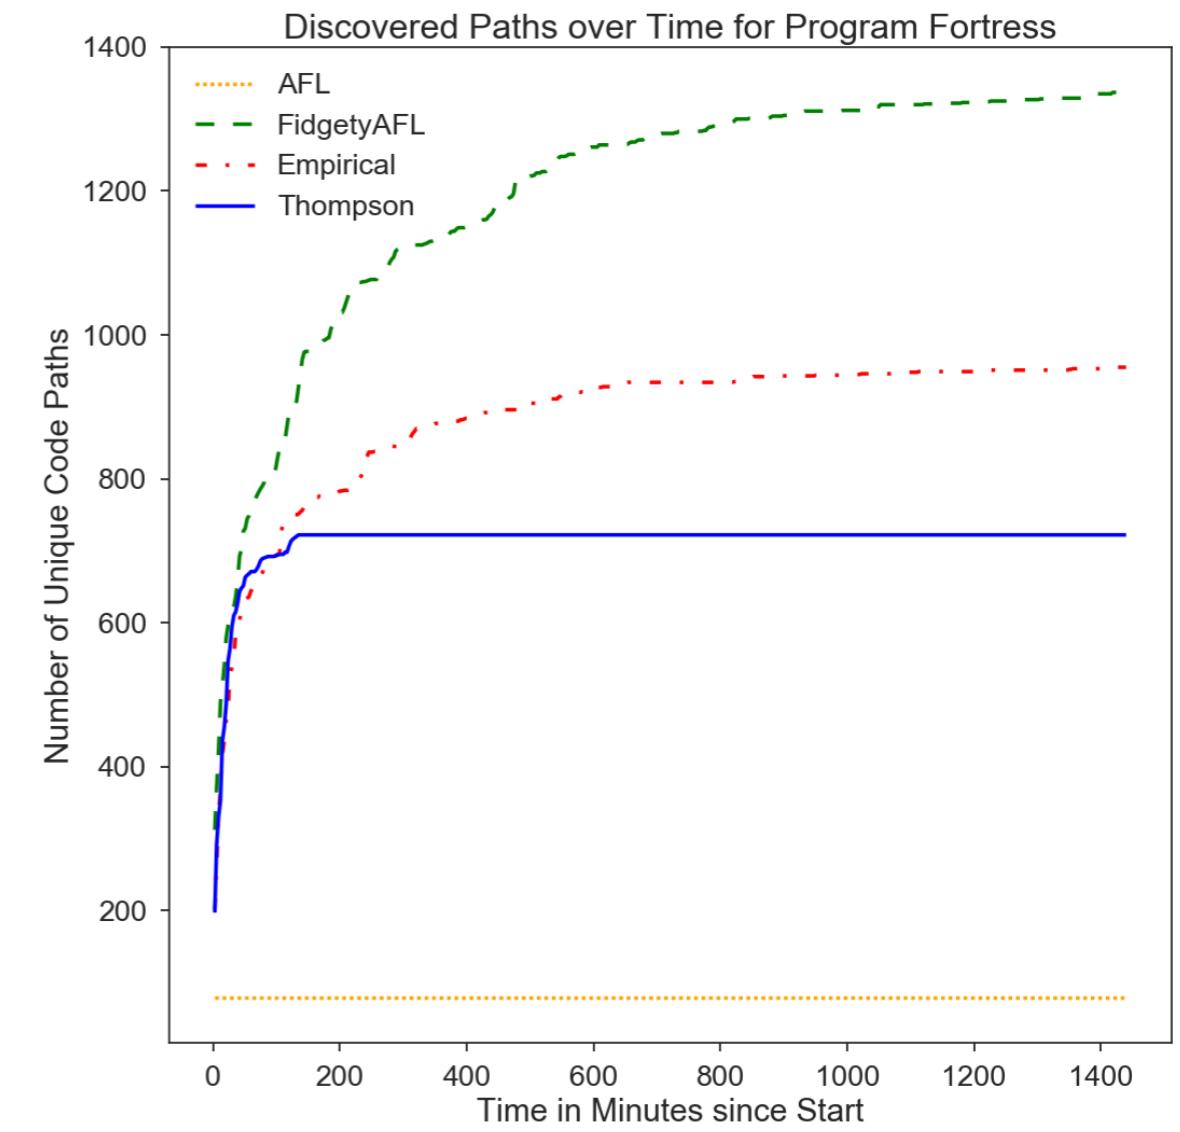
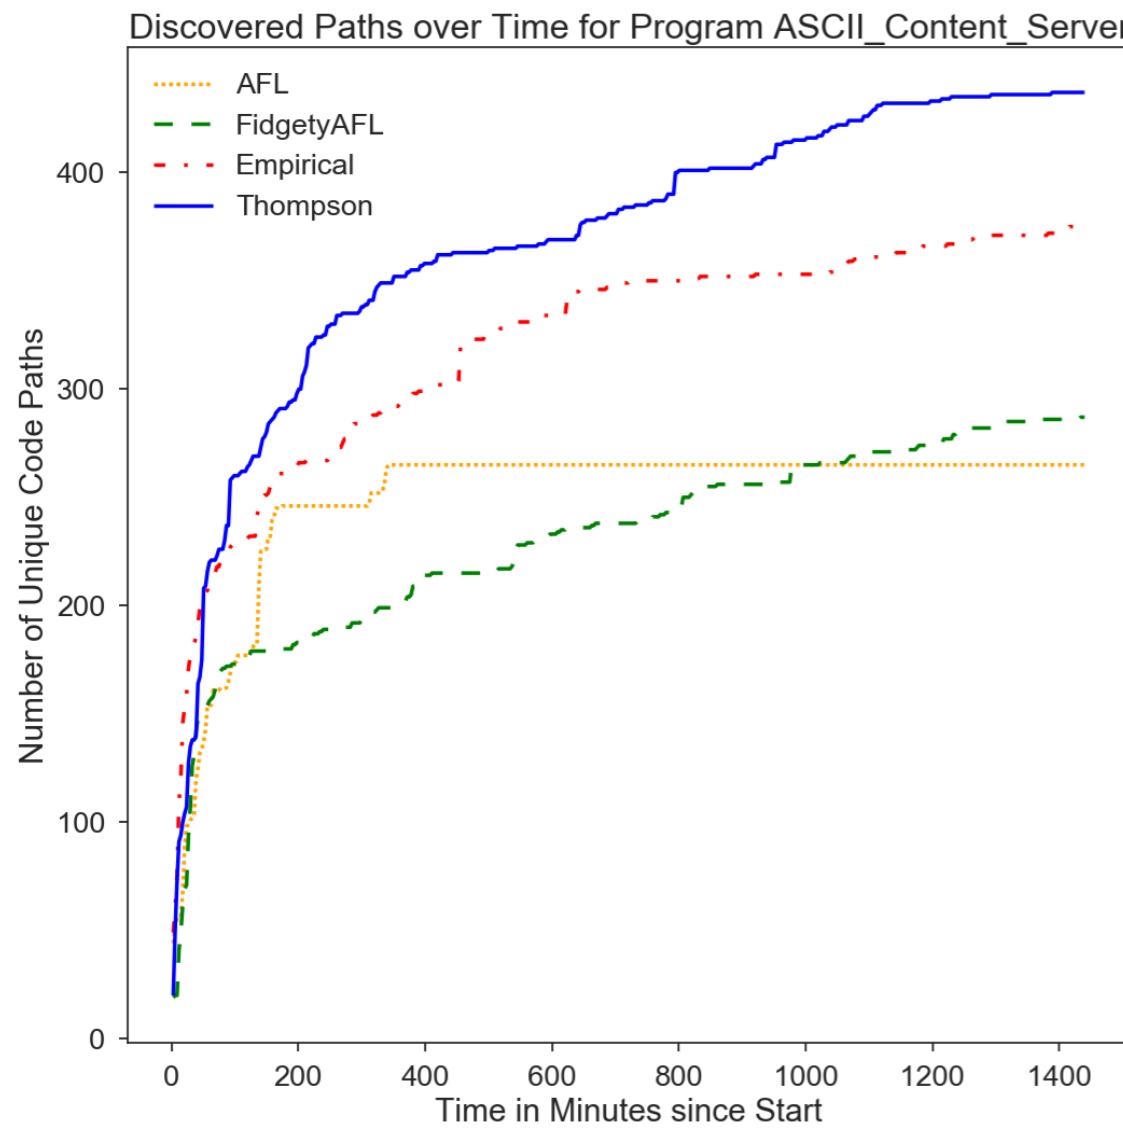
	base64	md5sum	uniq	who
AFL	117 ± 20	55 ± 2	13 ± 0	37 ± 1
FidgetyAFL	133 ± 10	340 ± 10	87 ± 1	372 ± 36
Empirical	134 ± 8	406 ± 22	80 ± 1	115 ± 9
Thompson	144 ± 14	405 ± 2	75 ± 2	106 ± 16

Number of unique code paths discovered

	base64	md5sum	uniq	who
AFL	15 ± 5	0 ± 0	0 ± 0	0 ± 0
FidgetyAFL	26 ± 9	4 ± 1	1 ± 1	201 ± 56
Empirical	22 ± 5	0 ± 0	0 ± 0	78 ± 17
Thompson	31 ± 8	1 ± 1	0 ± 0	106 ± 16

Number of unique crashes discovered

RESULTS - GRAPHS



FUTURE WORK

- Non-Stationary Distributions - need to adapt sampling
- Can we get Thompson Sampling to work together with other learned/data-driven heuristics?
 - Well...

	24 hr	Crashes	Wins / All
FairFuzz	0.88 ± 0.02	734	17
Thompson	0.95 ± 0.01	1287	49
<i>FairFuzz + Thompson</i>	<i>0.57 ± 0.03</i>	<i>245</i>	<i>1</i>

- Modeling Programs Directly - how inputs behave!
 - Talk to me after about entropy ranking!

SUMMARY

- Grey-box Mutational Fuzzing is good, but inefficient
 - Heuristics are non-optimal, can lead to redundant work
- Our contribution: improve grey-box fuzzing with data-driven learning!
 - Change distribution over mutators adaptively, via Thompson Sampling - focus on mutators that matter!
 - Results show huge gains over baselines, but not perfect

Questions: Email me @ sidd.karamchetti@gmail.com

Thank You!